

Decoding from Neural Text Generation Models

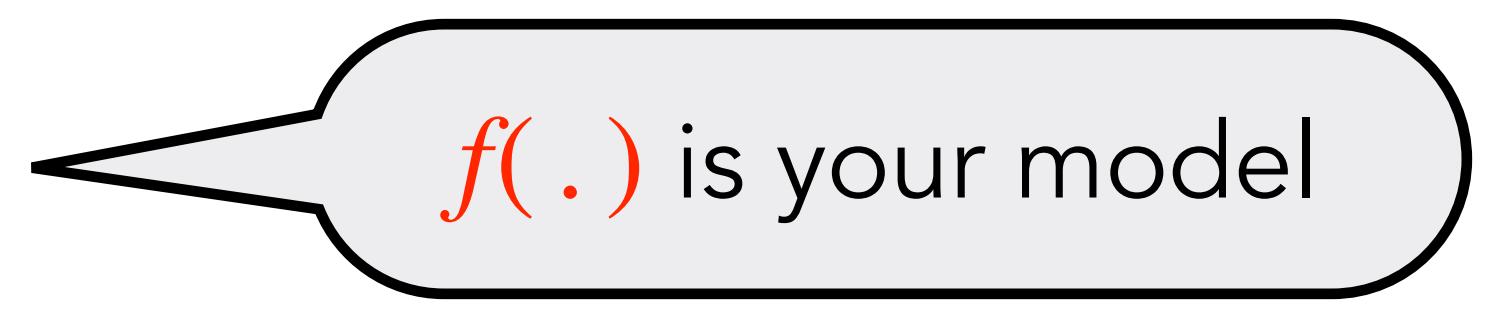
Antoine Bosselut



Generation Model Basics

1. At each time step, model computes a score o_n for each token in our vocabulary, $w_n \in V$

$$O_n = f(\{y\}_{<t})$$



$f(\cdot)$ is your model

Generation Model Basics

1. At each time step, model computes a score o_n for each token in our vocabulary, $w_n \in V$

$$O_n = f(\{y\}_{<t})$$

f(.) is your model

2. Compute a probability distribution over these scores (usually softmax)

P(.) is your distribution
over tokens

$$P(y_t = w_n | \{y\}_{<t}) = \frac{e^{o_n}}{\sum_{m=1}^M e^{o_m}}$$

Generation Model Basics

1. At each time step, model computes a score o_n for each token in our vocabulary, $w_n \in V$

$$O_n = f(\{y\}_{<t})$$

f(.) is your model

2. Compute a probability distribution over these scores (usually softmax)

P(.) is your distribution
over tokens

$$P(y_t = w_n | \{y\}_{<t}) = \frac{e^{o_n}}{\sum_{m=1}^M e^{o_m}}$$

3. Define a function to select a token from this distribution

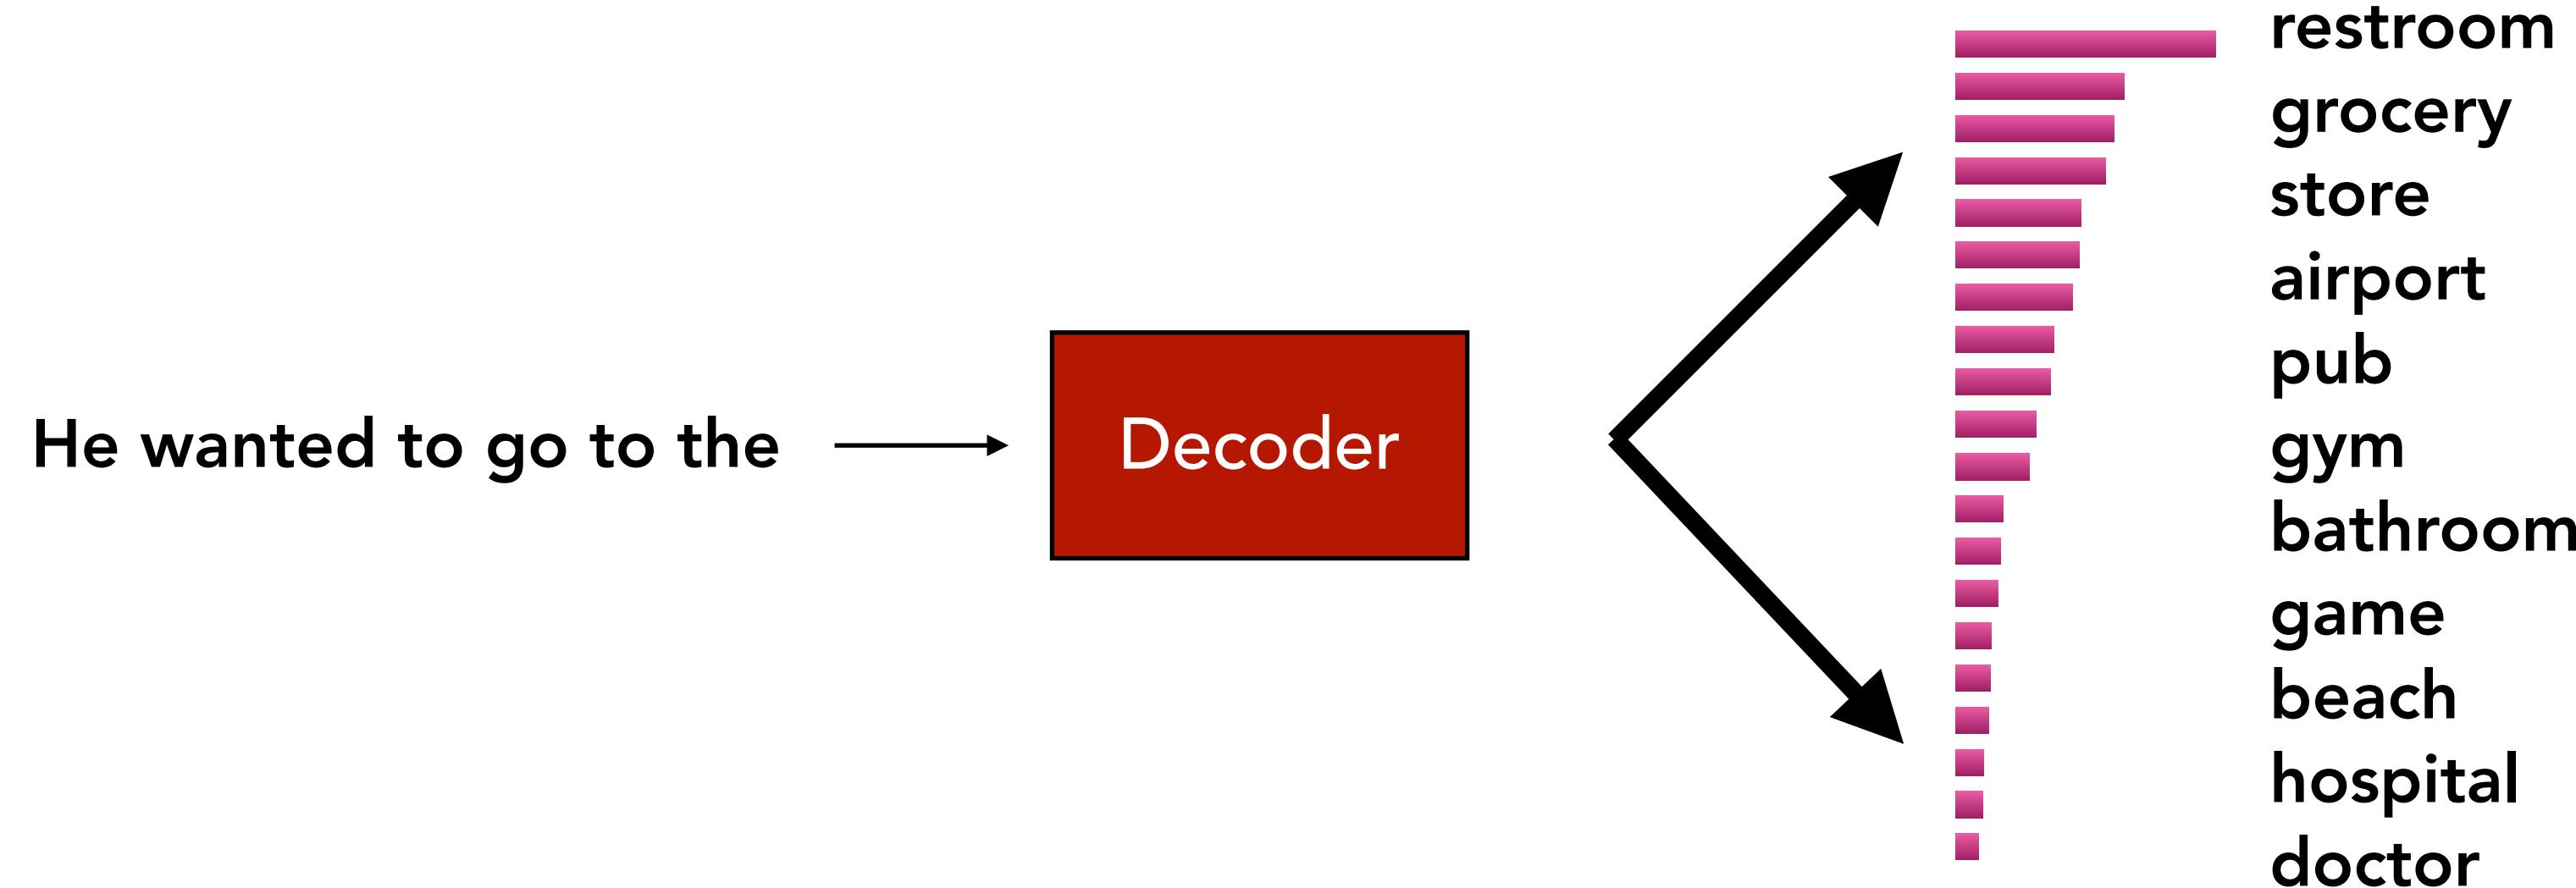
$$\hat{y}_t = g(P(y_t | \{y\}_{<t}))$$

g(.) is your decoding
algorithm

Simplest approach: Argmax Decoding

- g = select the token with the highest probability:

$$\hat{y}_t = \underset{w \in V}{\operatorname{argmax}} P(y_t = w | \{y\}_{<t})$$



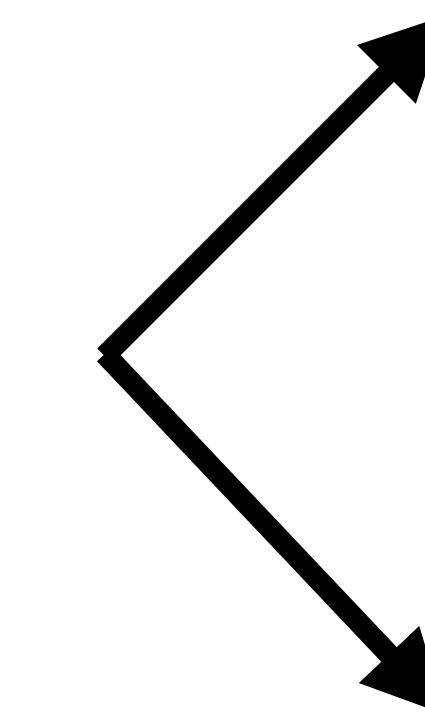
Simplest approach: Argmax Decoding

- g = select the token with the highest probability:

$$\hat{y}_t = \underset{w \in V}{\operatorname{argmax}} P(y_t = w | \{y\}_{<t})$$

Select highest scoring token

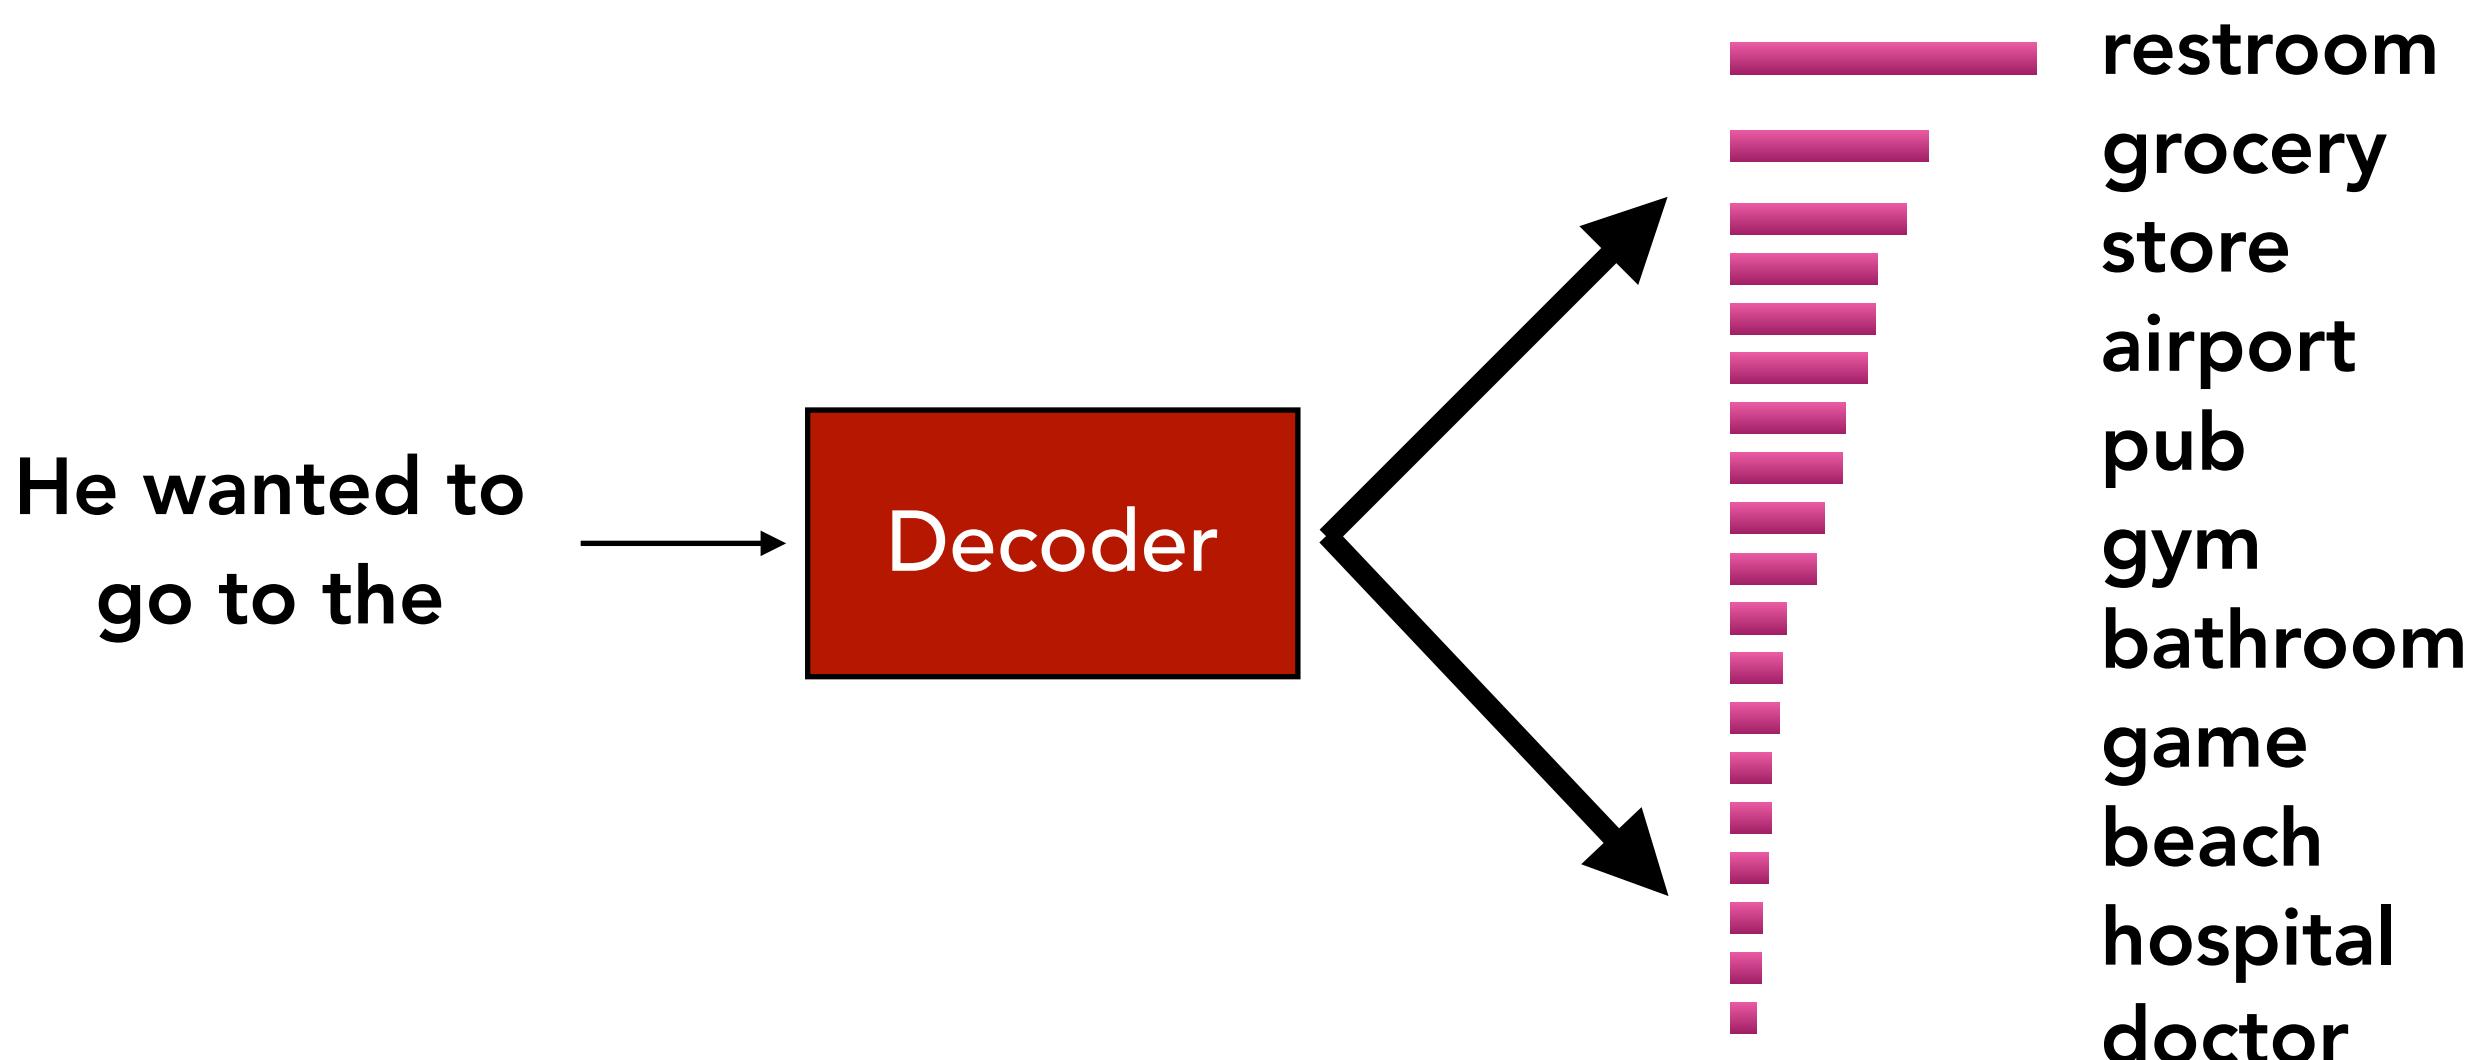
He wanted to go to the →



restroom
grocery
store
airport
pub
gym
bathroom
game
beach
hospital
doctor

Maybe we need more options: Beam Search

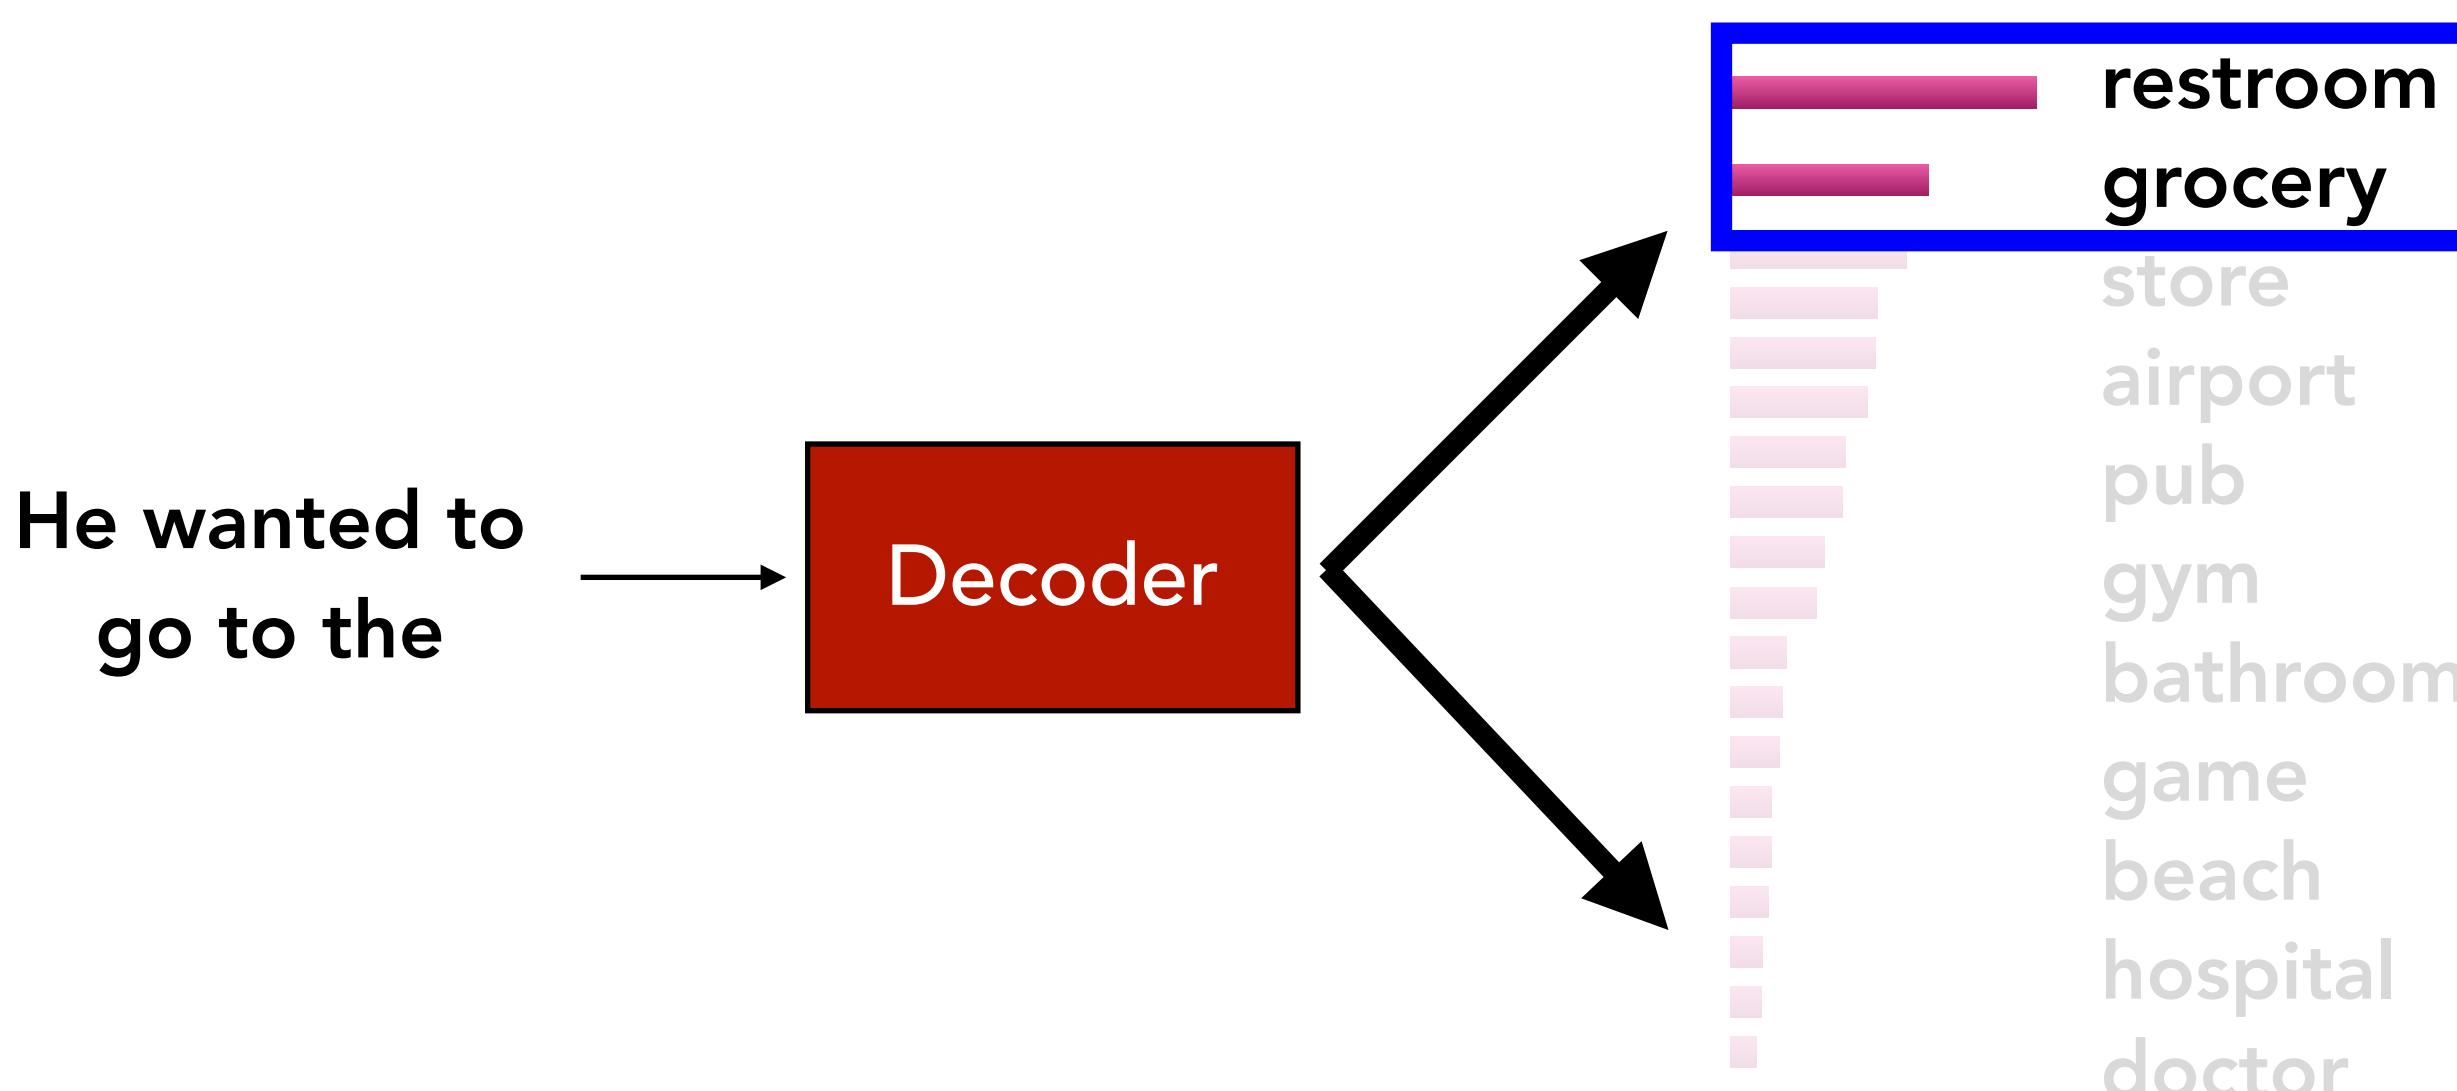
- g = cache b paths for two steps



Maybe we need more options: Beam Search

- g = cache b paths for two steps

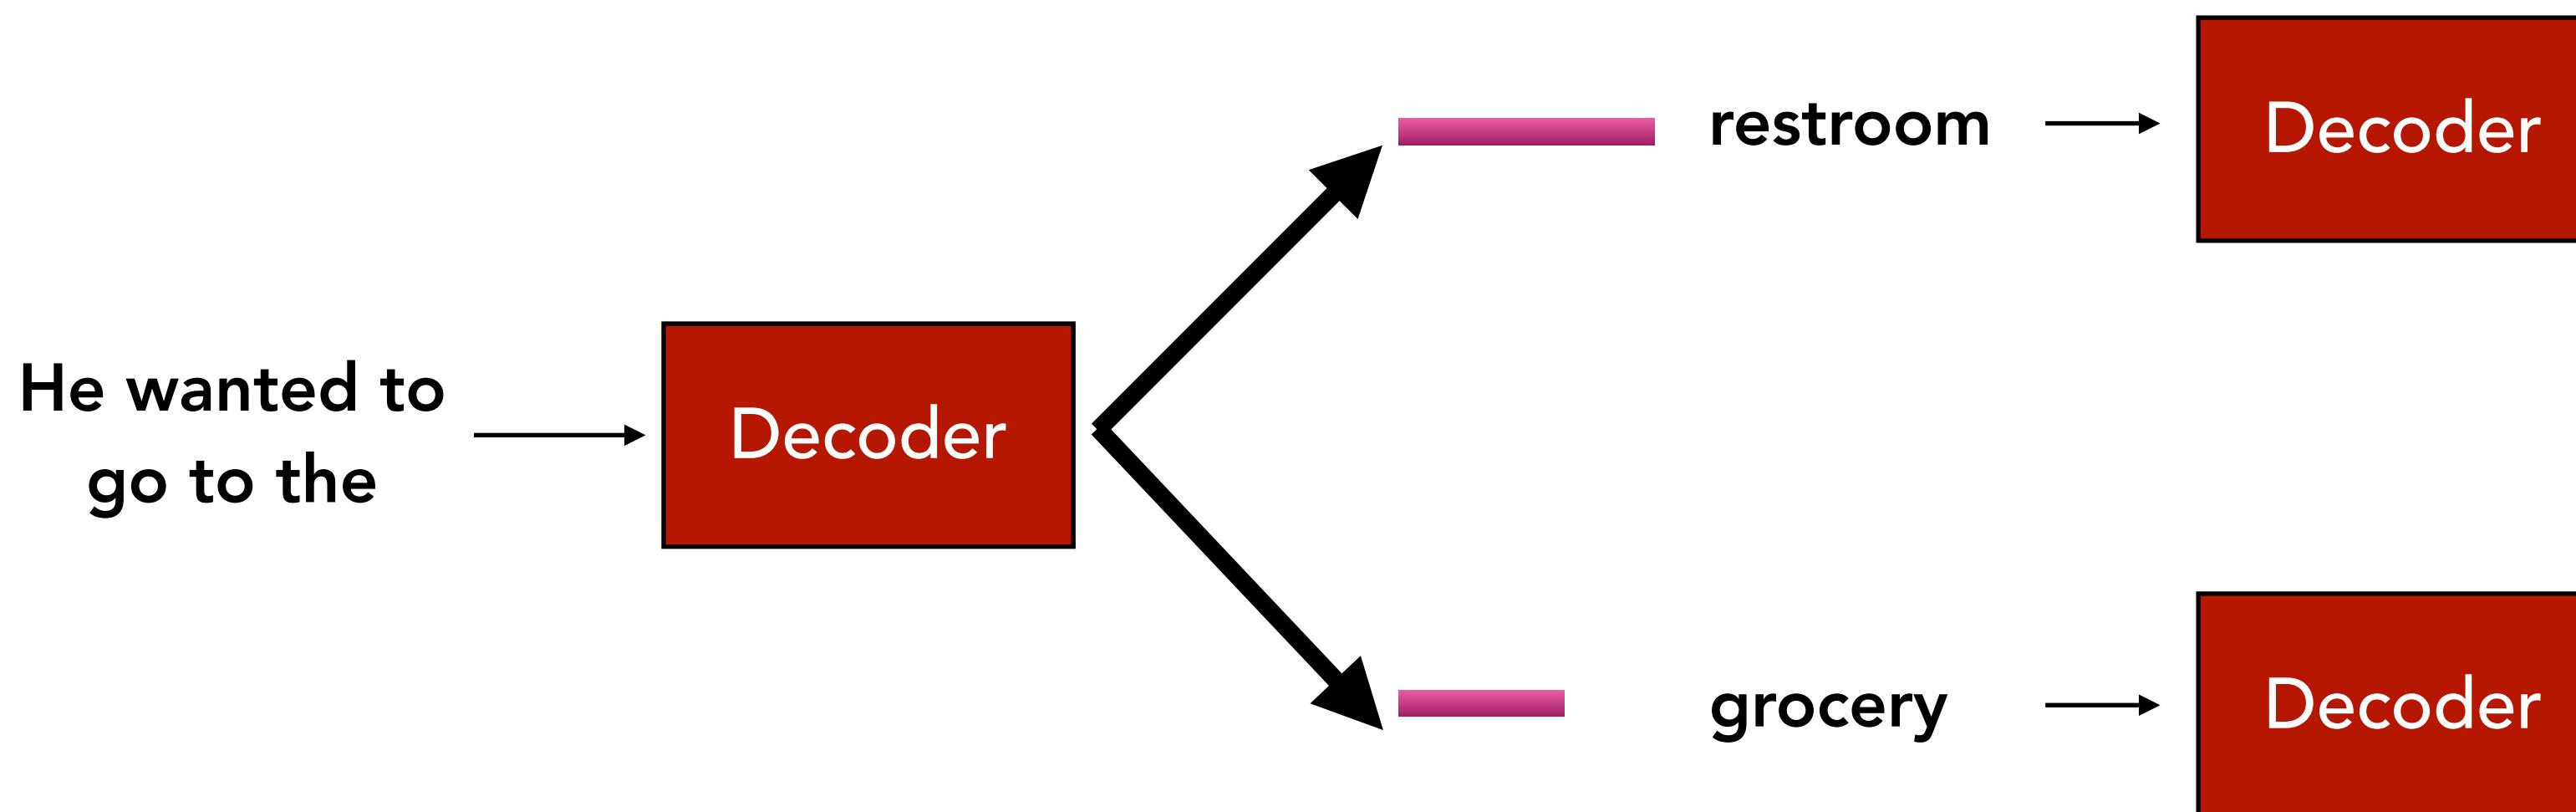
If $b = 2$, select
top two tokens



Maybe we need more options: Beam Search

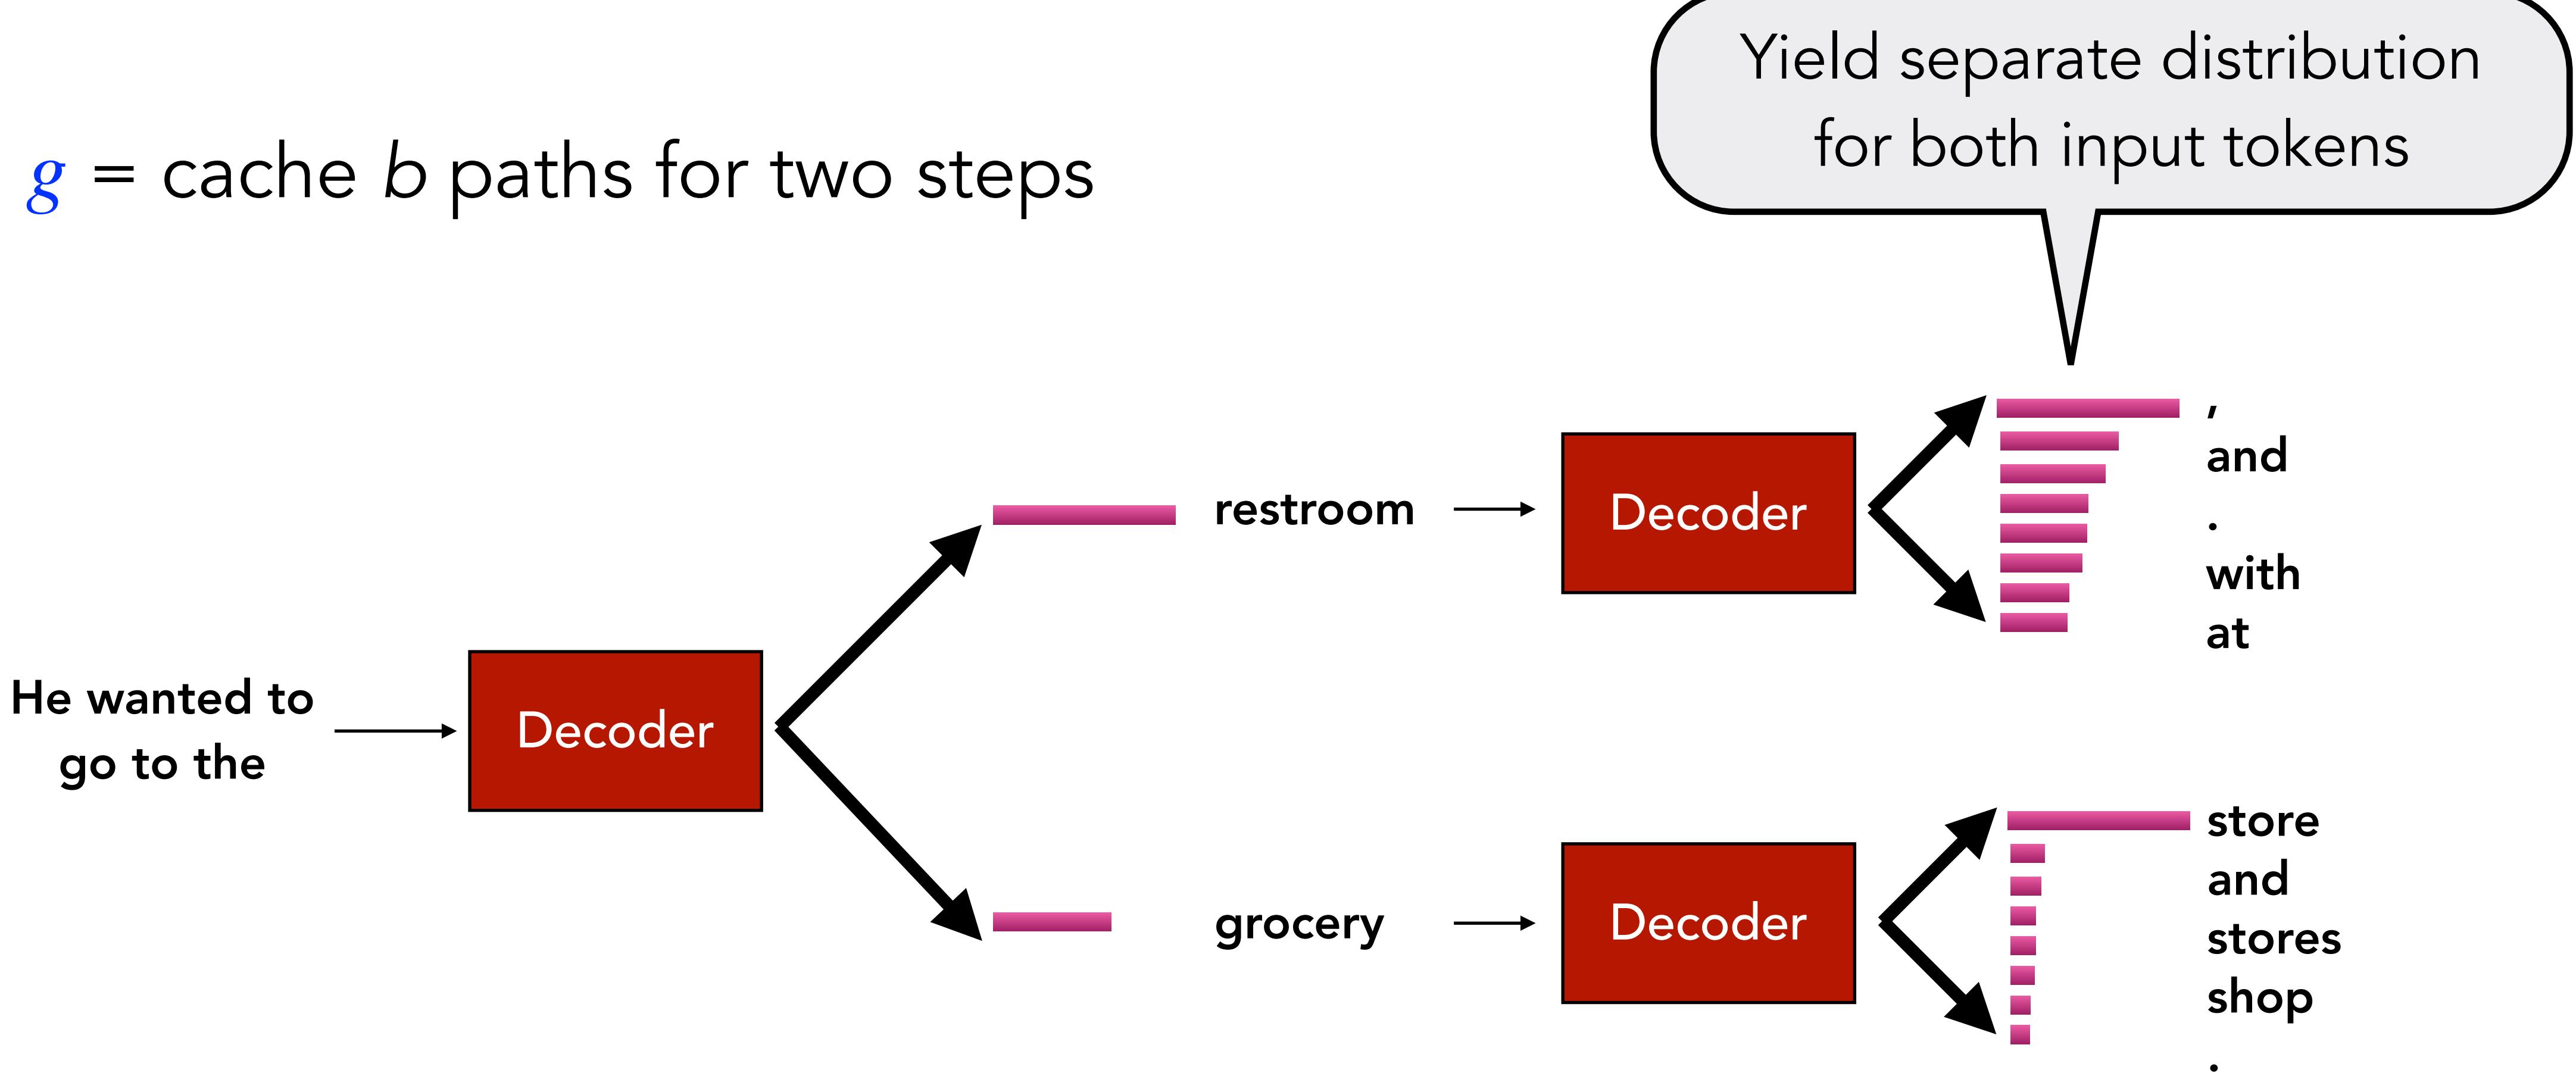
- g = cache b paths for two steps

Use them both as inputs to
the decoder at next step



Maybe we need more options: Beam Search

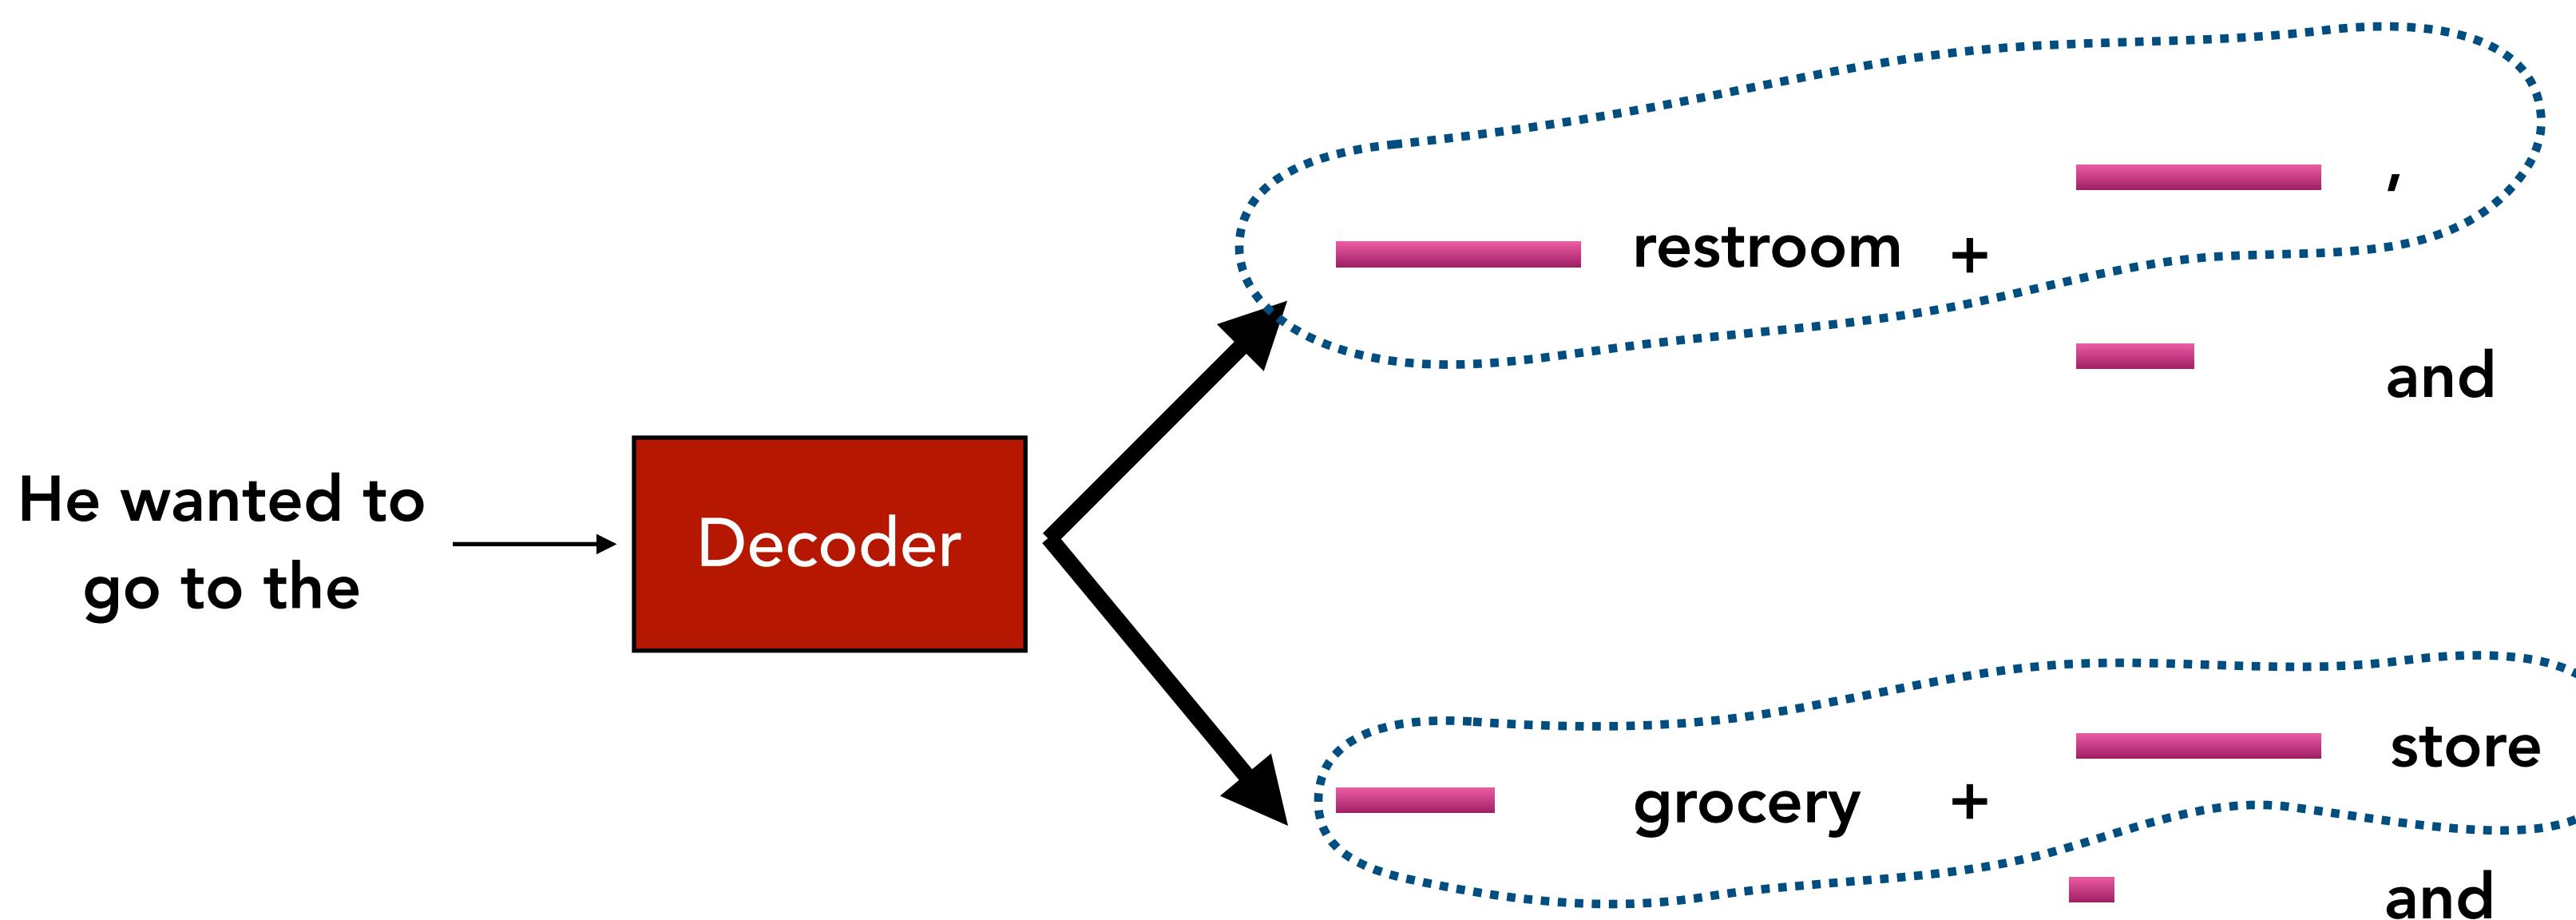
- g = cache b paths for two steps



Maybe we need more options: Beam Search

- g = cache b paths for two steps

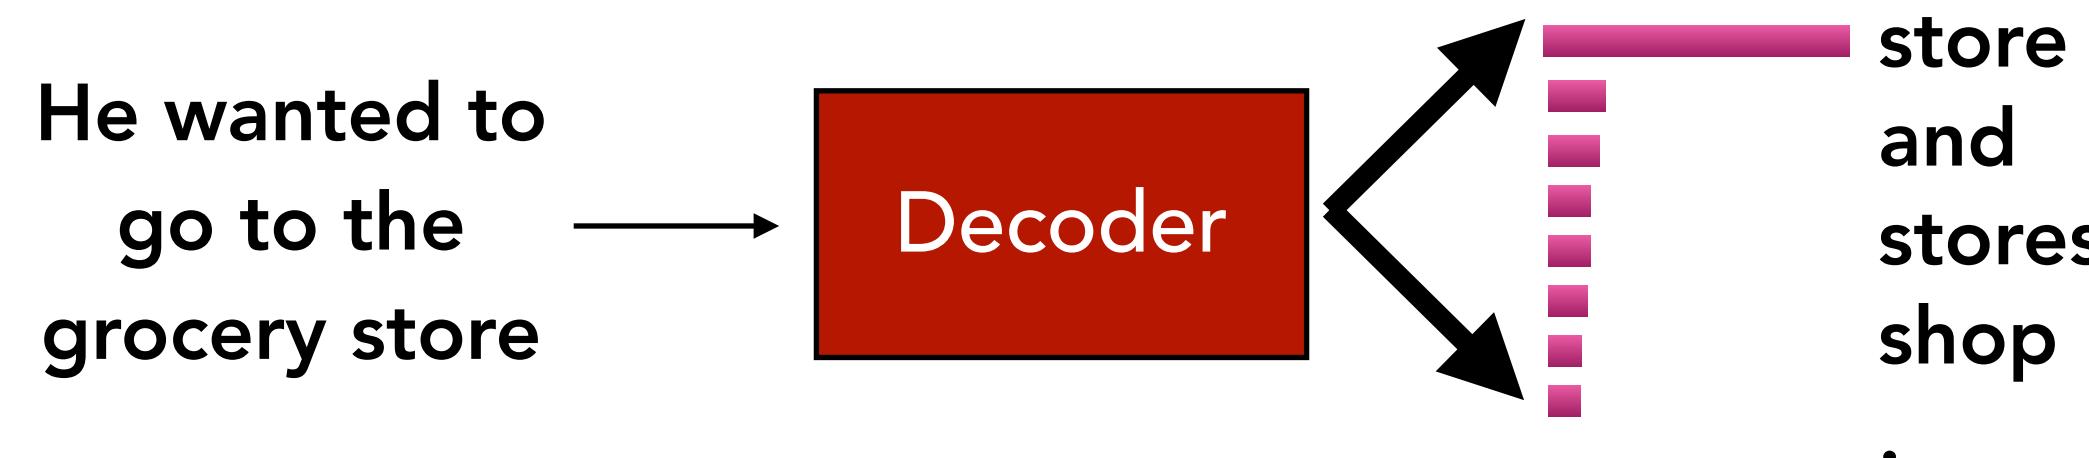
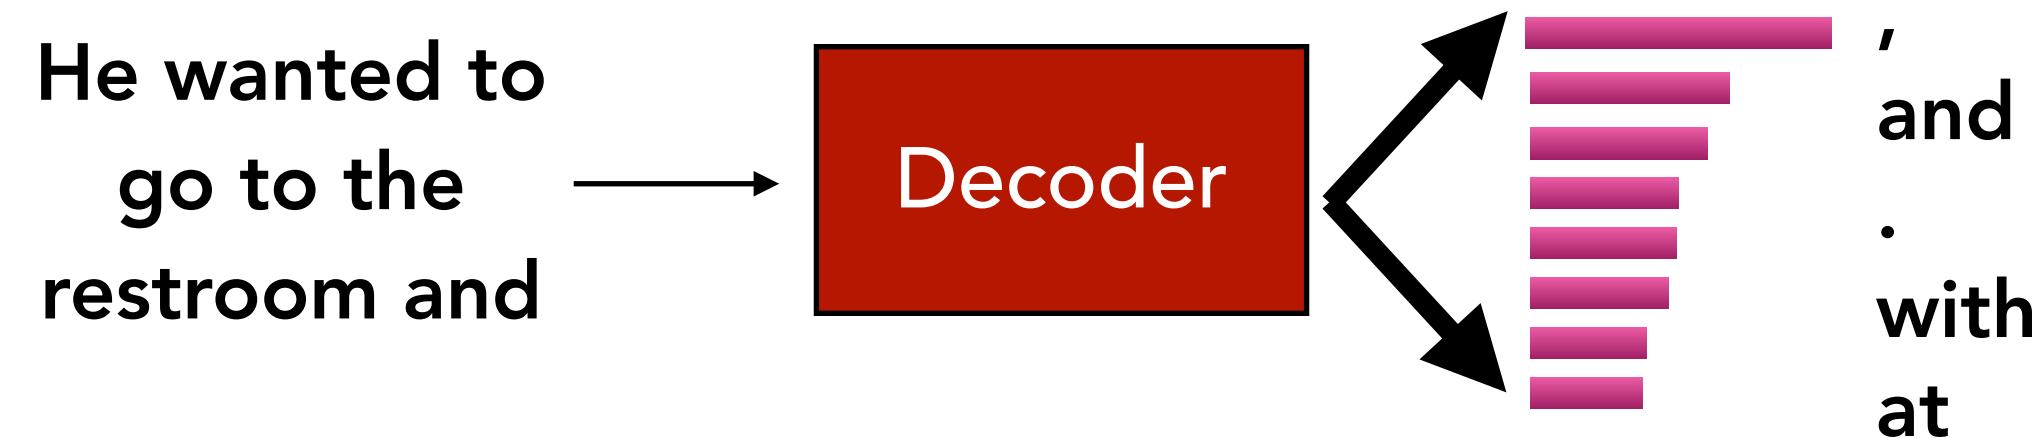
Select top b sequence continuations
across both distributions



Maybe we need more options: Beam Search

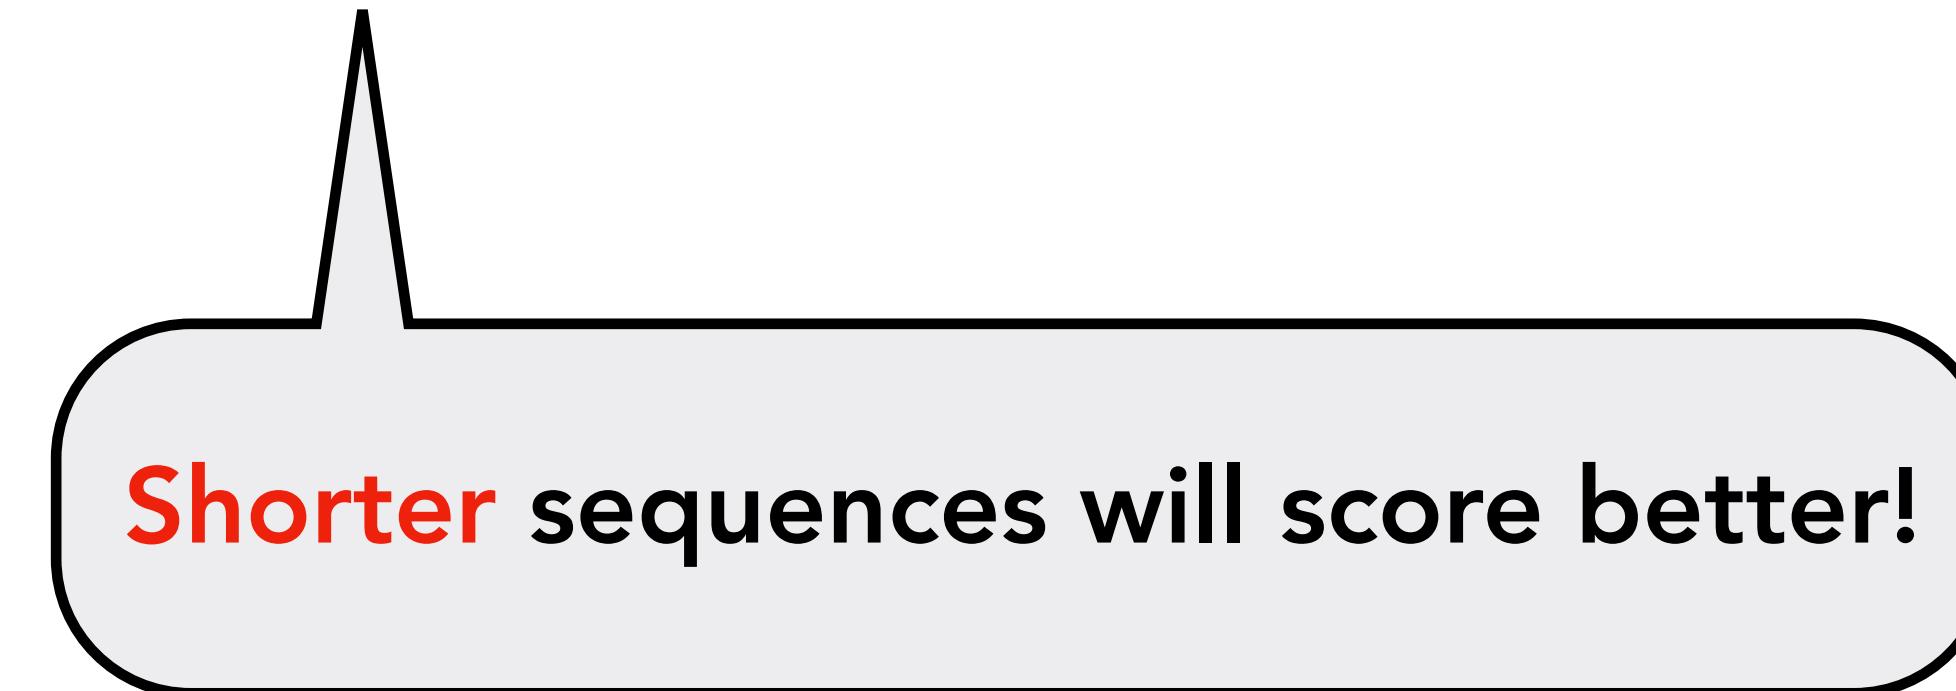
- g = cache b paths for two steps

Repeat!



Does this penalize longer sequences?

$$s(Y) = \sum_{t=1}^T \log P(y_t | \{y\}_{<t})$$



Does this penalize longer sequences?

- **Solution:** Normalize by token length of sequence

$$s(Y) = \frac{1}{T} \sum_{t=1}^T \log P(y_t | \{y\}_{<t})$$

- **Solution:** Normalize by token length relative to reference sequence

$$s(Y) = \frac{1}{lp(Y)} \sum_{t=1}^T \log P(y_t | \{y\}_{<t}) \quad lp(Y) = \frac{(5 + |Y|)^\alpha}{(5 + 1)^\alpha}$$

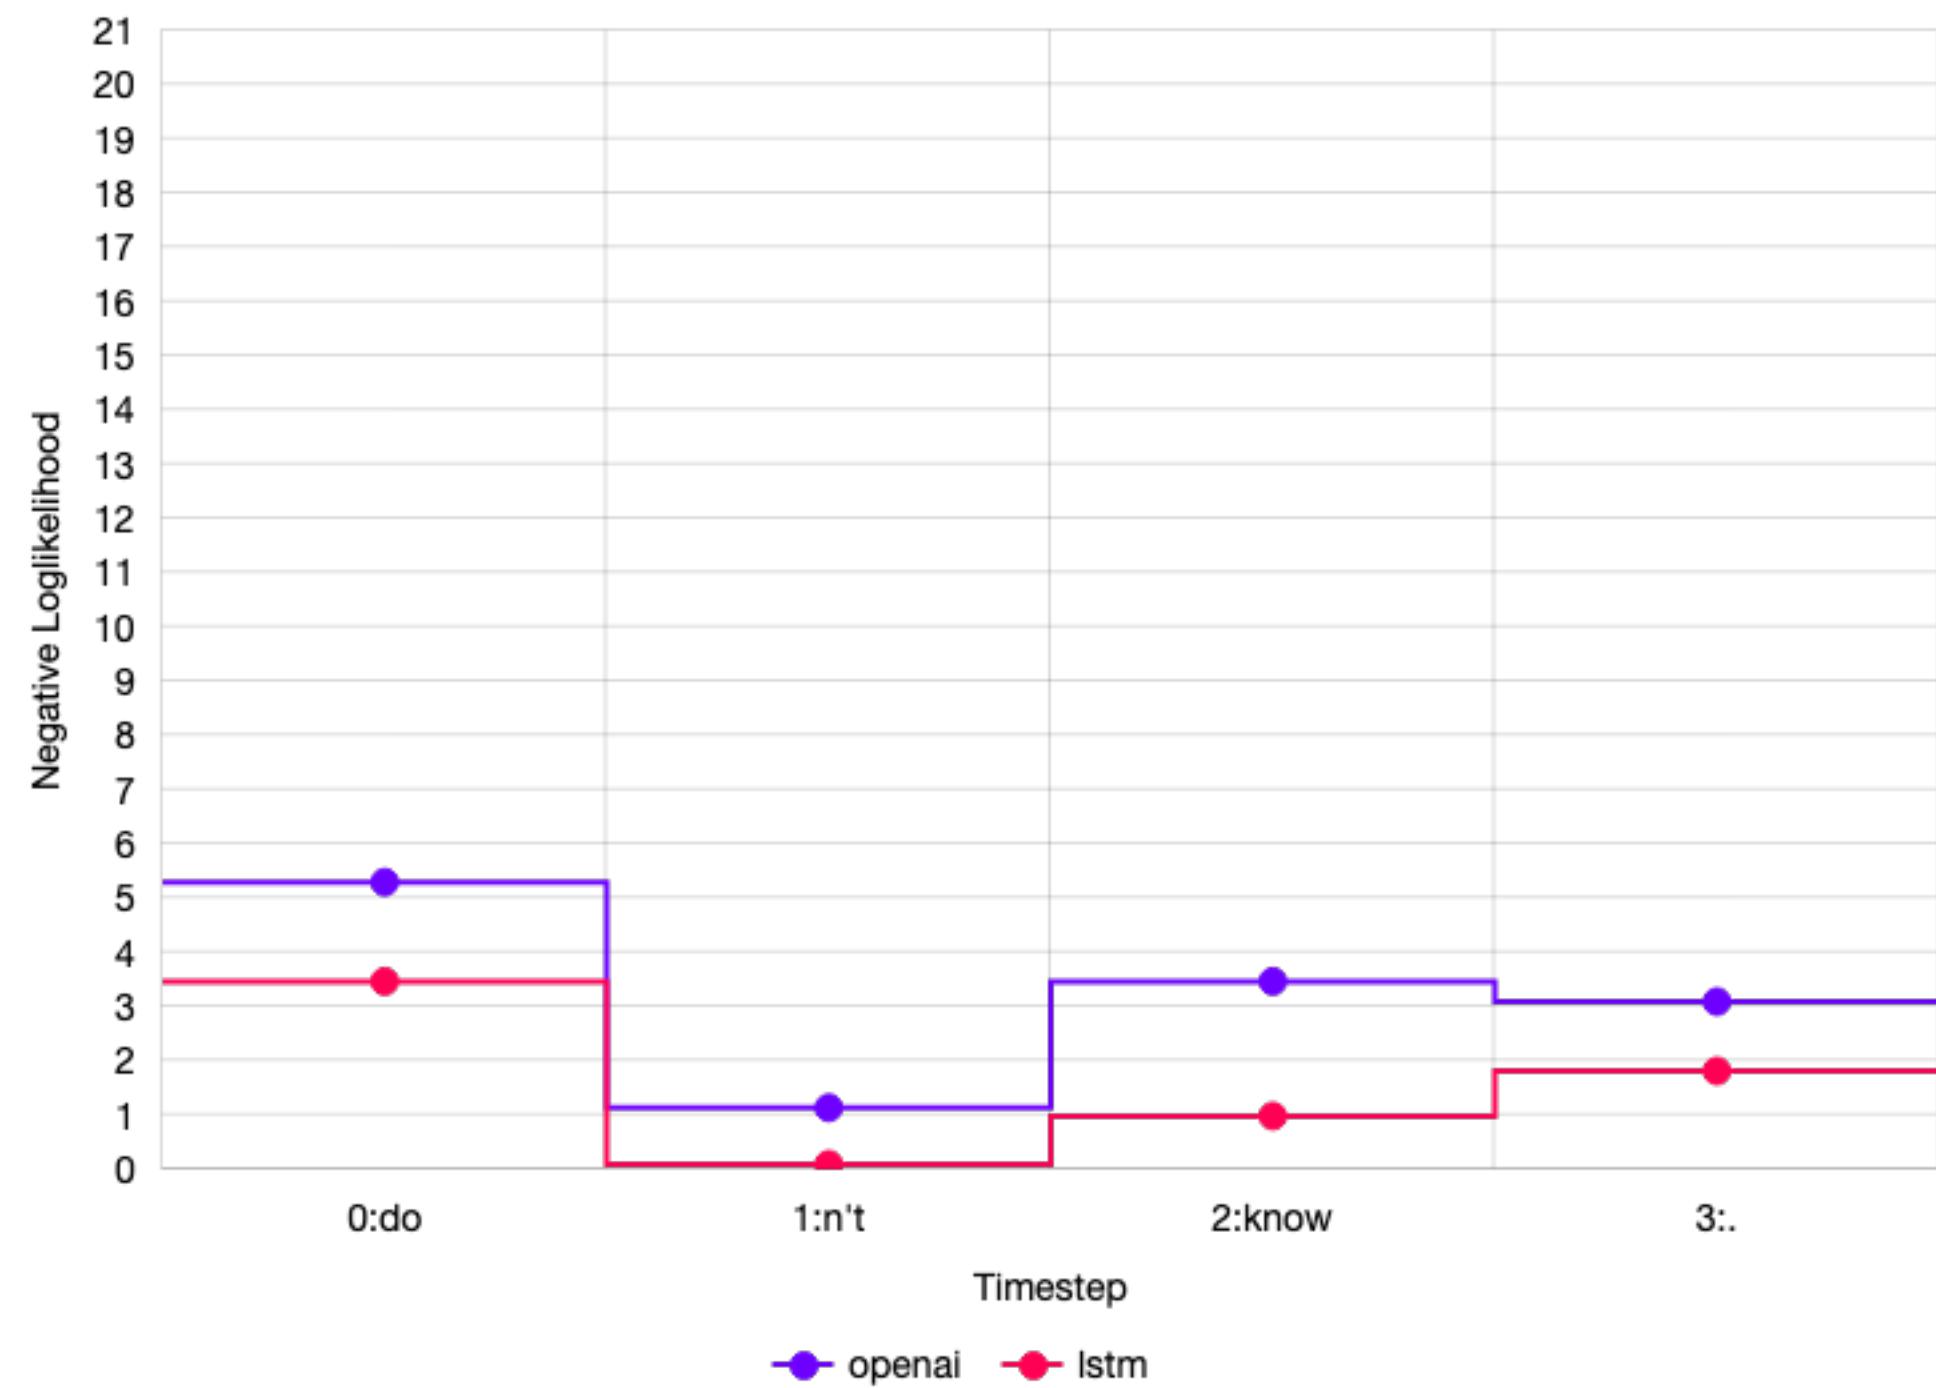
Beam search gets repetitive and repetitive

Context: In a shocking finding, scientist discovered a herd of unicorns living in a remote, previously unexplored valley, in the Andes Mountains. Even more surprising to the researchers was the fact that the unicorns spoke perfect English.

Continuation: The study, published in the Proceedings of the National Academy of Sciences of the United States of America (PNAS), was conducted by researchers from the **Universidad Nacional Autónoma de México (UNAM)** and **the Universidad Nacional Autónoma de México (UNAM/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México...)**

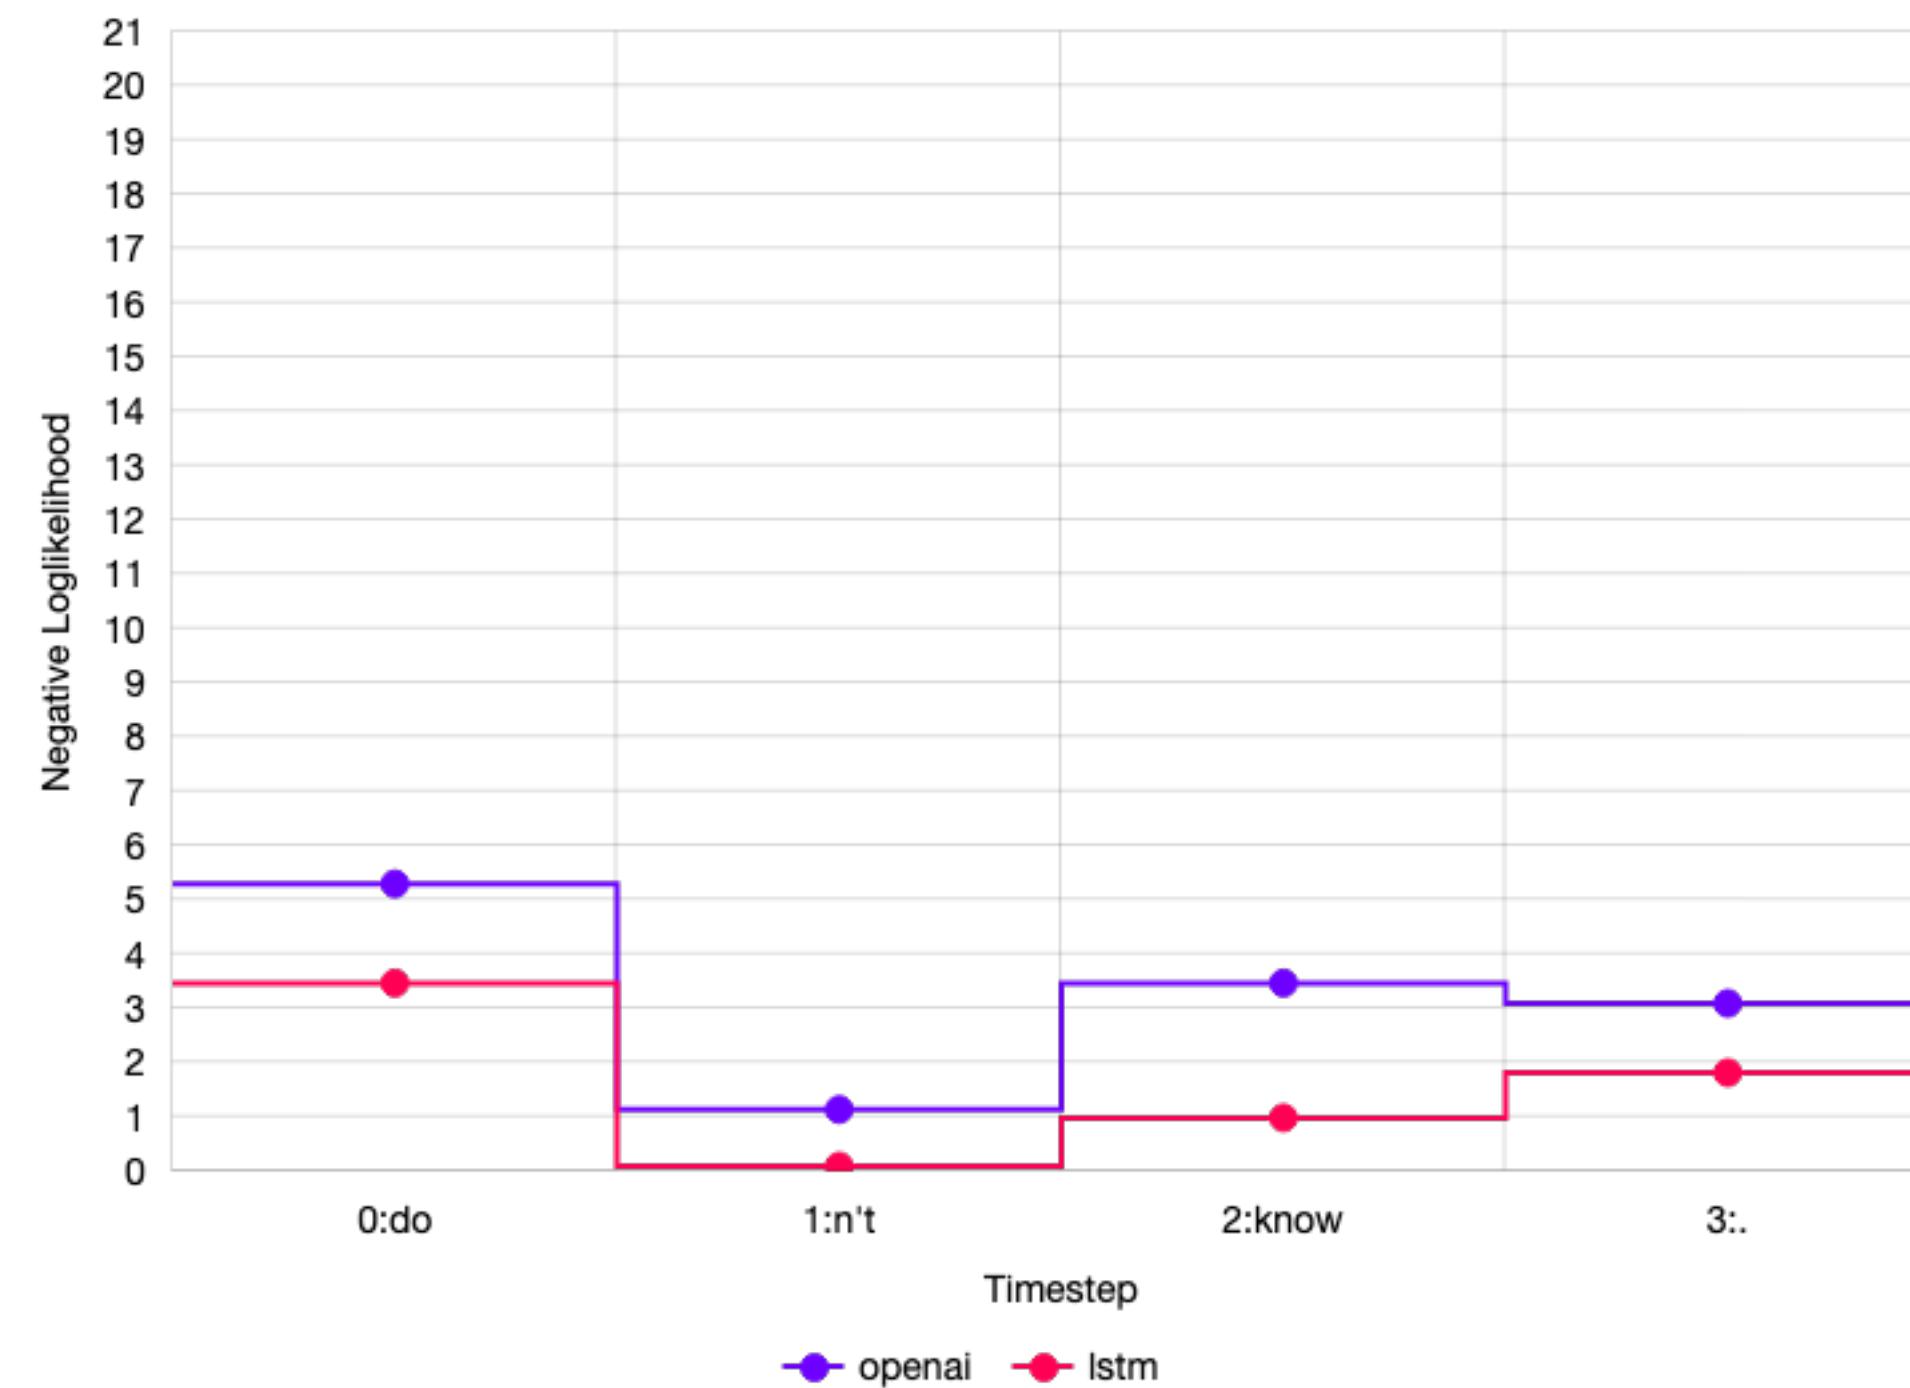
Why does this happen?

I don't know.

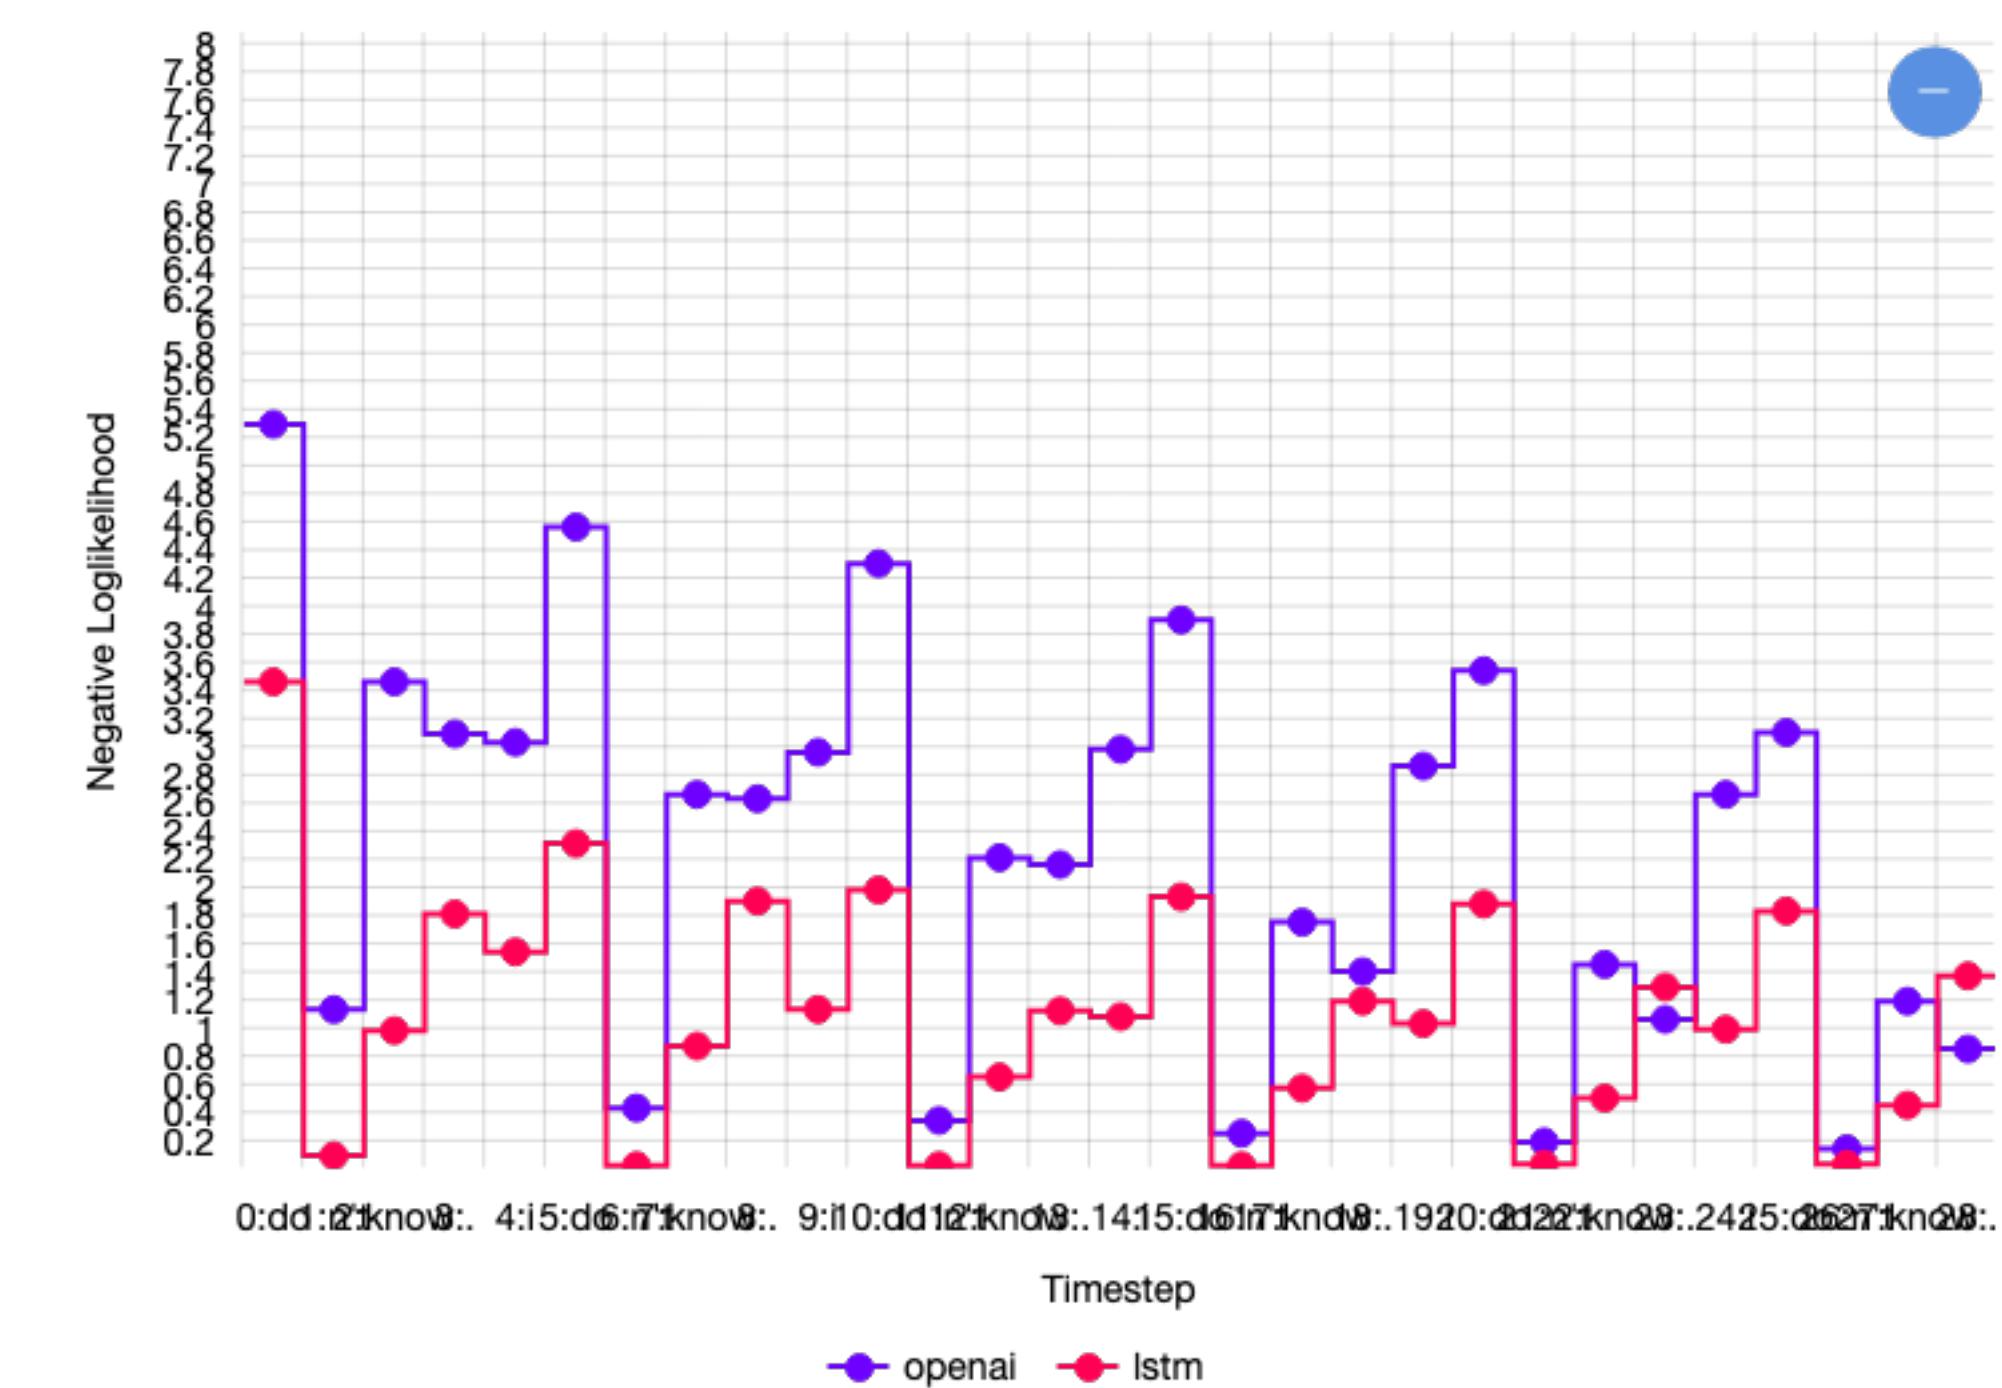


Artifact of Maximum Likelihood Training

I don't know.

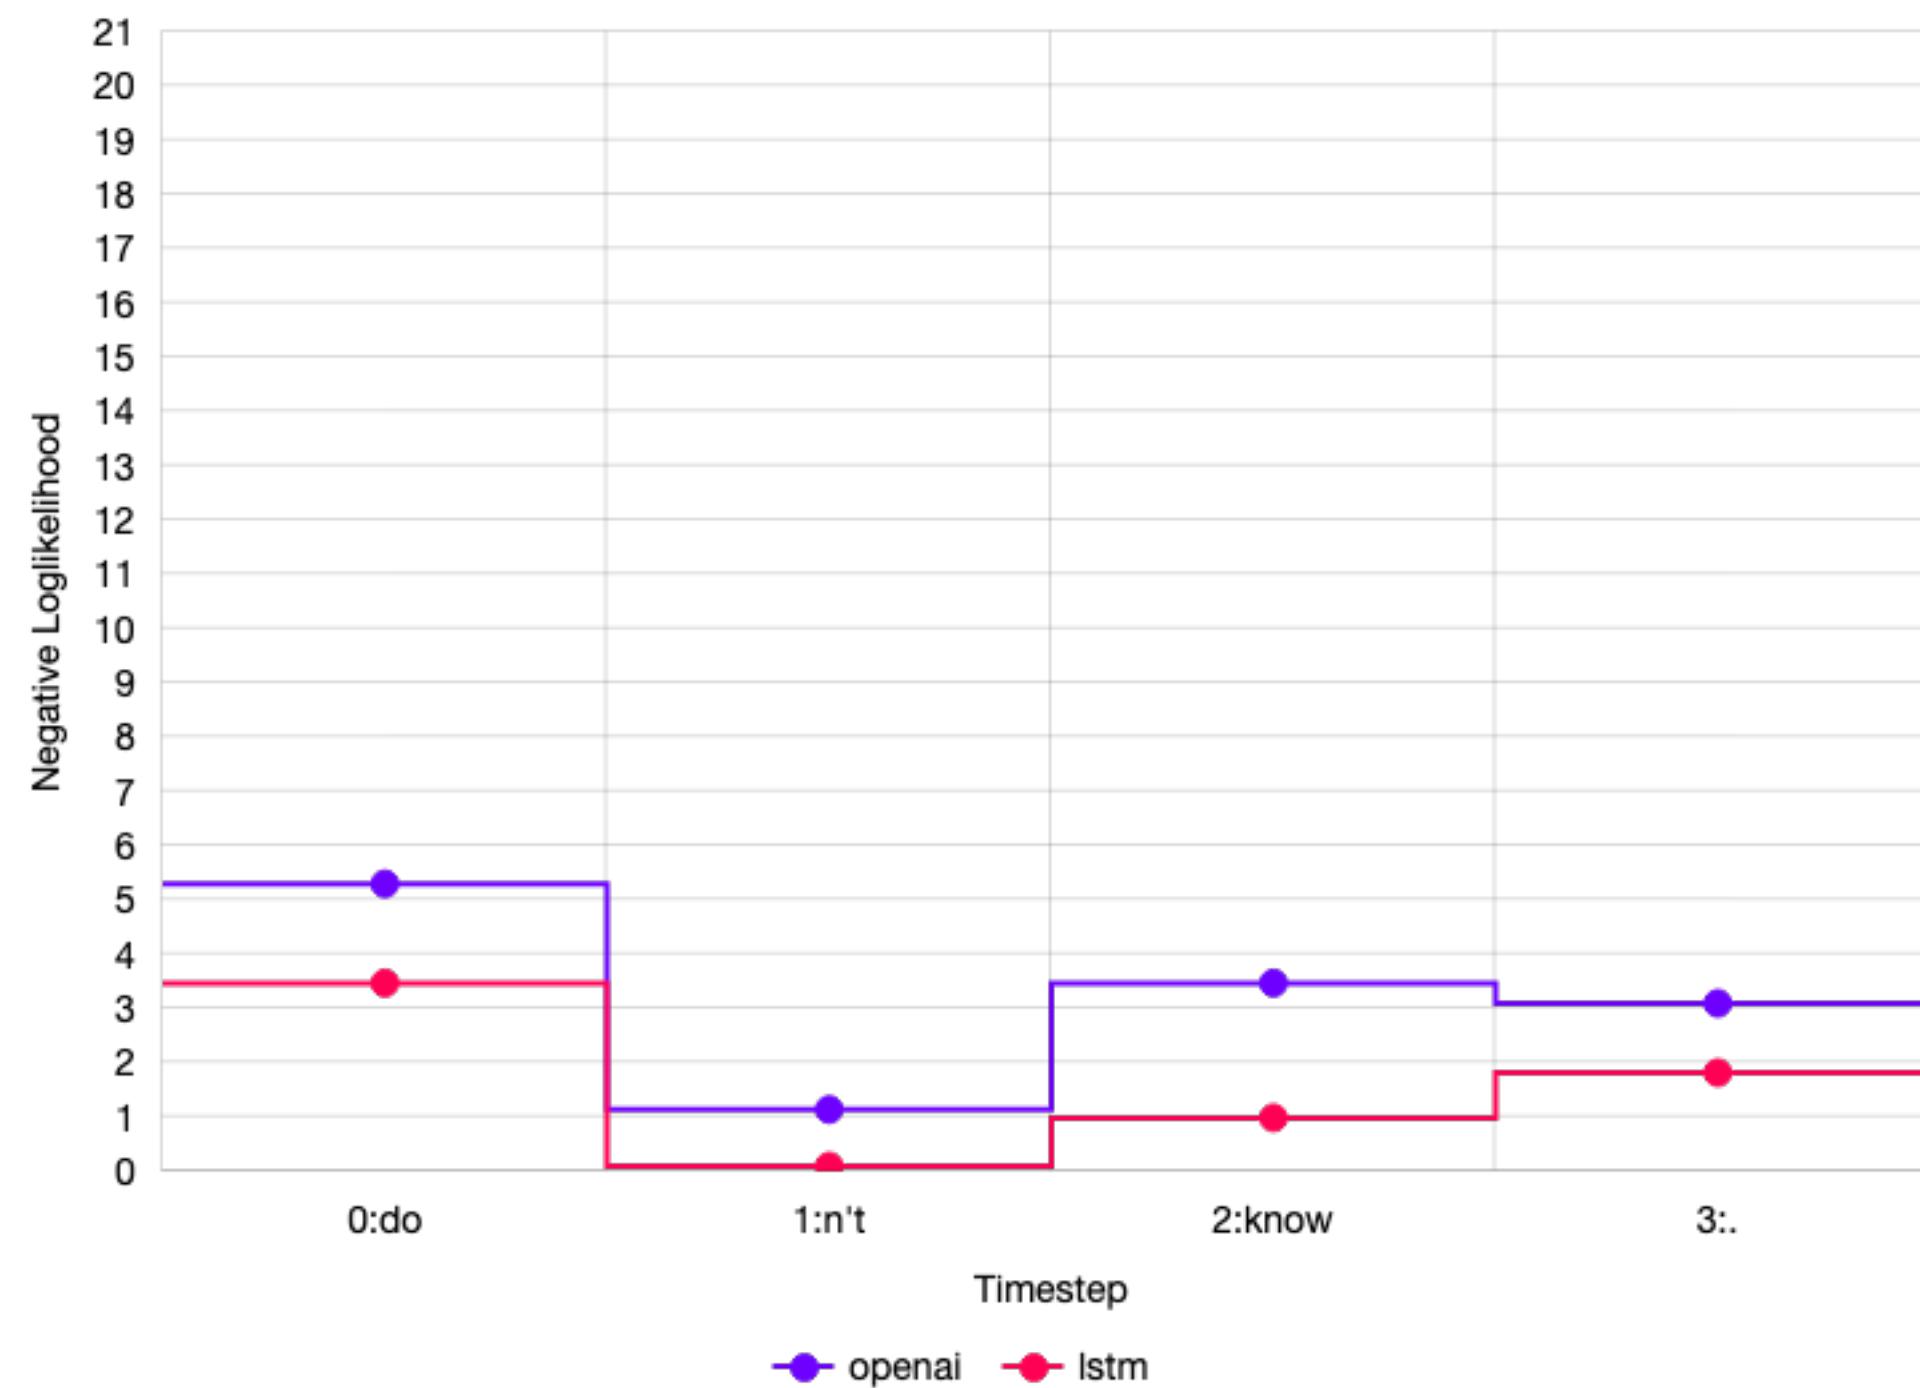


I don't know. I don't know.

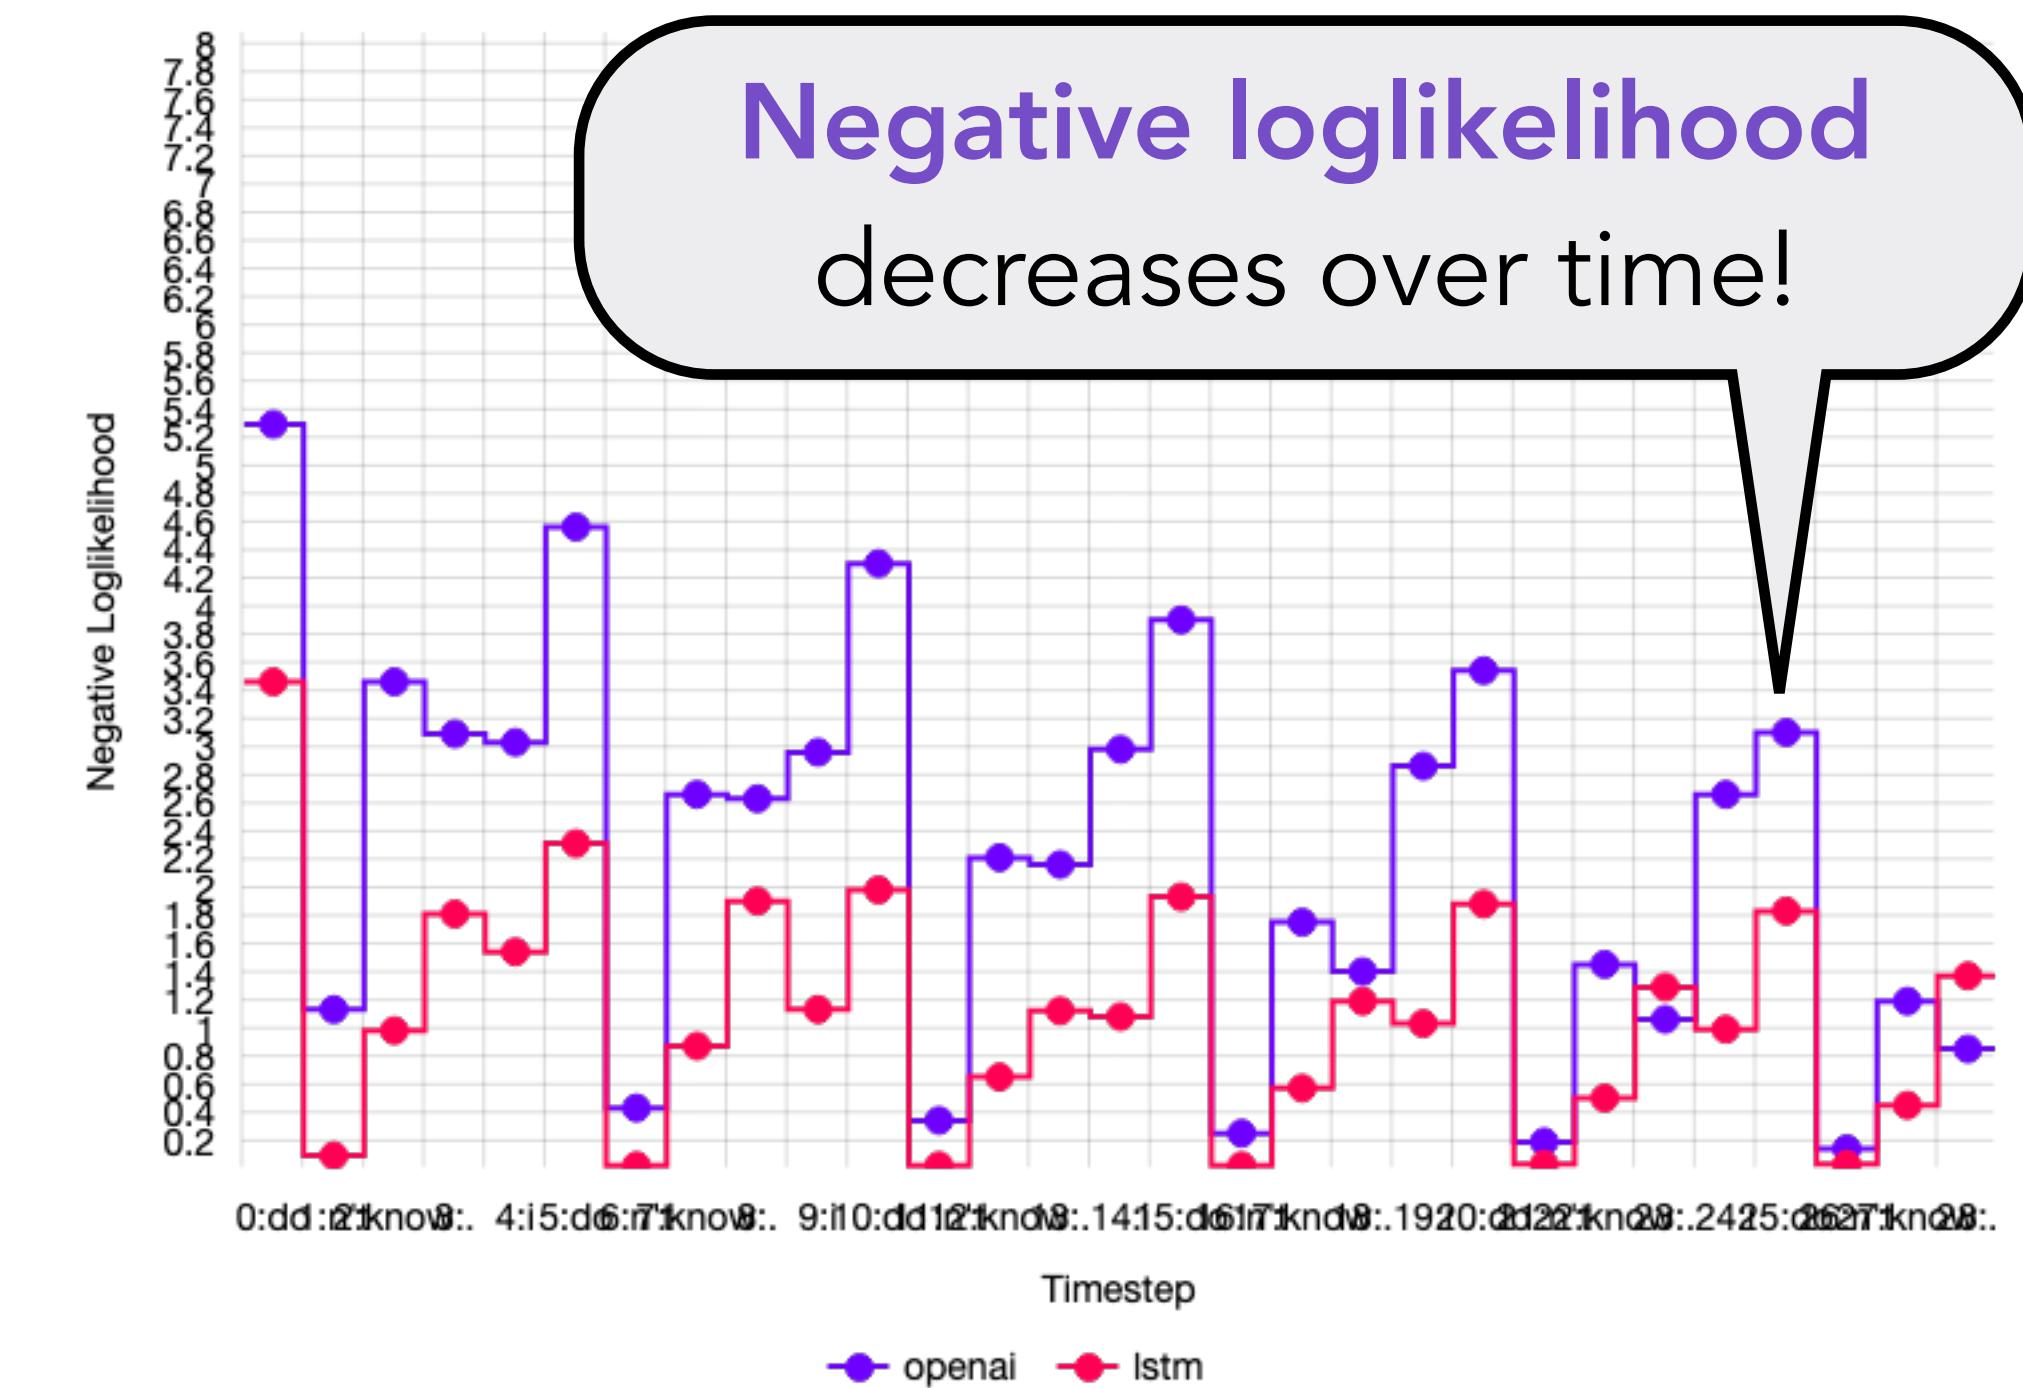


Artifact of Maximum Likelihood Training

I don't know.

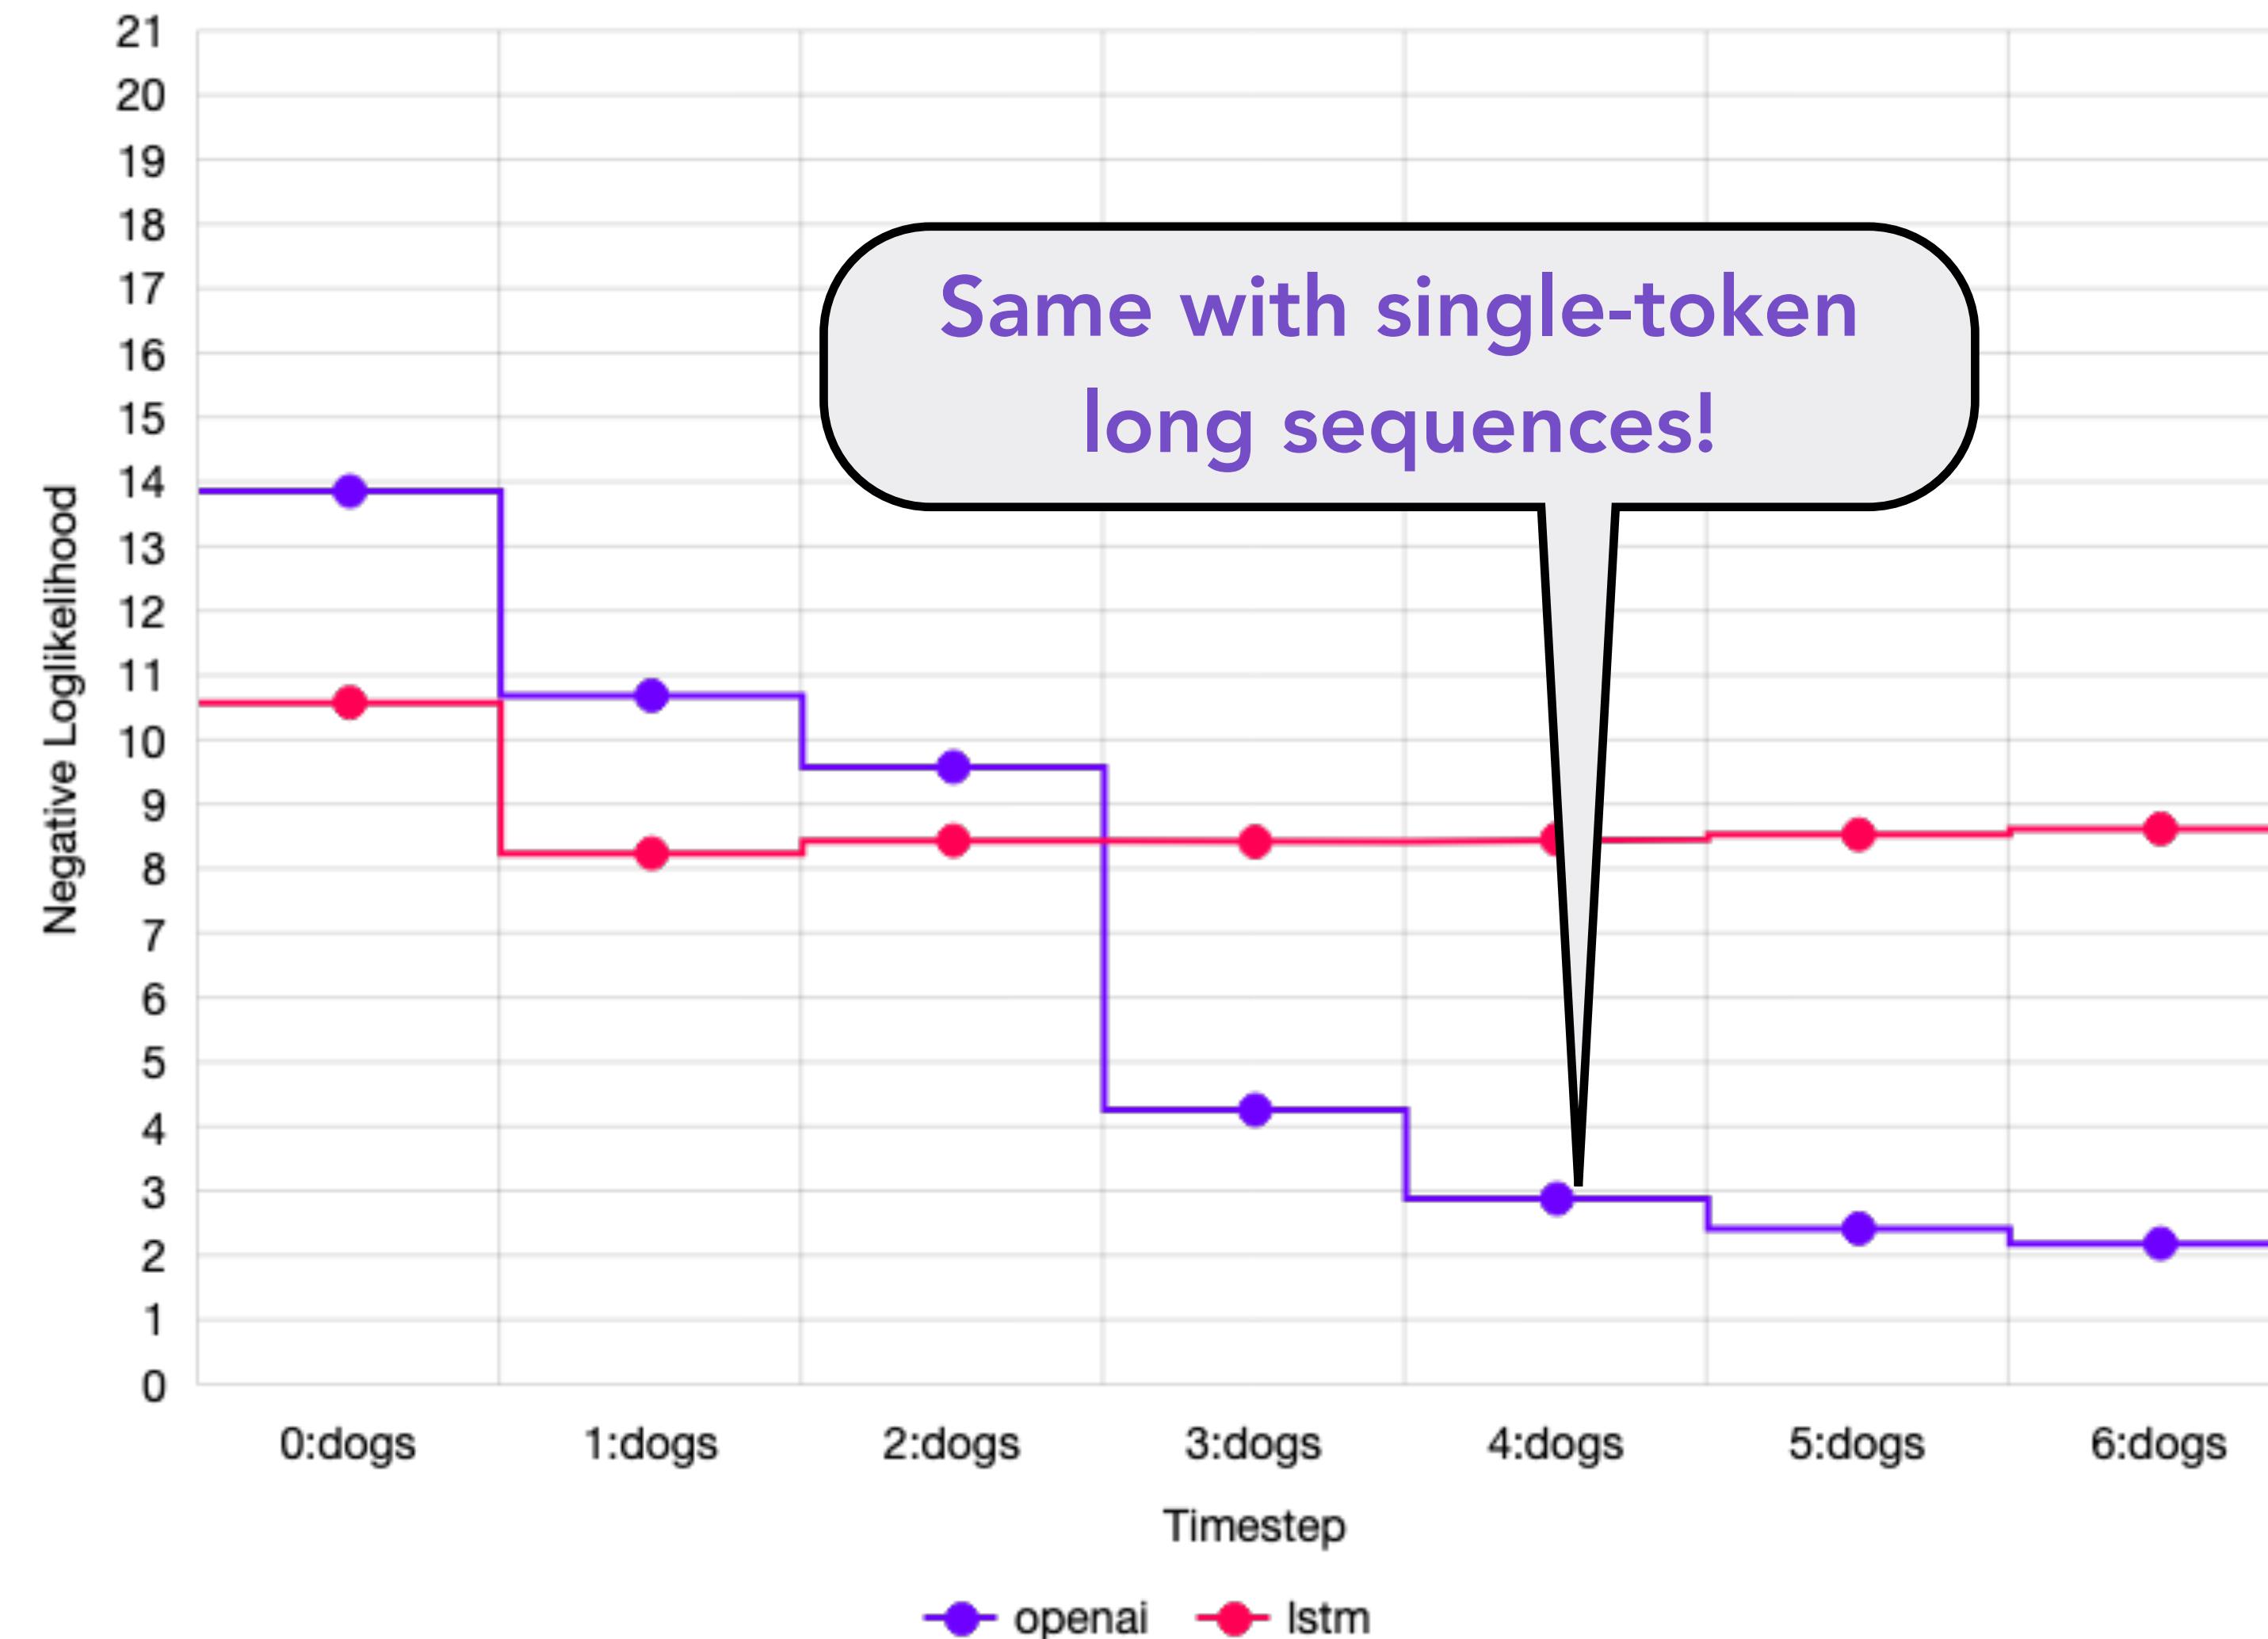


I don't know. I don't know.



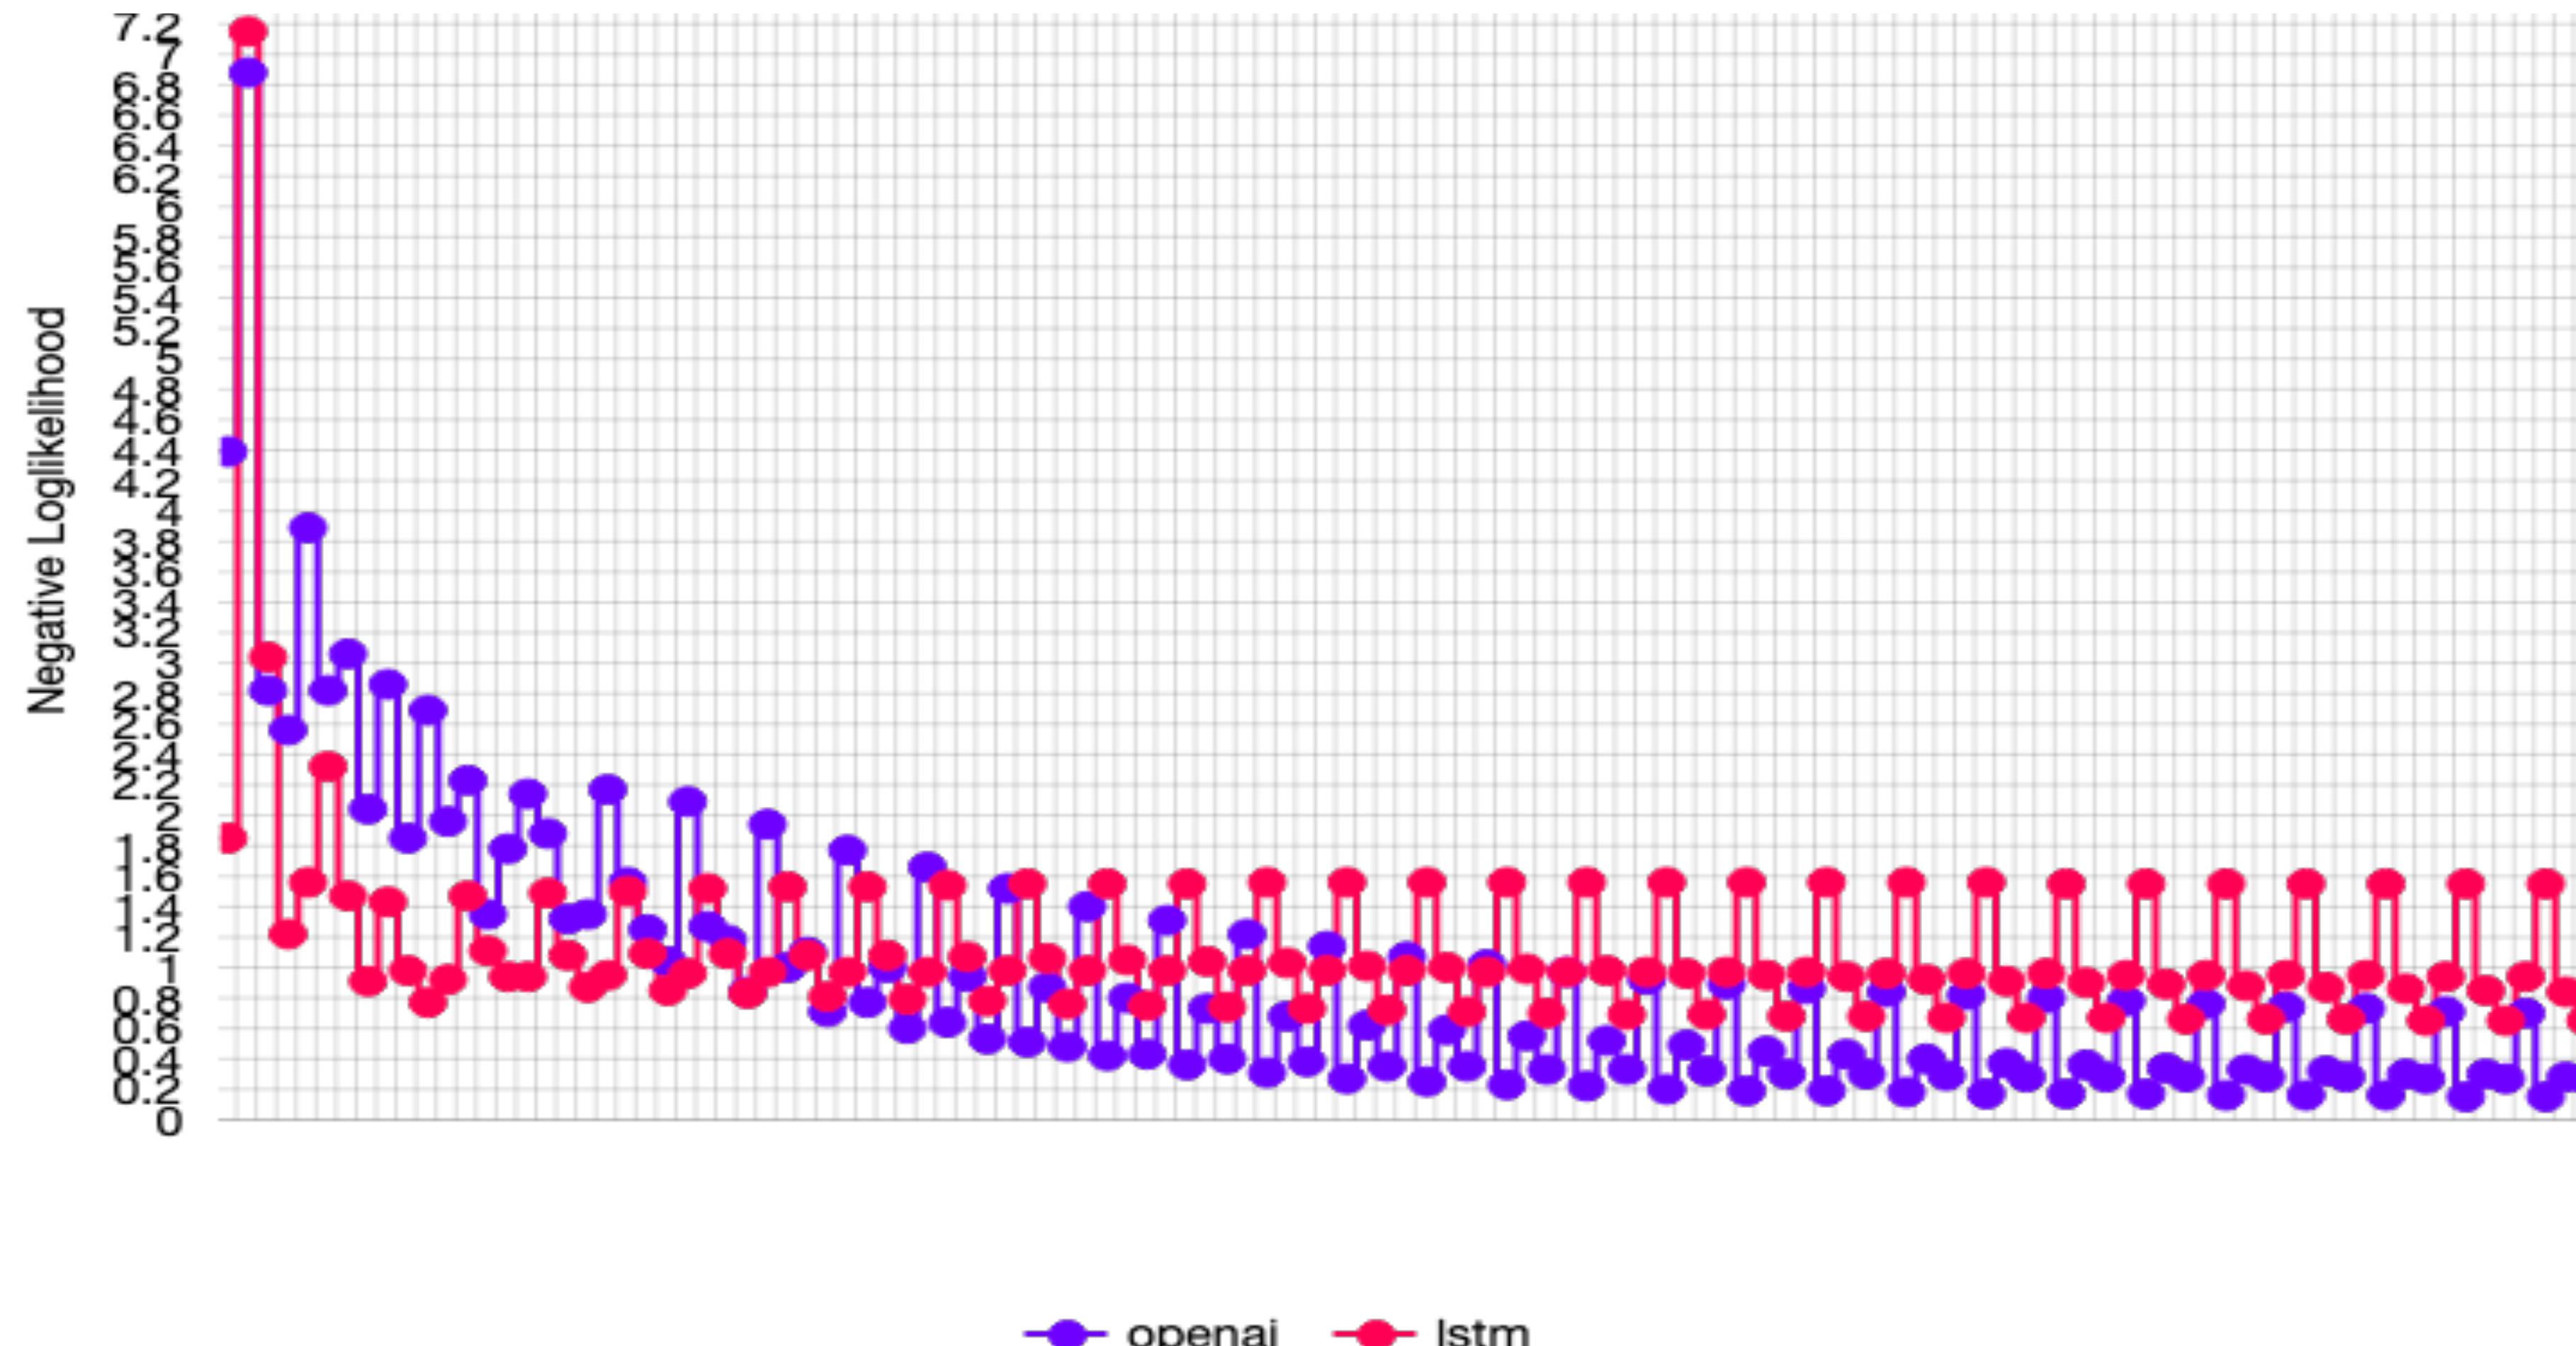
Beam search gets repetitive and repetitive

dogs dogs dogs dogs dogs dogs dogs



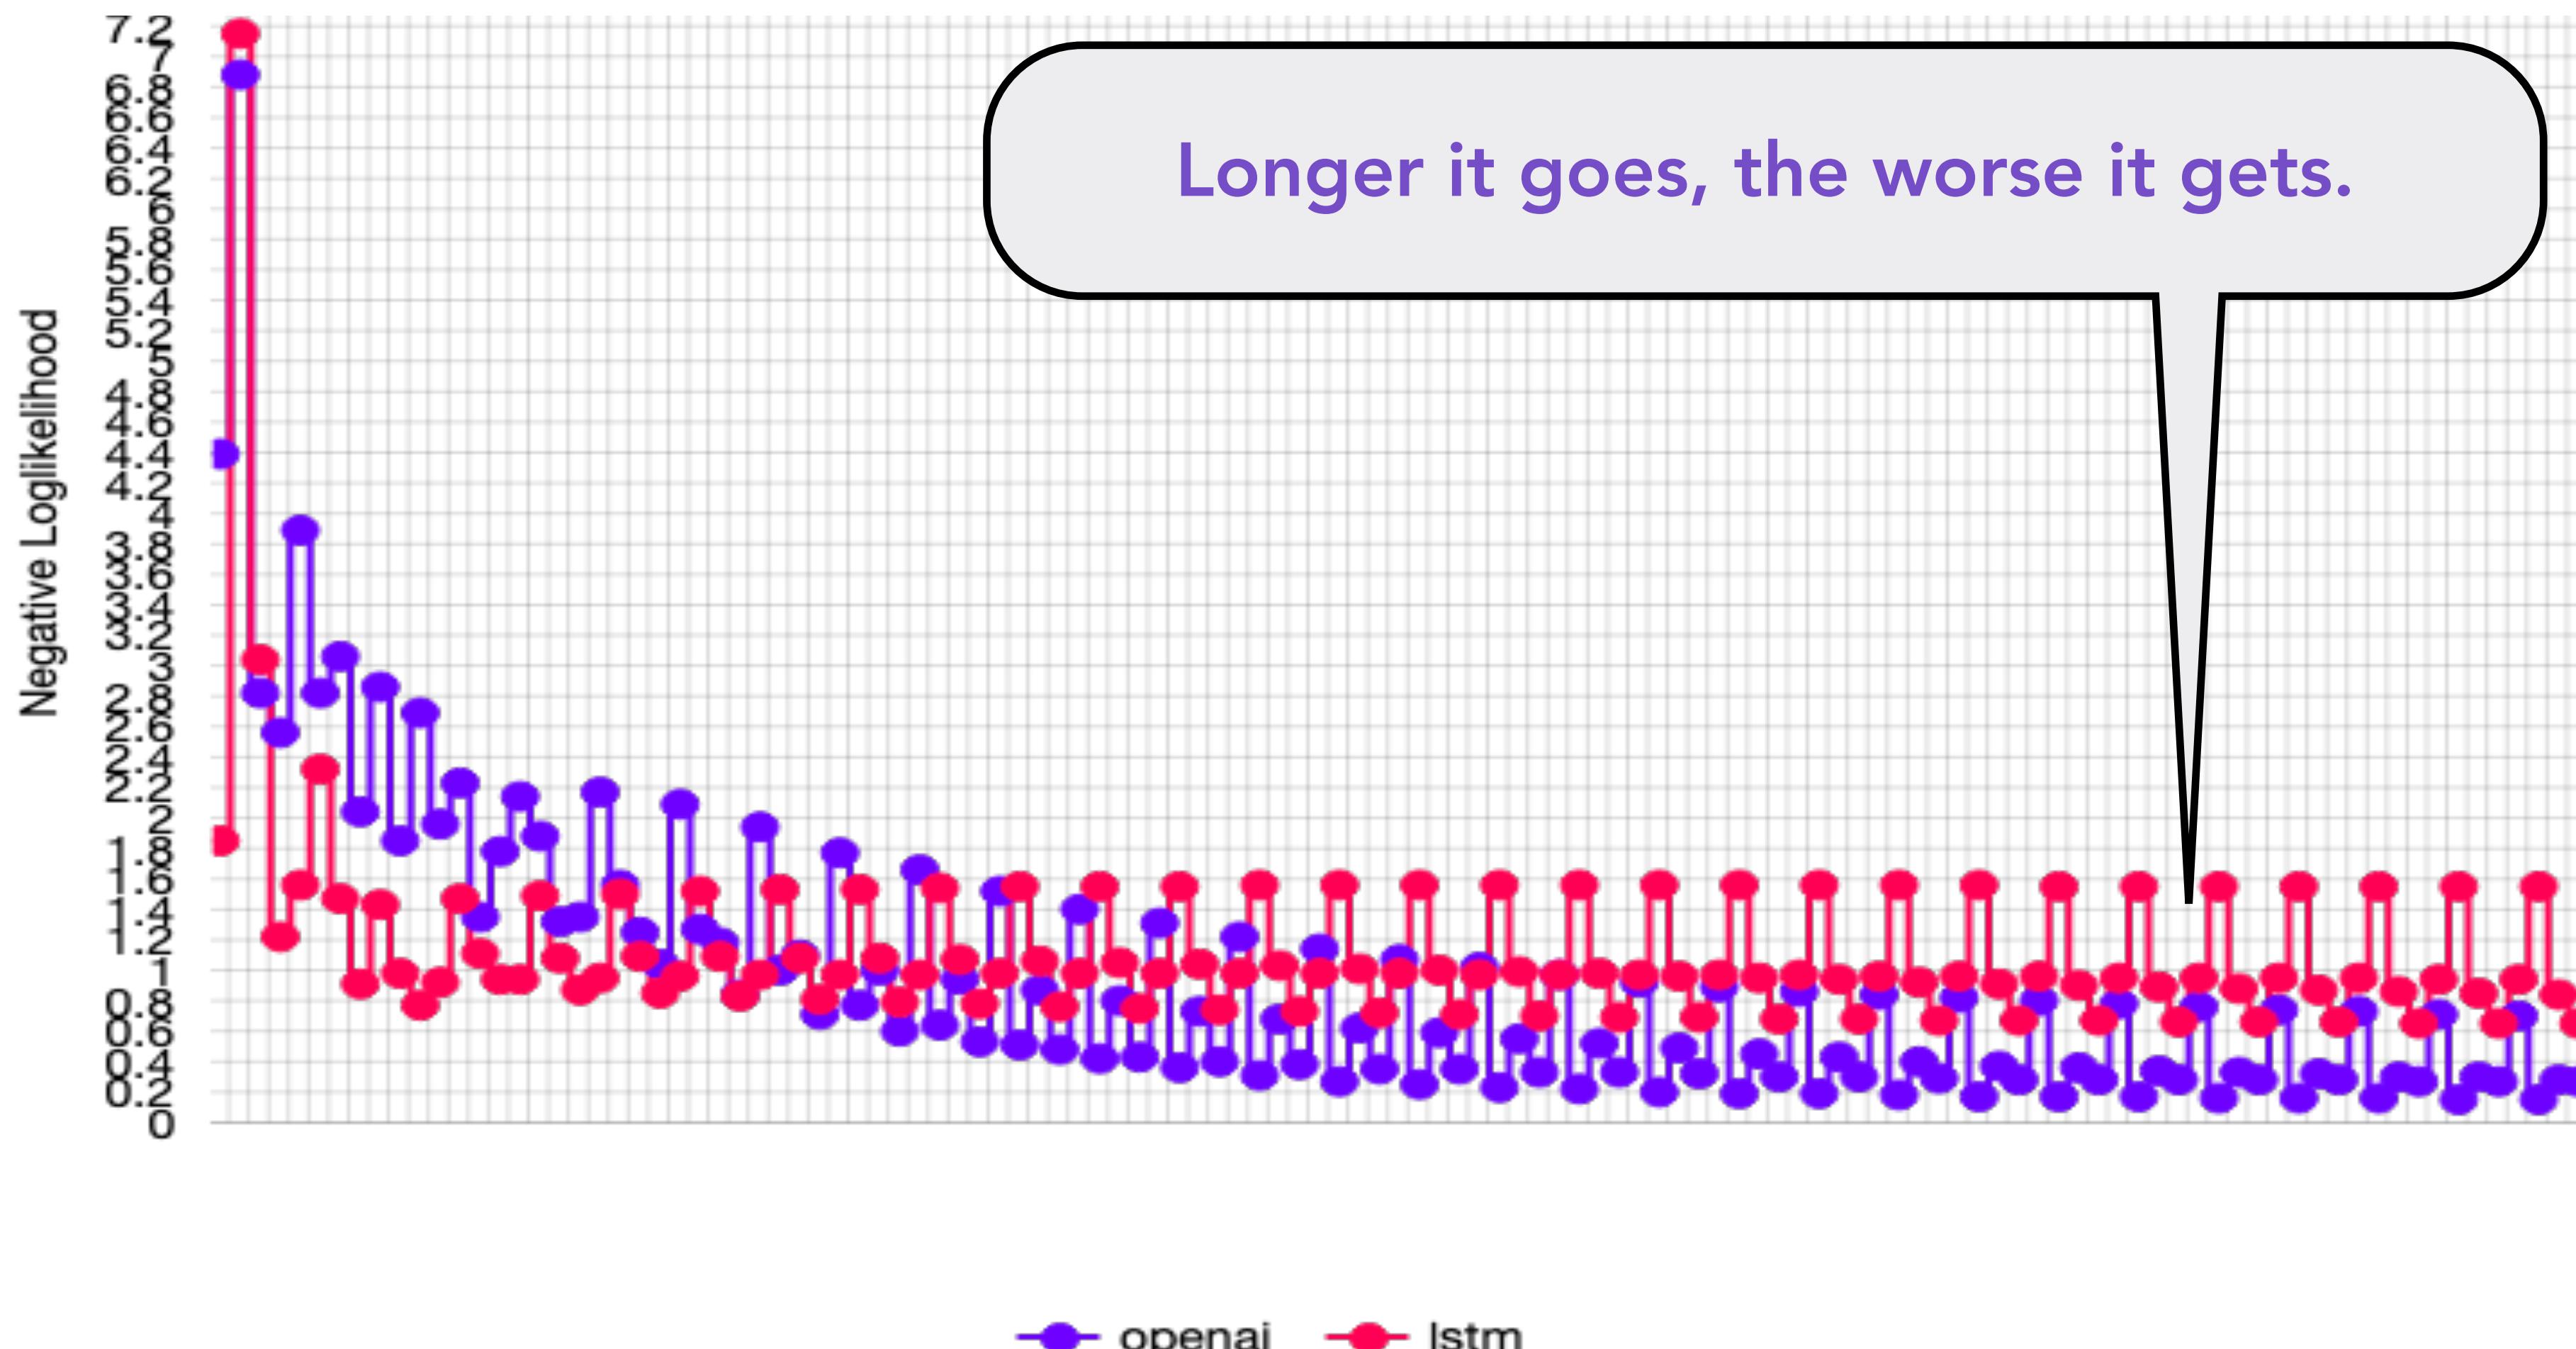
Beam search gets repetitive and repetitive

I'm tired. I'm tired.



Beam search gets repetitive and repetitive

I'm tired. I'm tired.



How can we reduce repetition?

- Don't repeat n -grams (Hacky, but works!)
- Minimize additional loss term for minimizing hidden state similarity (LSTMs)

$$\hat{y}_t = g \left(\log P(y_t | \{y\}_{<t}) - s(h_t, h_{t-m}) \right)$$

- Select token that maximizes mutual information between target and history

$$\hat{y}_t = g \left(\log P(y_t | \{y\}_{<t}) - \log P(y_t) \right)$$

How can we reduce repetition?

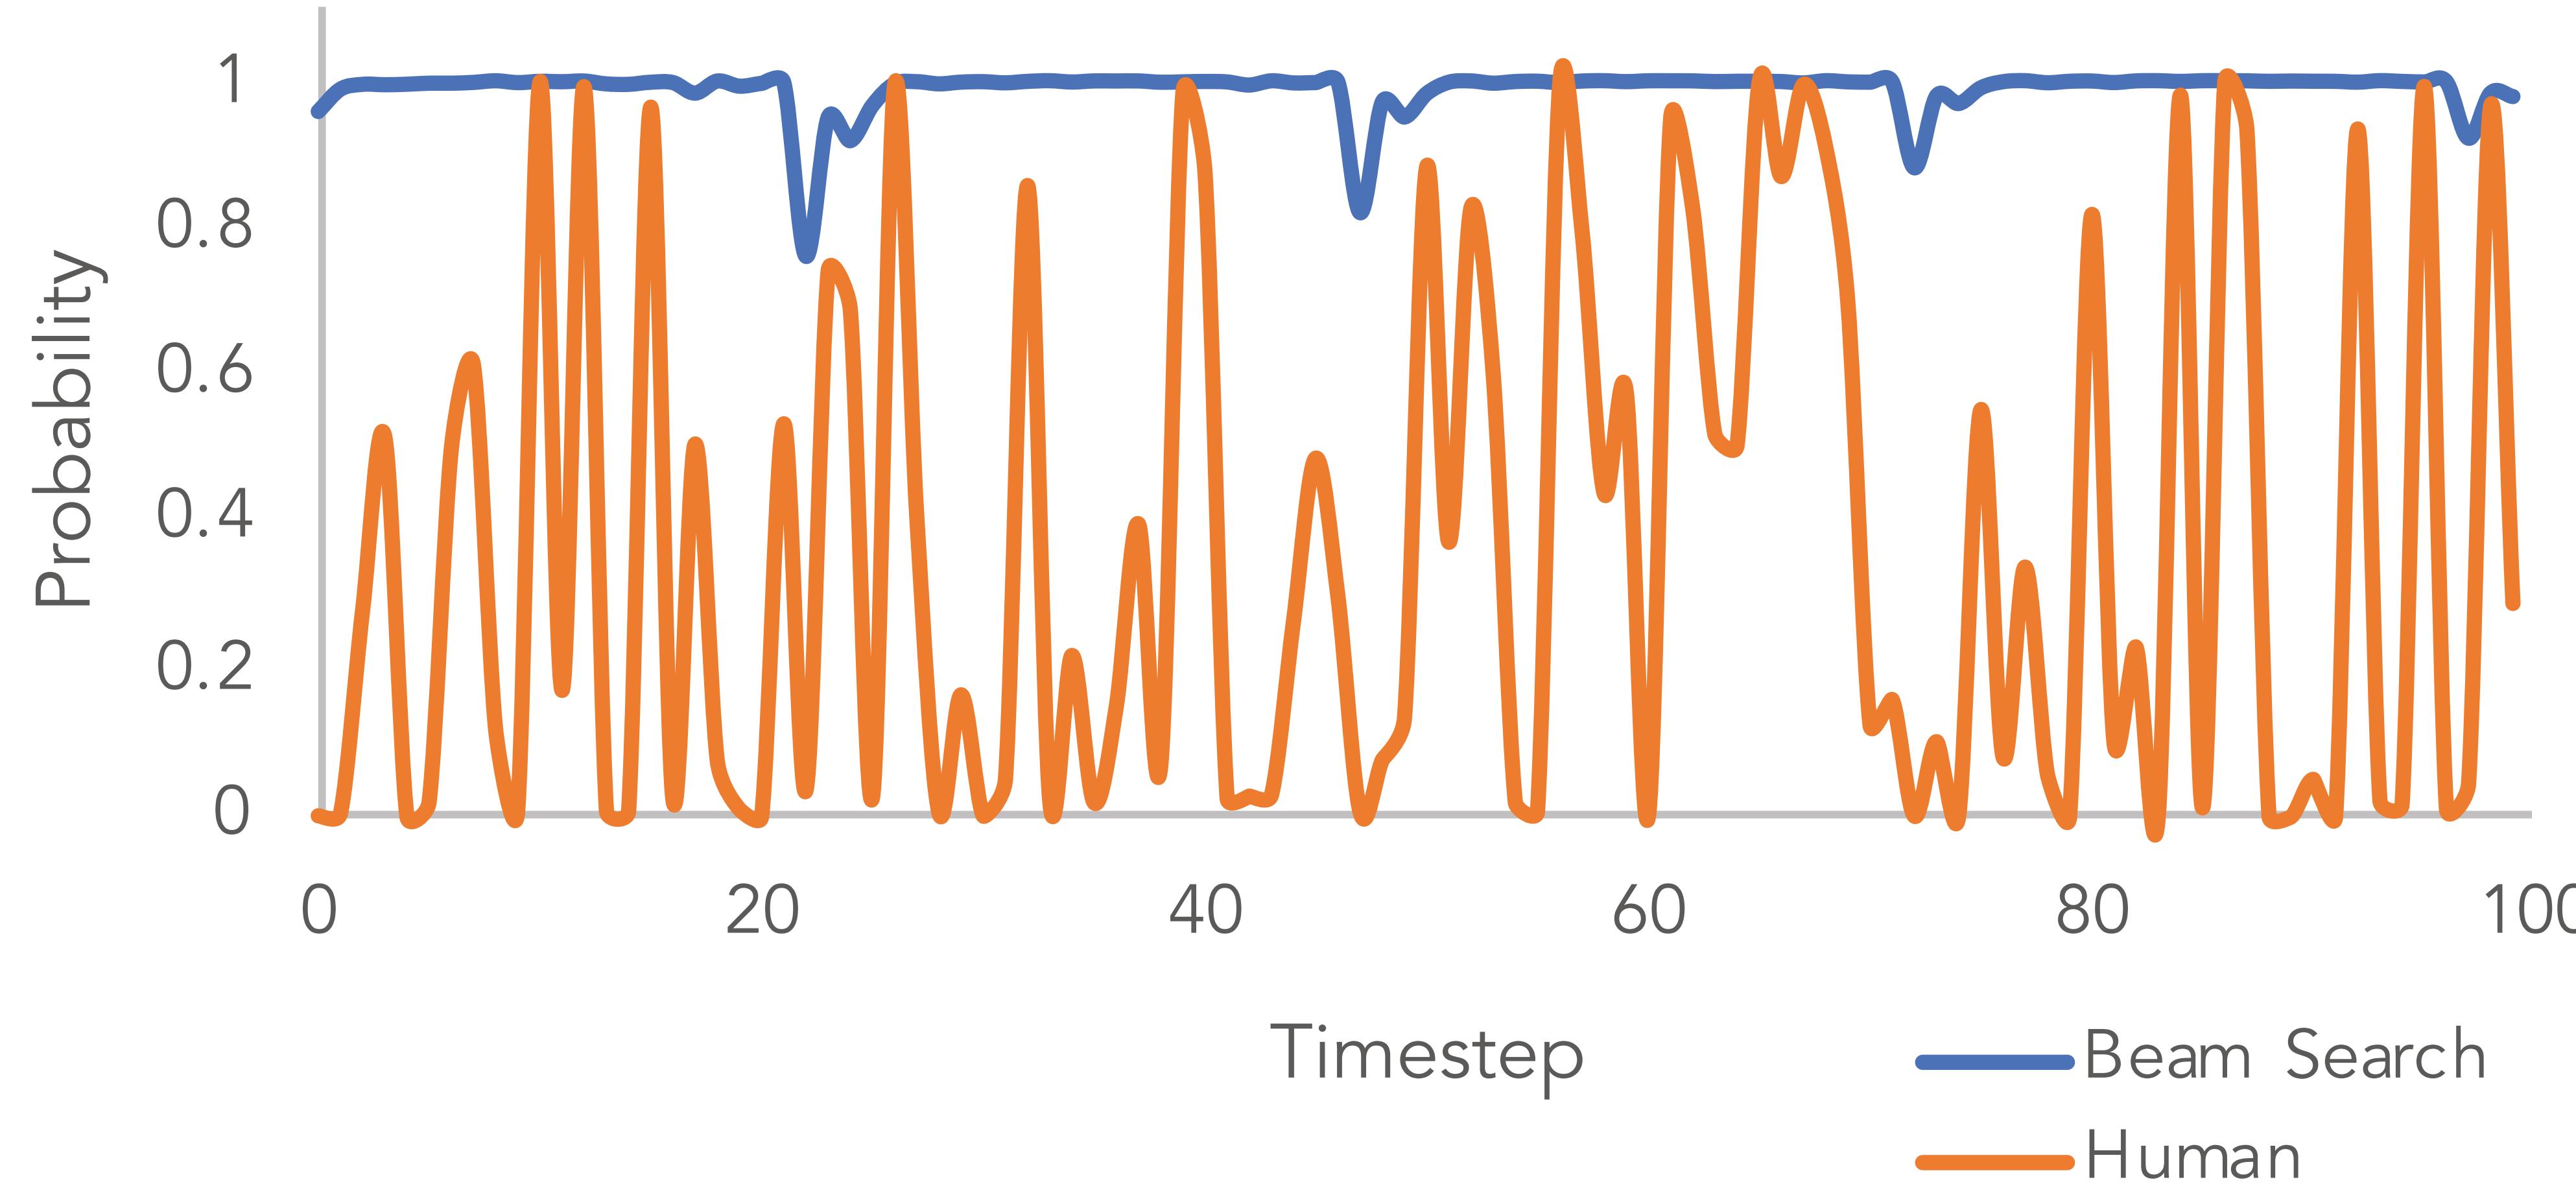
- Don't repeat n -grams (Hacky, but ✓)
- Minimize additional loss term for minimizing hidden state similarity (LSTMs)
- Select token that maximizes mutual information between target and history

Minimize **hidden state similarity** in LSTMs

$$\hat{y}_t = g \left(\log P(y_t | \{y\}_{<t}) - s(h_t, h_{t-m}) \right)$$

$$\hat{y}_t = g \left(\log P(y_t | \{y\}_{<t}) - \log P(y_t) \right)$$

Step-by-step Maximization



Time to get *random*: Sampling

- g = sample a token from the distribution of tokens

$$\hat{y}_t \sim P(y_t = w | \{y\}_{<t})$$



Whoa, too *random*: Temperature Scaling

- Too much randomness: distribution has too much entropy

Whoa, too *random*: Temperature Scaling

- Too much randomness: distribution has too much entropy
- **Solution:** Make the distribution more “peaky” with temperature scaling

Recall: $P(y_t | \{y\}_{<t}) = \frac{e^{o_n}}{\sum_{m=1}^M e^{o_m}}$

Whoa, too *random*: Temperature Scaling

- Too much randomness: distribution has too much entropy
- **Solution:** Make the distribution more “peaky” with temperature scaling

Recall: $P(y_t | \{y\}_{<t}) = \frac{e^{o_n}}{\sum_{m=1}^M e^{o_m}}$

→ $P(y_t | \{y\}_{<t}) = \frac{e^{o_n / \tau}}{\sum_{m=1}^M e^{o_m / \tau}}$

Whoa, too *random*: Temperature Scaling

- Too much randomness: distribution has too much entropy
- **Solution:** Make the distribution more “peaky” with temperature scaling

$$P(y_t | \{y\}_{<t}) = \frac{e^{o_n / \tau}}{\sum_{m=1}^M e^{o_m / \tau}}$$

$\tau > 1$ “flatter” distribution

$\tau < 1$ “peakier” distribution

Maybe we need fewer options: Top- k sampling

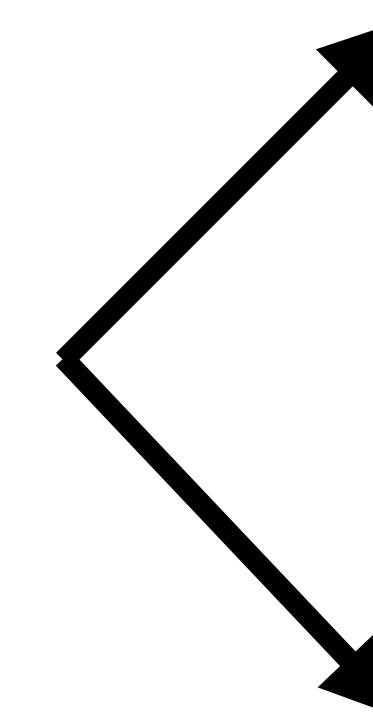
- The entire distribution over tokens is not needed at every step
- Many token choices should have no chance of being selected

Maybe we need fewer options: Top- k sampling

- The entire distribution over tokens is not needed at every step
- Many token choices should have no chance of being selected
- Only sample from the top k tokens in the distribution

$$\hat{y}_t \sim P^*(y_t = w | \{y\}_{<t})$$

He wanted to go to the

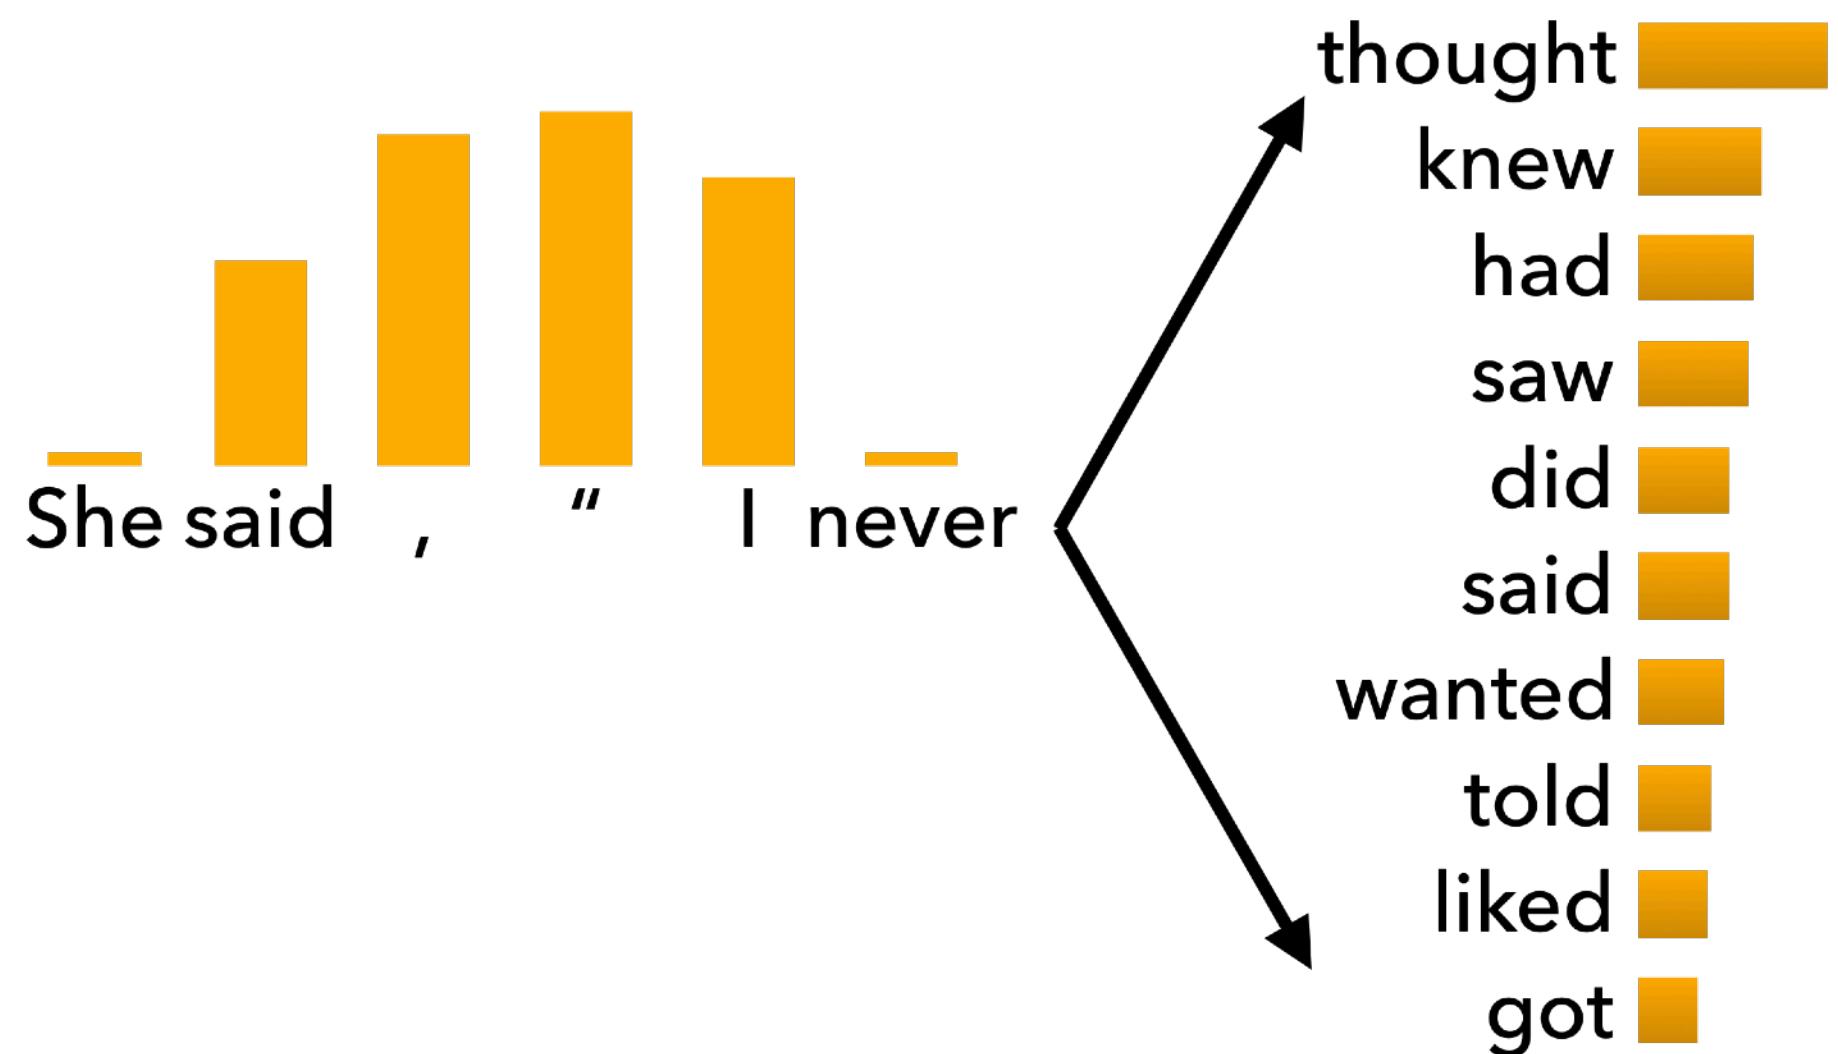


restroom
grocery
store
airport
pub
gym
bathroom
game
beach
hospital
doctor

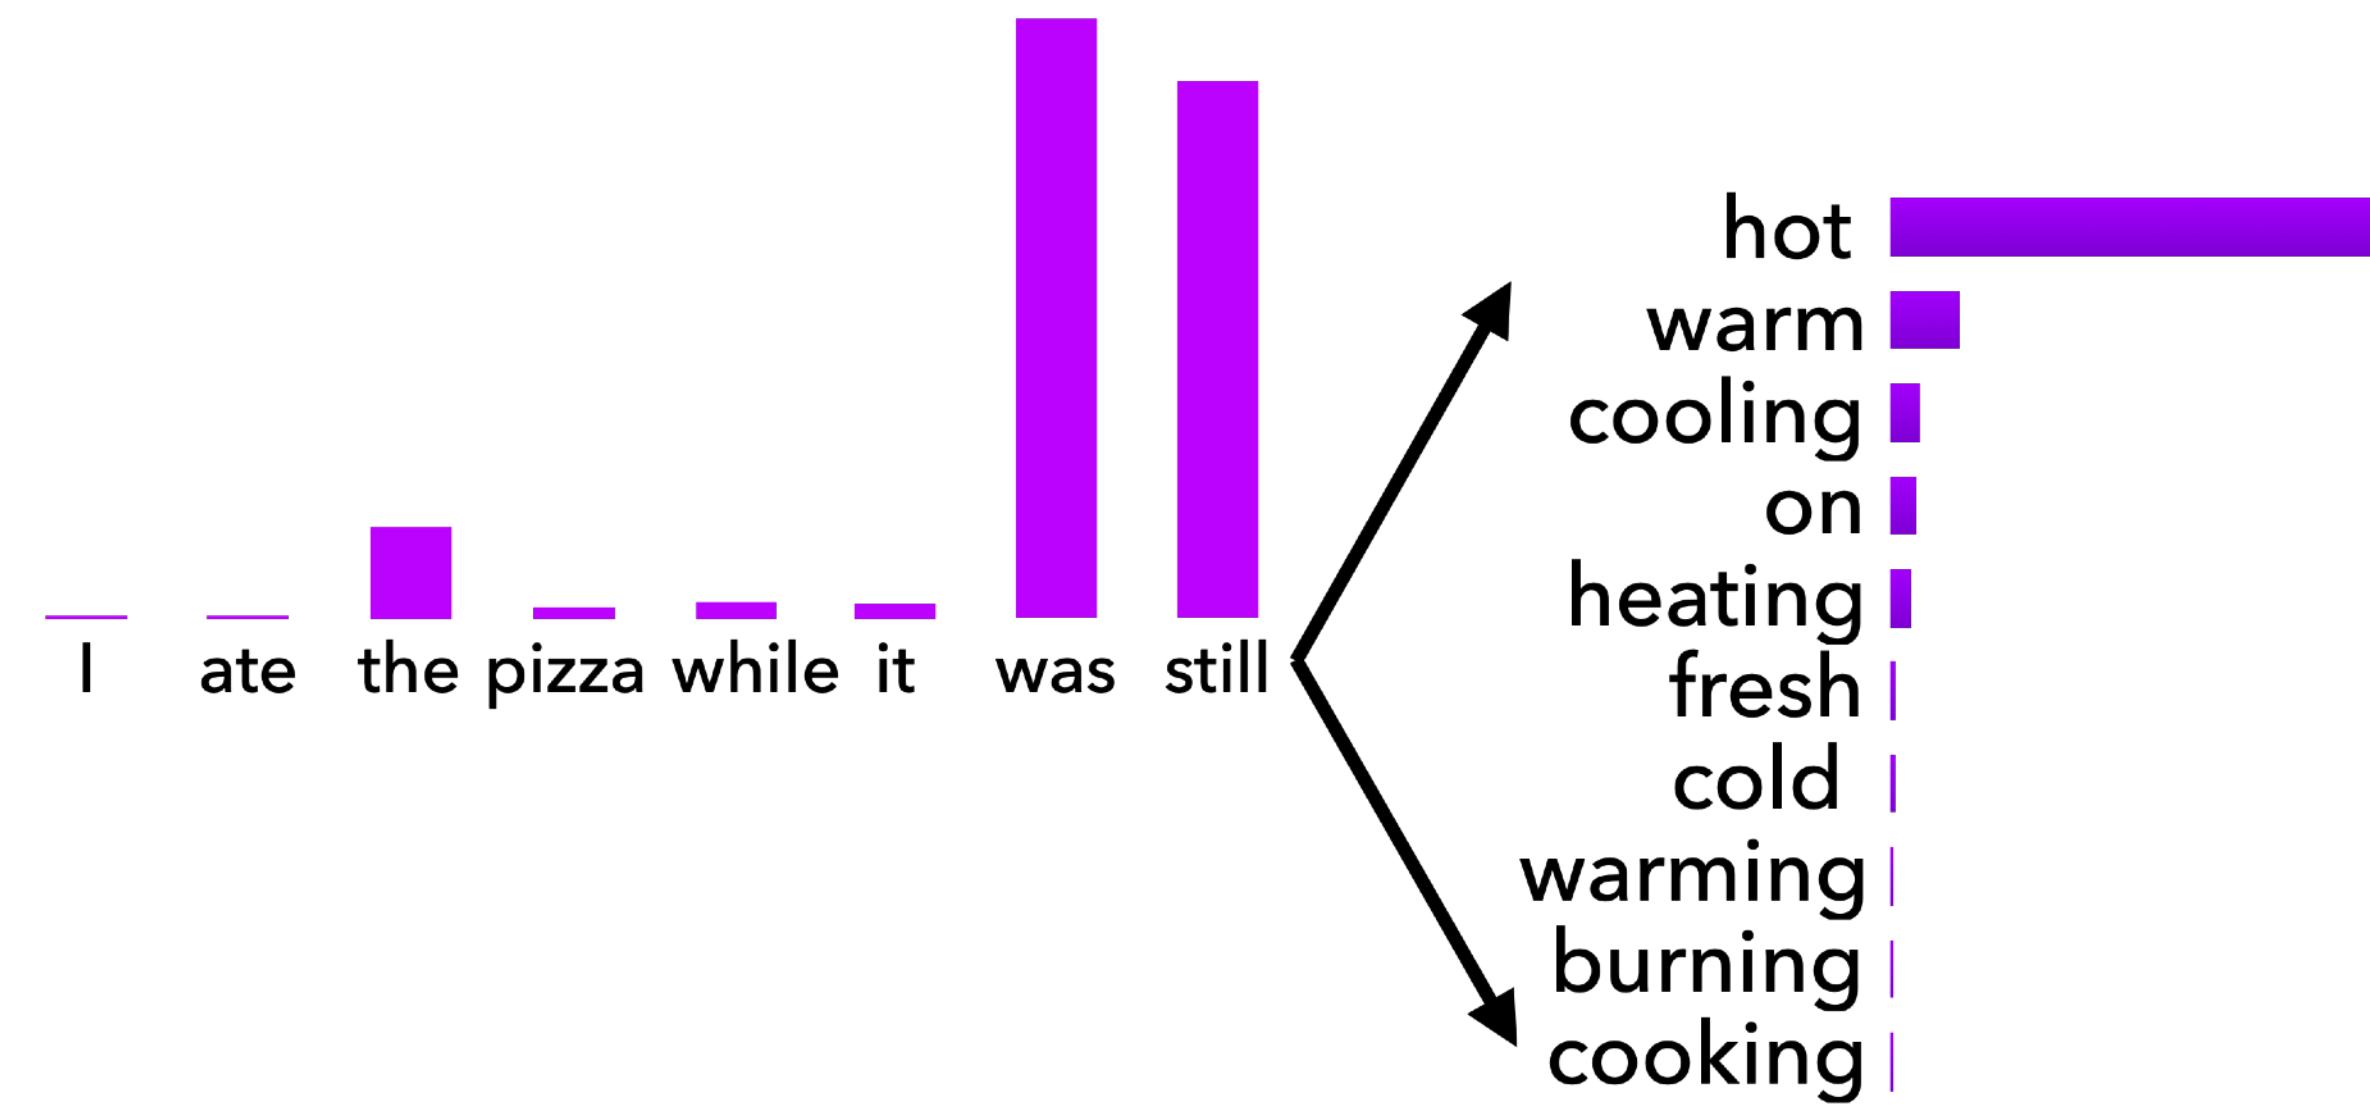
Randomly sample
token from top k
highest probability
tokens in $P(.)$

Issues with top-k sampling

Top- k can cut-off too *quickly*



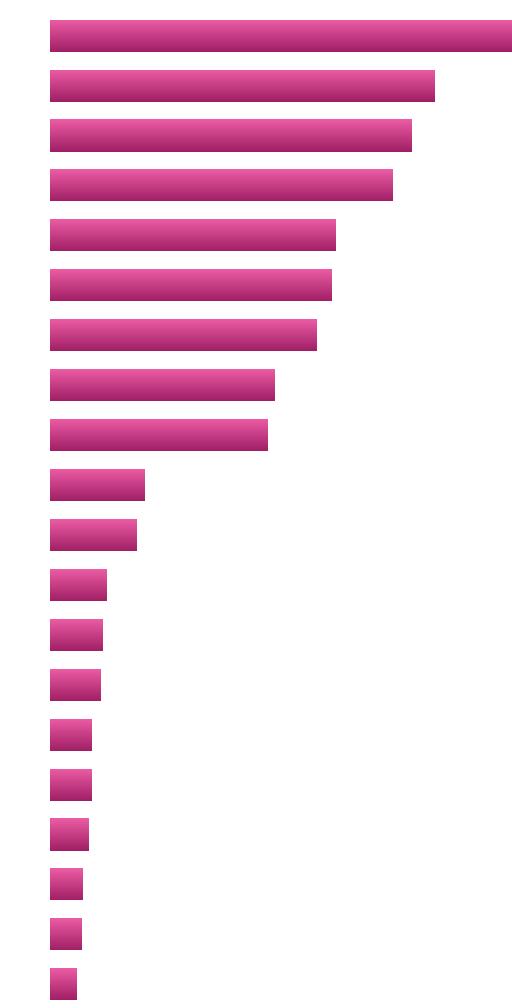
Top- k can cut-off too *slowly*



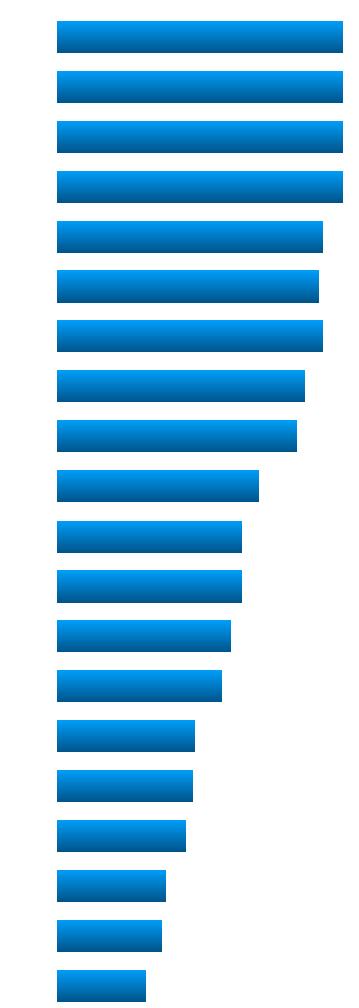
I don't know how many options I need: Top-p sampling

- Also known as **nucleus** sampling
- Sample from subset of vocabulary where probability mass is concentrated

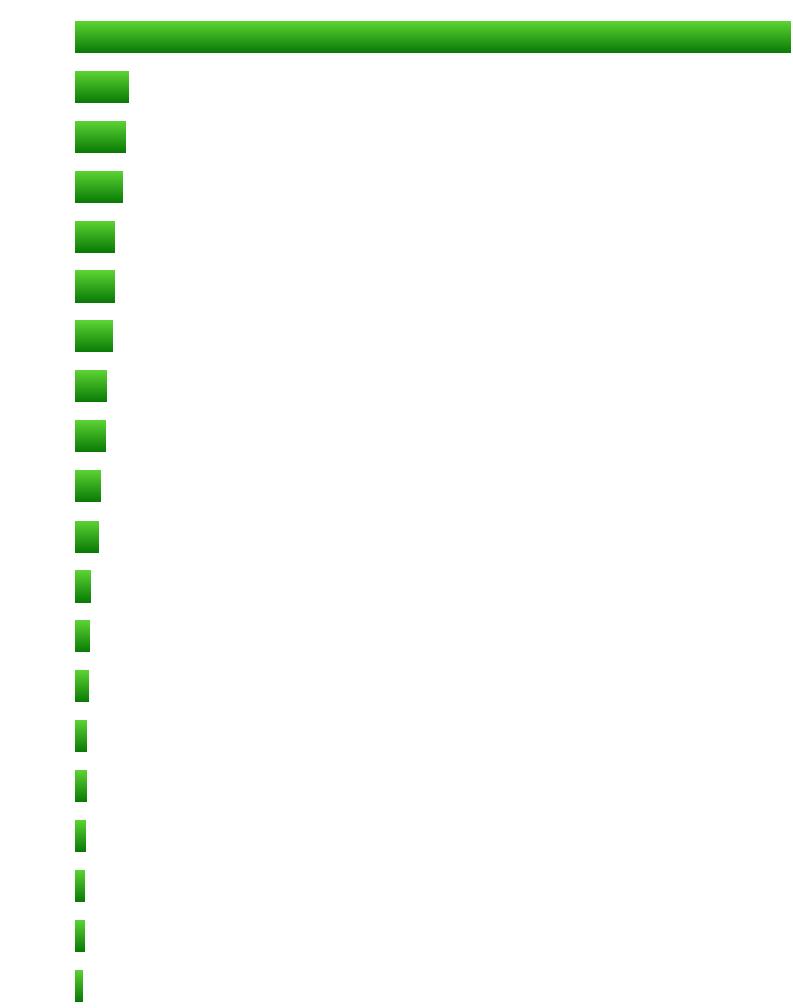
$$P_1(y_t | \{y\}_{<t})$$



$$P_2(y_t | \{y\}_{<t})$$



$$P_3(y_t | \{y\}_{<t})$$

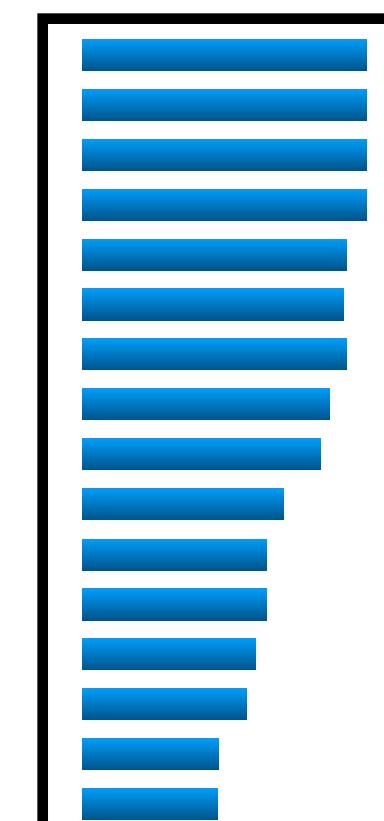


I don't know how many options I need: Top- p sampling

- Also known as **nucleus** sampling:
- Sample from subset of vocabulary where probability mass is concentrated
- Probability mass has a dynamically changing **nucleus**



$$P_1^*(y_t | \{y\}_{<t})$$



$$P_2^*(y_t | \{y\}_{<t})$$

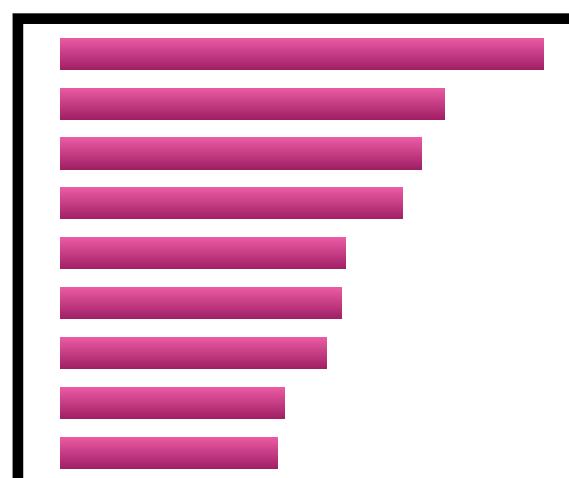


$$P_3^*(y_t | \{y\}_{<t})$$

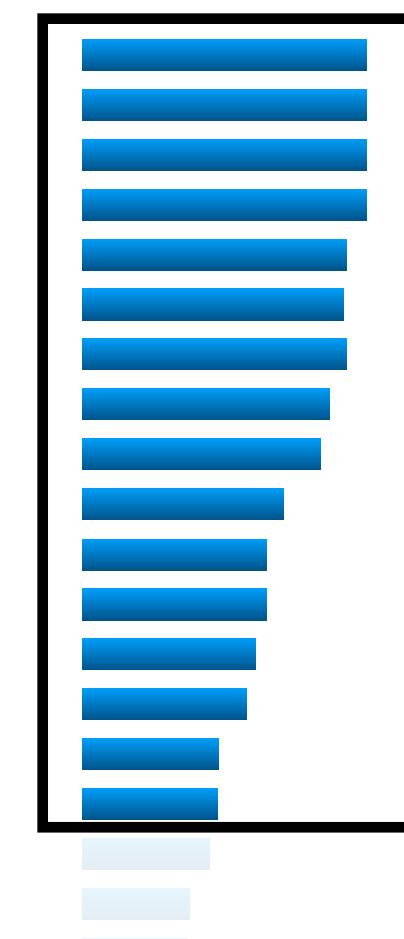
I don't know how many options I need: Top-p sampling

- Also known as **nucleus** sampling:
- Sample from subset of vocabulary where probability mass is concentrated
- Probability mass has a dynamically changing **nucleus**

Only sample from the
nucleus of the **distribution**



$$P_1^*(y_t | \{y\}_{<t})$$



$$P_2^*(y_t | \{y\}_{<t})$$



$$P_3^*(y_t | \{y\}_{<t})$$

This all sounds a bit *risky*

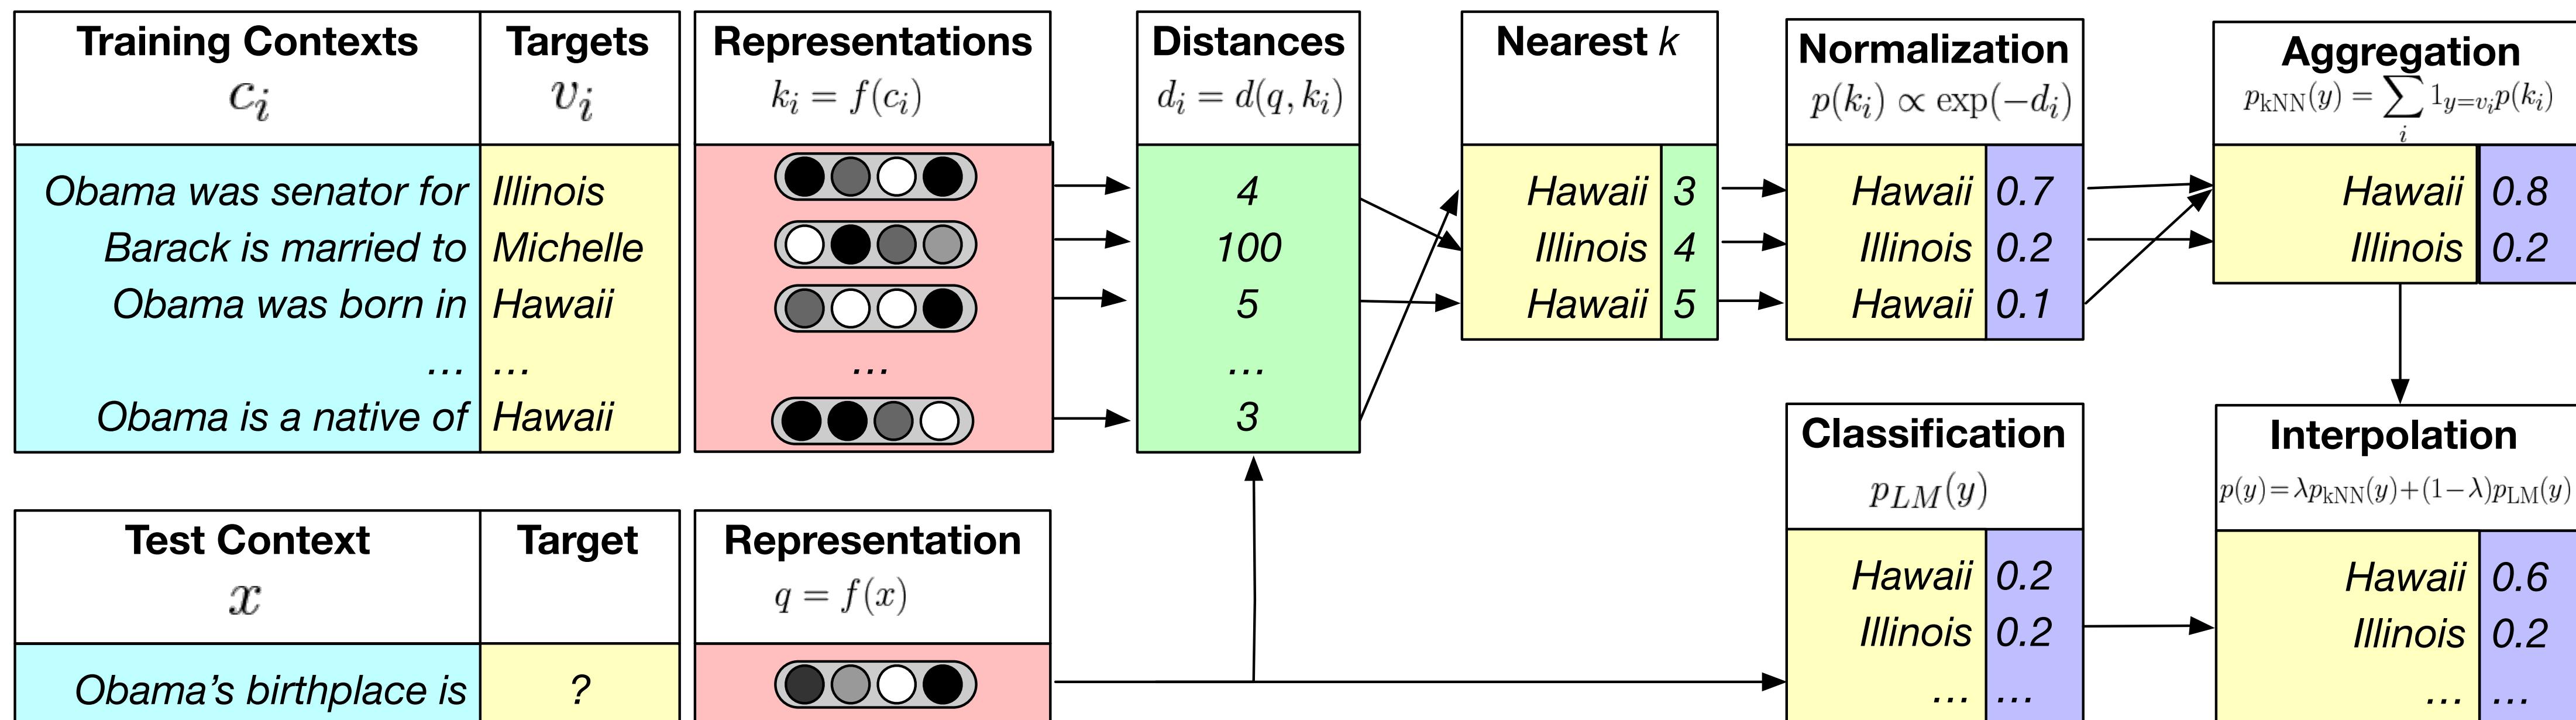
- What if my distribution isn't great to begin with? Can I mitigate this?

This all sounds a bit *risky*

- What if my distribution isn't great to begin with? Can I mitigate this?
 1. Make your word-level distributions more robust

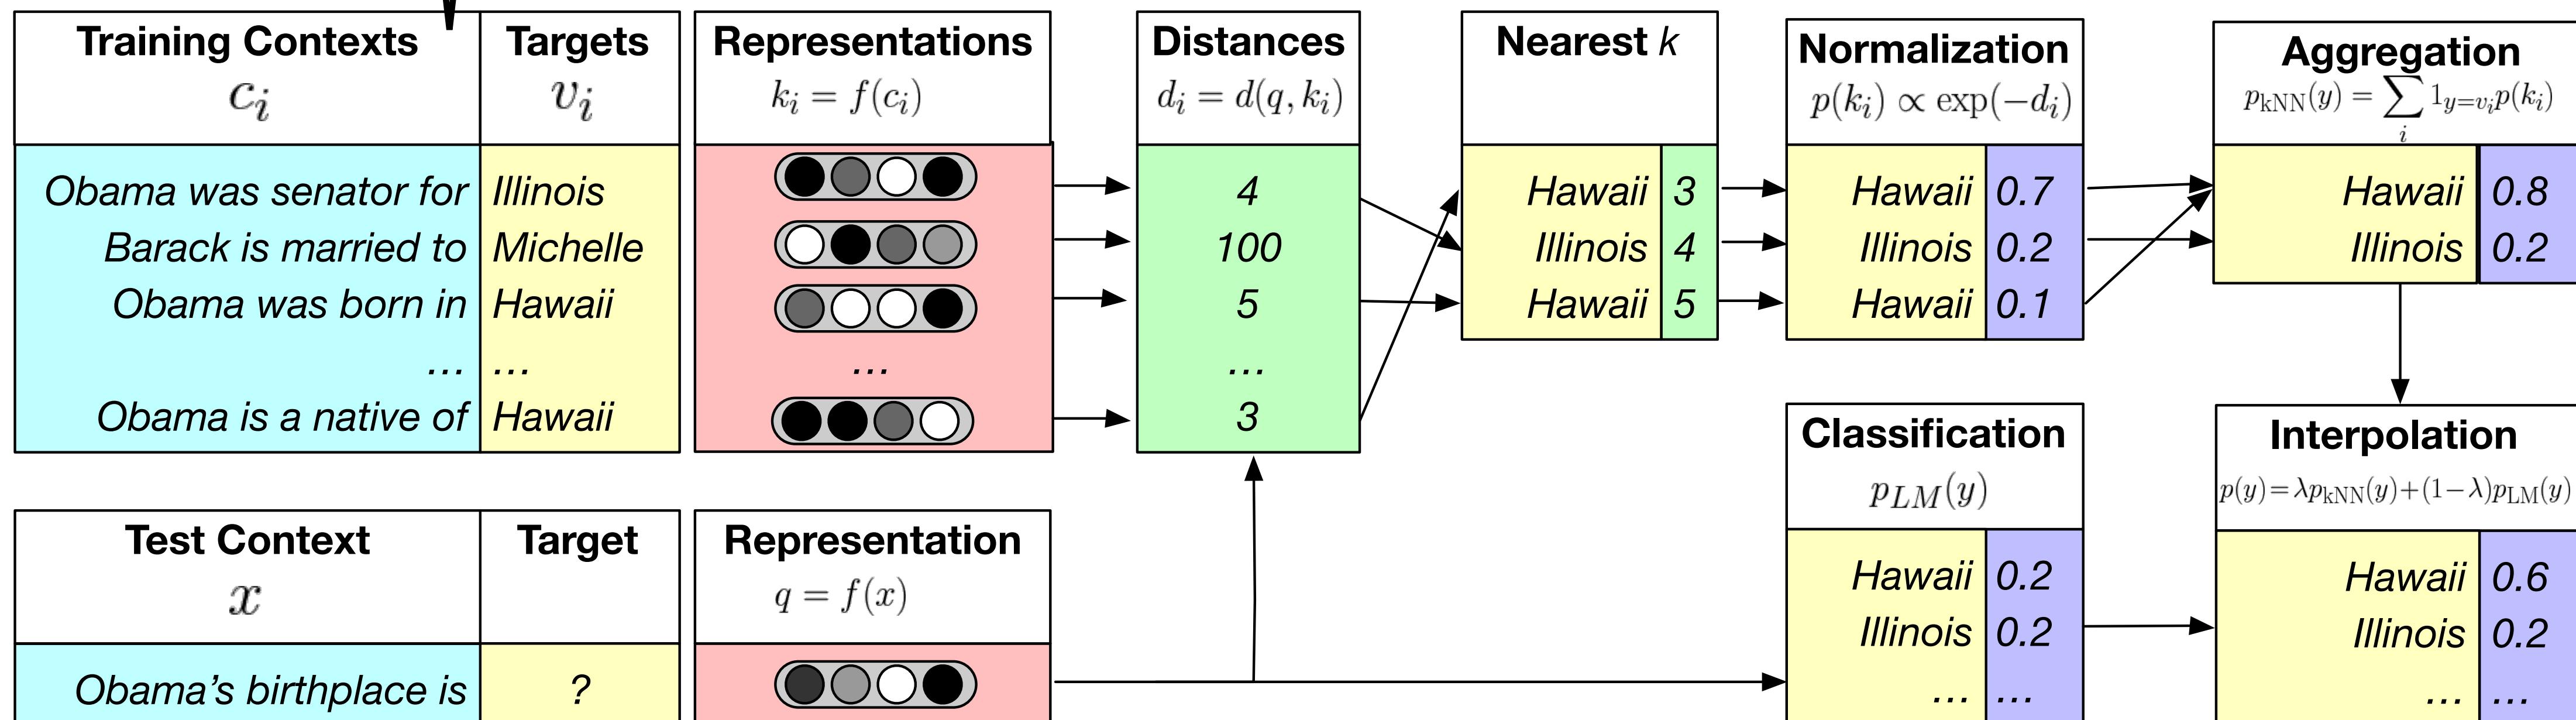
kNN Language Models

- Don't just rely on your trained model to generate a distribution over tokens
- Use knowledge of similar contexts from another corpus



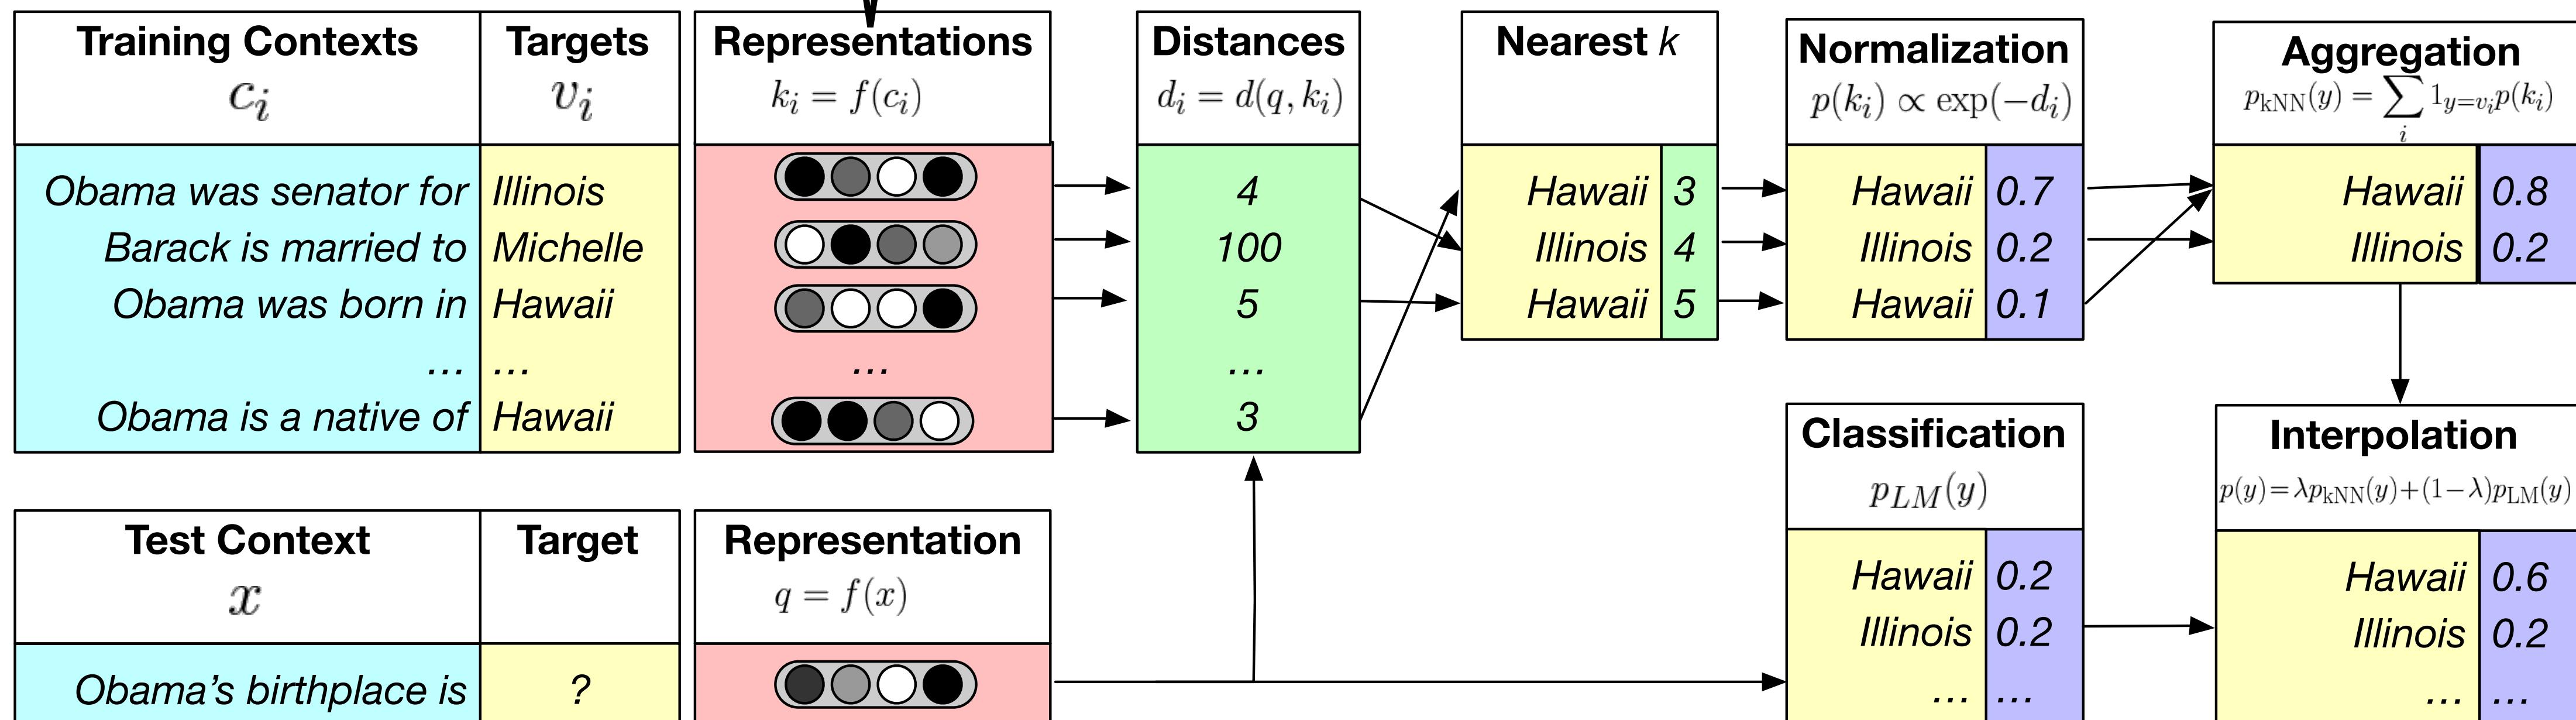
kNN Language Models

- Don't implement from scratch, use your trained model to generate a distribution over contexts
- Initialize a **database** of contexts
- Use knowledge of similar contexts from another corpus



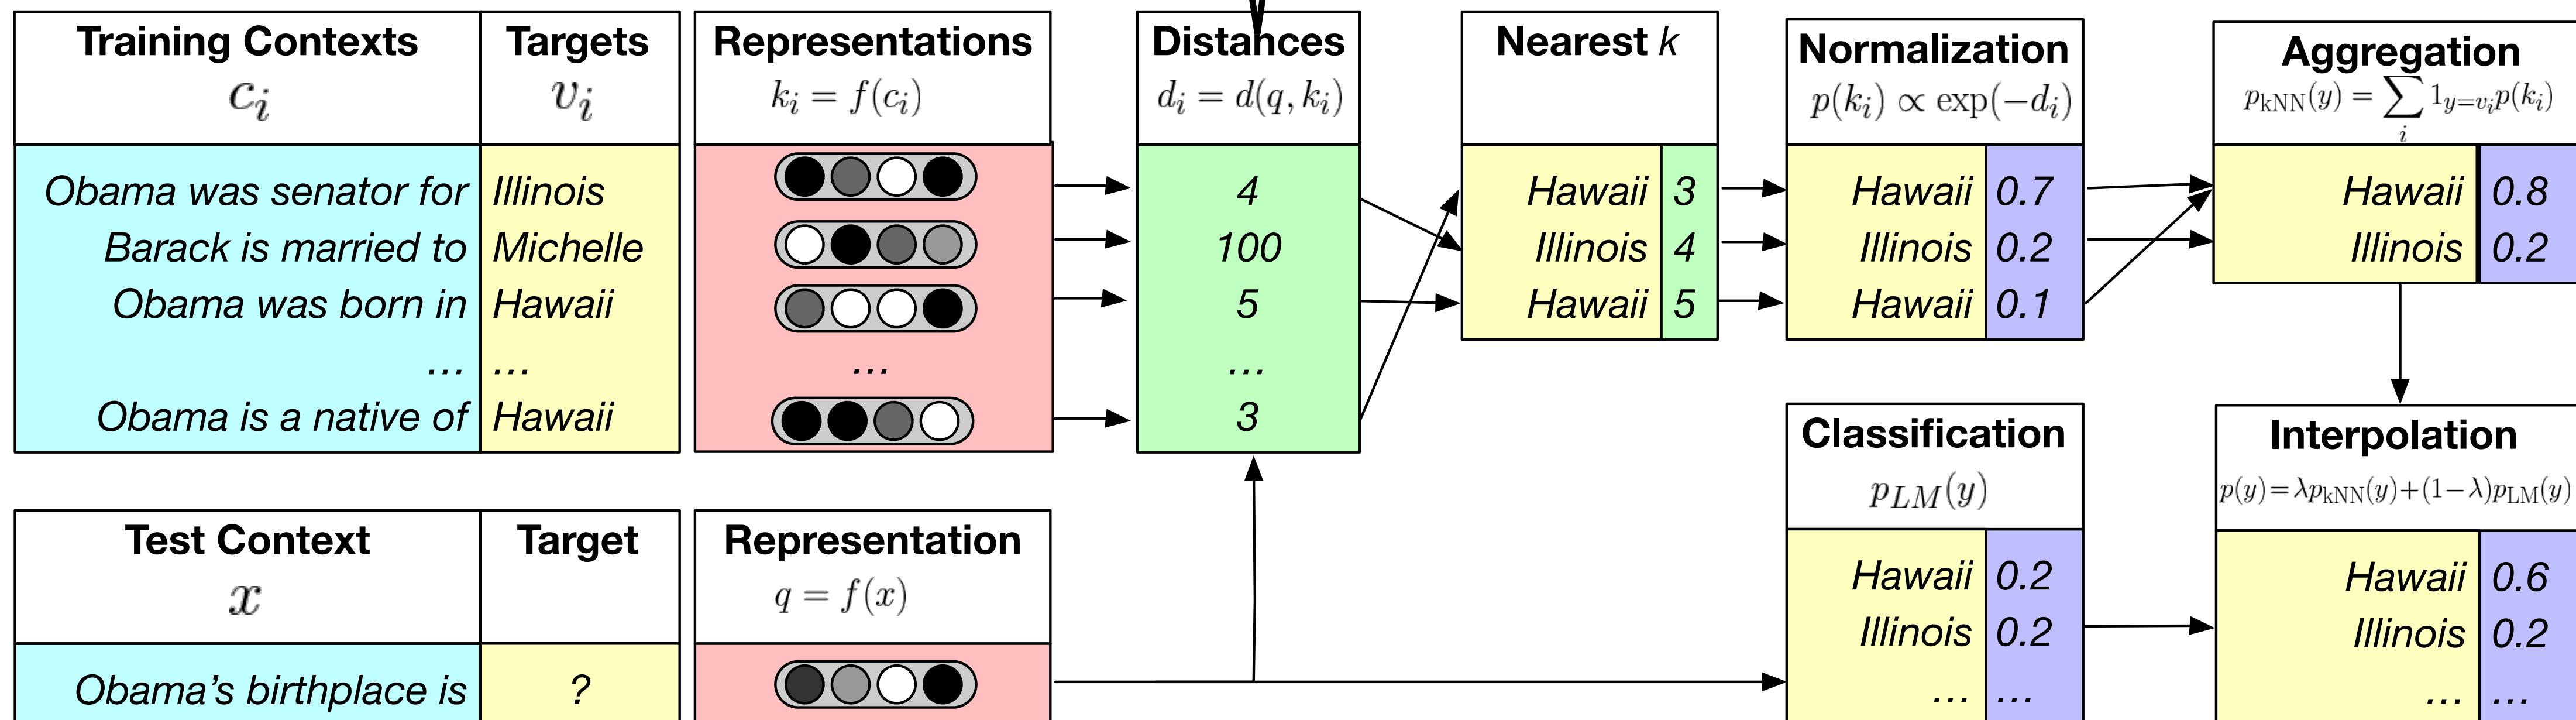
kNN Language Models

- Don't just rely on your trained model to generate a distribution over tokens.
- Compute **representations** for each context in DB
- Use knowledge of similar contexts from another corpus



kNN Language Models

- Don't just rely on your trained model to generate a distribution over to Efficiently compute **distance** between each context in DB and current sequence
- Use knowledge of similar contexts from another corpus

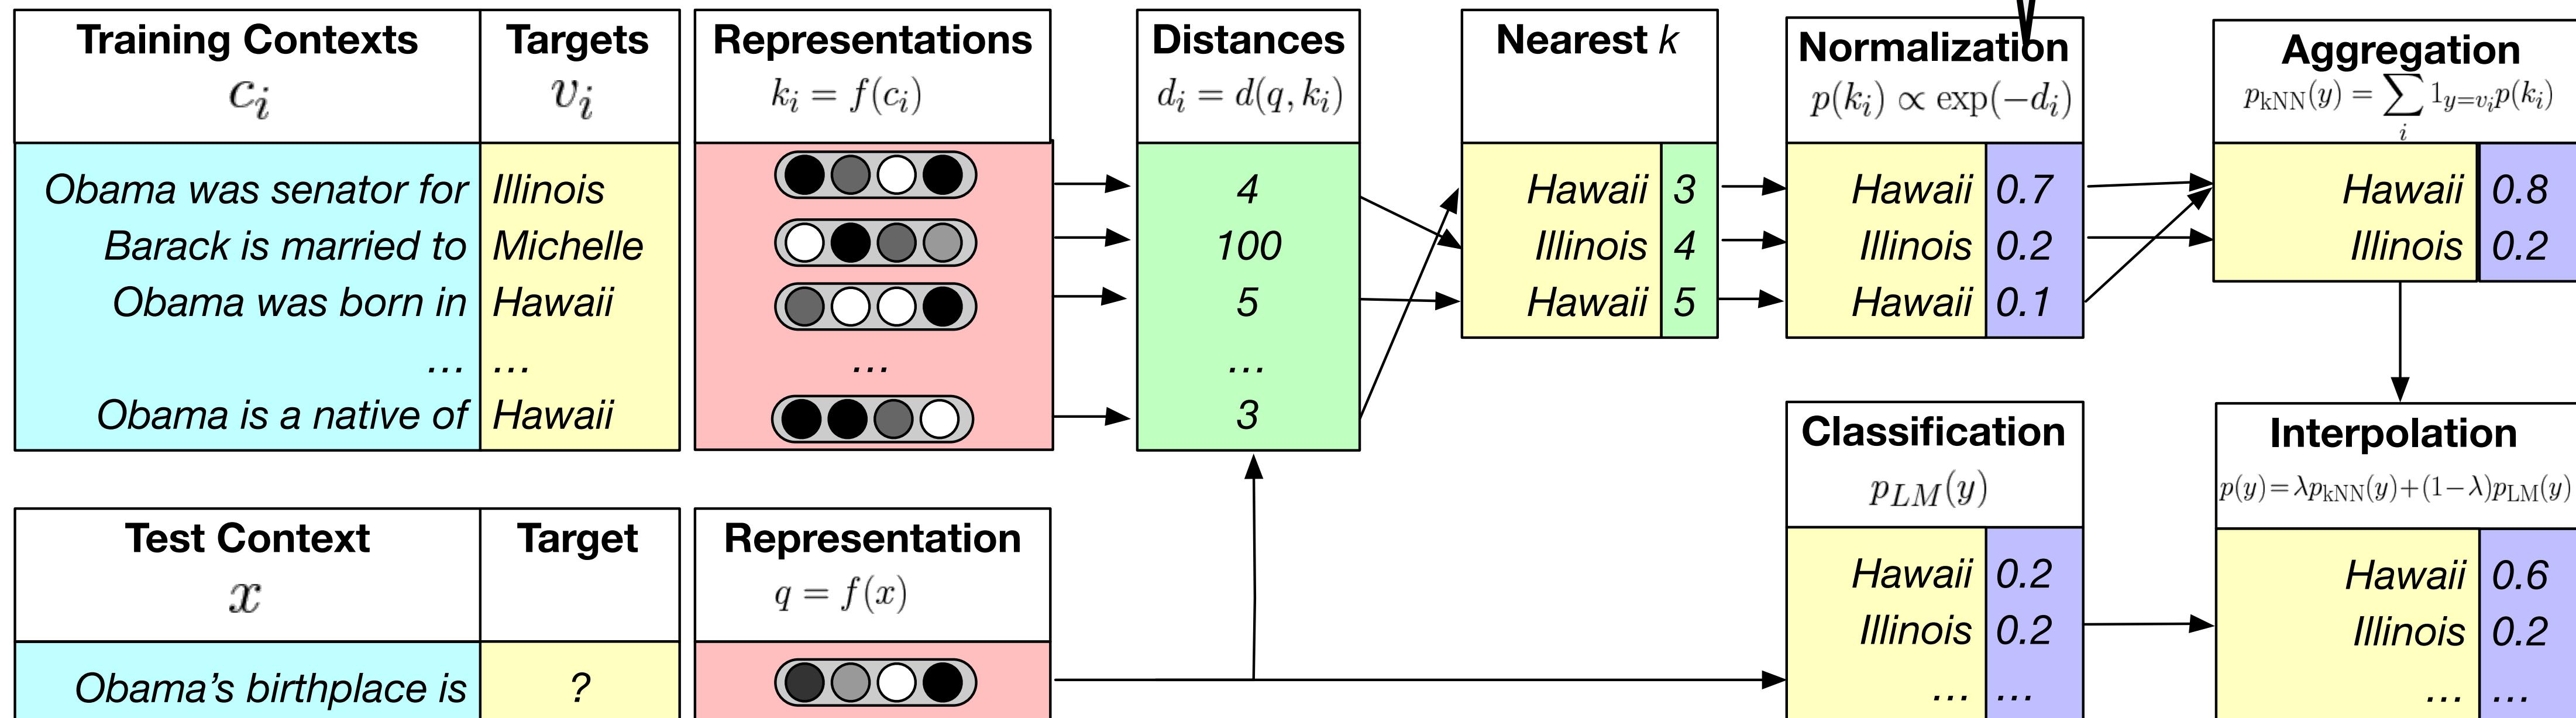


kNN Language Models

- Don't just rely on your trained model to generate a distribution over tokens

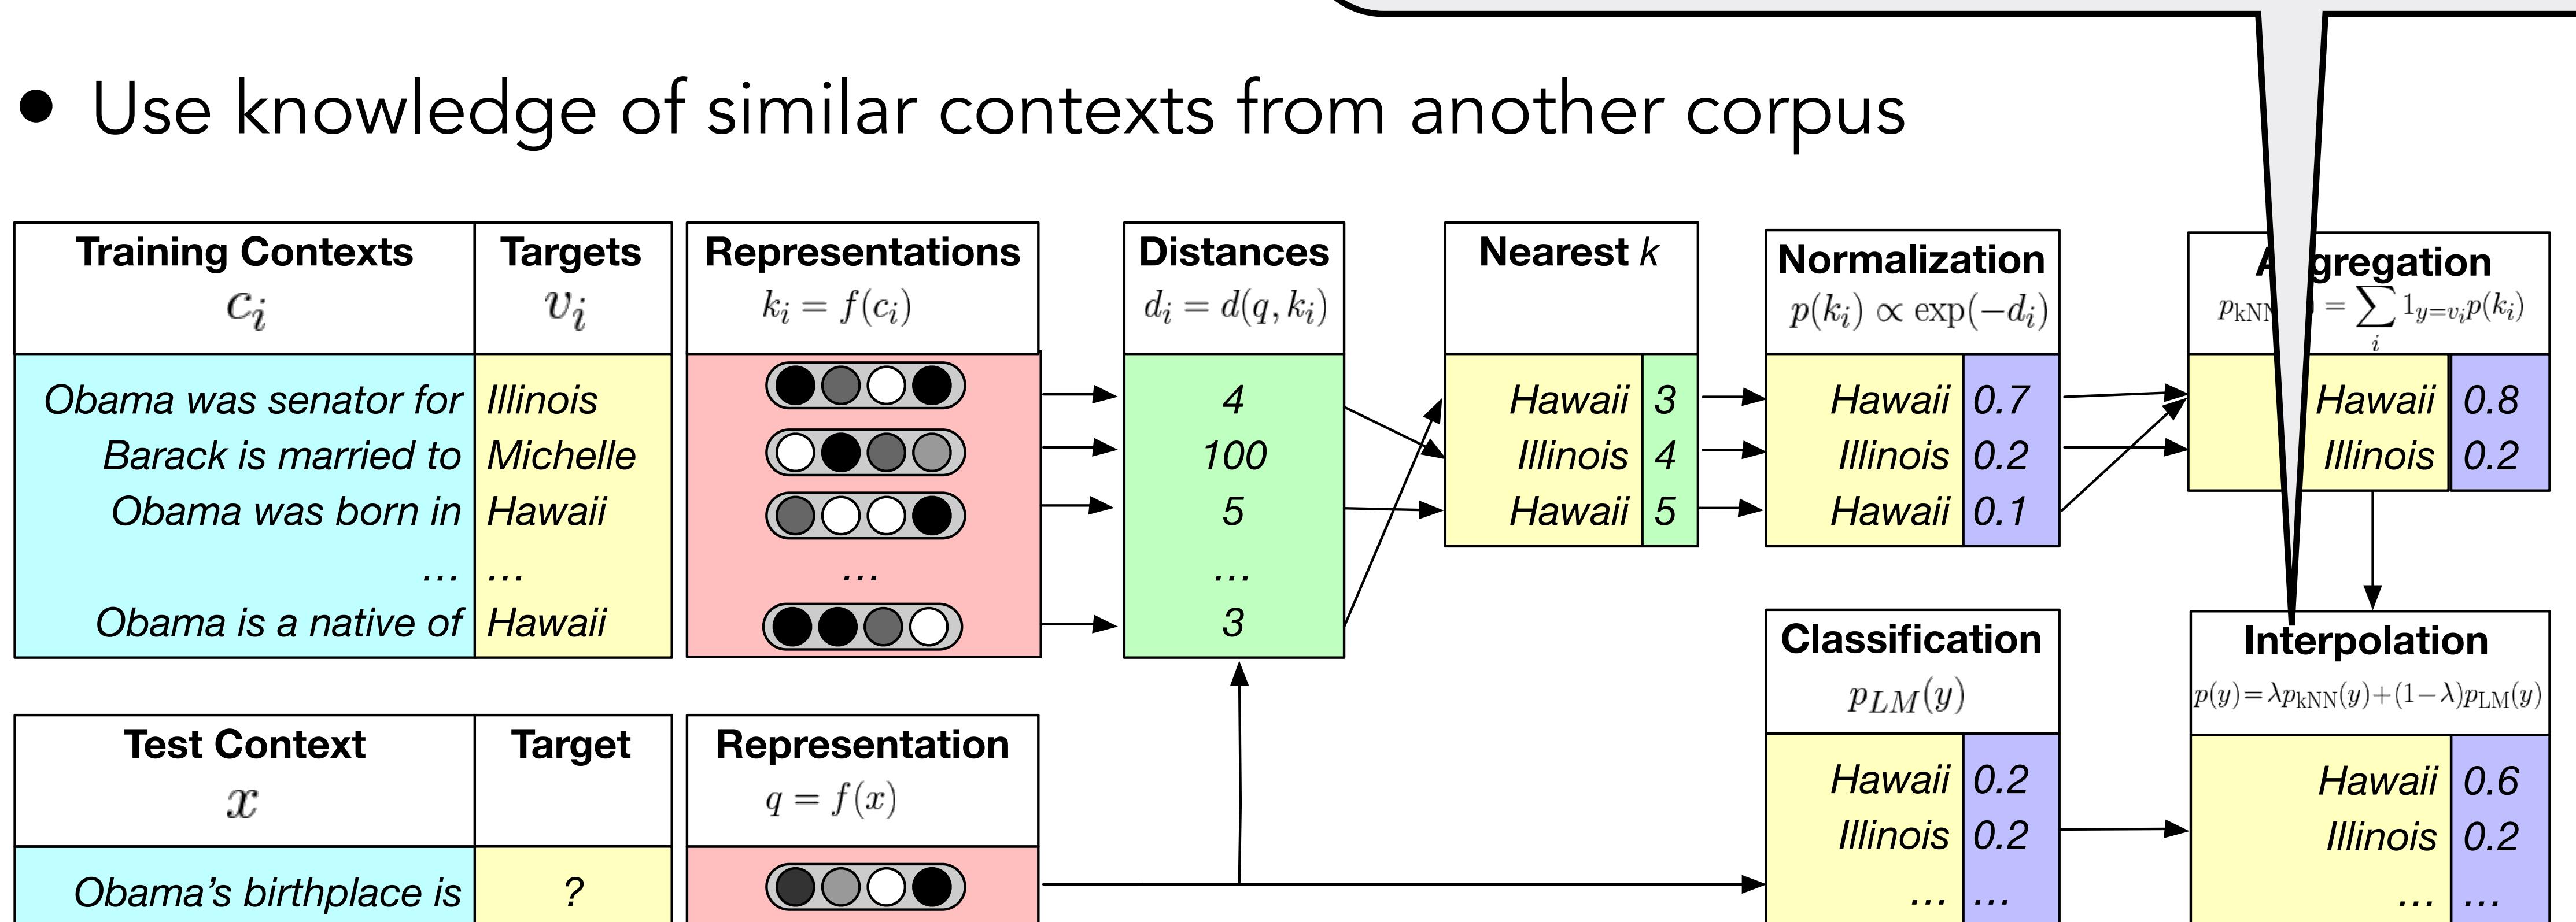
Compute **distribution** over possible **targets** from context DB sequences using **distance of history**

- Use knowledge of similar contexts from another corpus



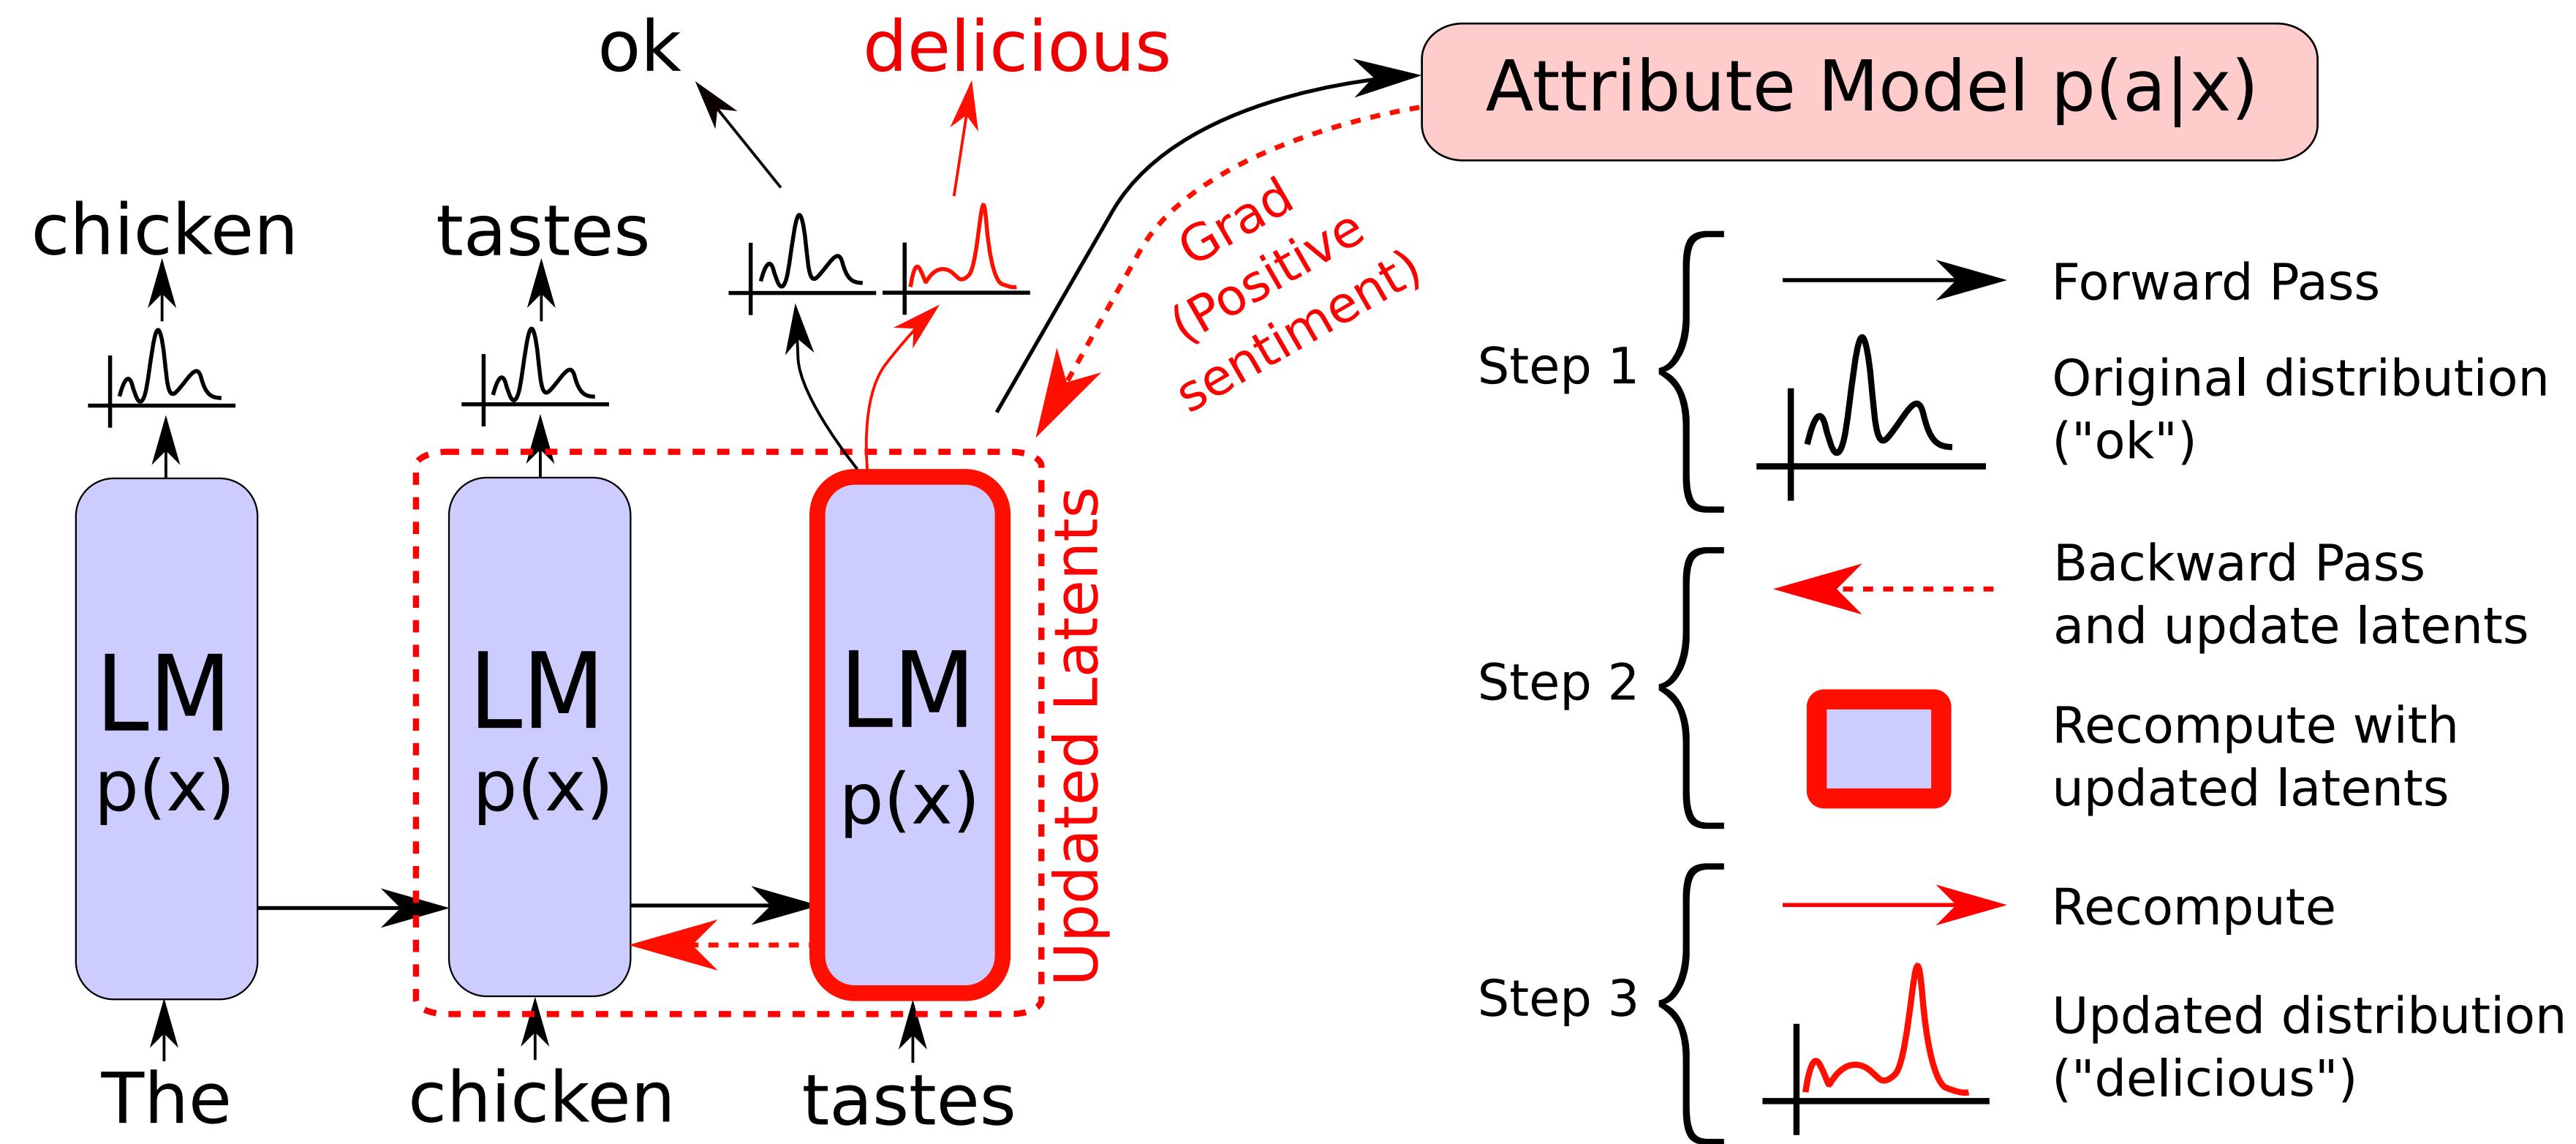
kNN Language Models

- Don't just rely on your trained model to generate a distribution over tokens
- Use knowledge of similar contexts from another corpus



Plug and Play Language Models!

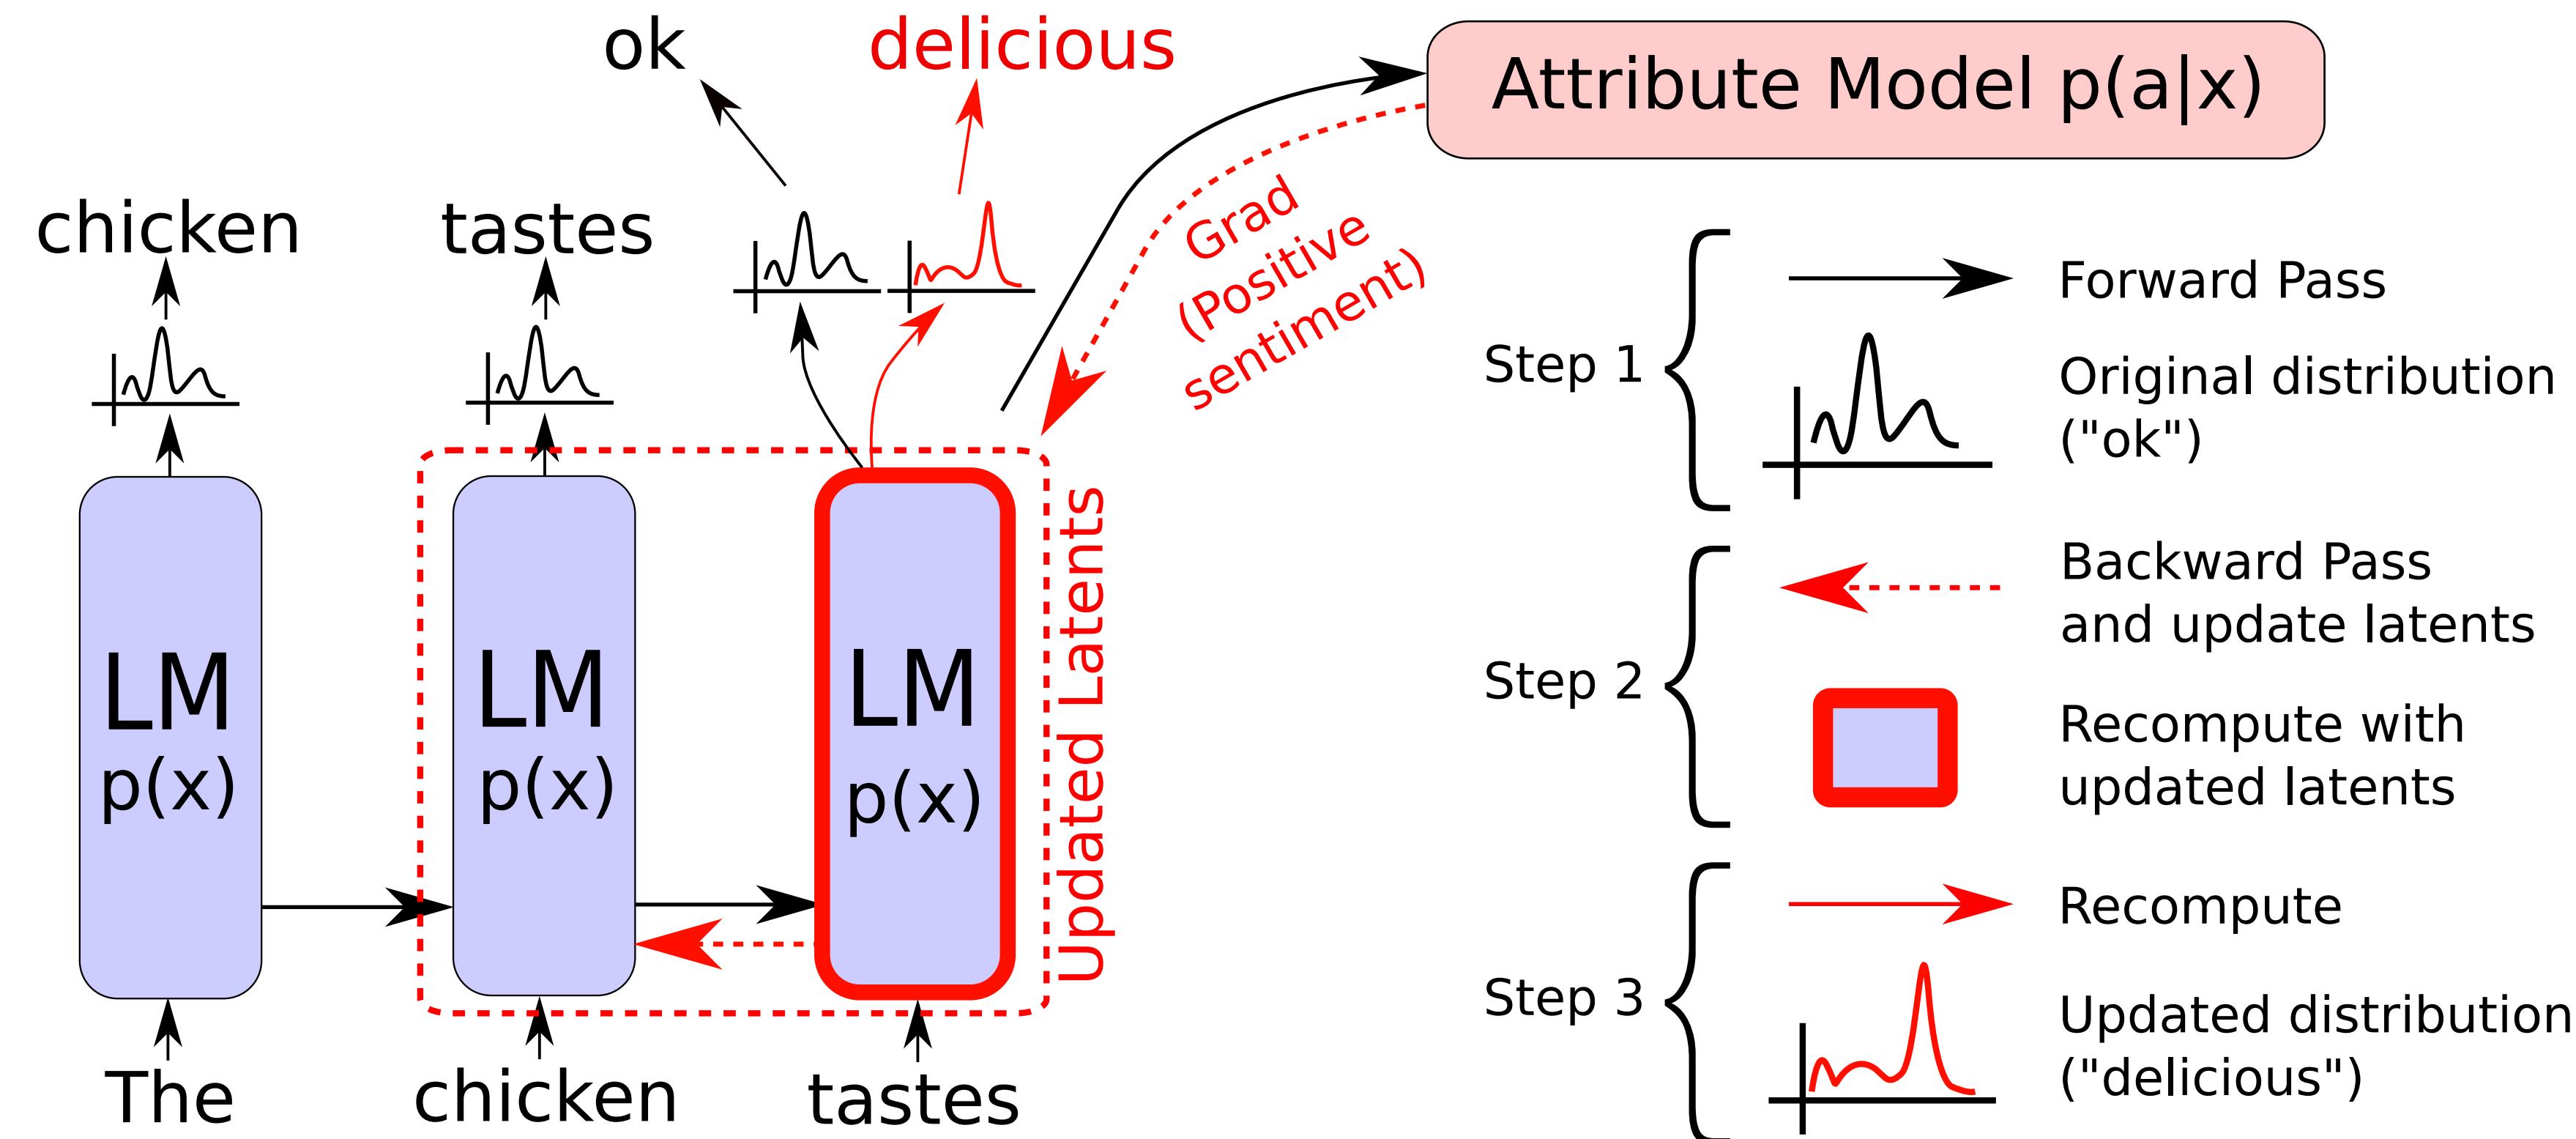
- What if I want to encourage a tough to formalize behavior at inference time?



Plug and Play

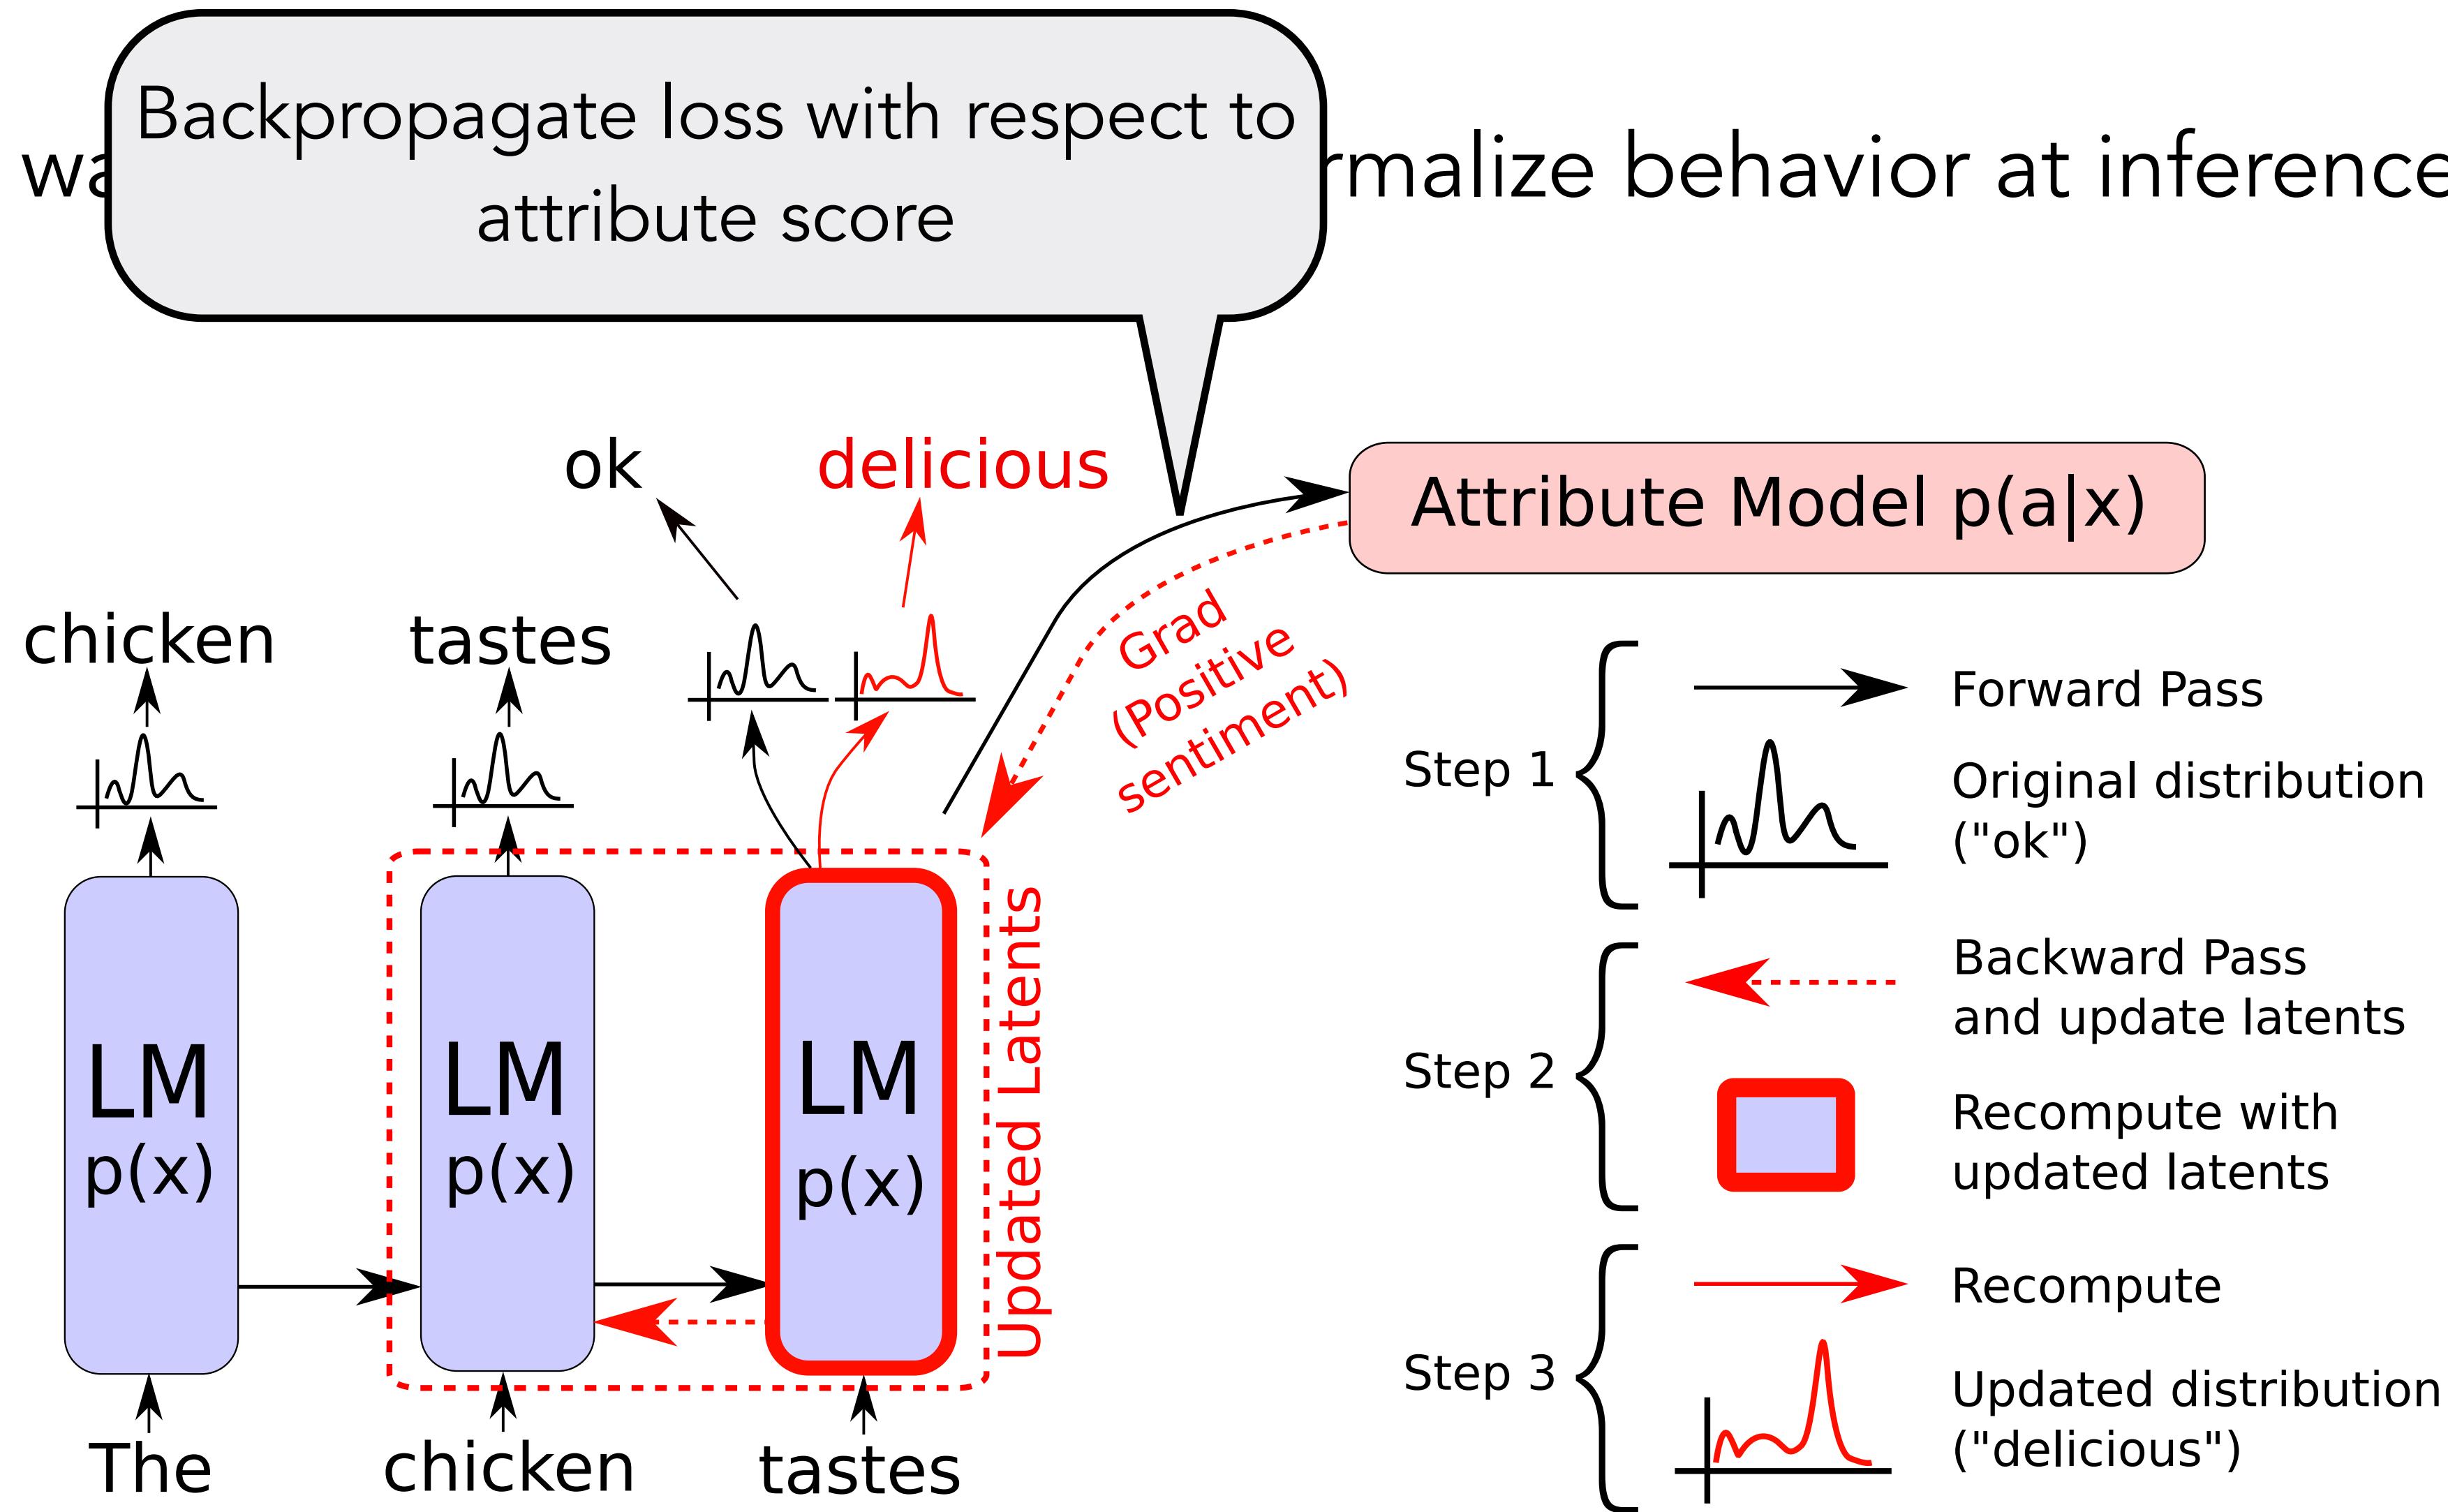
Define an attribute model that scores the generated sequence. Each generated token must try to increase the score given to the sequence by the attribute model

- What if I want to encourage positive sentiment over time?



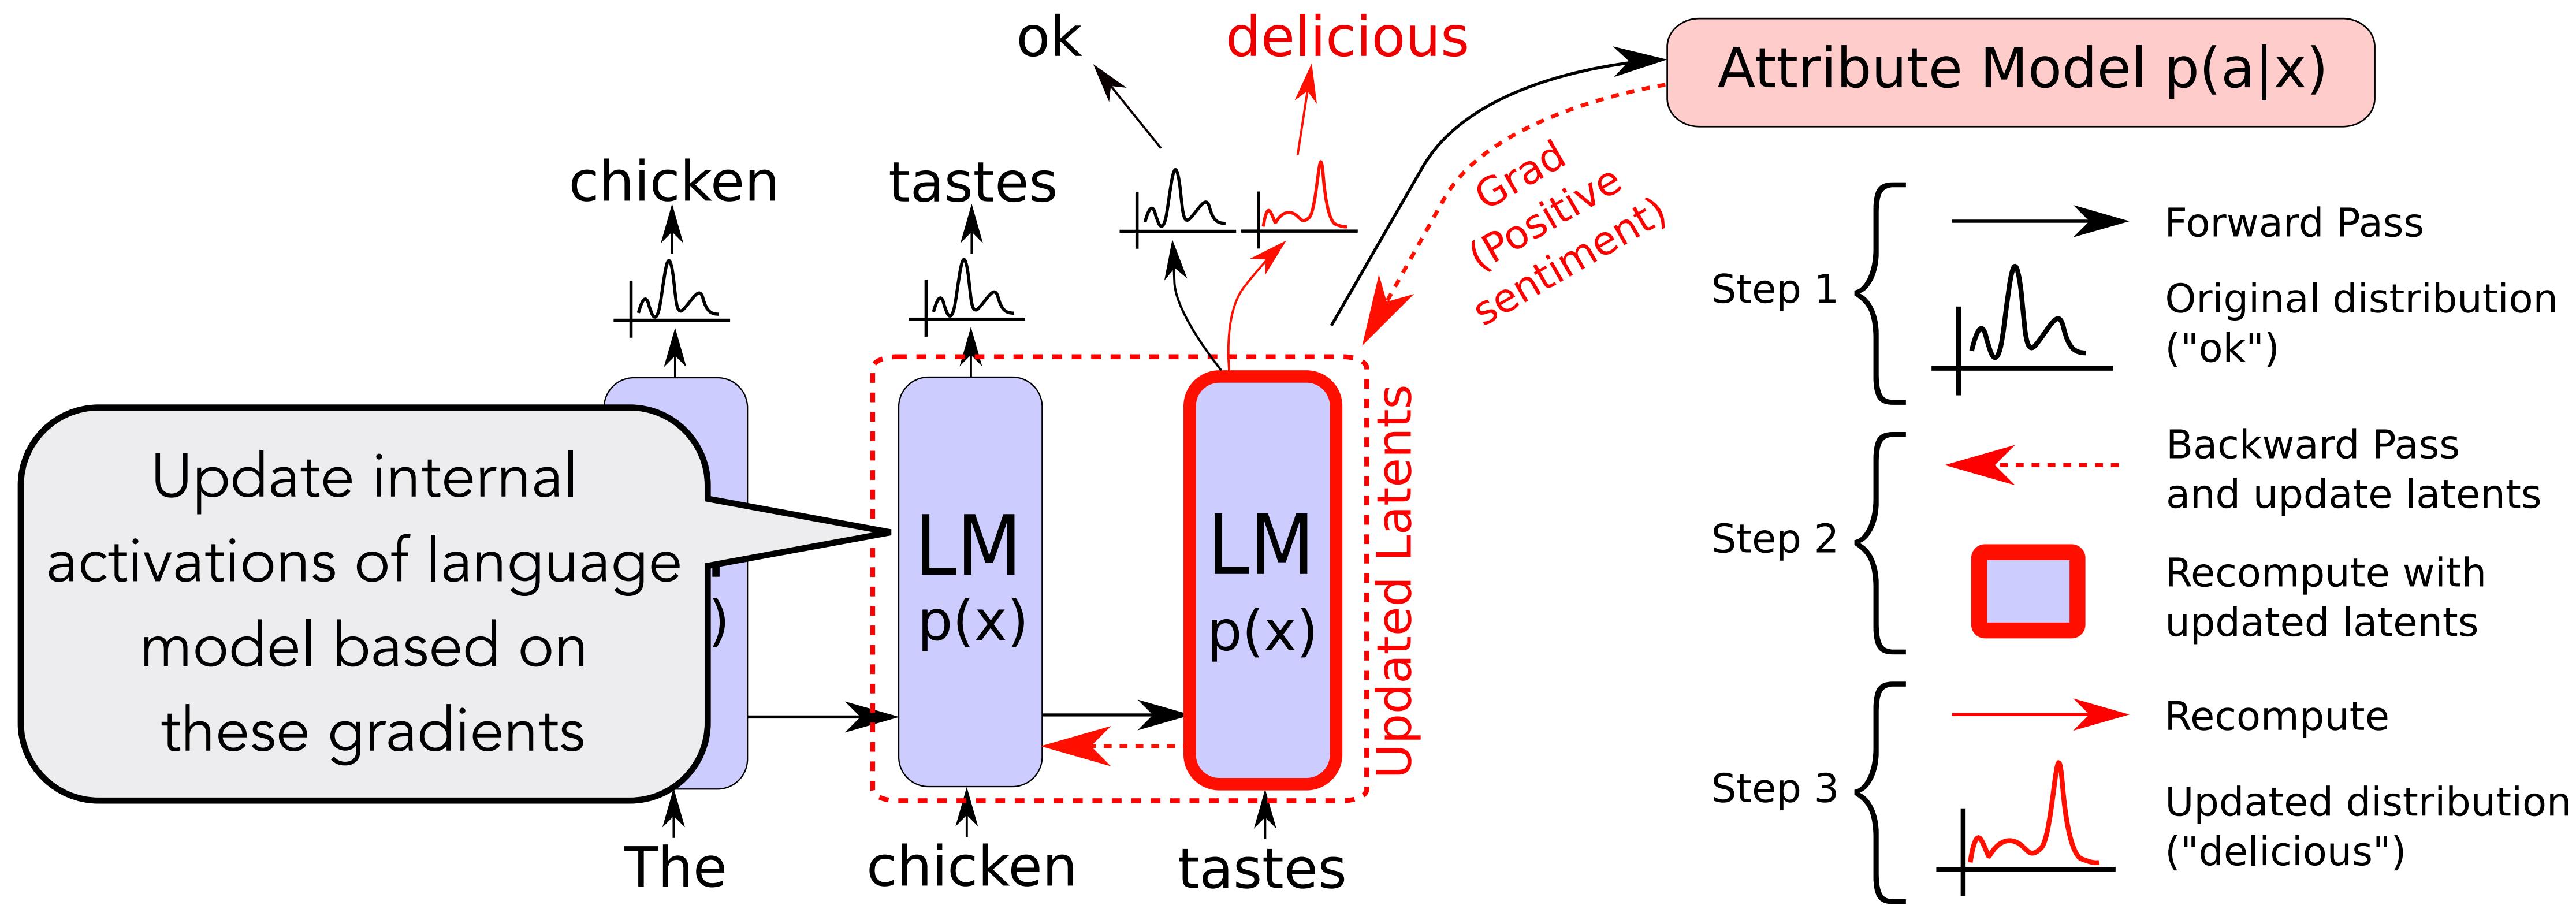
Plug and Play Language Models!

- What if I want to normalize behavior at inference time?



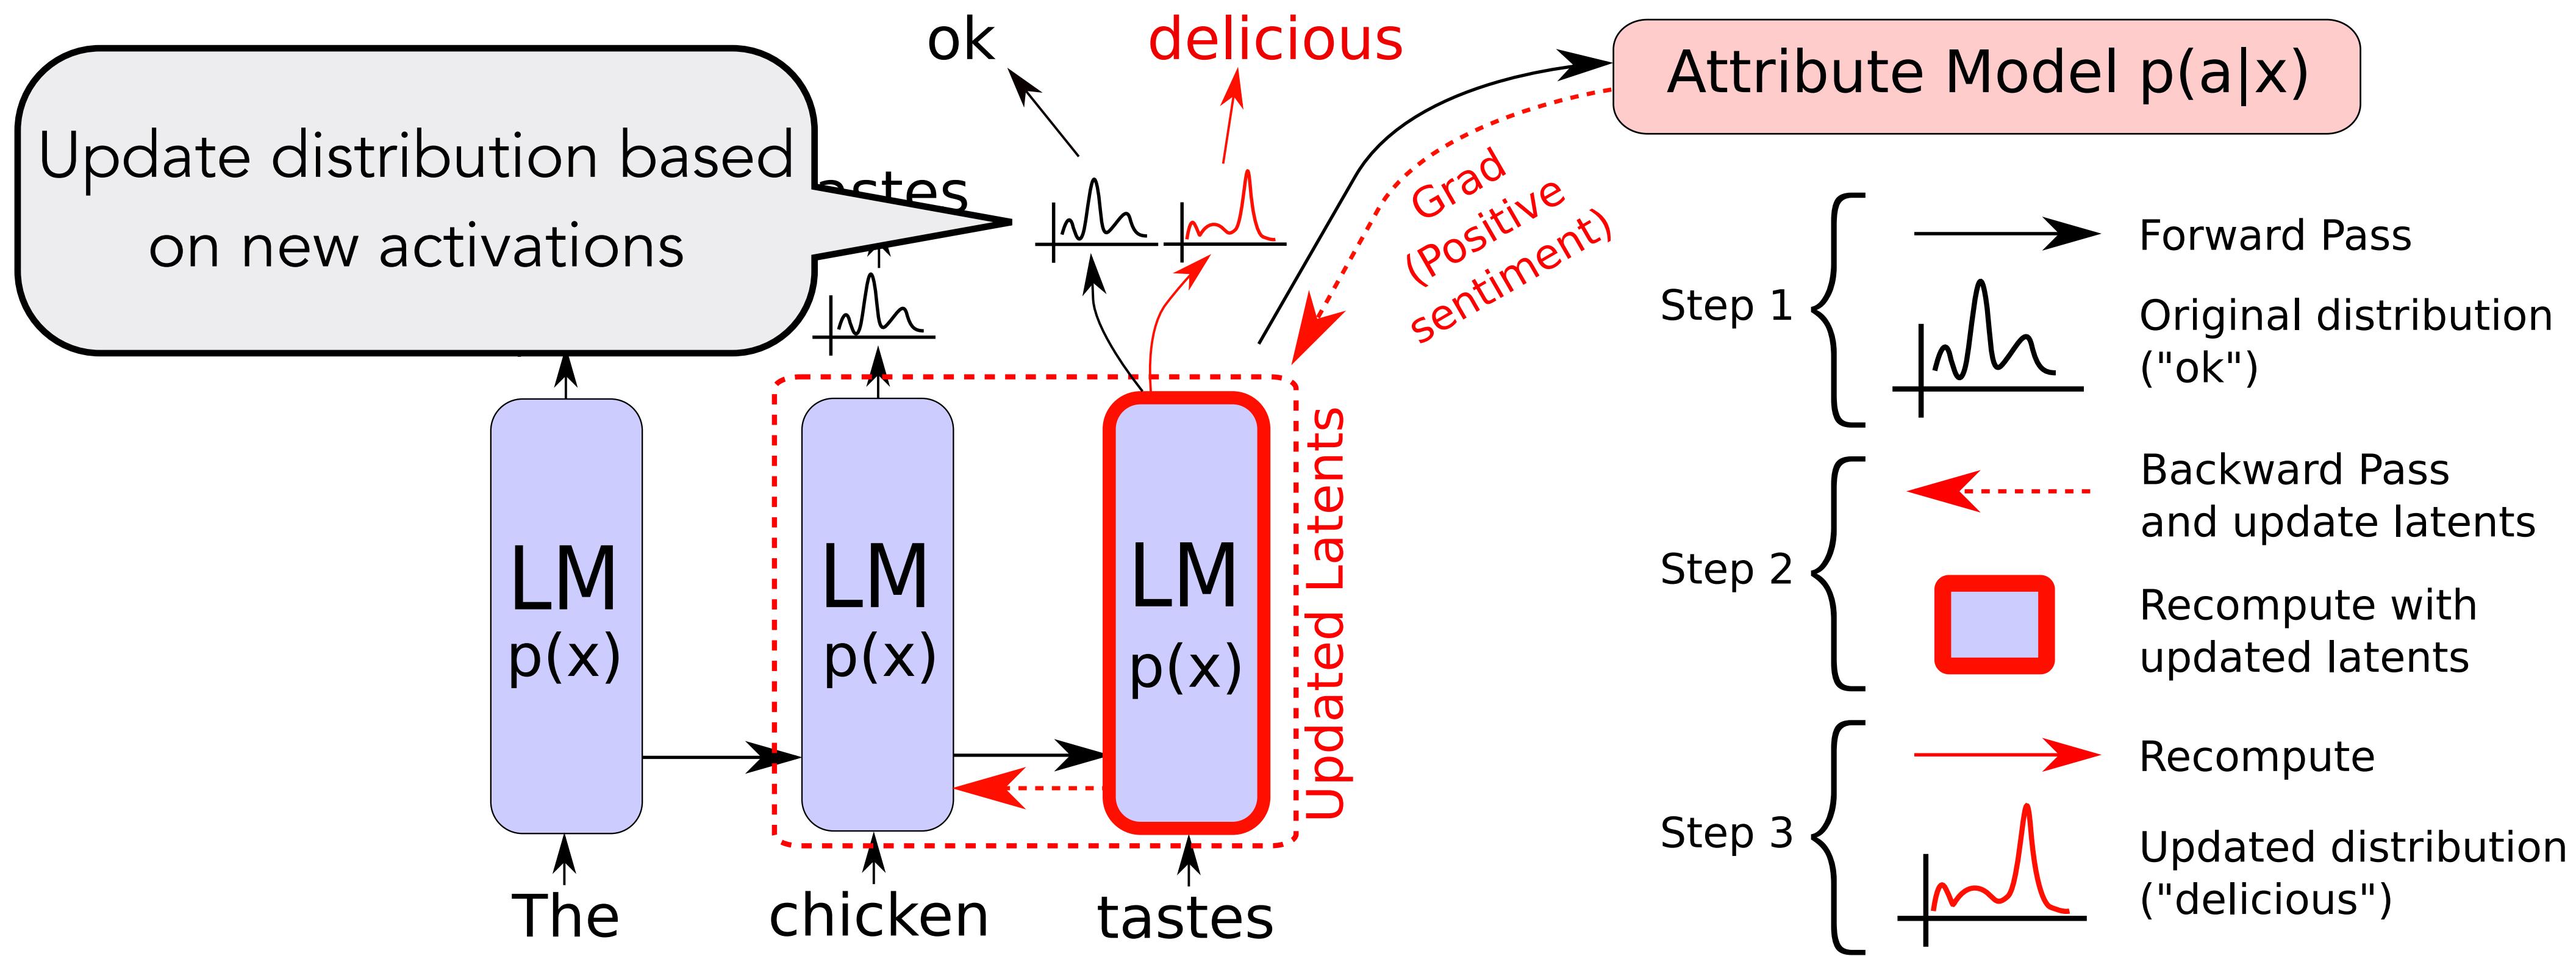
Plug and Play Language Models!

- What if I want to encourage a tough to formalize behavior at inference time?



Plug and Play Language Models!

- What if I want to encourage a tough to formalize behavior at inference time?



This all sounds a bit *risky*

- What if my distribution isn't great to begin with? Can I mitigate this?
 1. Make your word-level distributions more robust
 - kNN LMs
 - Plug and play LMs
 2. Optimize other sequence-level scores
 - Re-ranking
 - DeLorean

Re-ranking

- Generate a bunch of sequences with a sampling-based procedure
- Define a score to approximate the quality of your sequence.
- Simplest is to just use perplexity!
- However, re-rankers can be used to score a variety of properties: style (Holtzman et al., 2018), discourse (Gabriel et al., 2019), entailment (Pasunuru et al., 2018), logical consistency (Li et al., 2020; Lu et al., 2020)
- TODO: Make this more exciting



DELOREAN

(DEcoding for nonmonotonic LOgical REAsoNing)



Initialization

Any decoding
algorithm will do!

Y

$x_1 \quad x_2 \quad \dots \quad x_{N_X}$

Ray hung a tire on a
rope to make his
daughter a swing.

LM

Ray ran to his
daughter to make
sure she was okay.

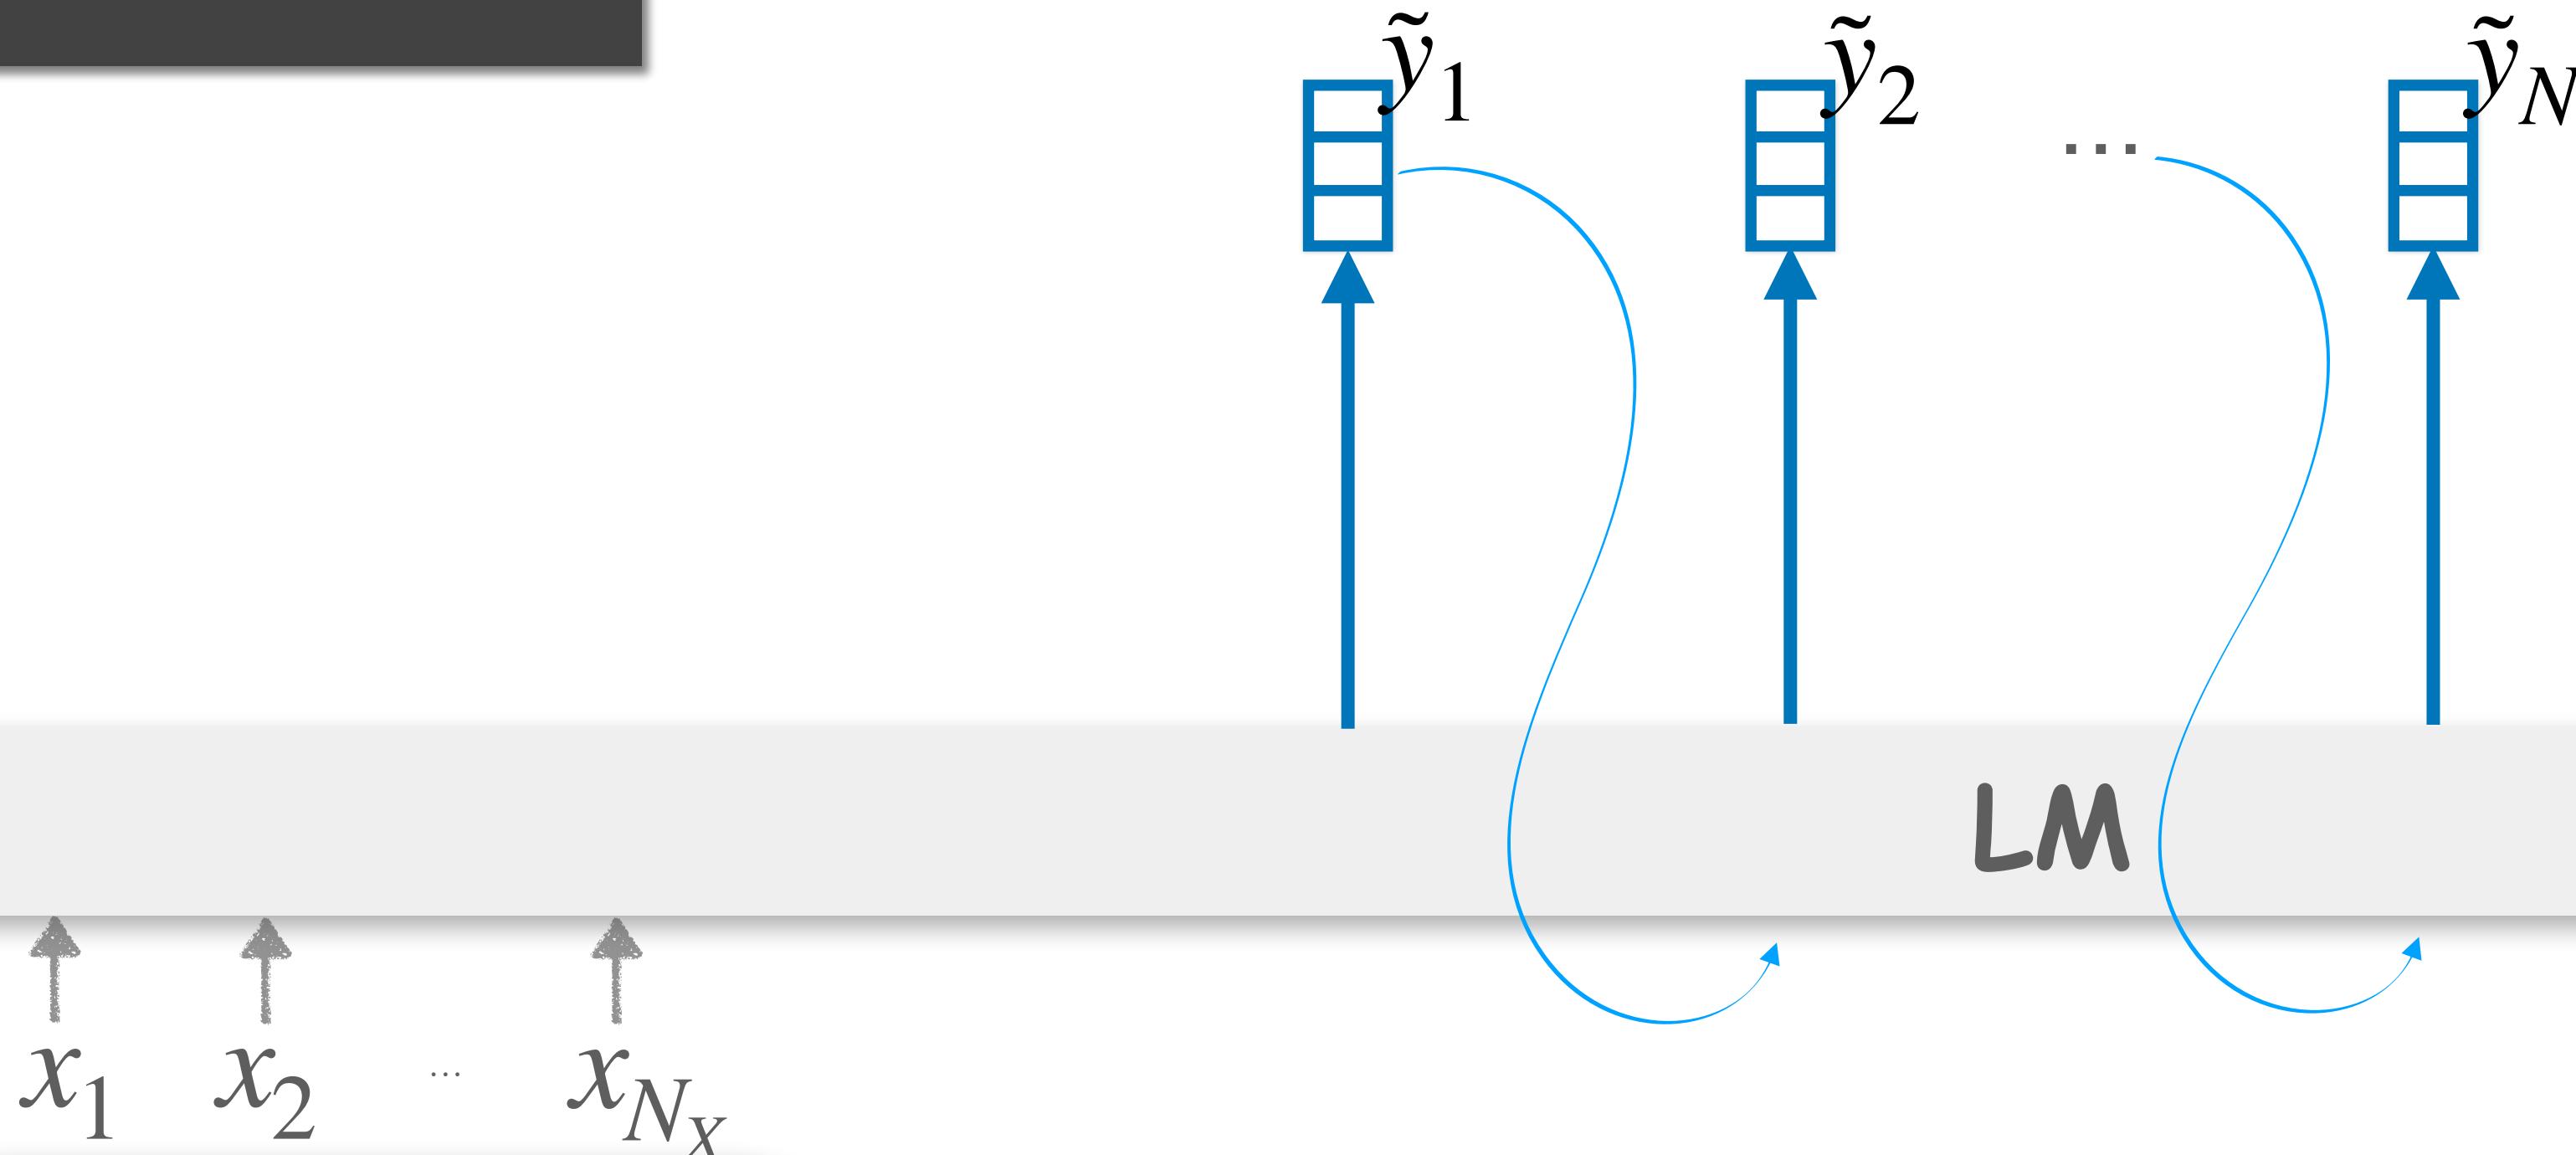
X

Z



Initialization

Any decoding algorithm will do!



Ray hung a tire on a
rope to make his
daughter a swing.

X

Y

Ray ran to his
daughter to make
sure she was okay.

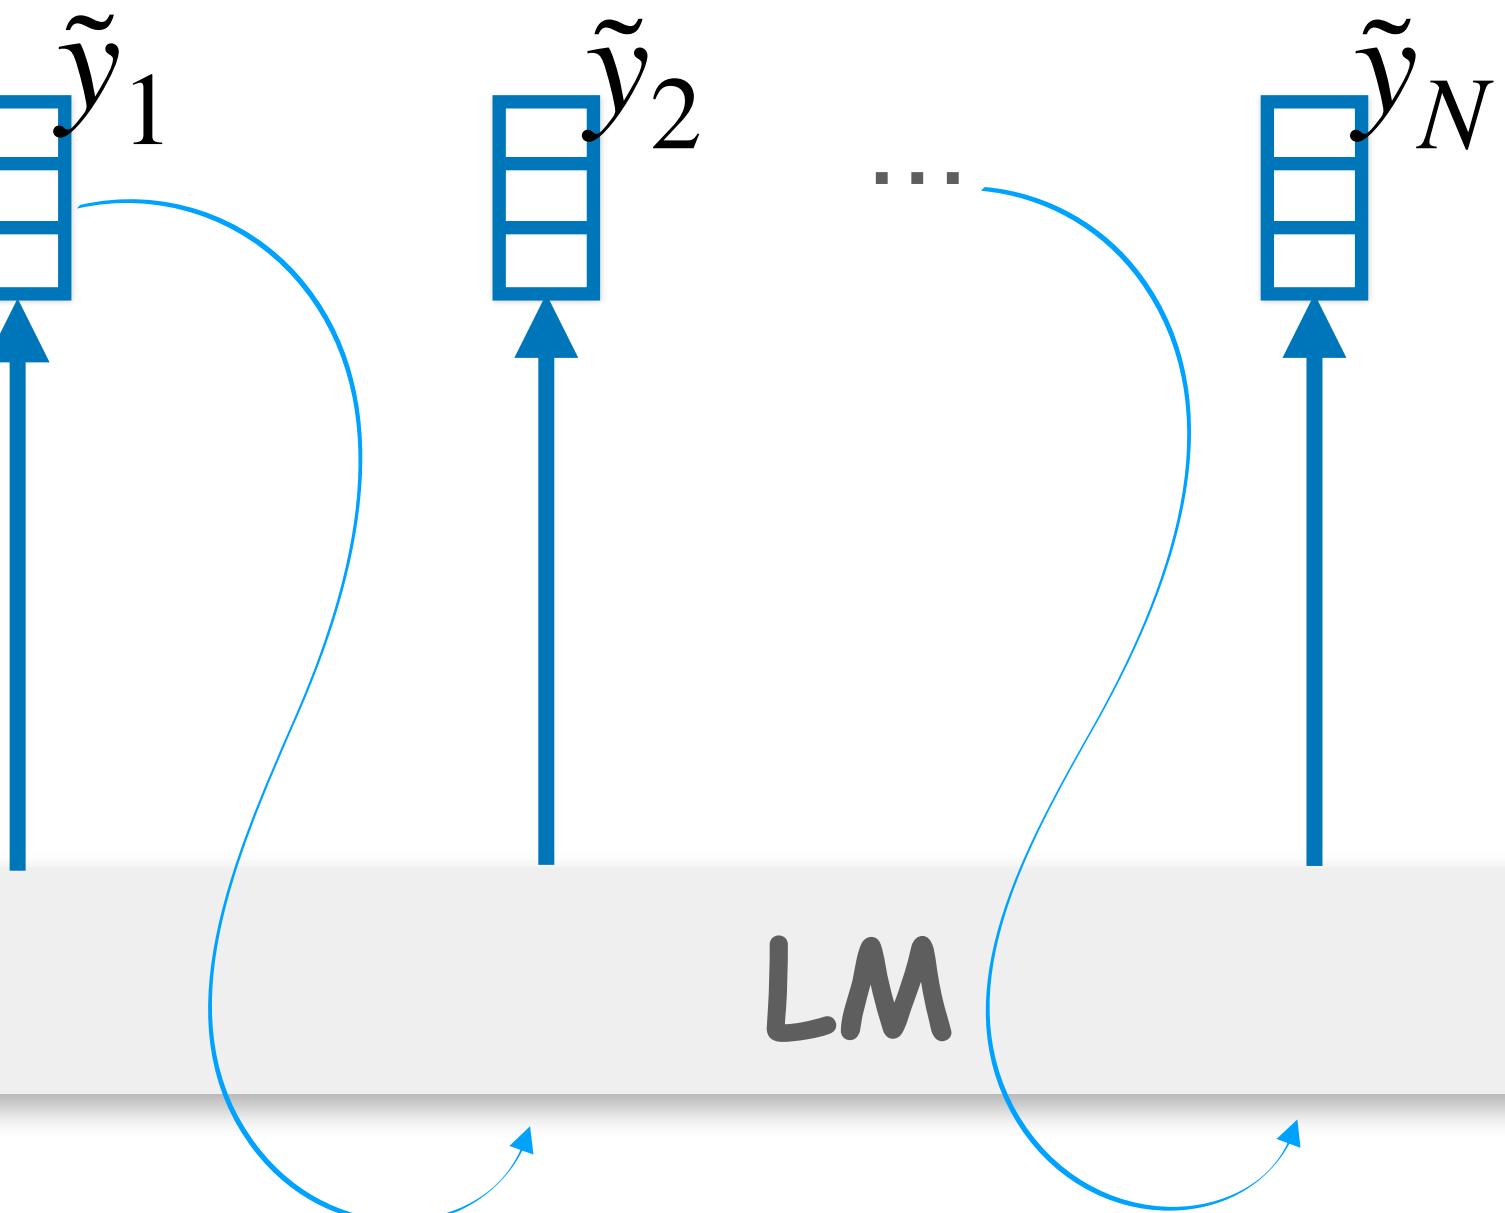
Z



Backward Pass

Backpropagate
future information

Y



Ray hung a tire on a
rope to make his
daughter a swing.

Ray ran to his
daughter to make
sure she was okay.

X

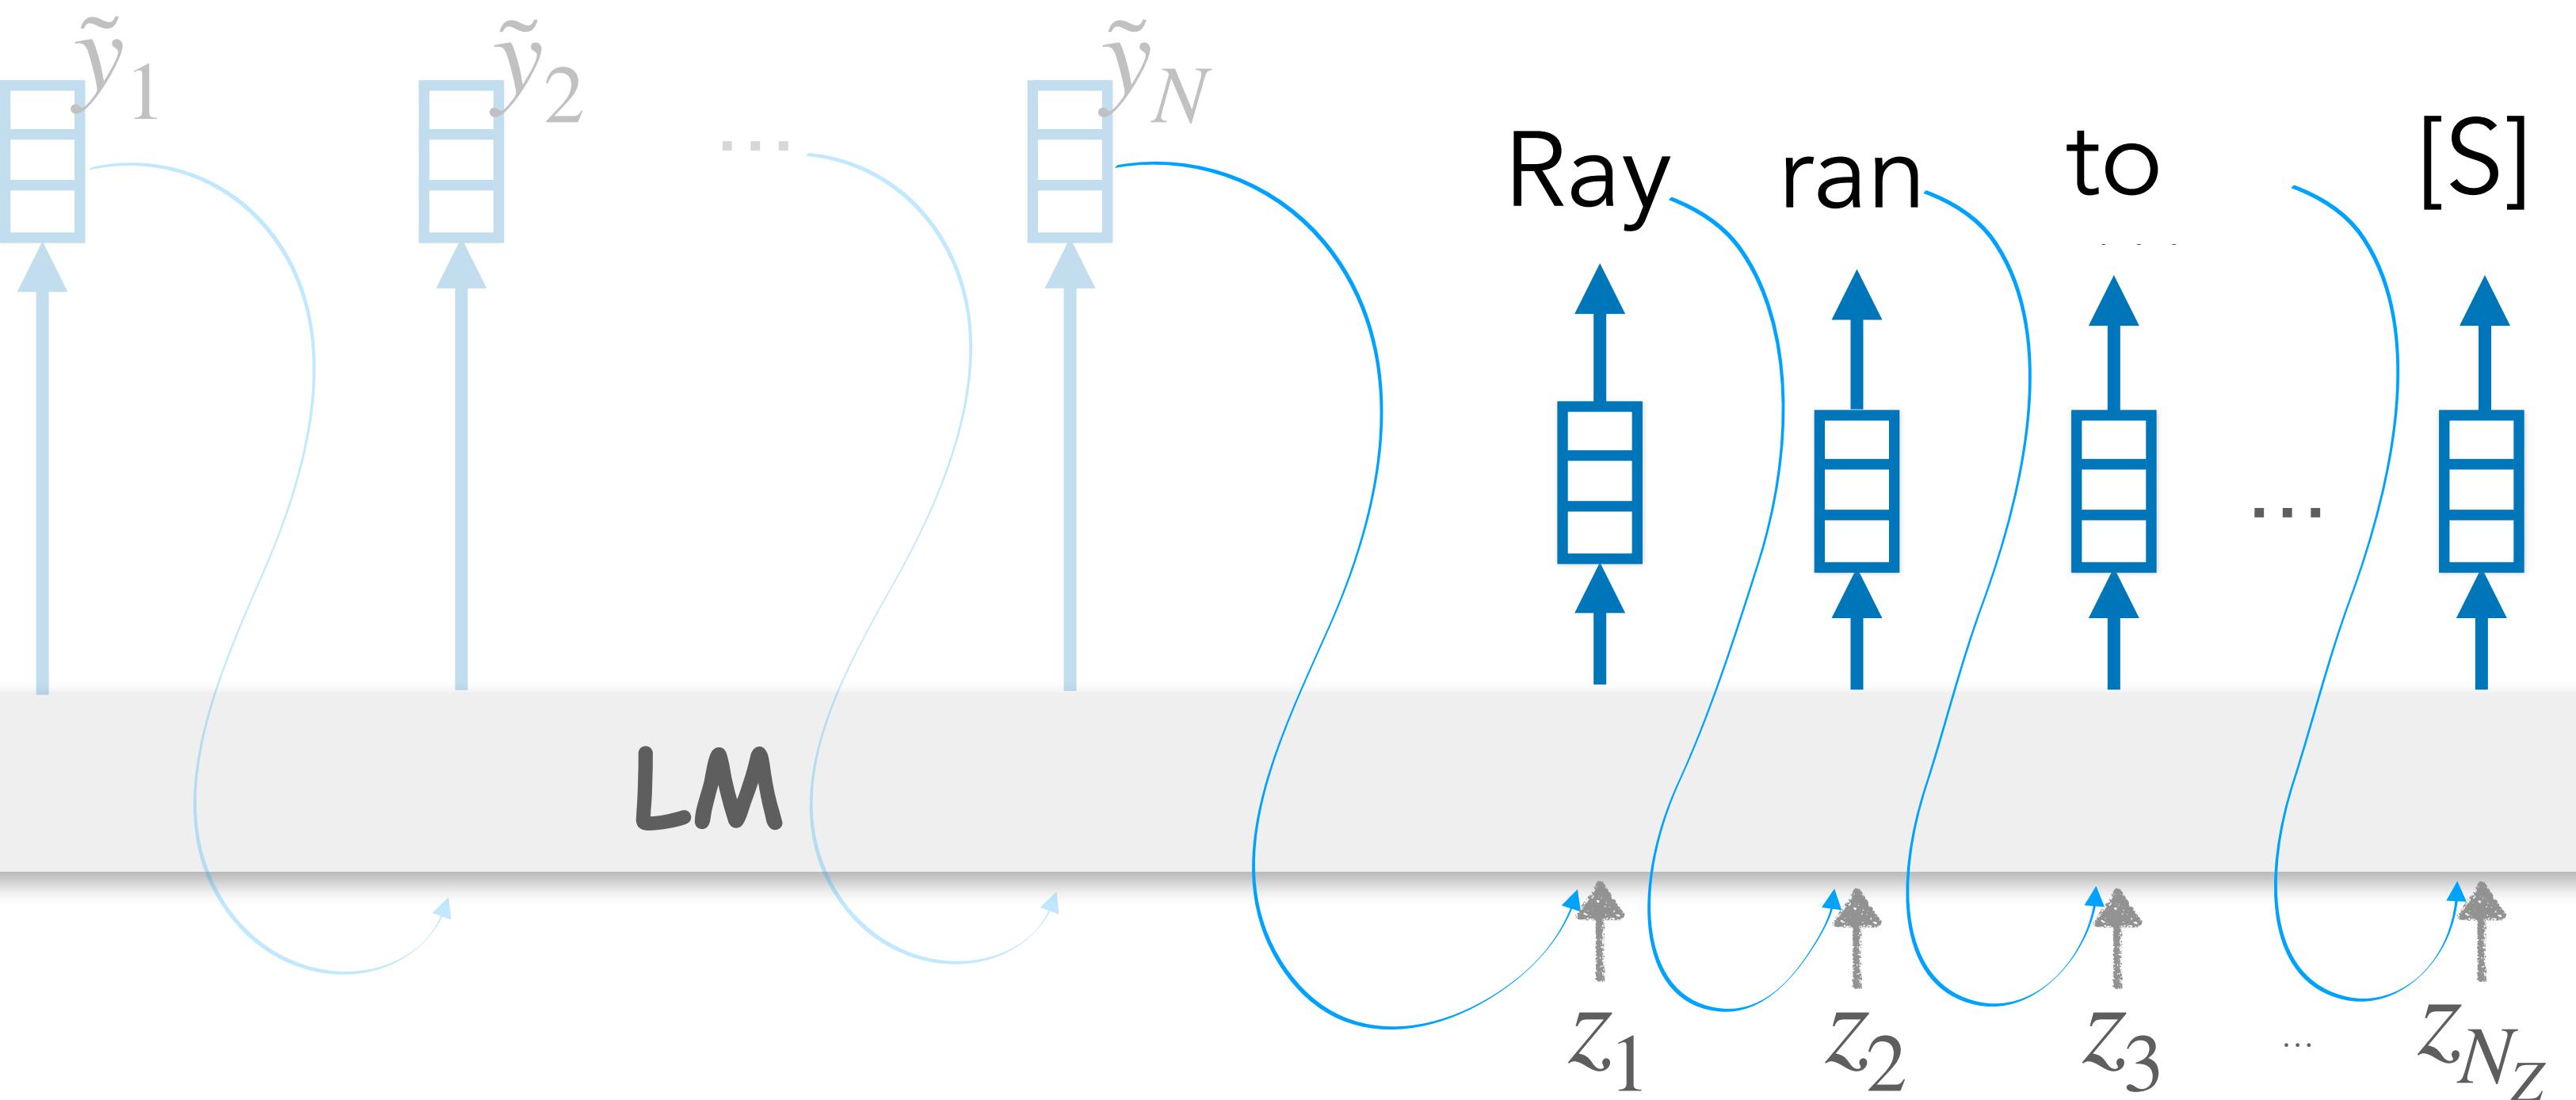
Z



Backward Pass

Backpropagate
future information

Y



Ray hung a tire on a
rope to make his
daughter a swing.

X

Ray ran to his
daughter to make
sure she was okay.

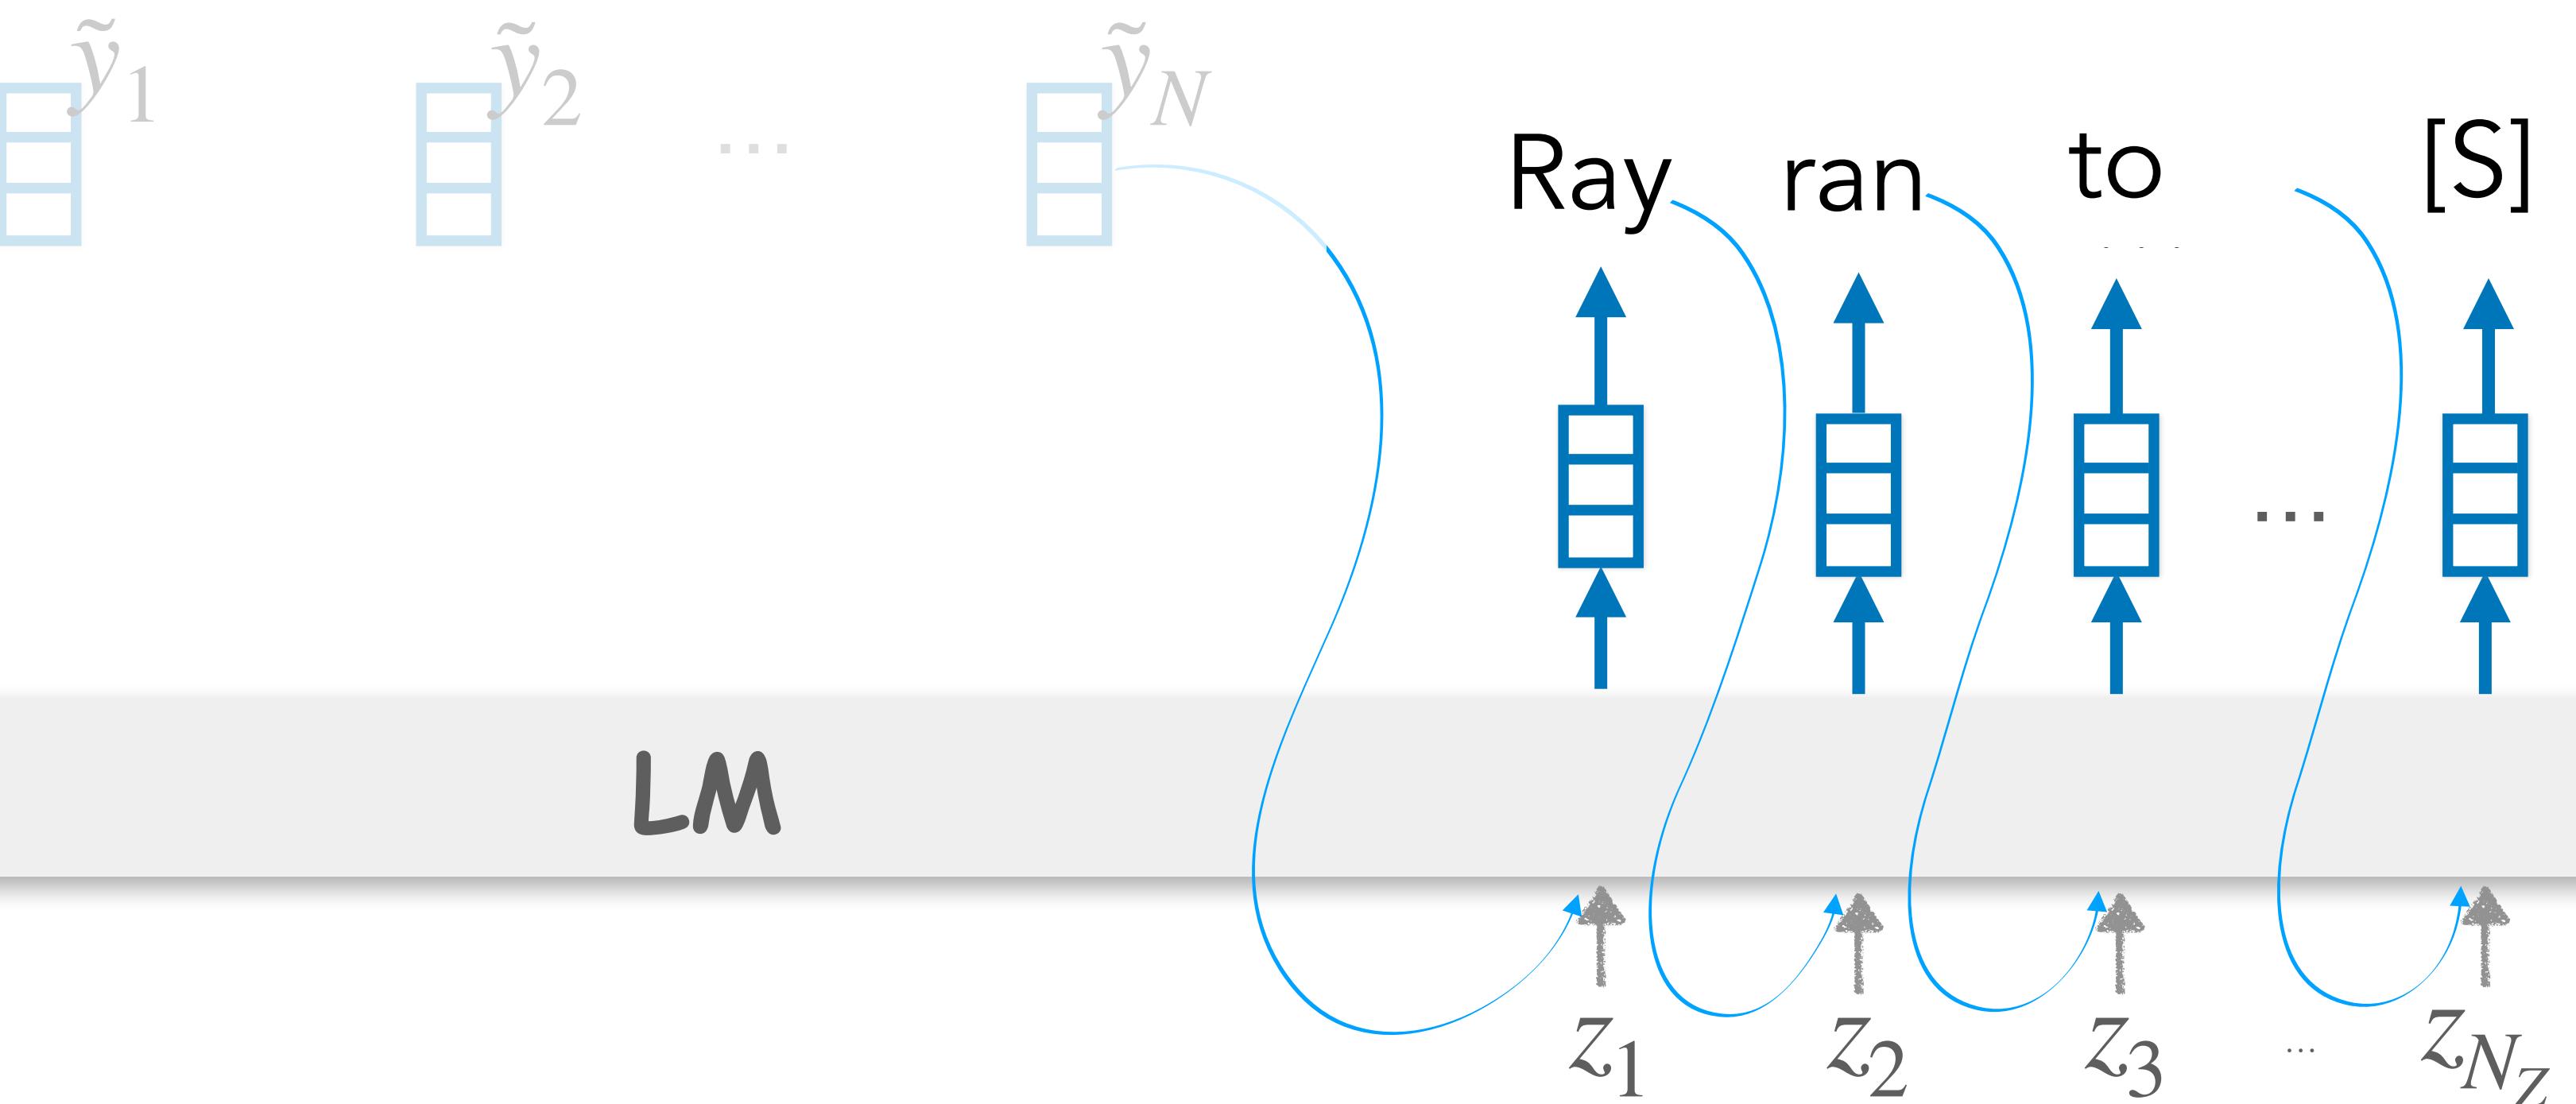
Z



Backward Pass

Backpropagate
future information

Y



Ray hung a tire on a
rope to make his
daughter a swing.

Ray ran to his
daughter to make
sure she was okay.

X

Z

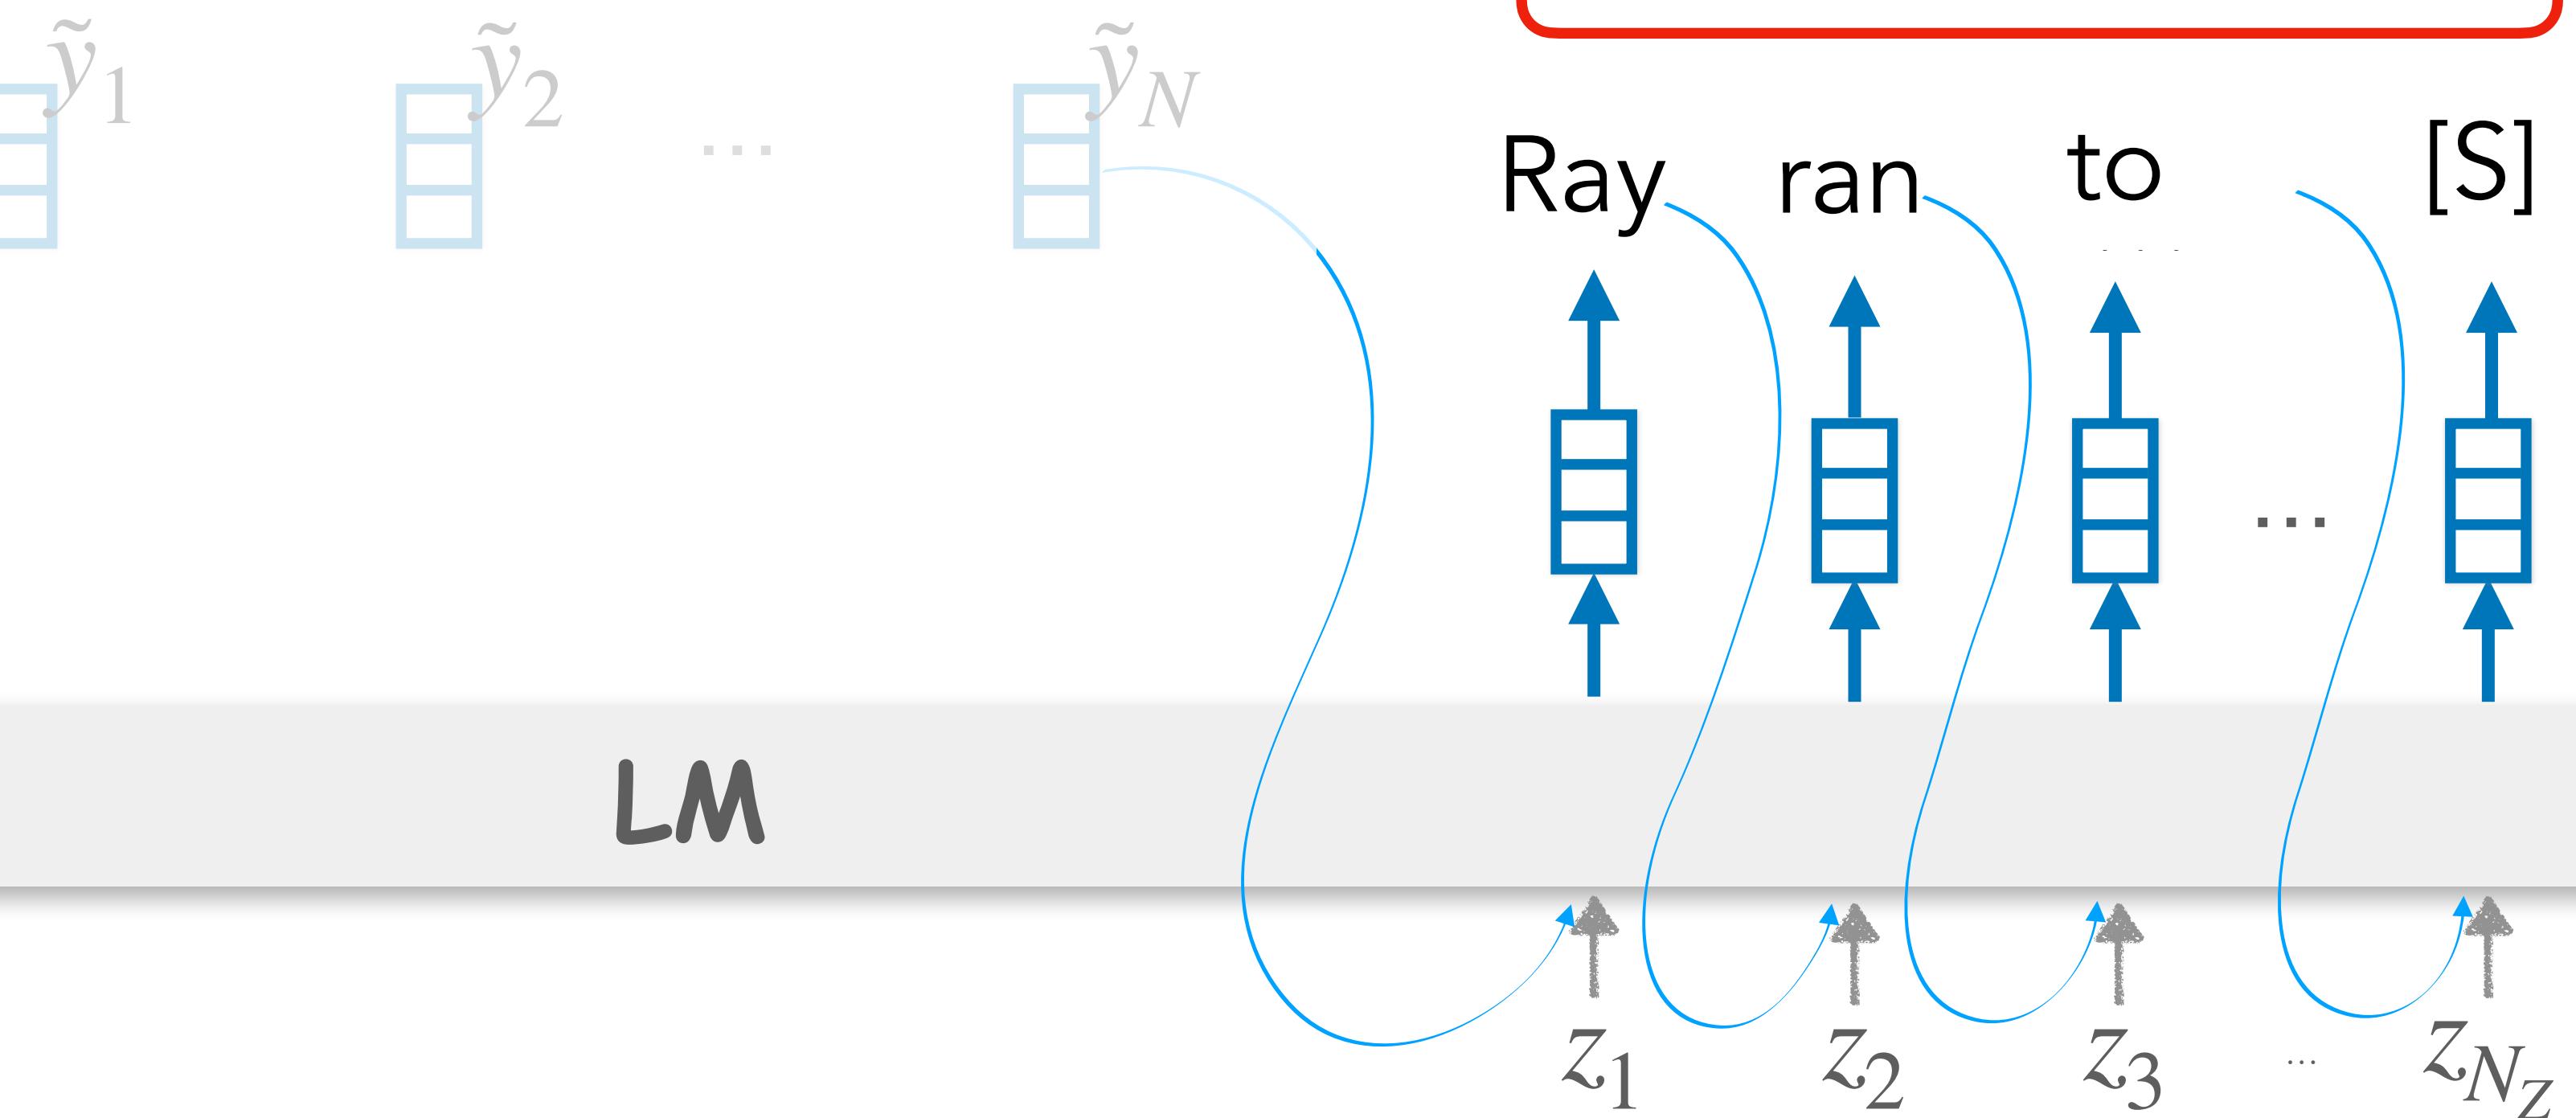


Backward Pass

Backpropagate
future information

Y

$$\mathcal{L}(X, \tilde{Y}, Z) := -\sum_{n=1}^{N_Z} \log P_{LM}(z_n | X, \tilde{Y}, Z_{1:n-1})$$



Ray hung a tire on a
rope to make his
daughter a swing.

X

Ray ran to his
daughter to make
sure she was okay.

Z



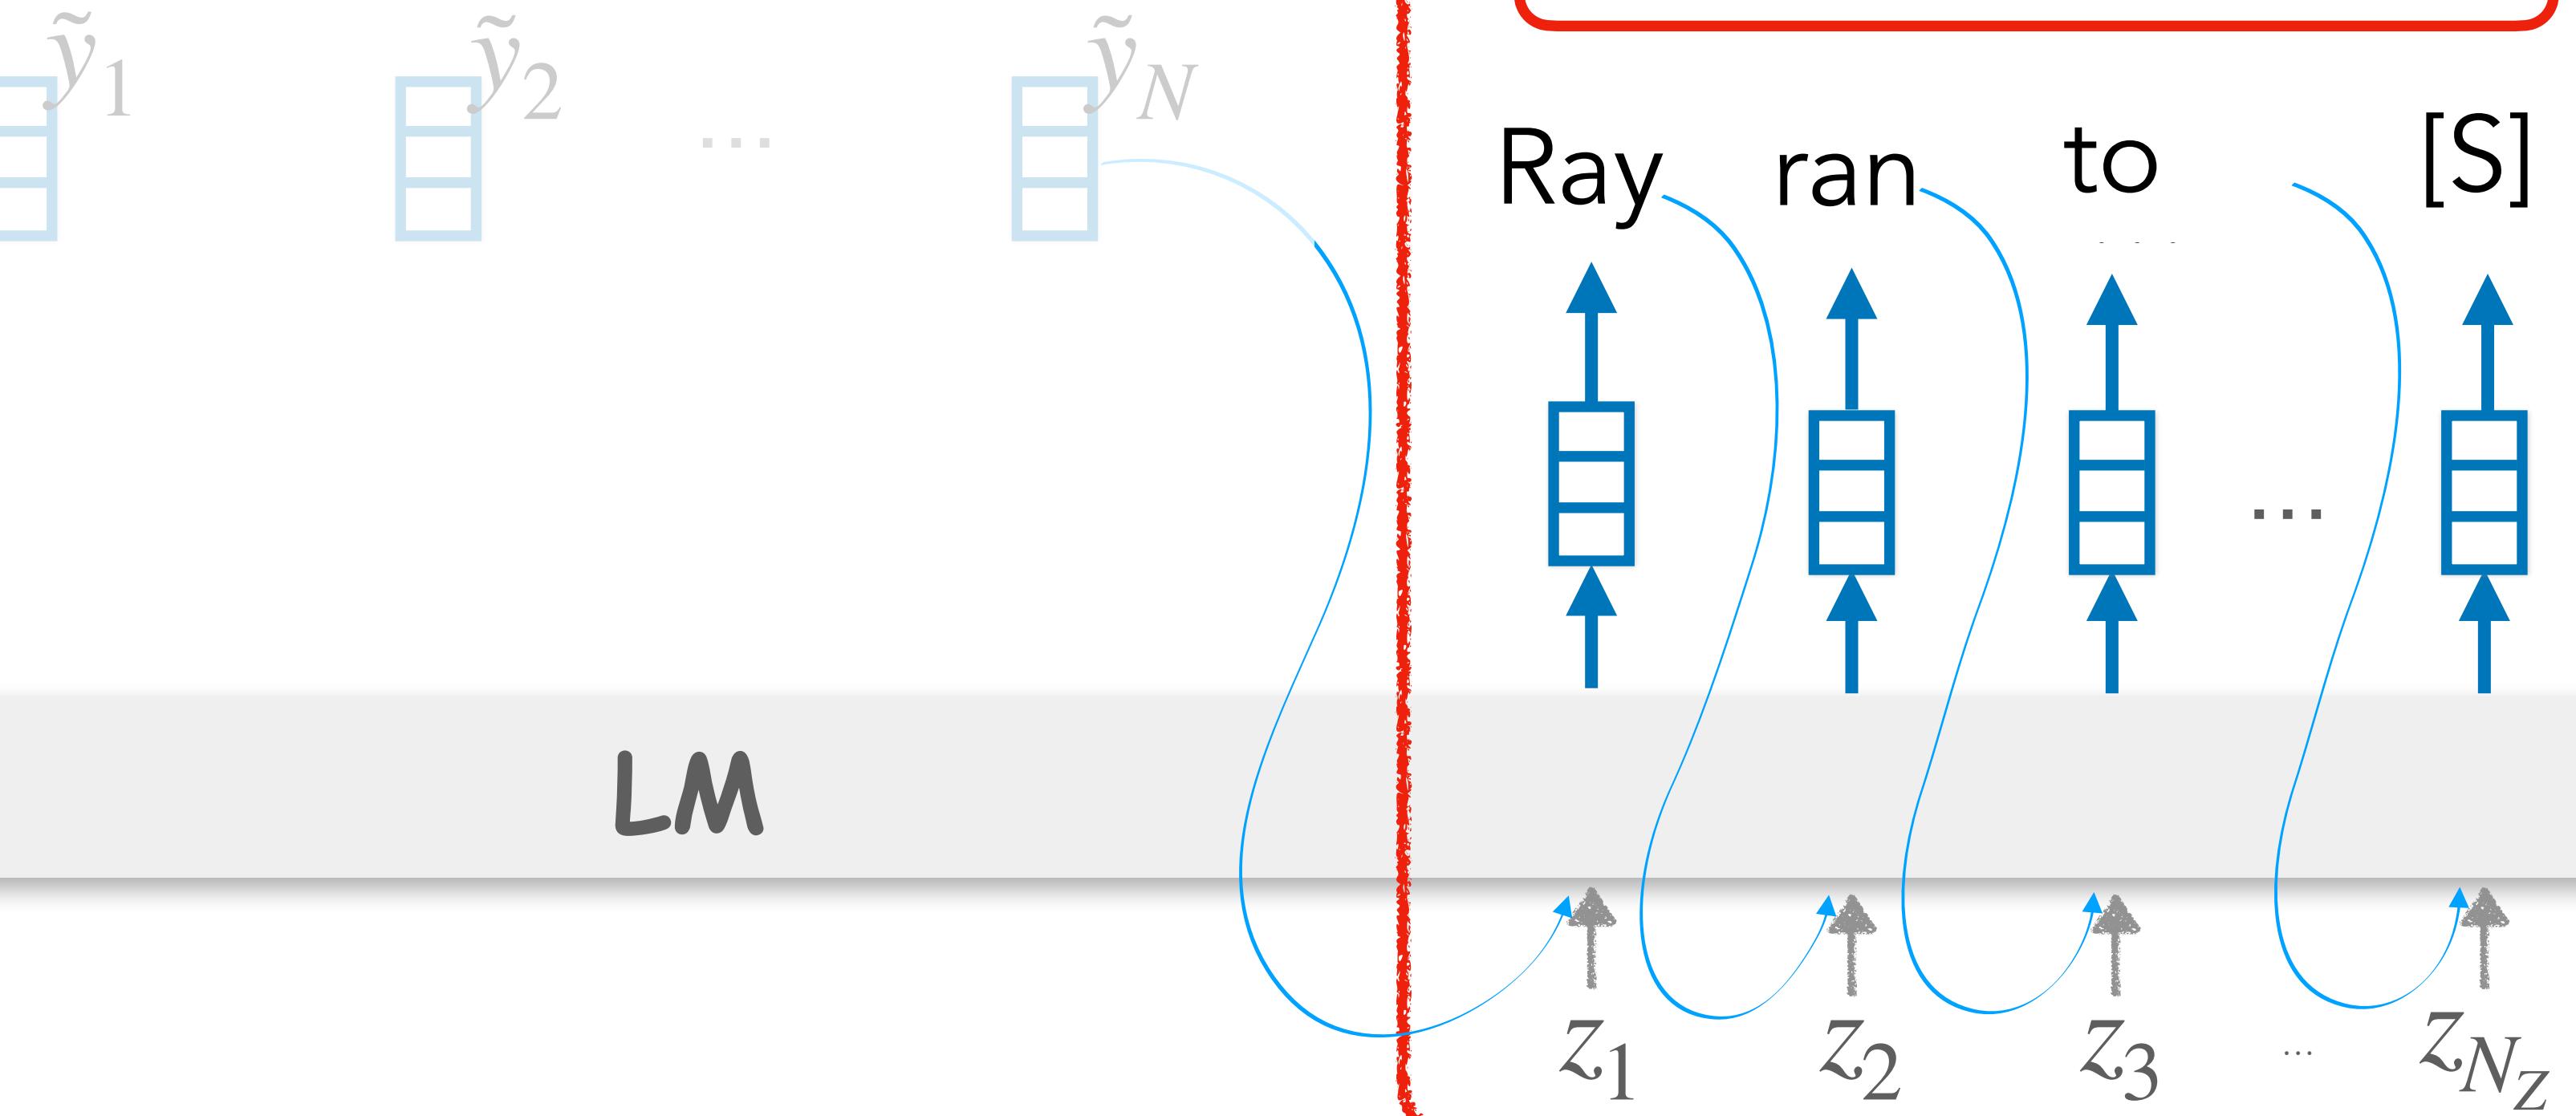
Backward Pass

Backpropagate
future information

Y

Backpropagation

$$\mathcal{L}(X, \tilde{Y}, Z) := -\sum_{n=1}^{N_Z} \log P_{LM}(z_n | X, \tilde{Y}, Z_{1:n-1})$$



Ray hung a tire on a
rope to make his
daughter a swing.

X

Z

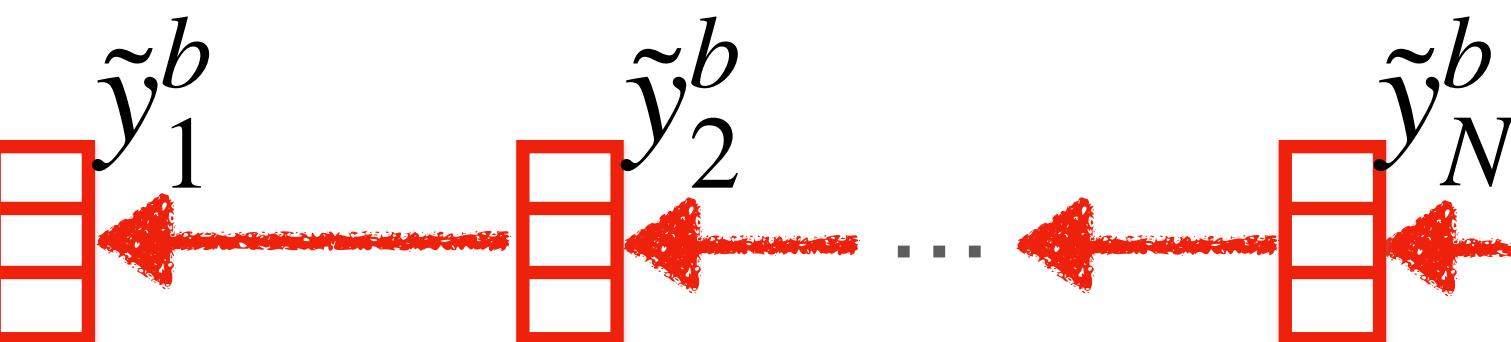
Ray ran to his
daughter to make
sure she was okay.



Backward Pass

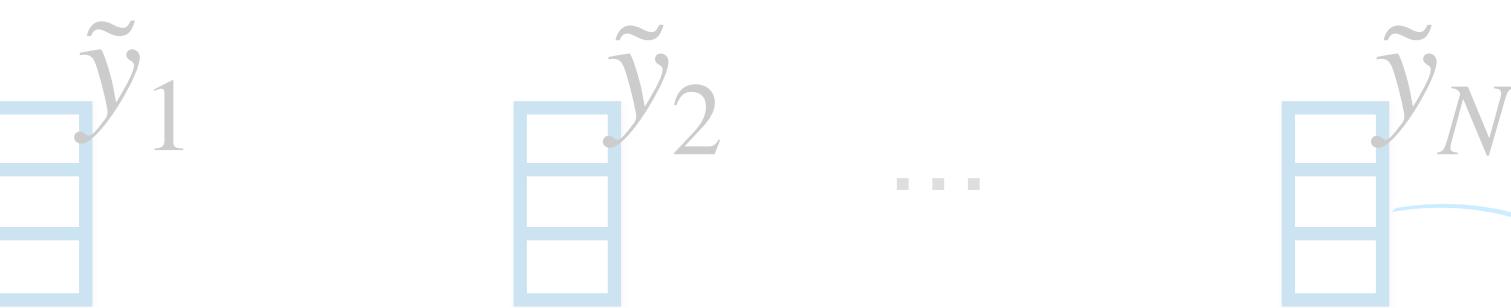
Backpropagate
future information

Y

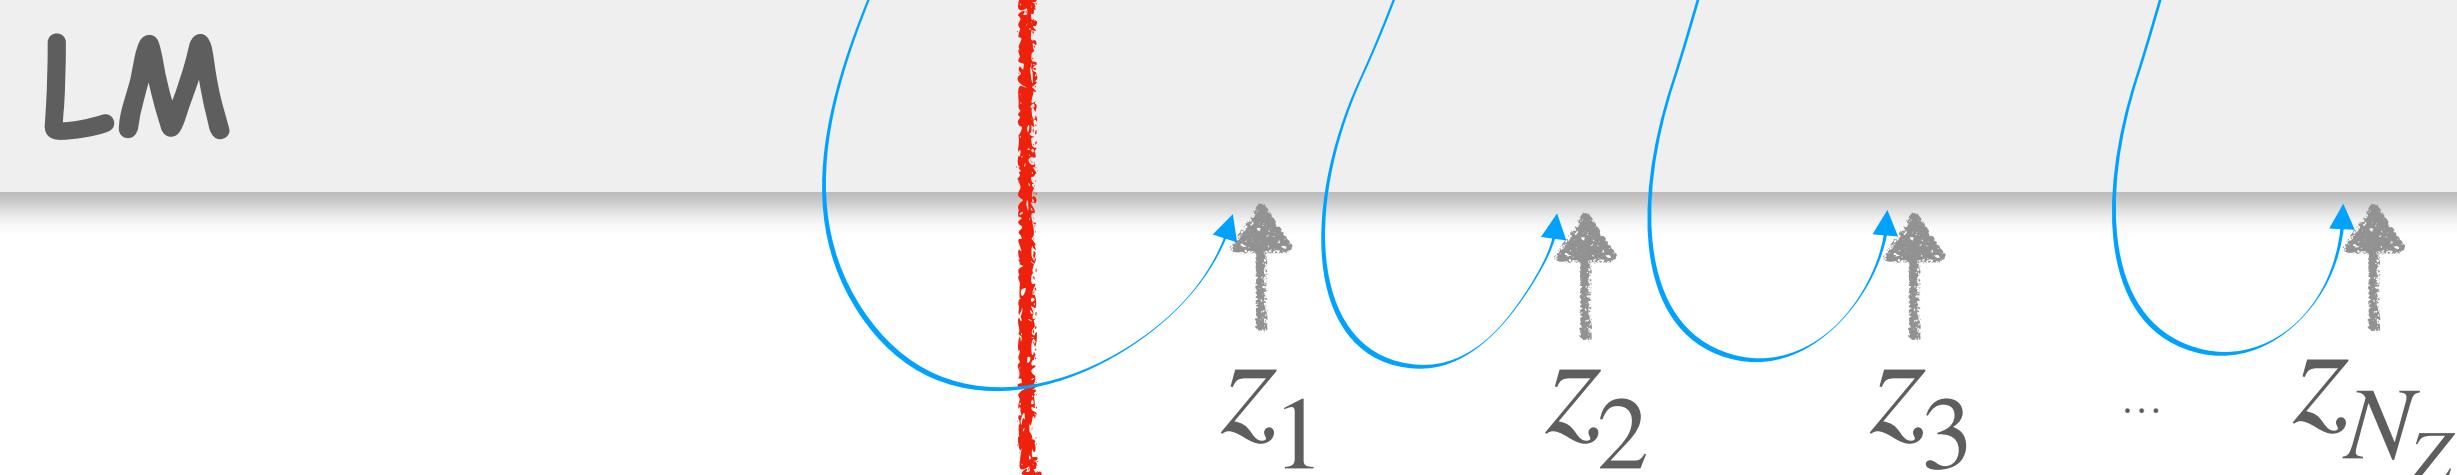


Backpropagation

$$\mathcal{L}(X, \tilde{Y}, Z) := -\sum_{n=1}^{N_Z} \log P_{LM}(z_n | X, \tilde{Y}, Z_{1:n-1})$$



LM



Ray hung a tire on a
rope to make his
daughter a swing.

Ray ran to his
daughter to make
sure she was okay.

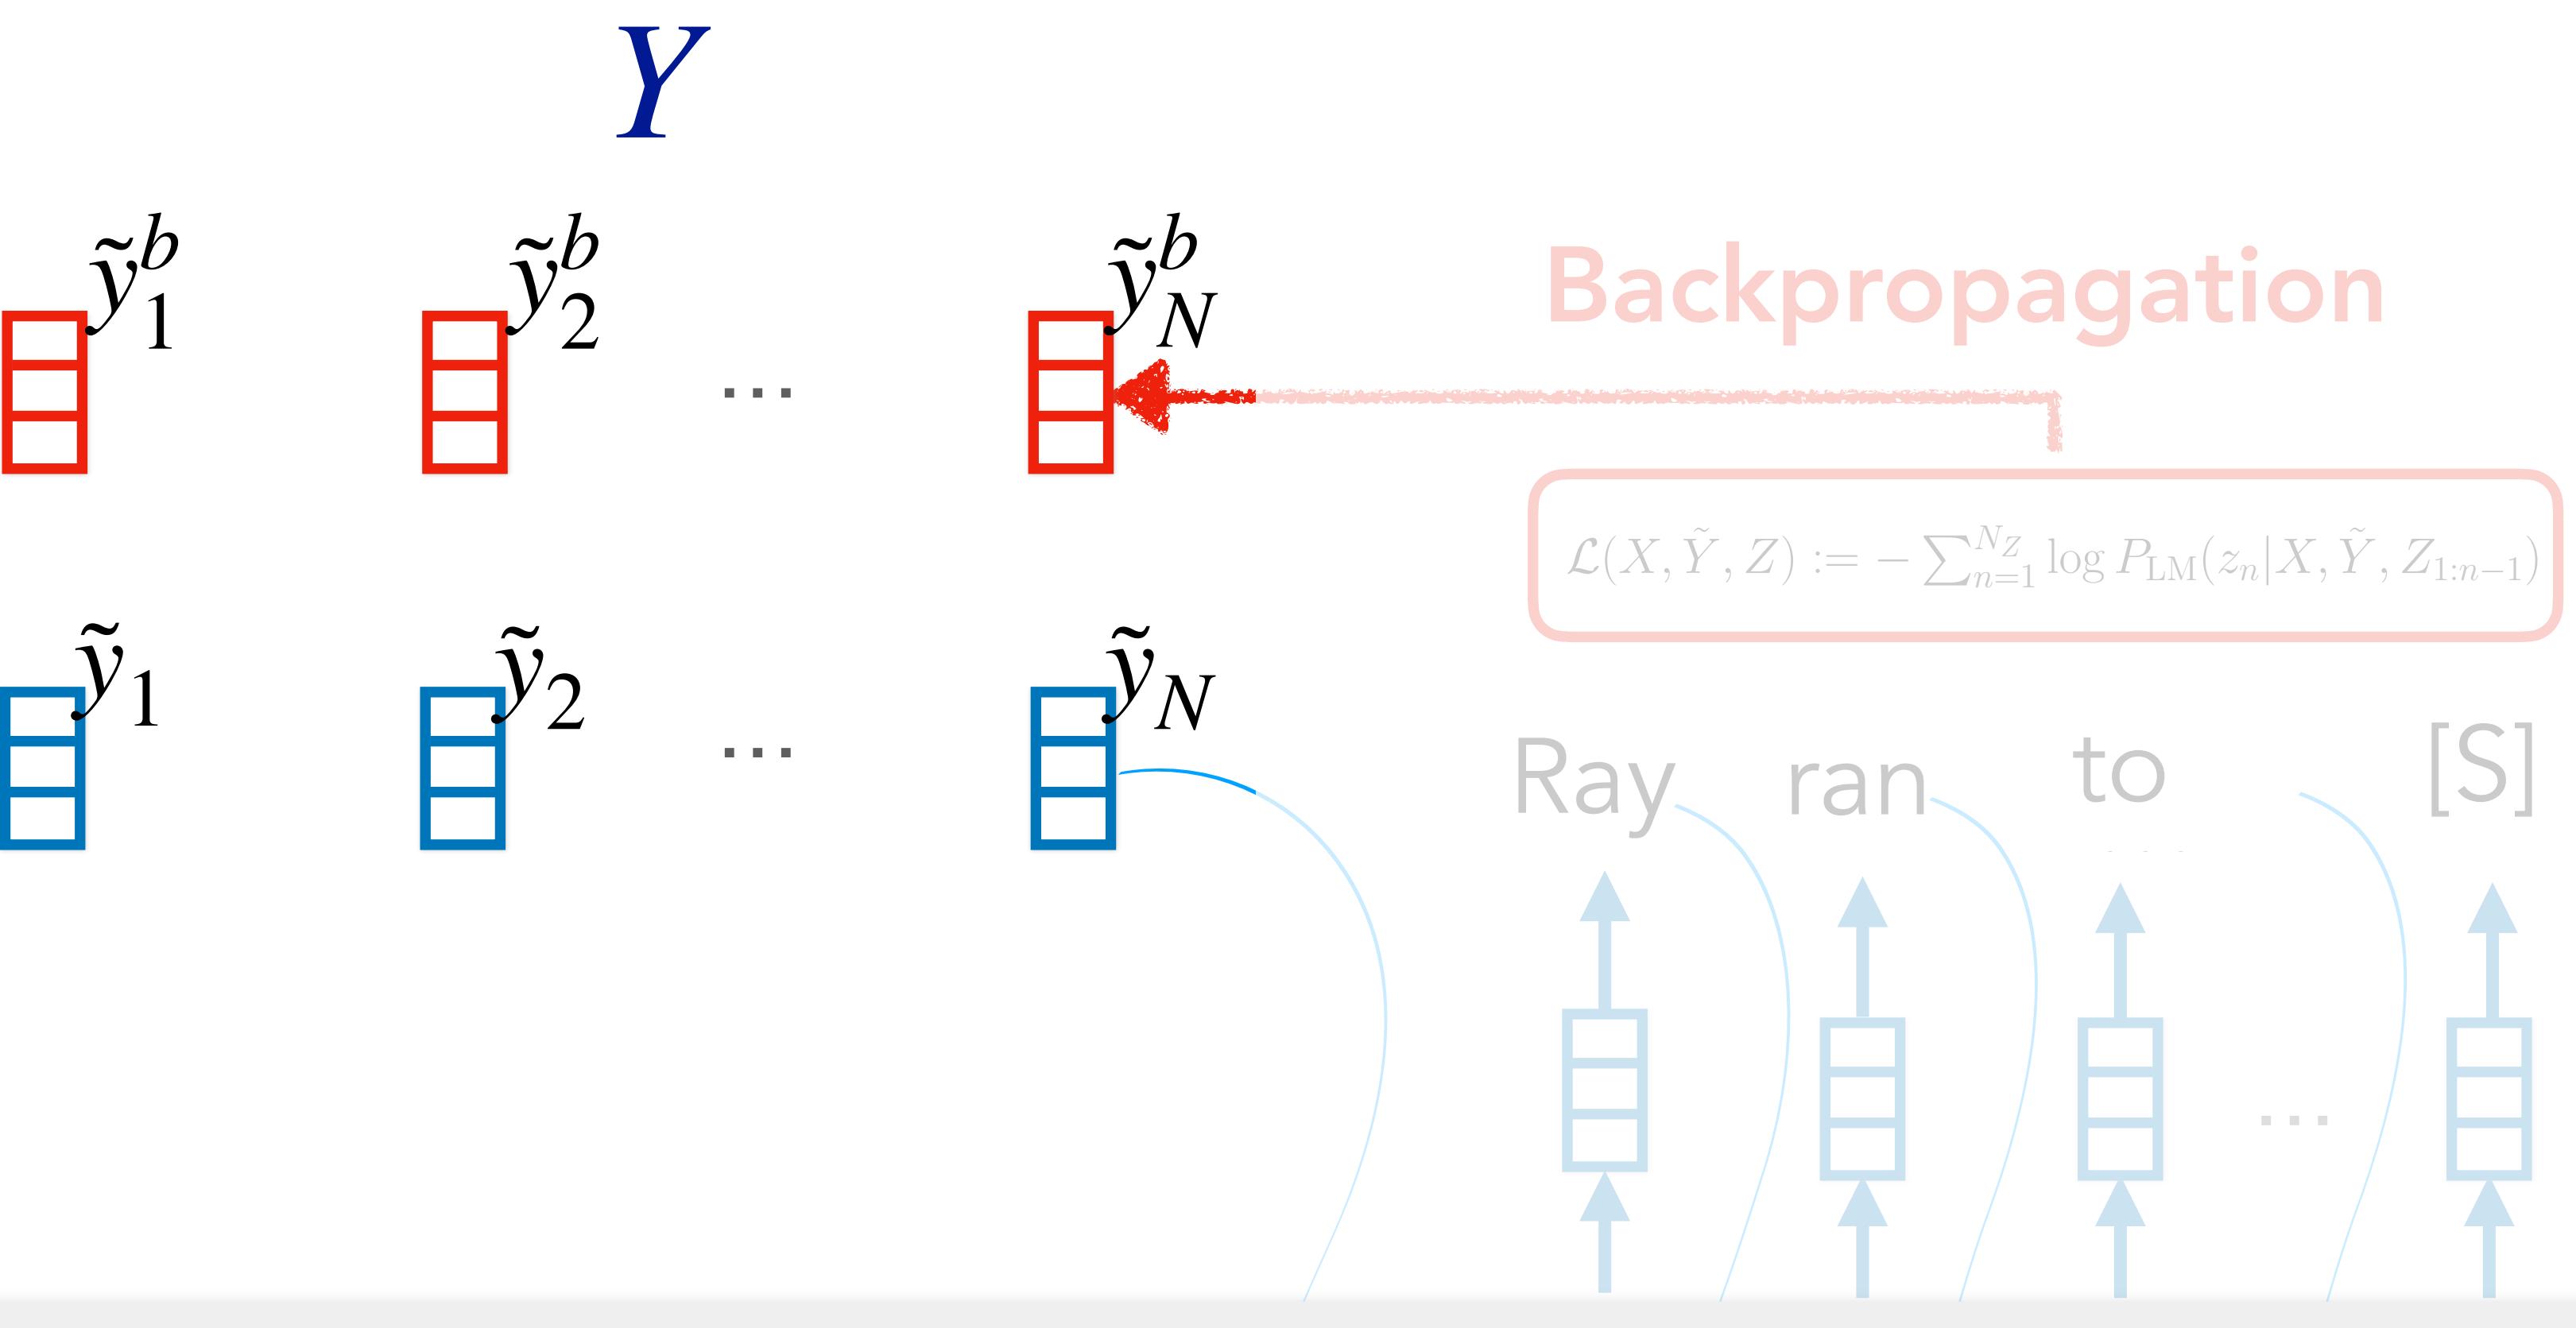
X

Z



Forward Pass

Mix both past and future information

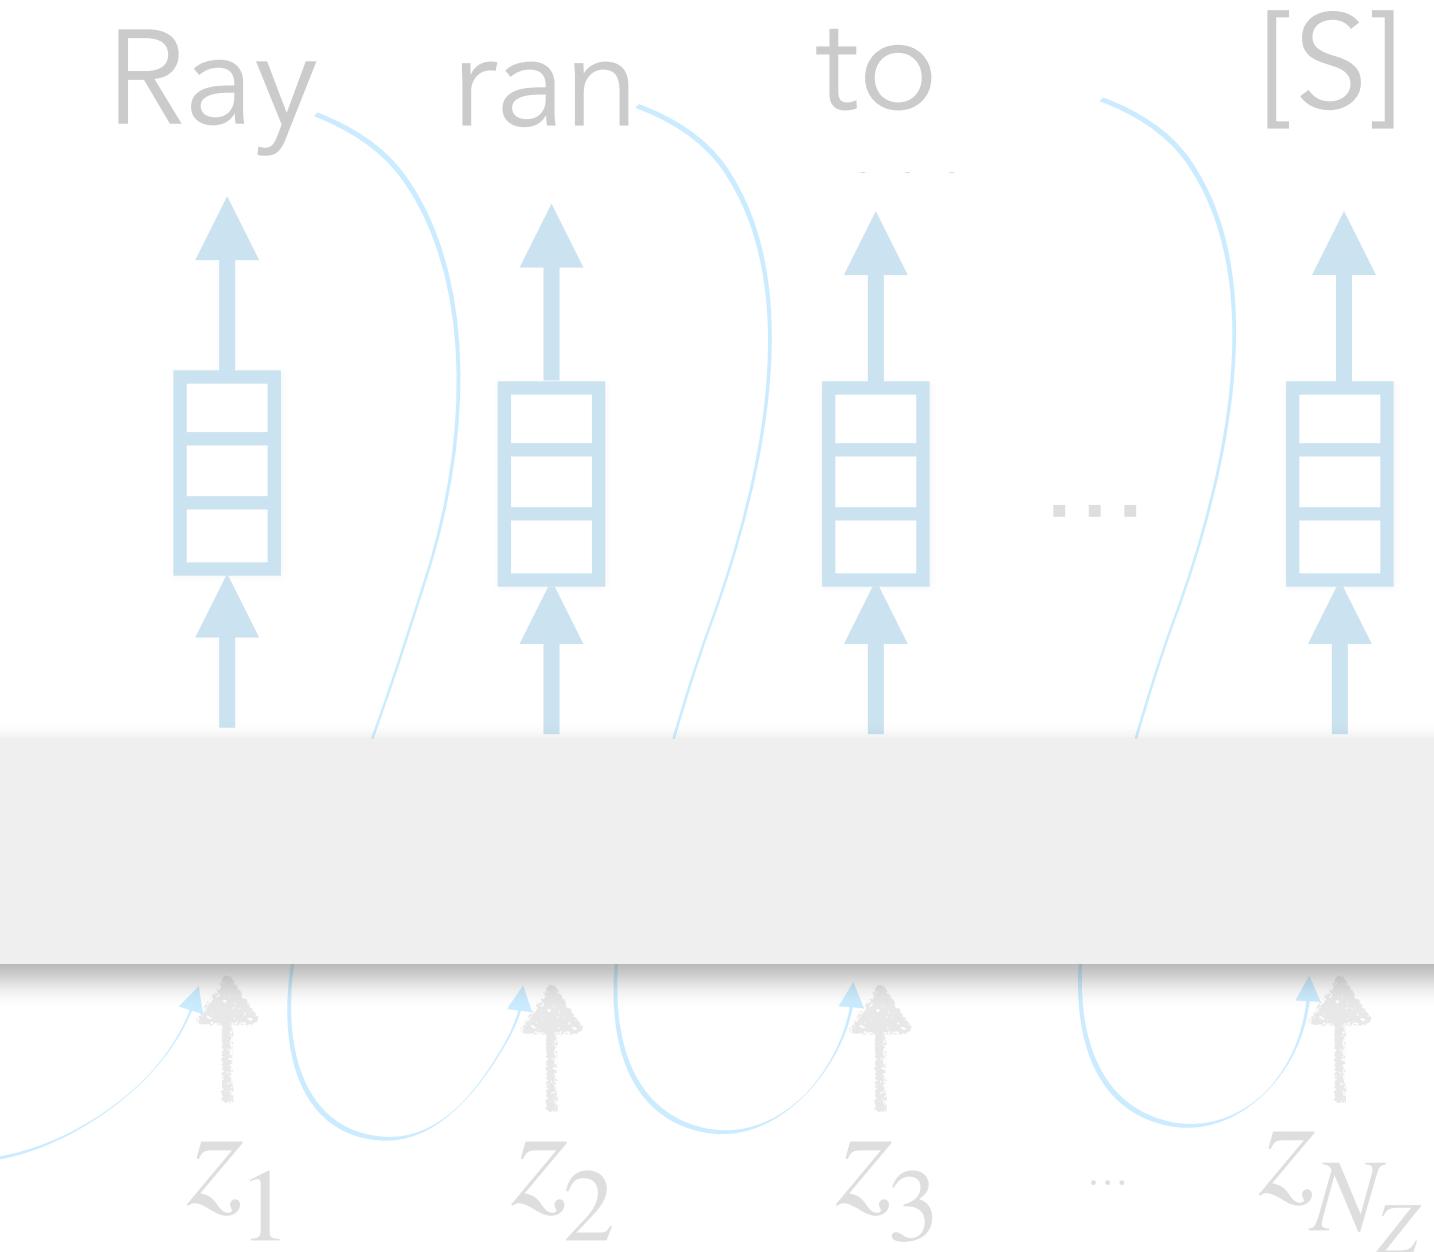


Ray hung a tire on a rope to make his daughter a swing.

X

Backpropagation

$$\mathcal{L}(X, \tilde{Y}, Z) := -\sum_{n=1}^{N_Z} \log P_{\text{LM}}(z_n | X, \tilde{Y}, Z_{1:n-1})$$



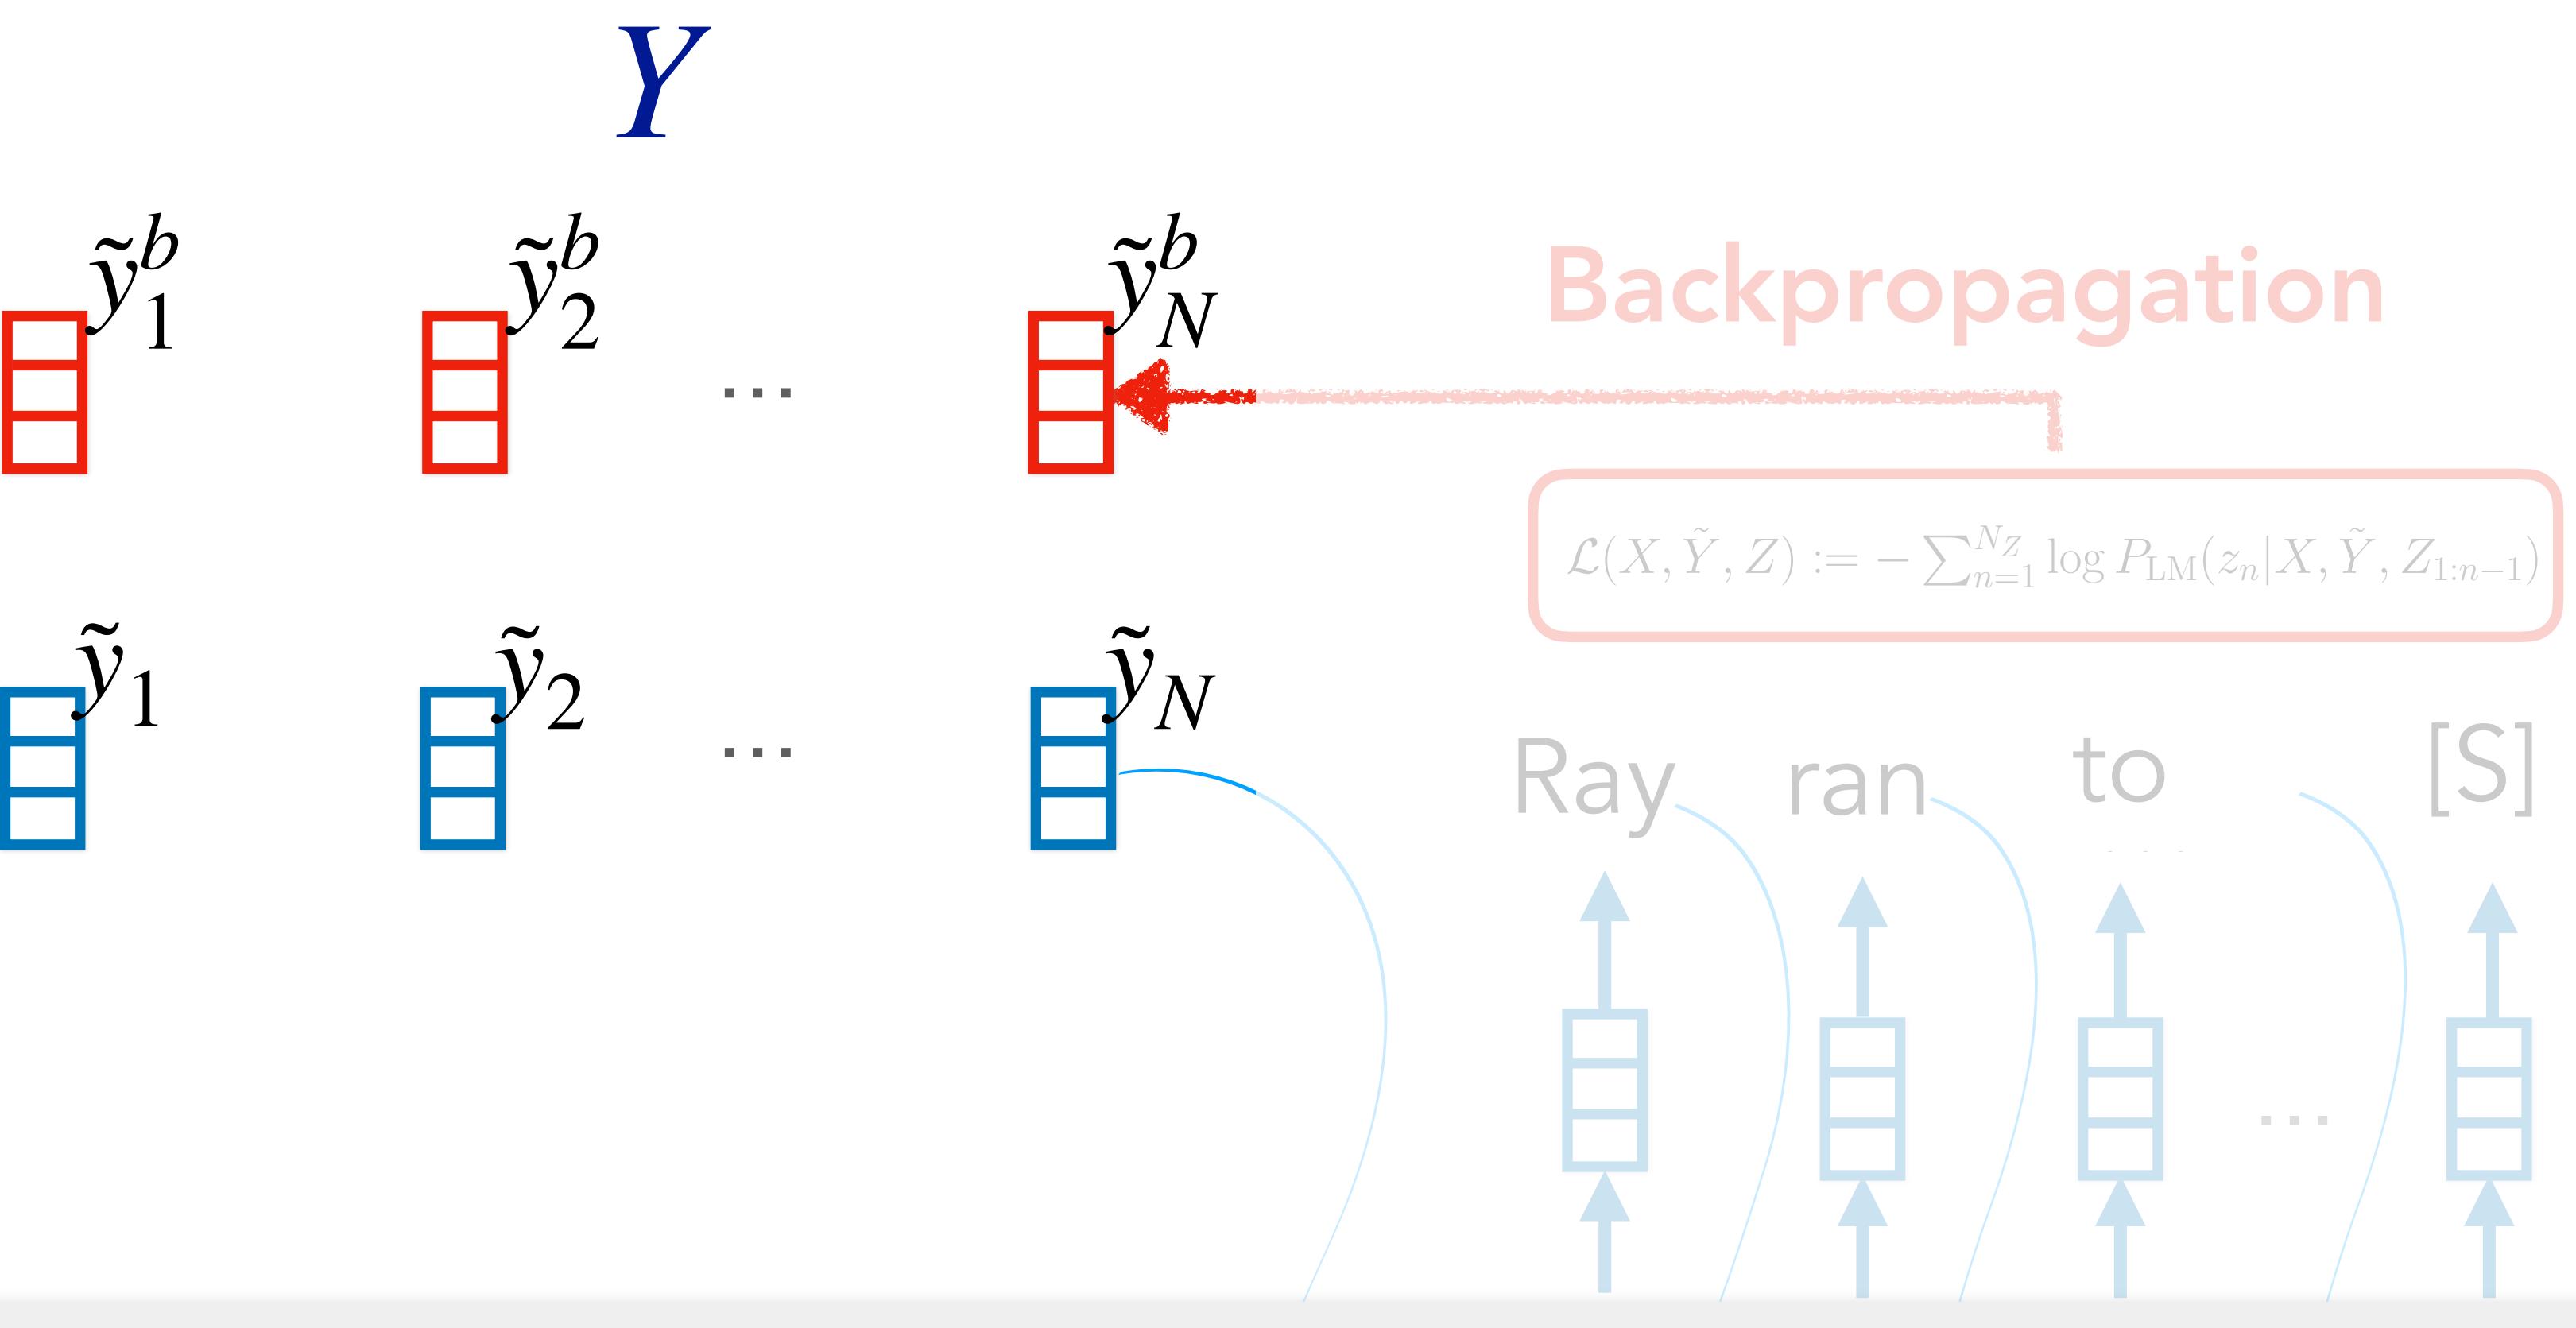
Ray ran to his daughter to make sure she was okay.

Z



Forward Pass

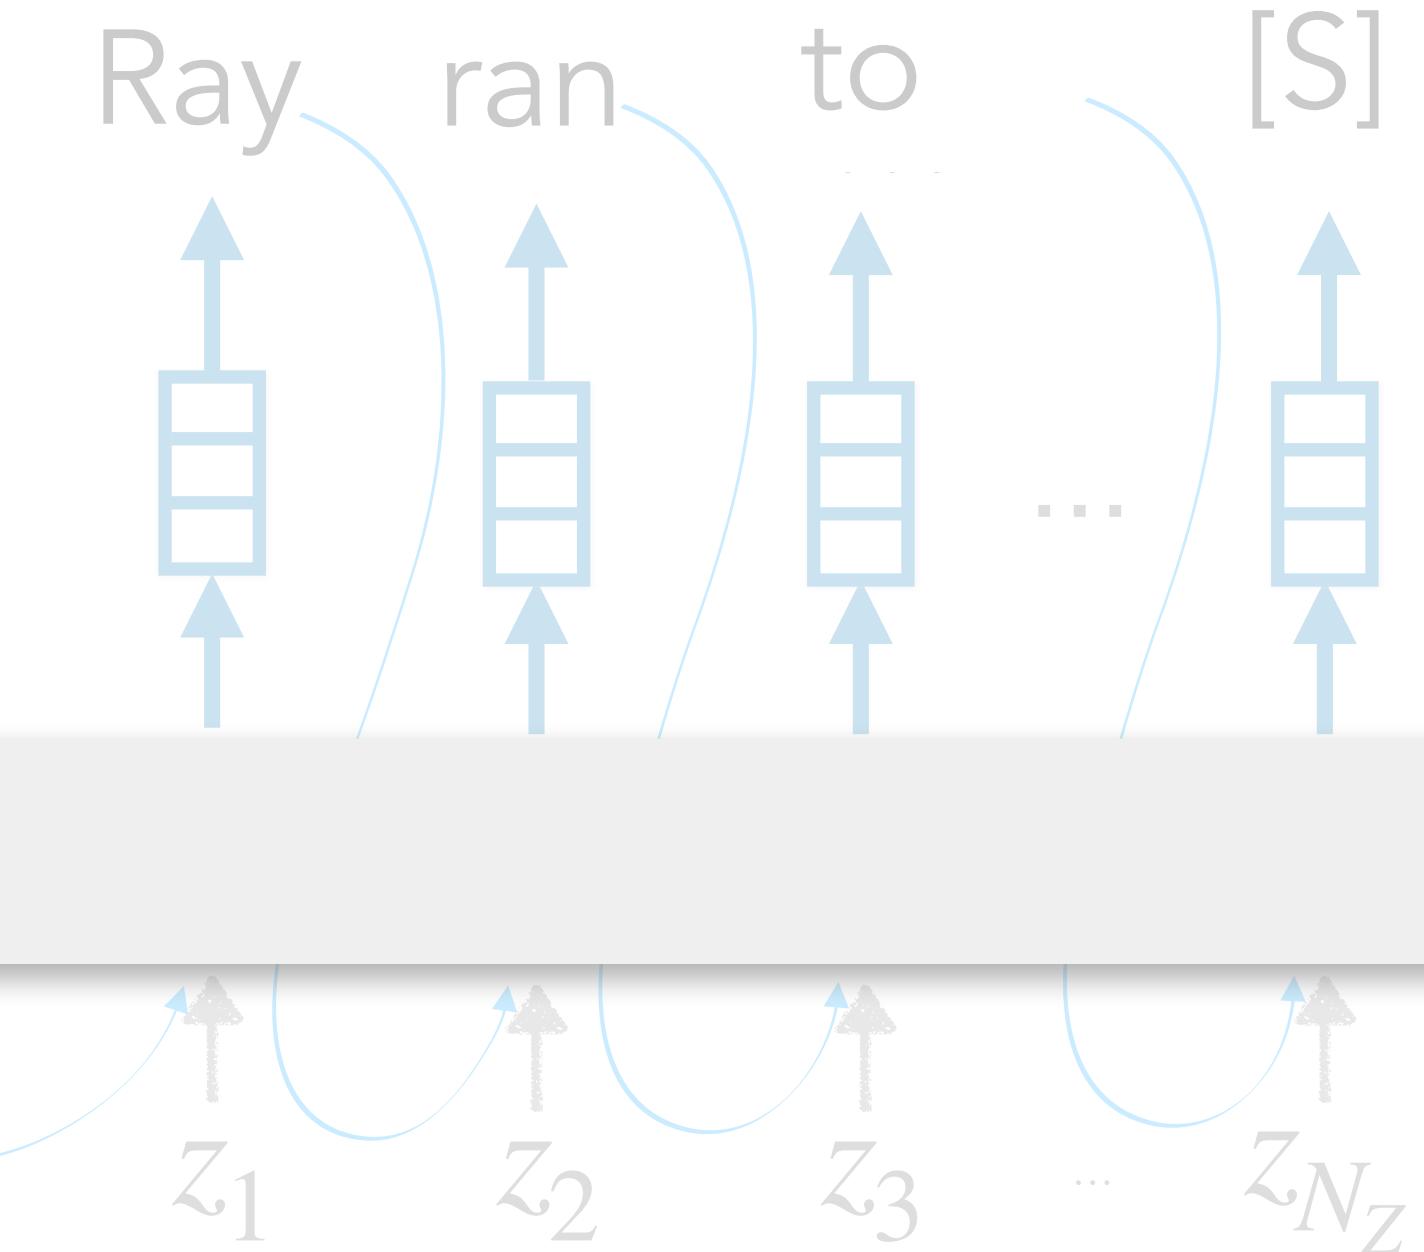
Mix both past and future information



Ray hung a tire on a rope to make his daughter a swing.

Backpropagation

$$\mathcal{L}(X, \tilde{Y}, Z) := -\sum_{n=1}^{N_Z} \log P_{\text{LM}}(z_n | X, \tilde{Y}, Z_{1:n-1})$$



Ray ran to his daughter to make sure she was okay.

X

Z



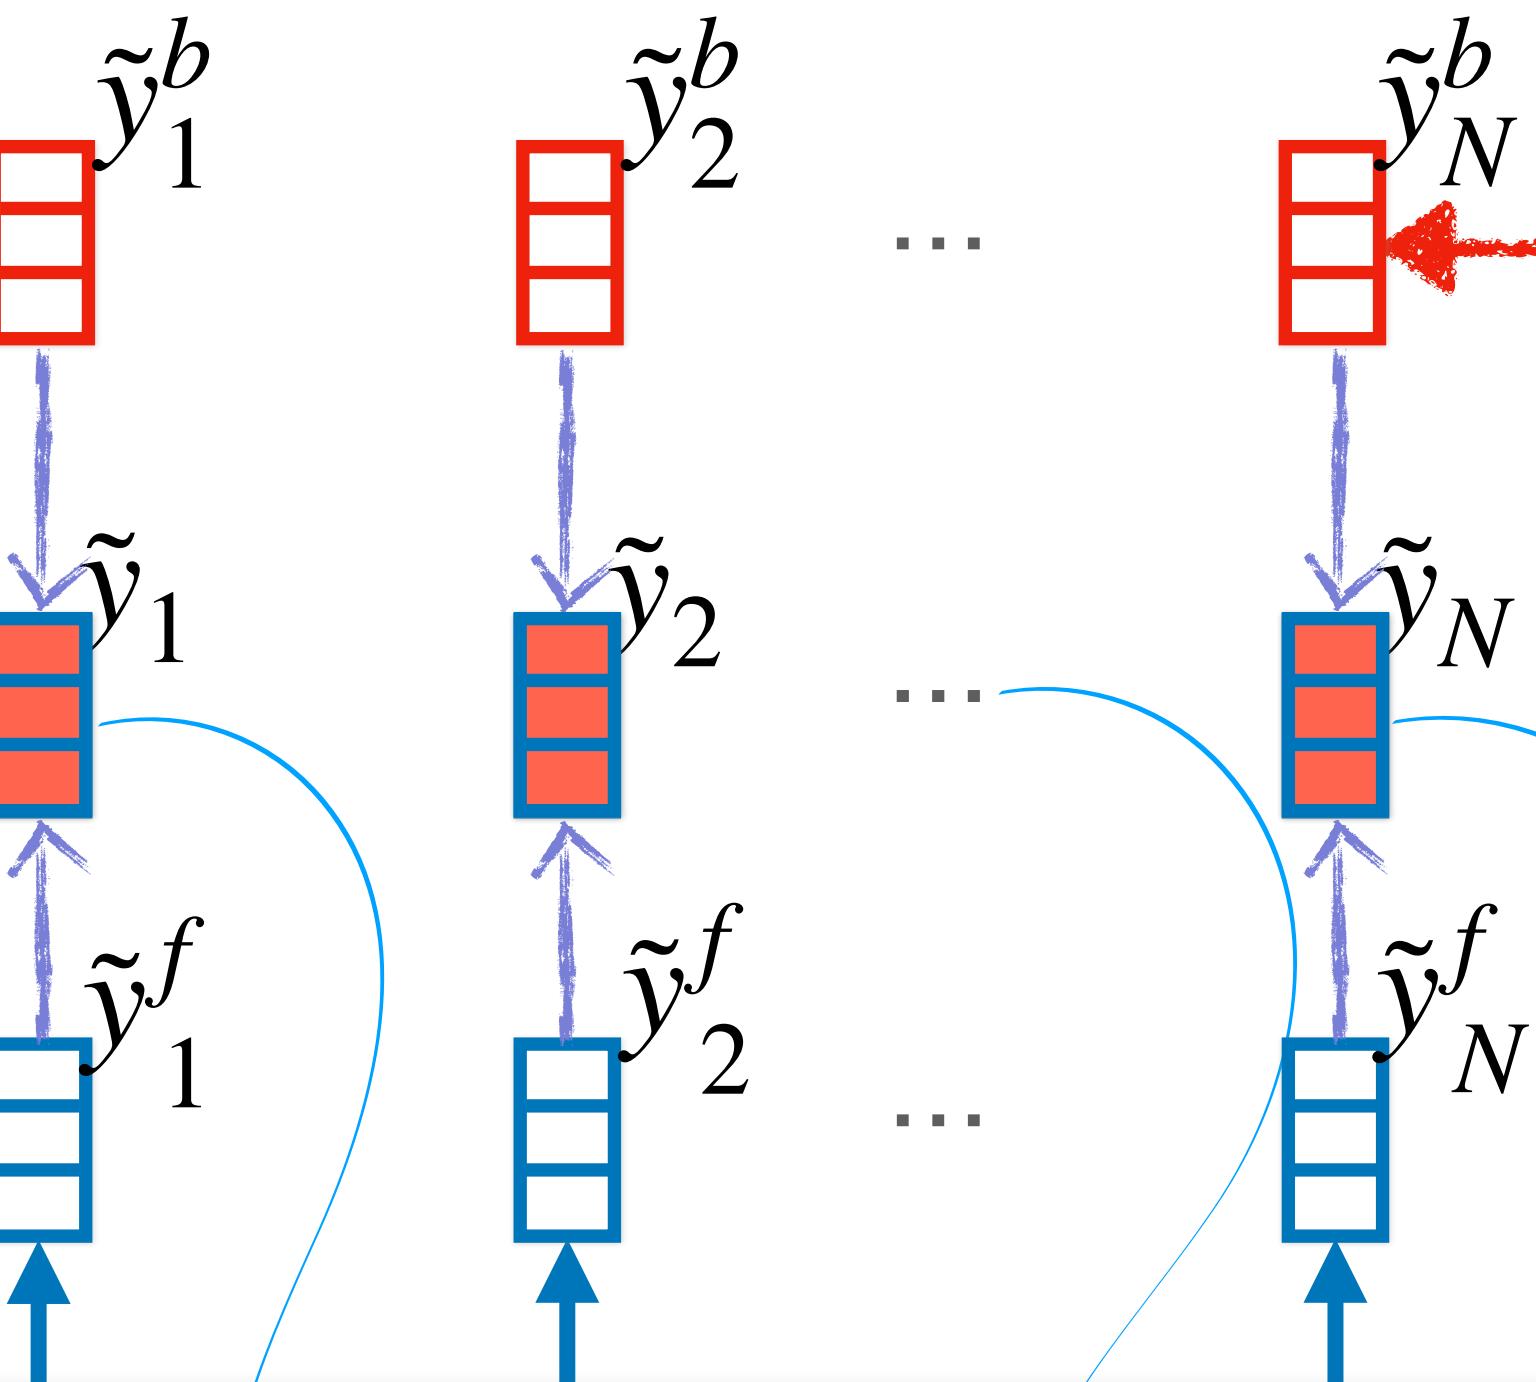
Forward Pass

Mix both past and future information

x_1 x_2 ... x_{N_X}

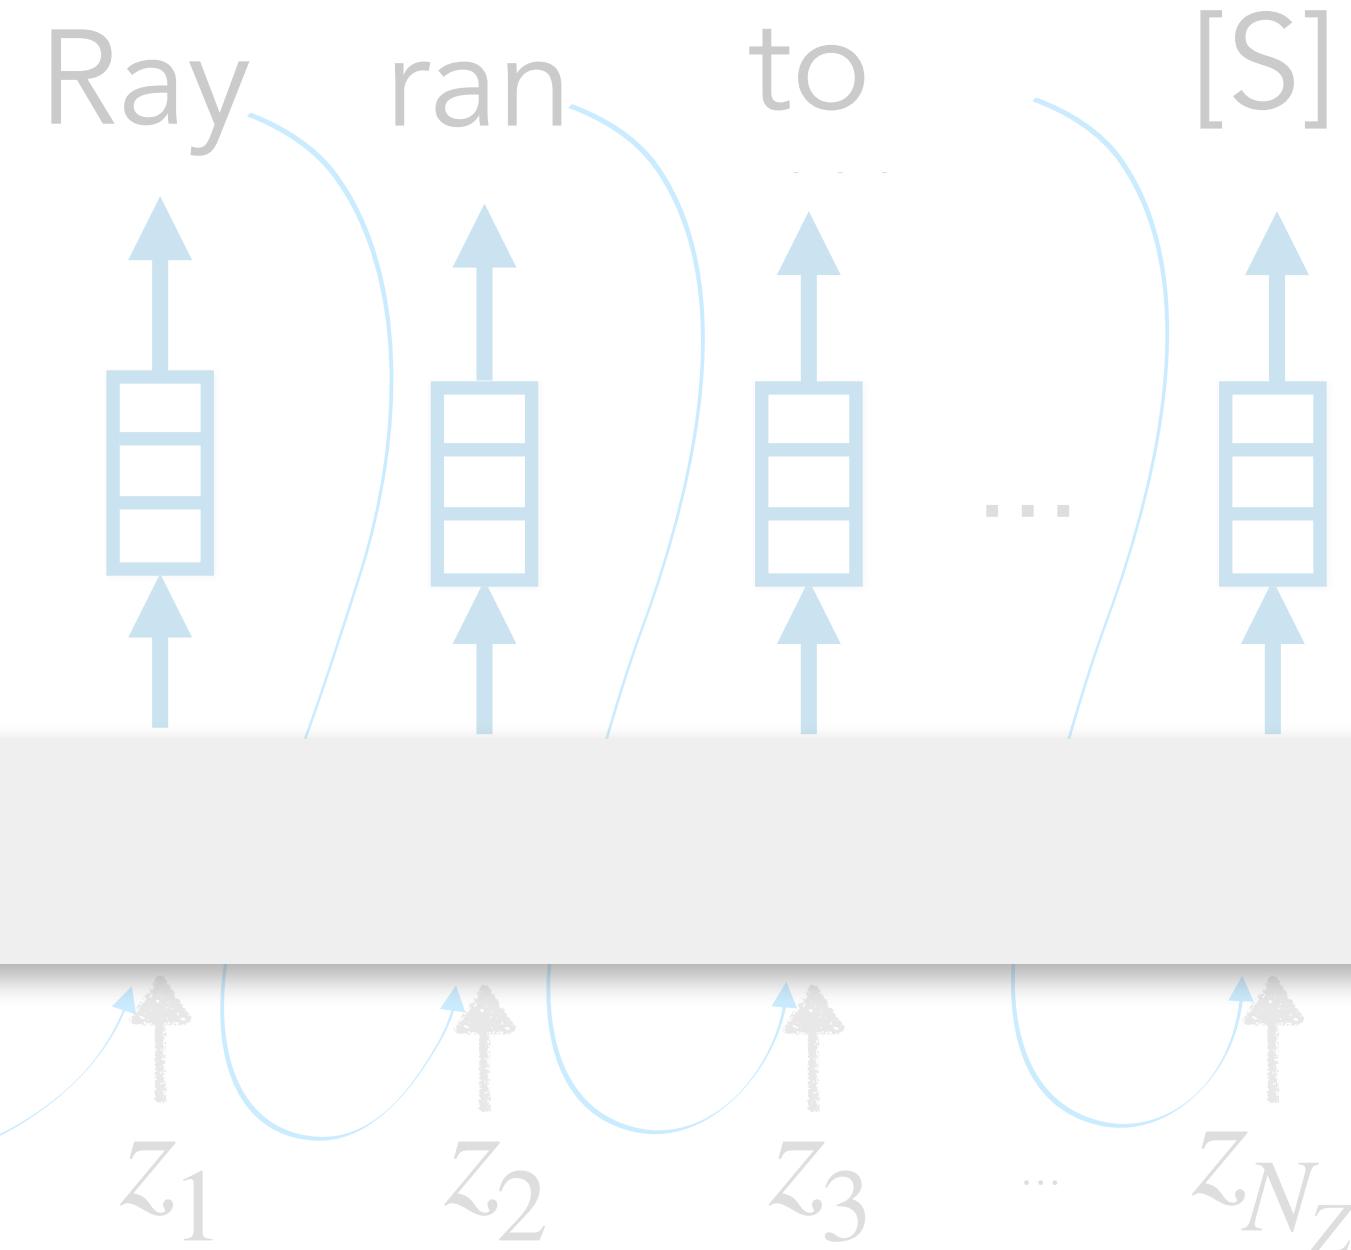
Ray hung a tire on a rope to make his daughter a swing.

Y



Backpropagation

$$\mathcal{L}(X, \tilde{Y}, Z) := -\sum_{n=1}^{N_Z} \log P_{LM}(z_n | X, \tilde{Y}, Z_{1:n-1})$$



X

Z

Ray ran to his daughter to make sure she was okay.



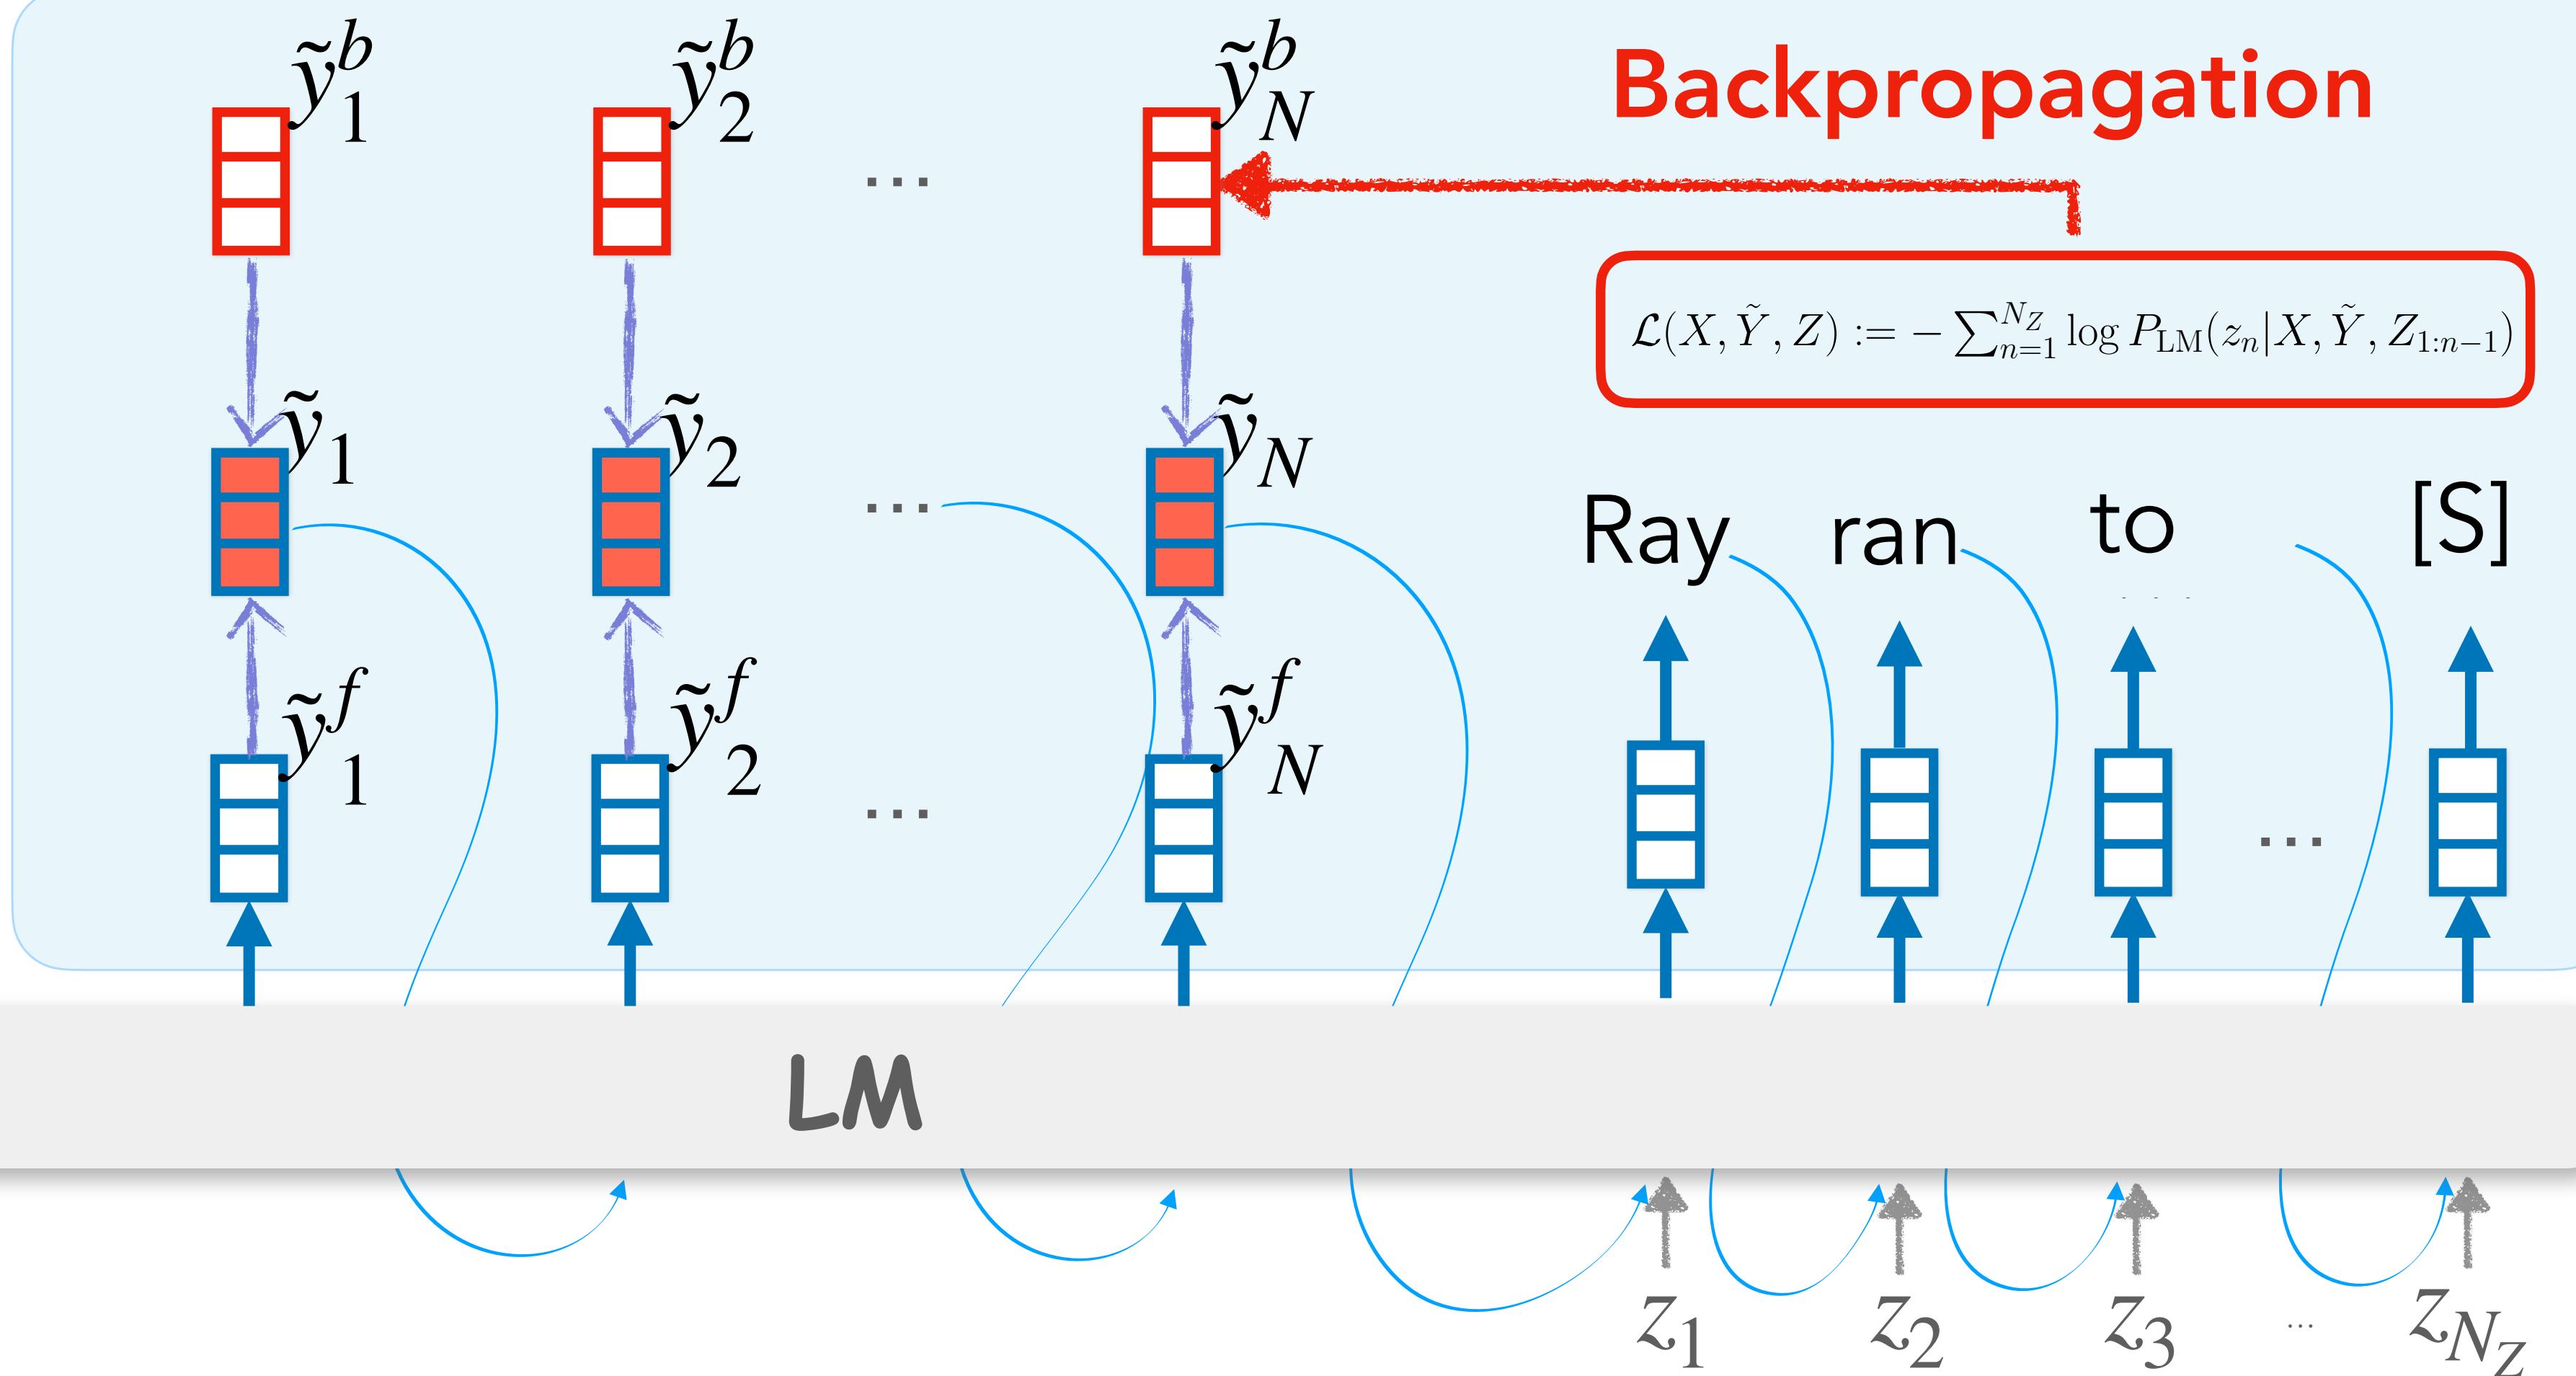
Iterate T
times

x_1 x_2 ... x_{N_X}

Ray hung a tire on a
rope to make his
daughter a swing.

X

Y



Z

Ray ran to his
daughter to make
sure she was okay.



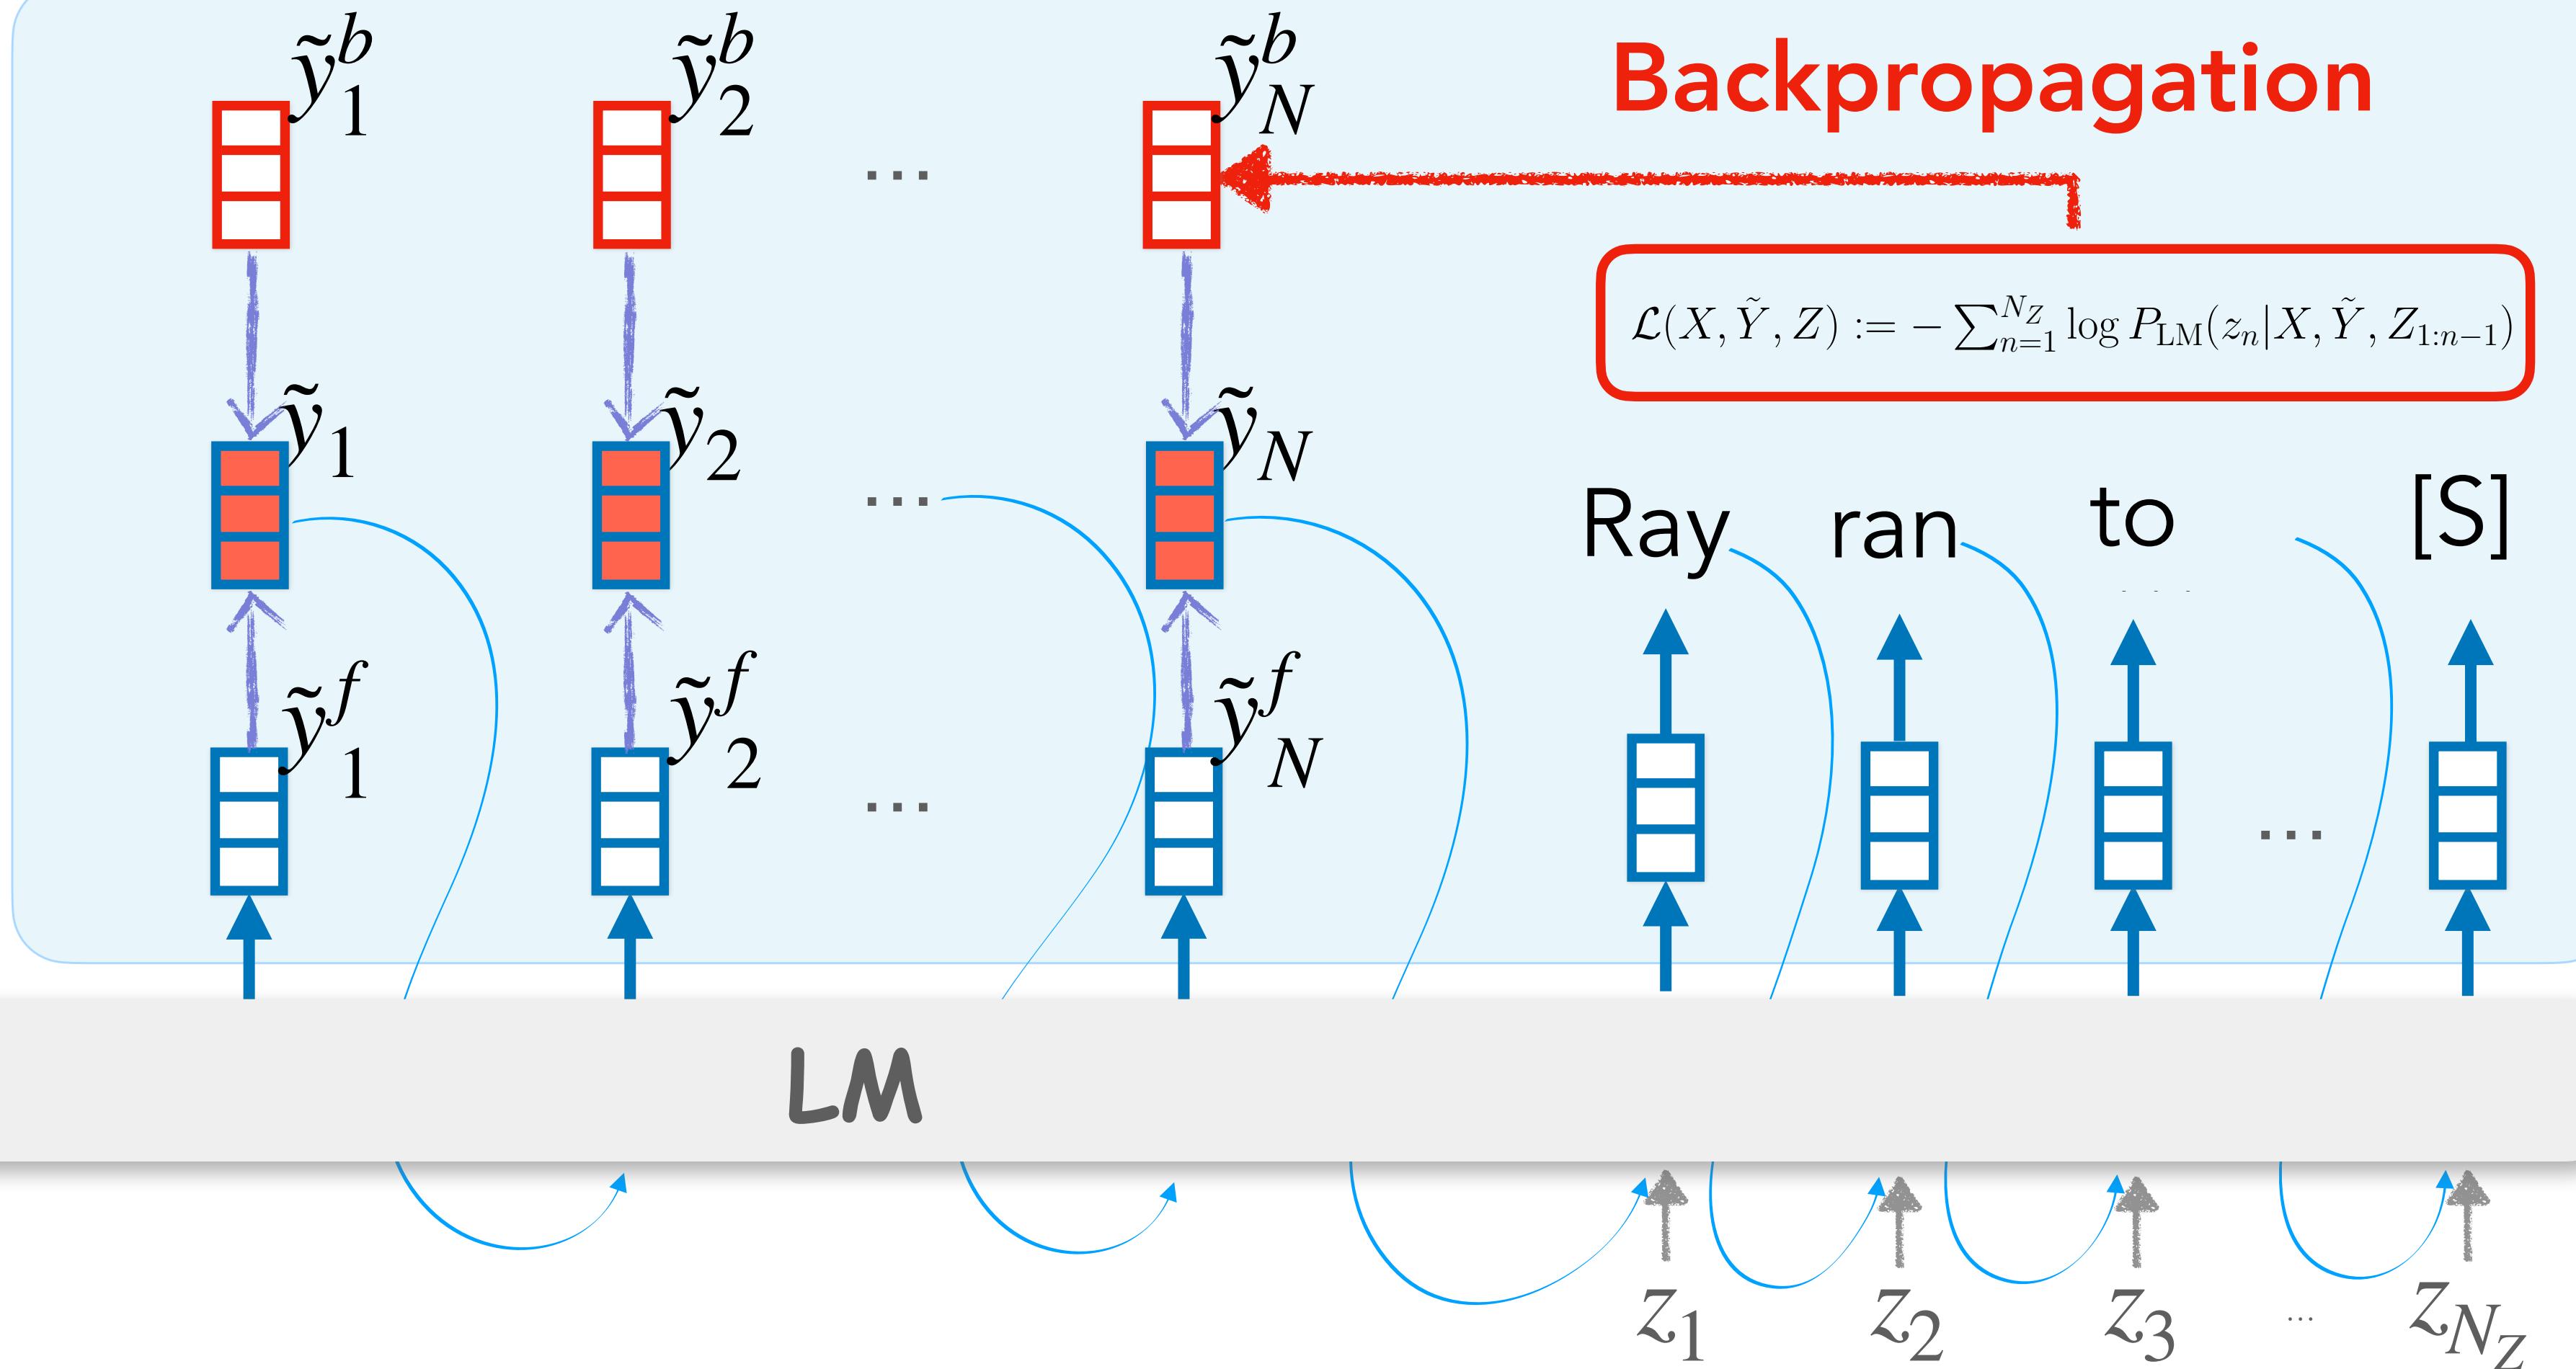
Sample final sequence

x_1 x_2 ... x_{N_X}

Ray hung a tire on a
rope to make his
daughter a swing.

X

Y



Ray ran to his
daughter to make
sure she was okay.

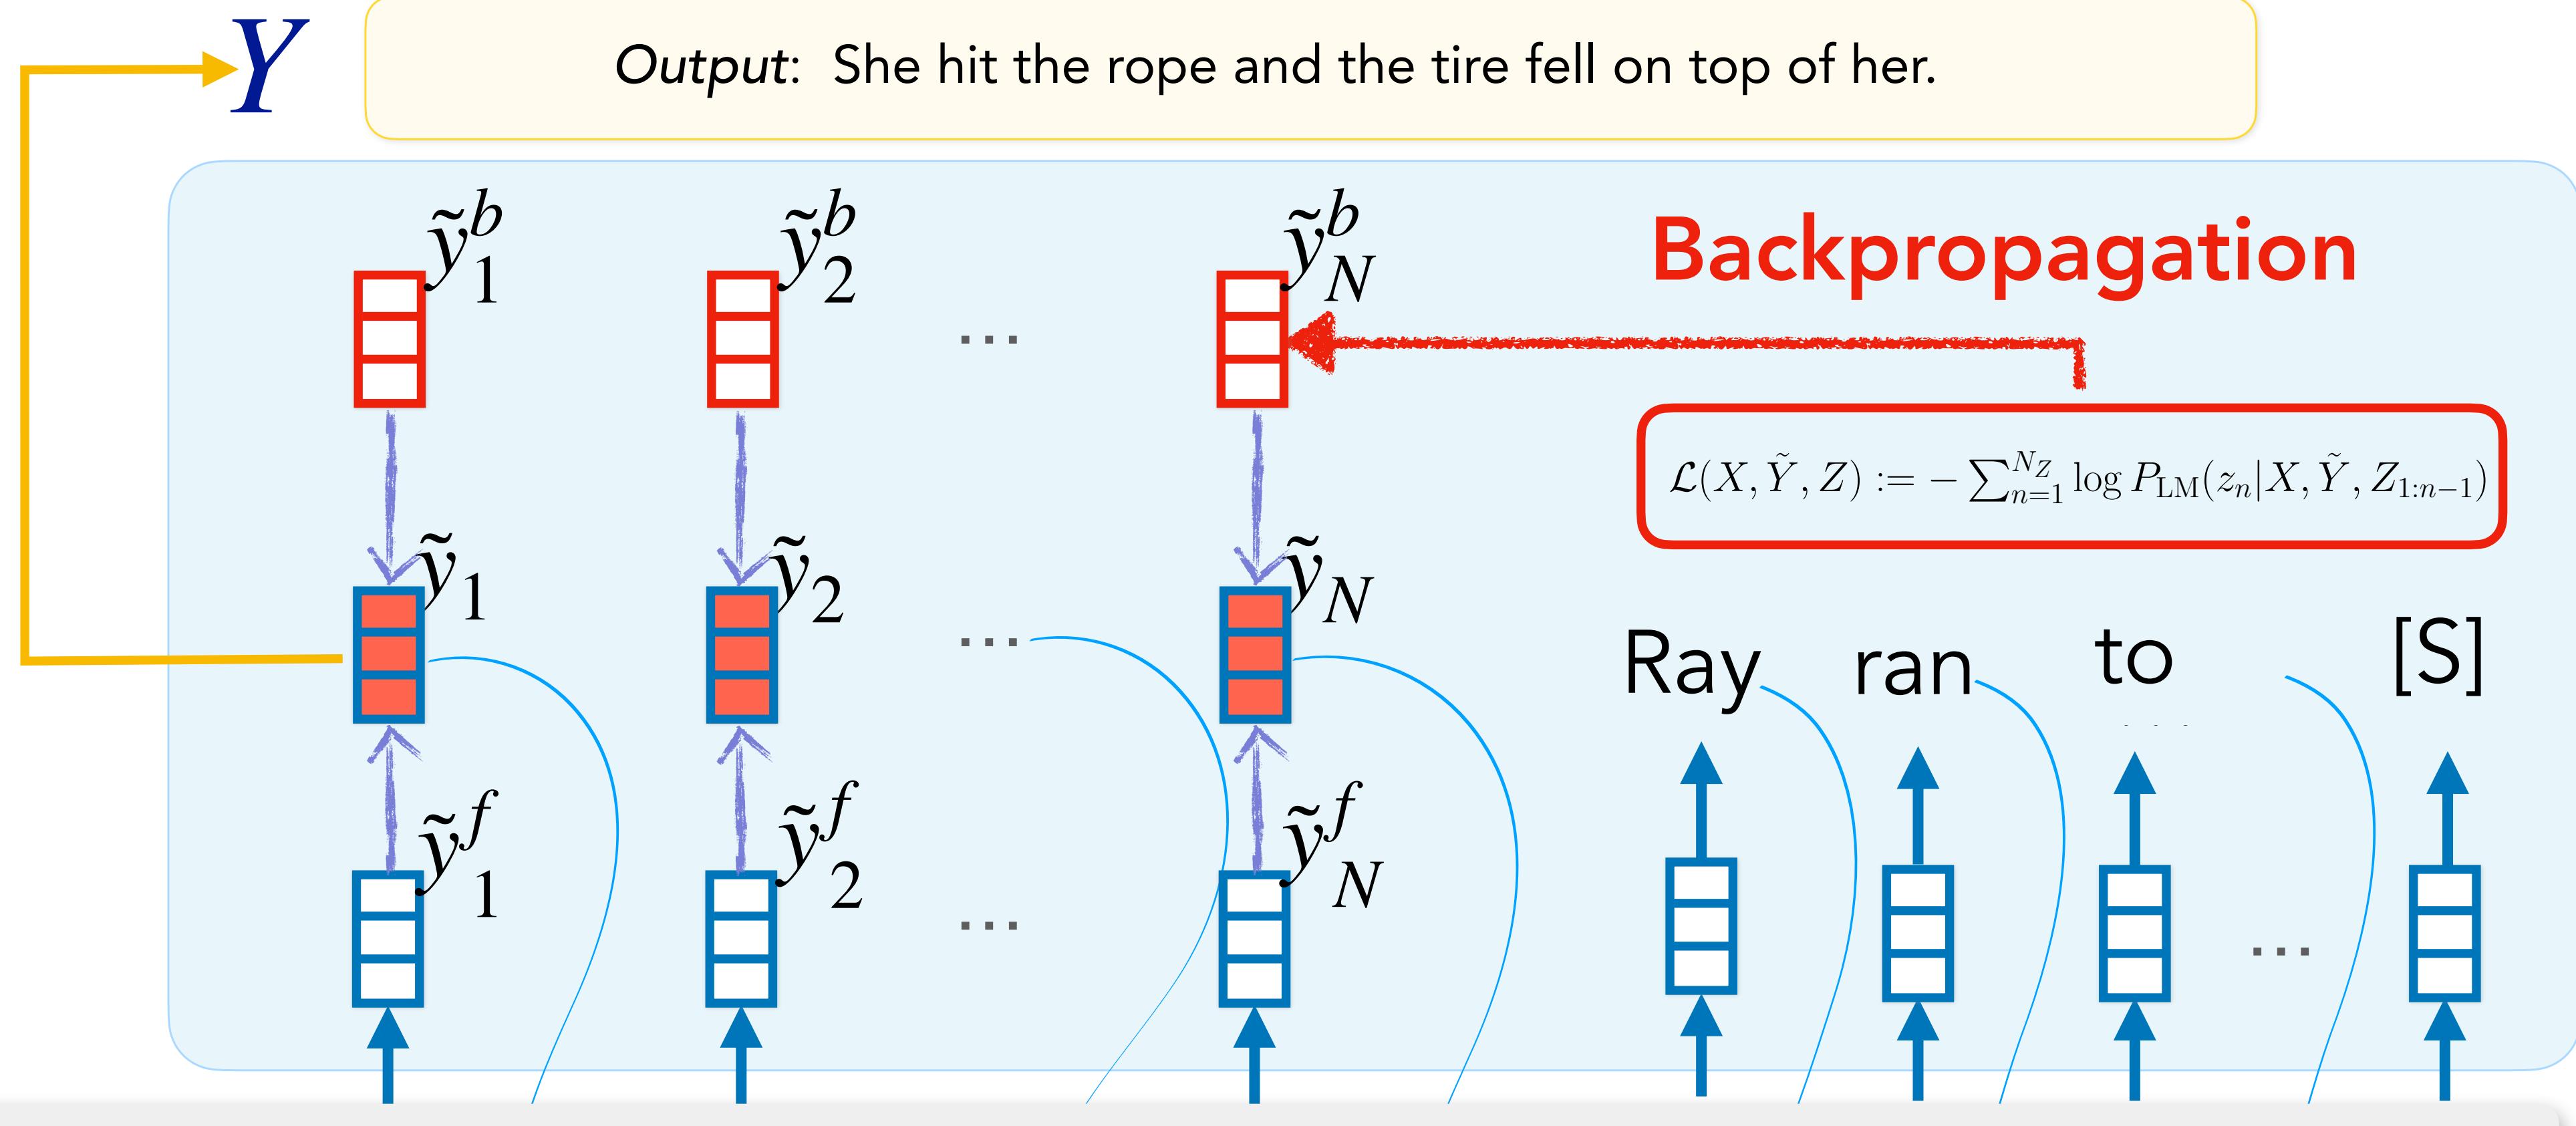
Z



Sample final sequence

$$x_1 \uparrow \quad x_2 \uparrow \quad \dots \quad x_{N_X} \uparrow$$

Ray hung a tire on a rope to make his daughter a swing.



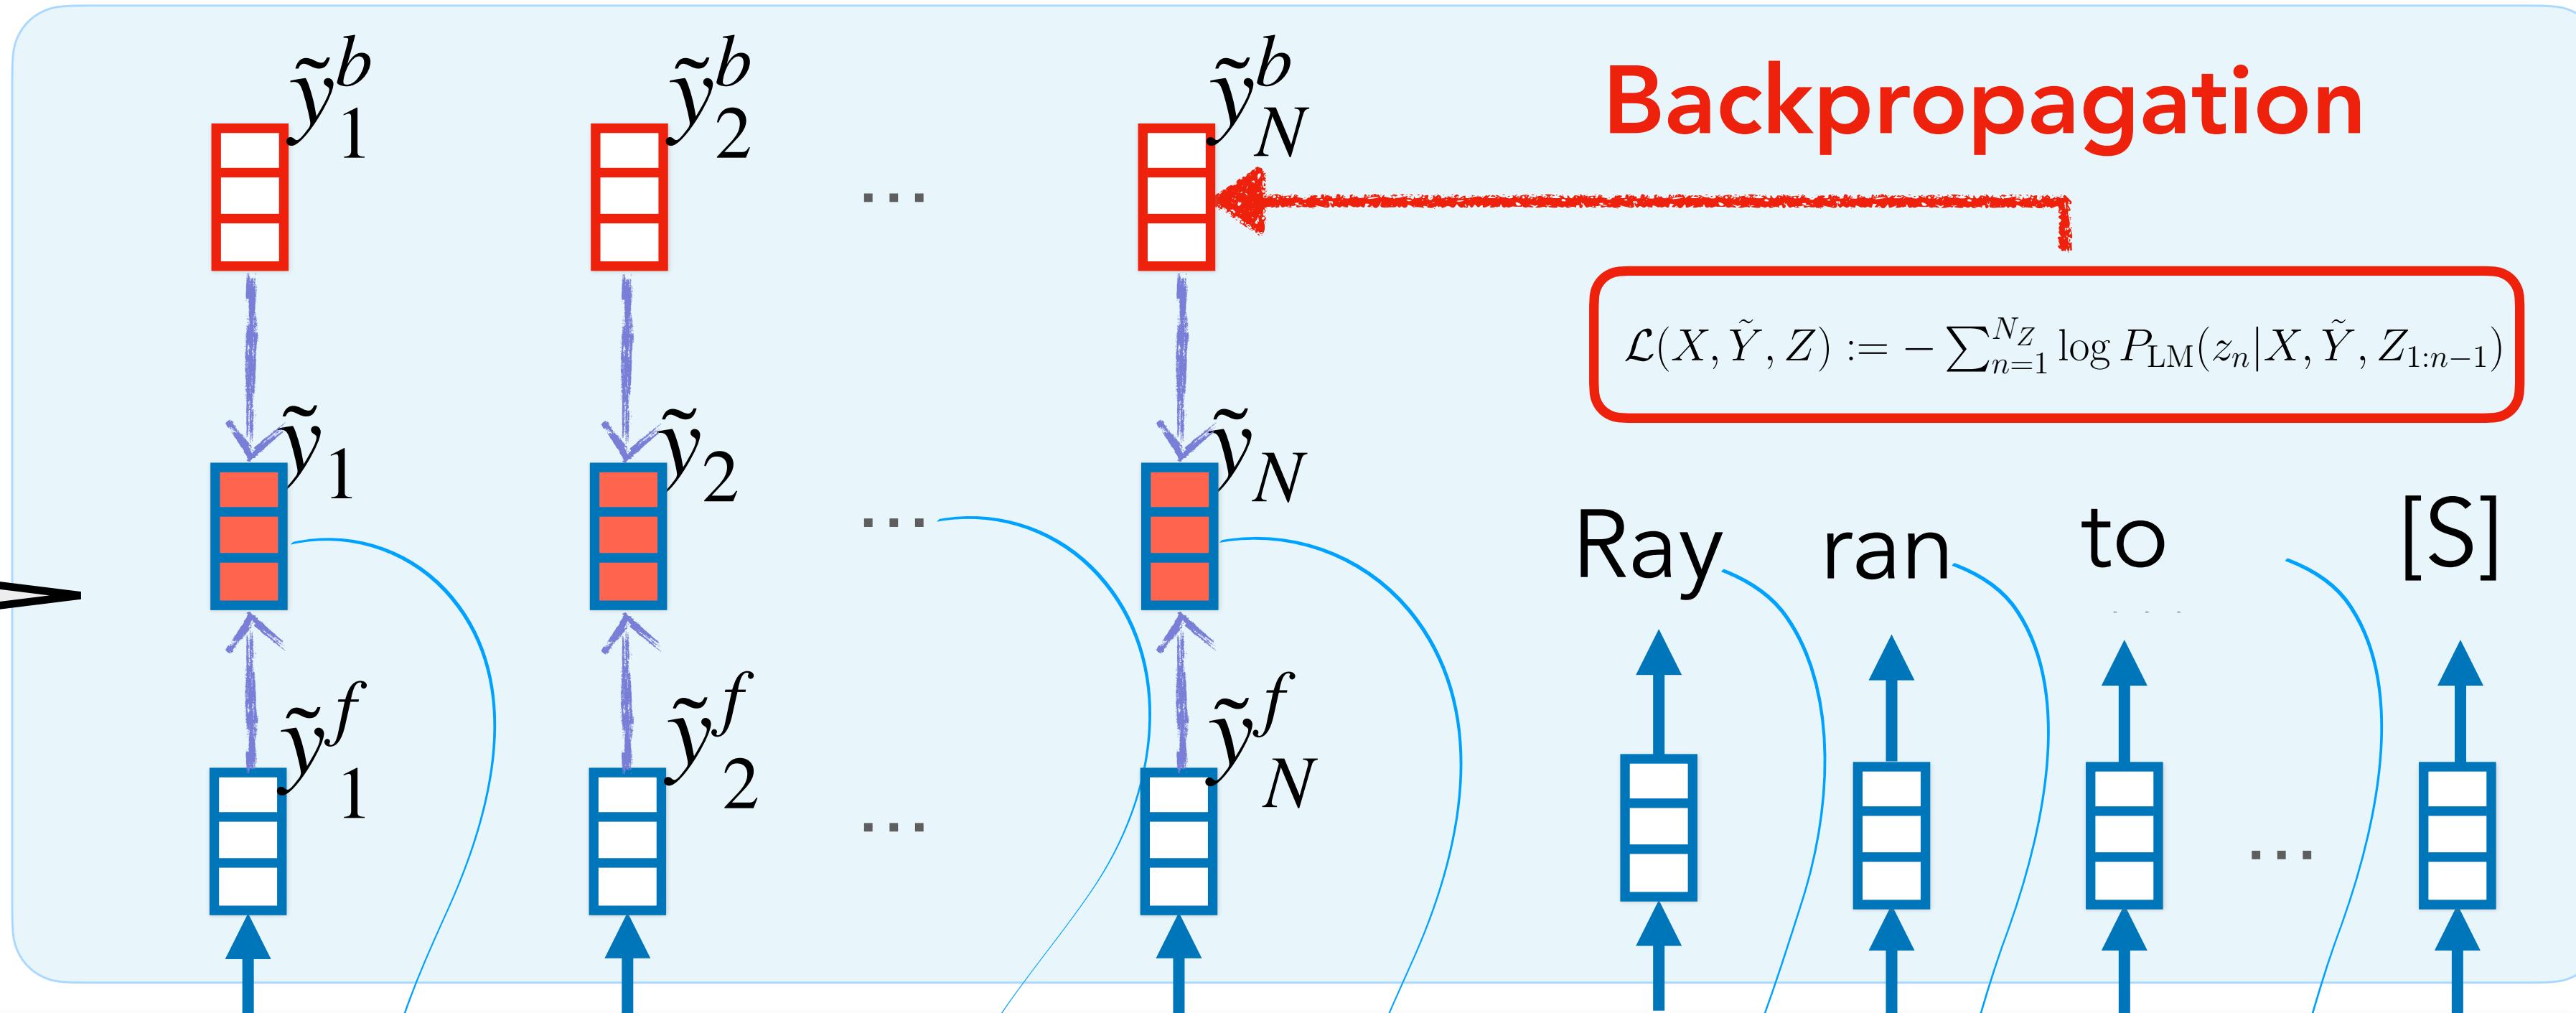
X

Z

Ray ran to his daughter to make sure she was okay.

Y

Inference time
optimization with
respect to future
constraints



x_1 x_2 ... x_{N_X}

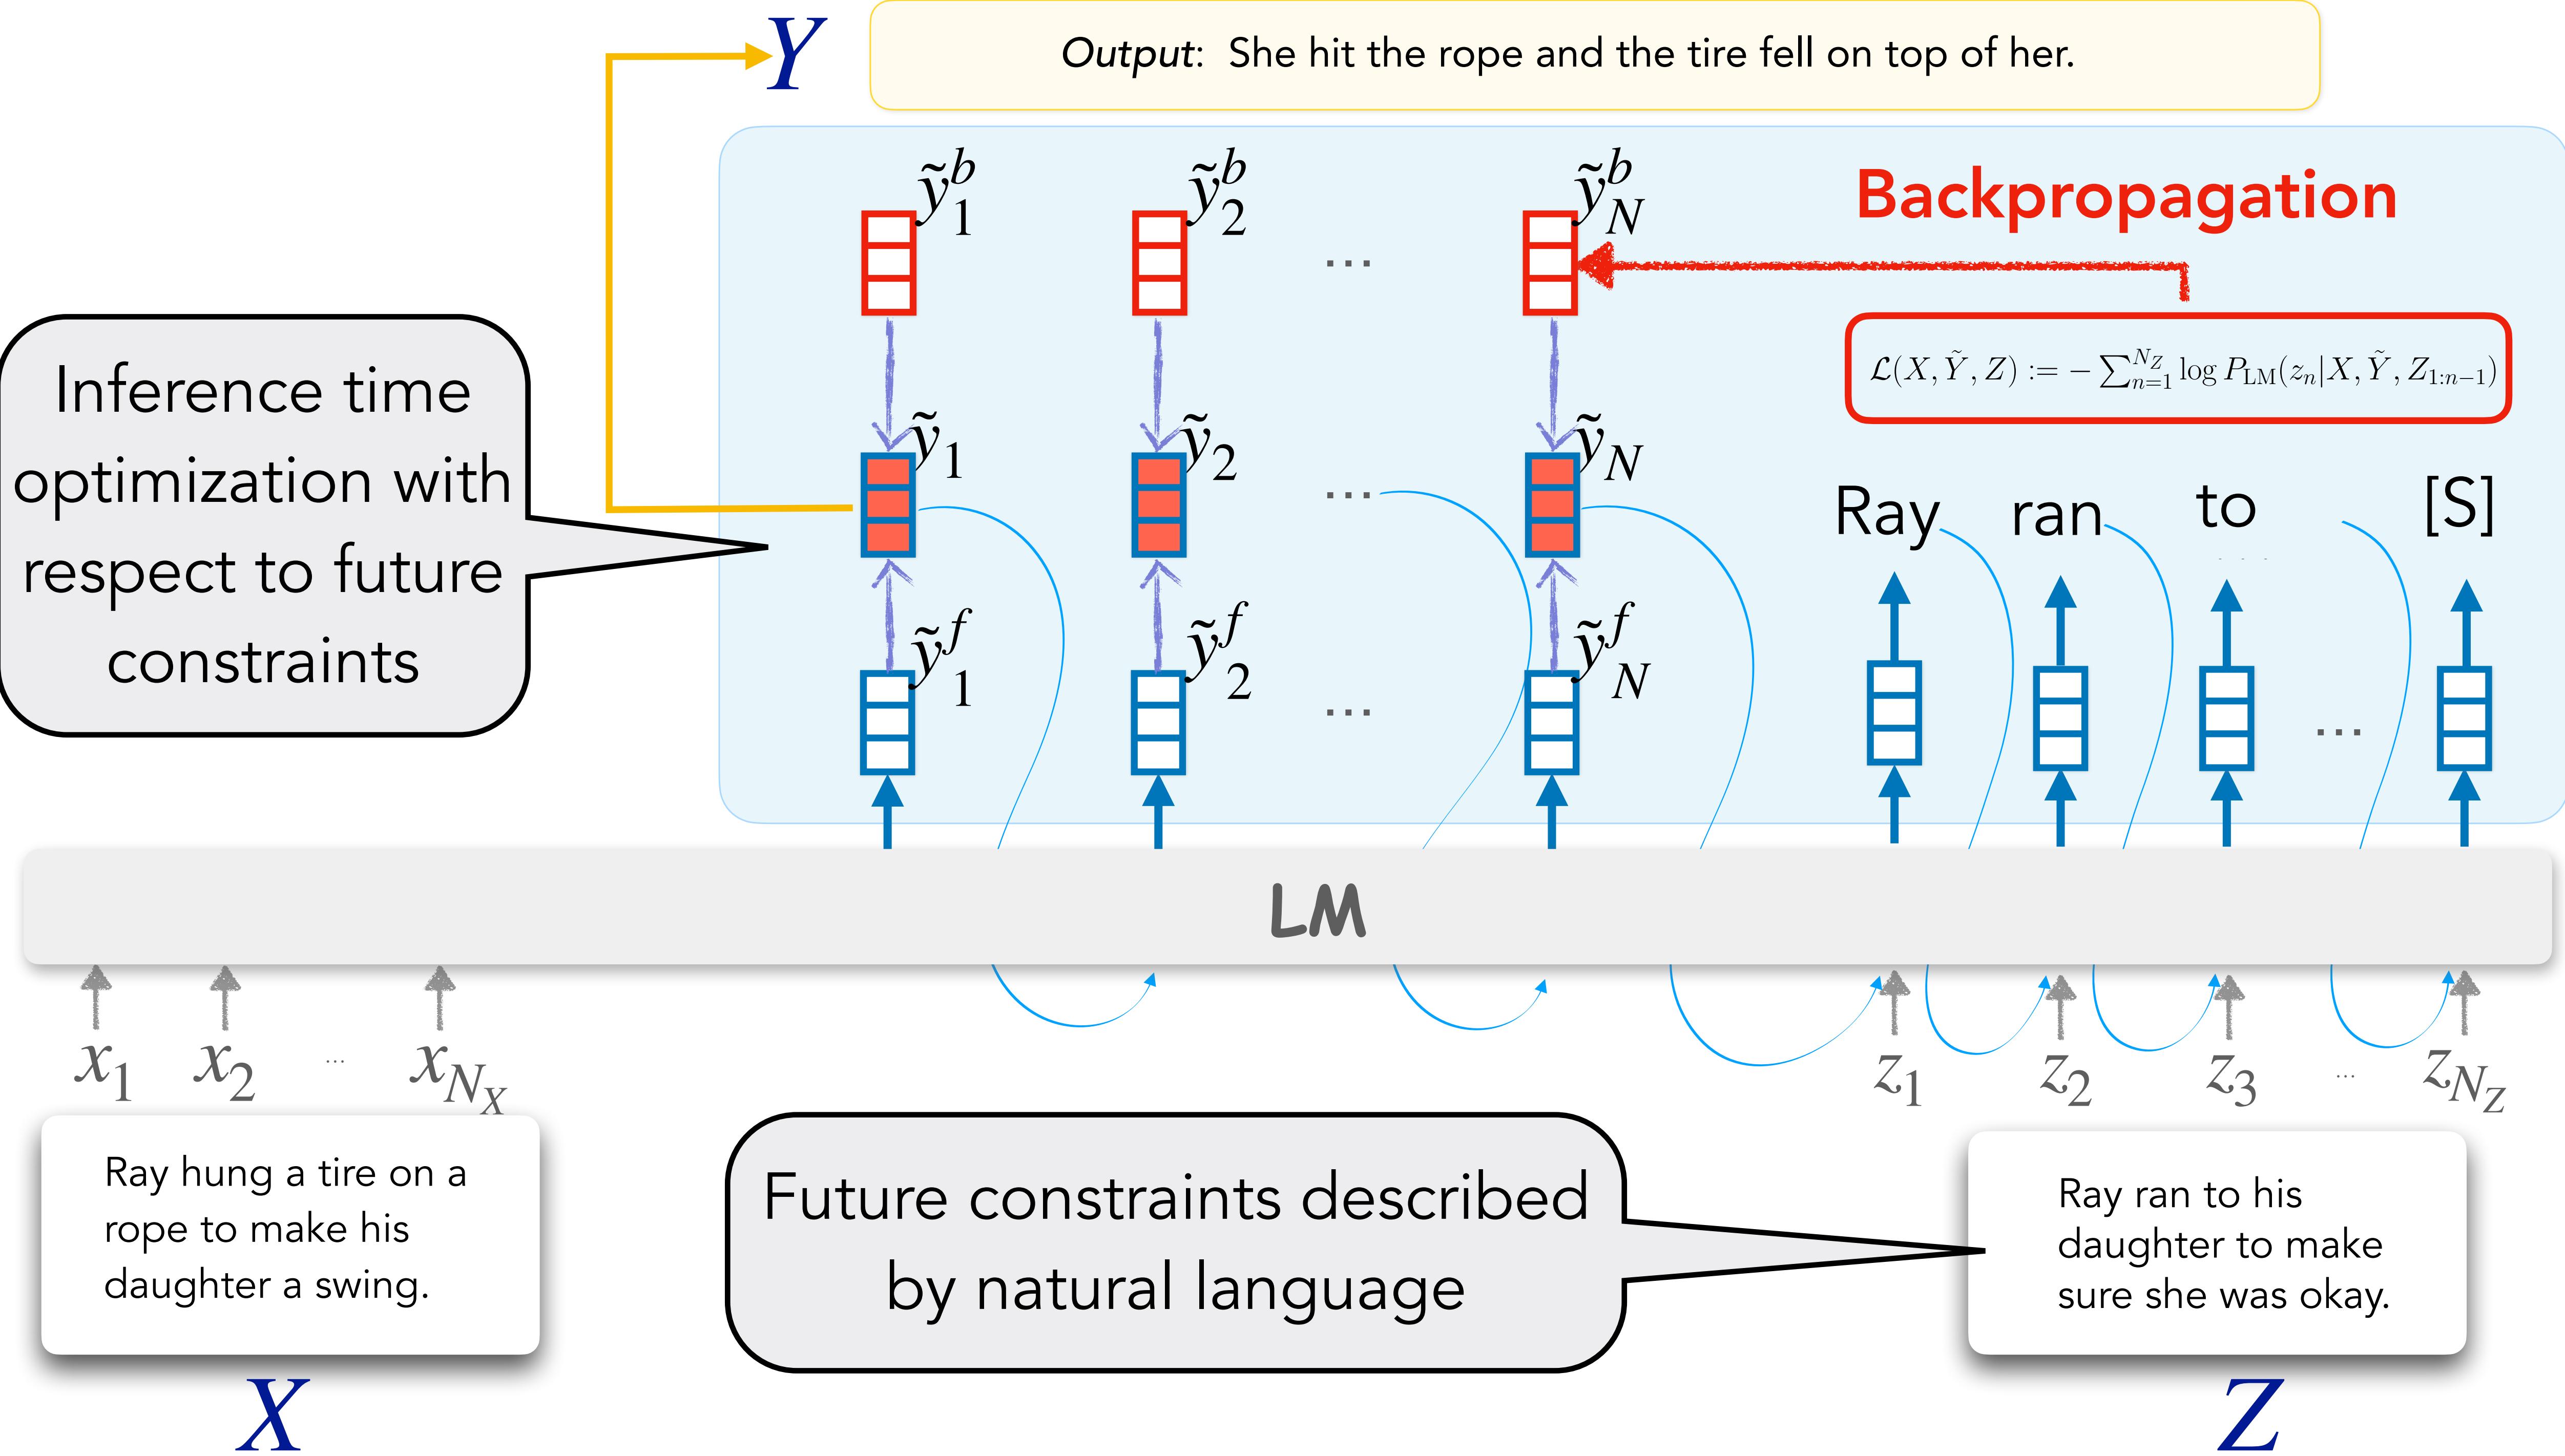
Ray hung a tire on a
rope to make his
daughter a swing.

Future constraints described
by natural language

X

Z

Ray ran to his
daughter to make
sure she was okay.



Training Neural Text Generation Models

Antoine Bosselut



Generation Model Basics

1. At each time step, model computes a score o_n for each token in our vocabulary, $w_n \in V$

$$O_n = f(\{y\}_{<t})$$

f(.) is your model

2. Compute a probability distribution over these scores (usually softmax)

P(.) is your distribution
over tokens

$$P(y_t = w_n | \{y\}_{<t}) = \frac{e^{o_n}}{\sum_{m=1}^M e^{o_m}}$$

Maximizing Likelihood

- Trained to generate the next word given a set of preceding words



↑
Allen

y_0^*

Maximizing Likelihood

- Trained to generate the next word given a set of preceding words



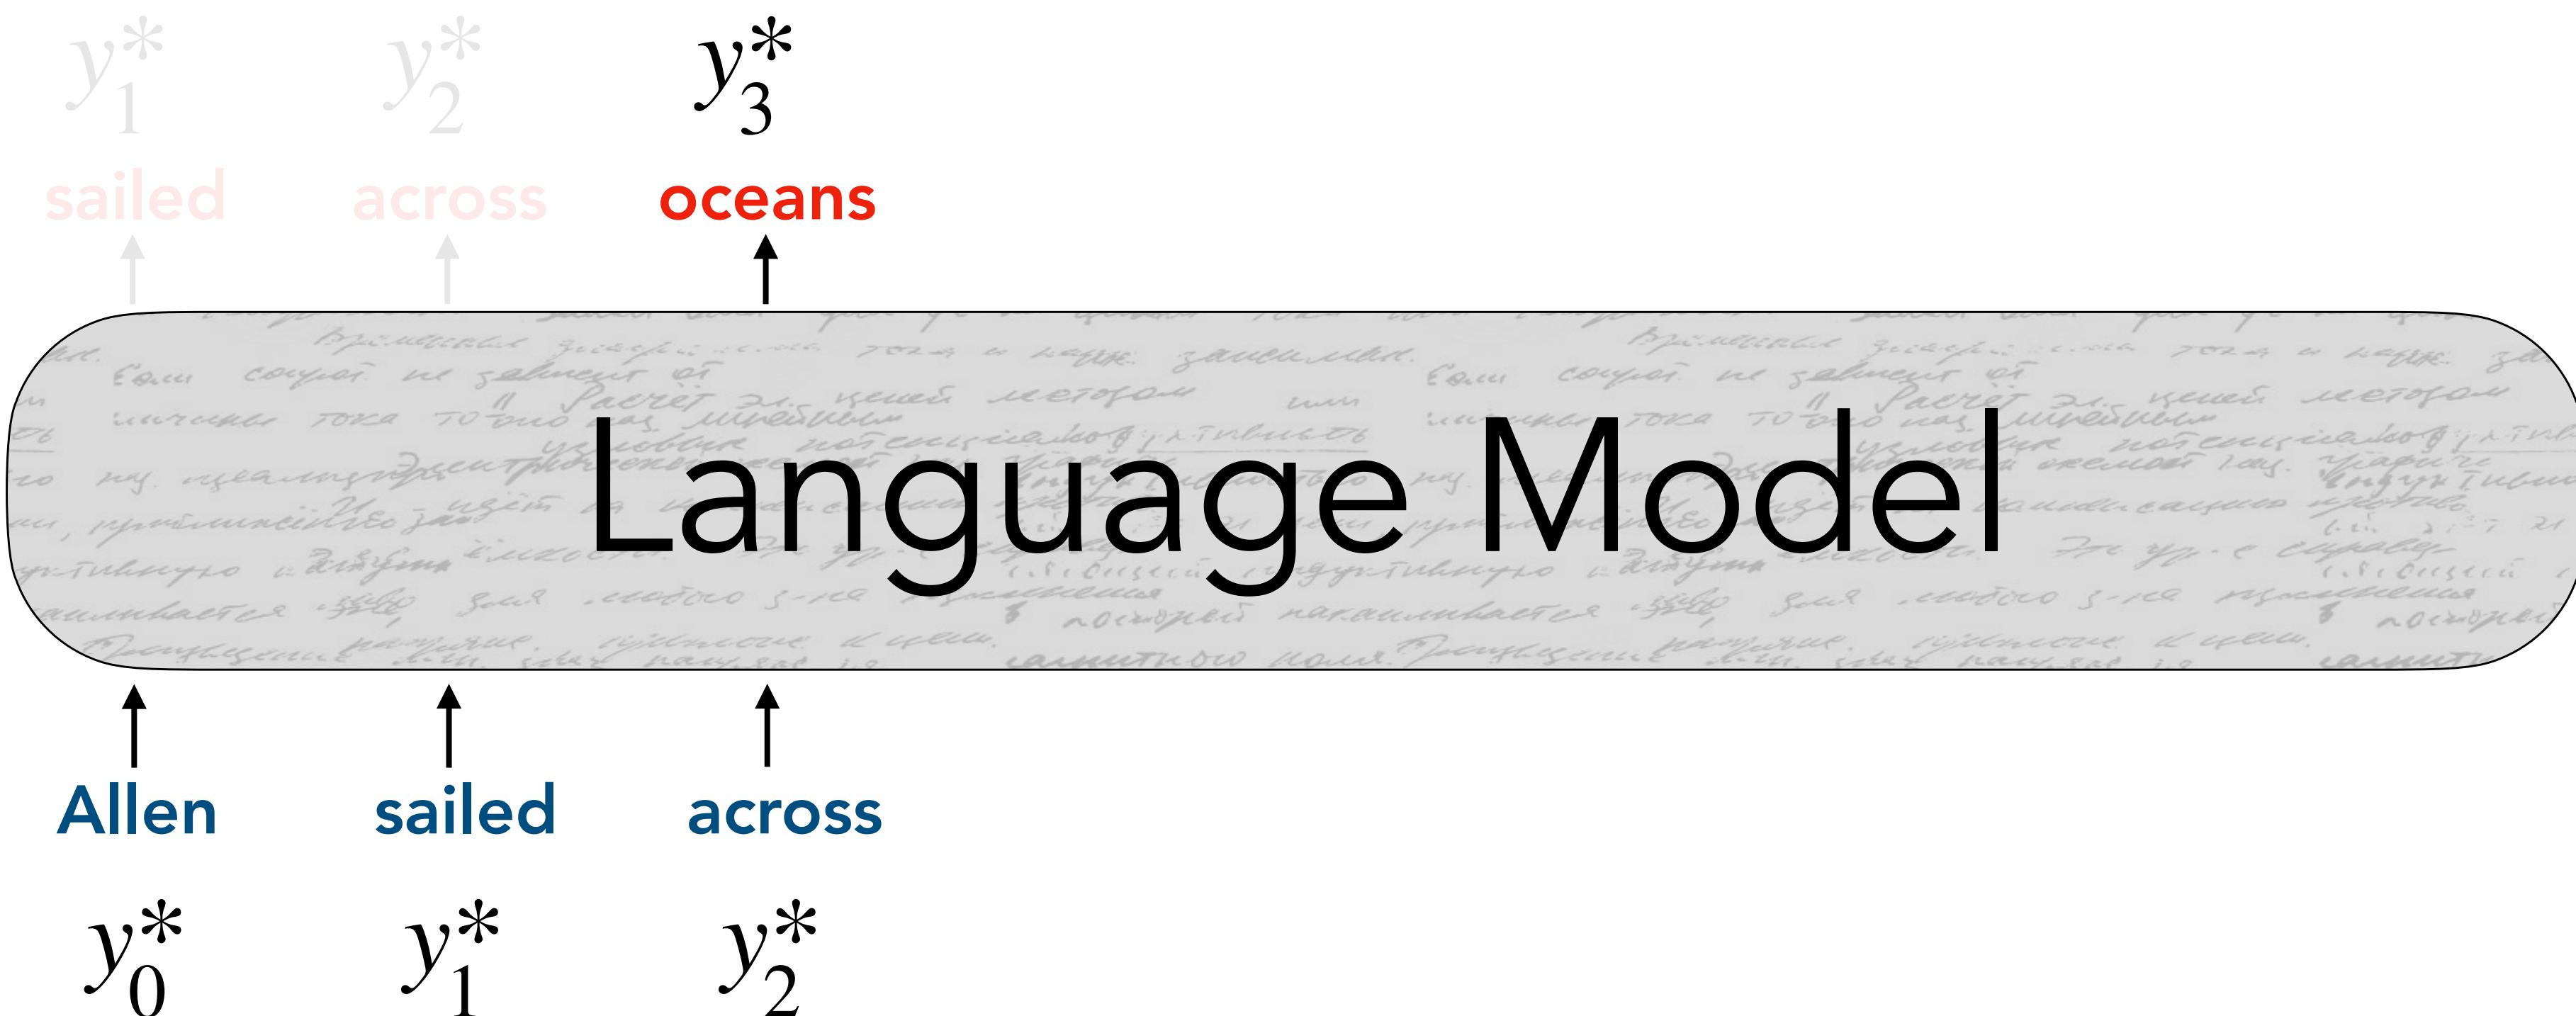
Maximizing Likelihood

- Trained to generate the next word given a set of preceding words



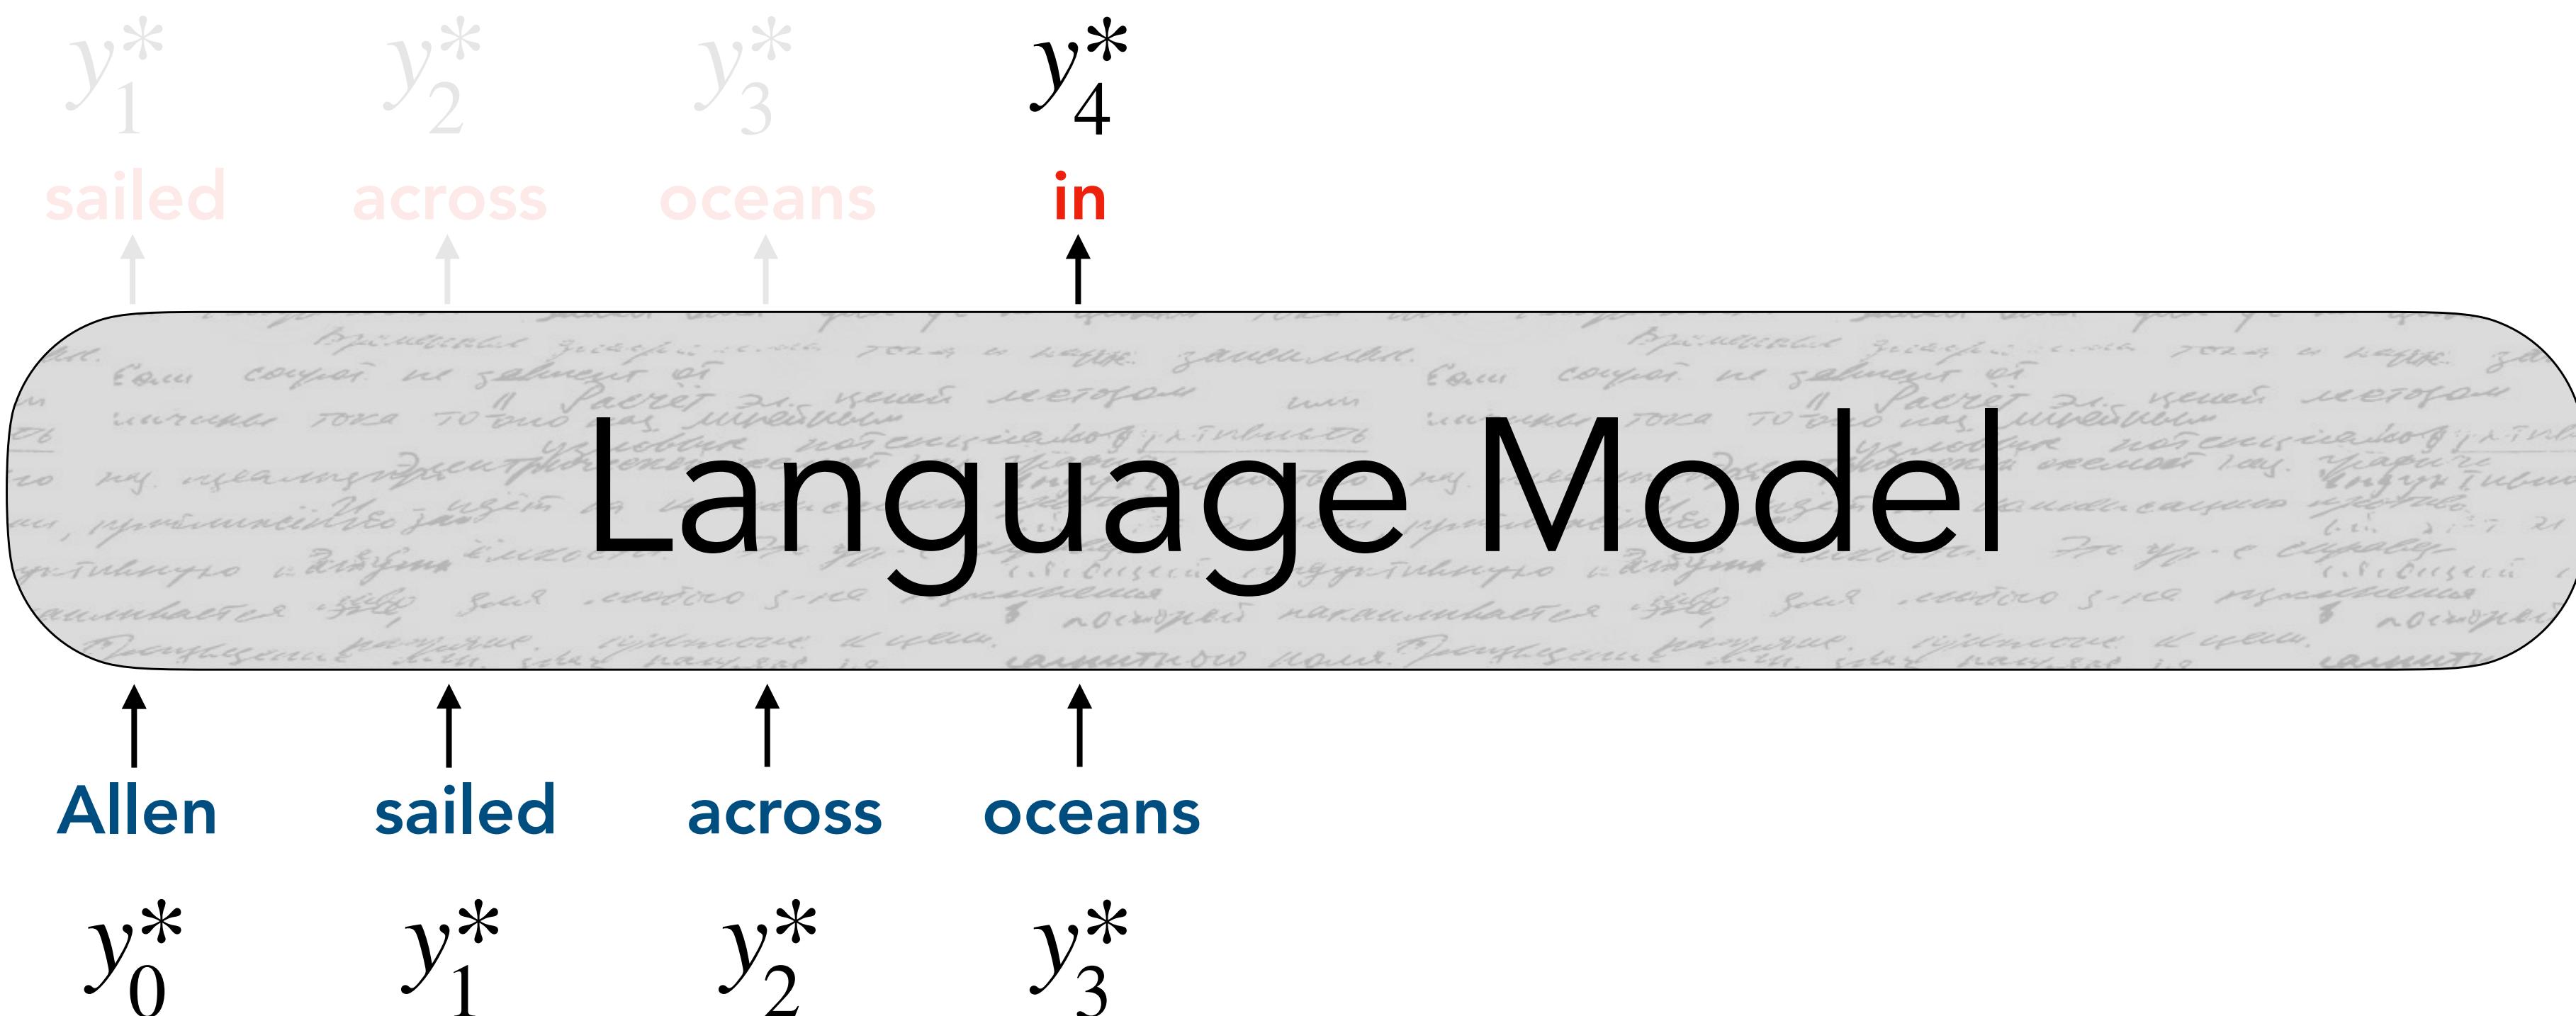
Maximizing Likelihood

- Trained to generate the next word given a set of preceding words



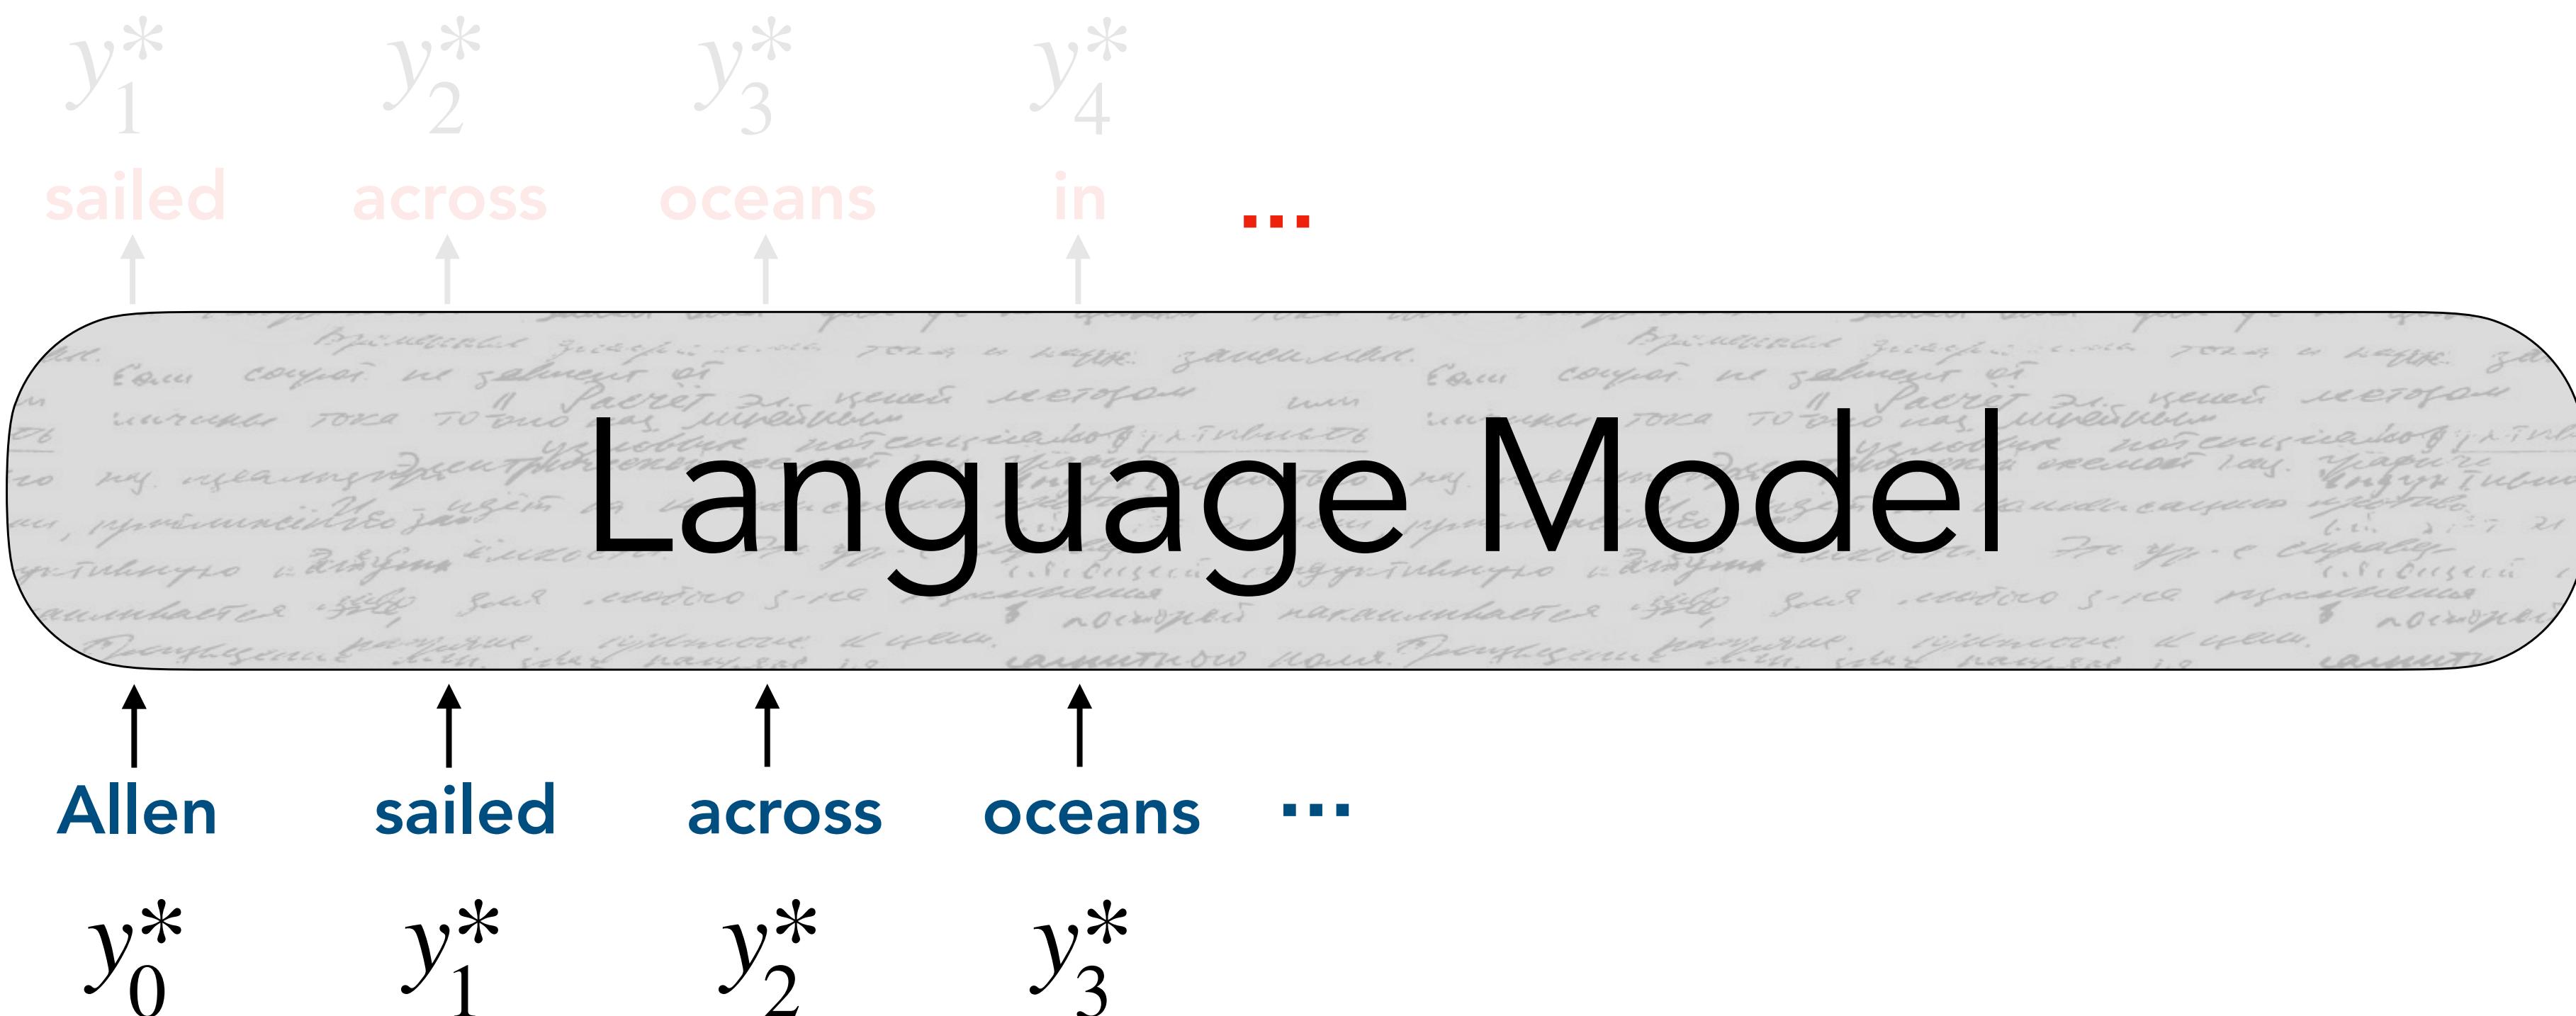
Maximizing Likelihood

- Trained to generate the next word given a set of preceding words



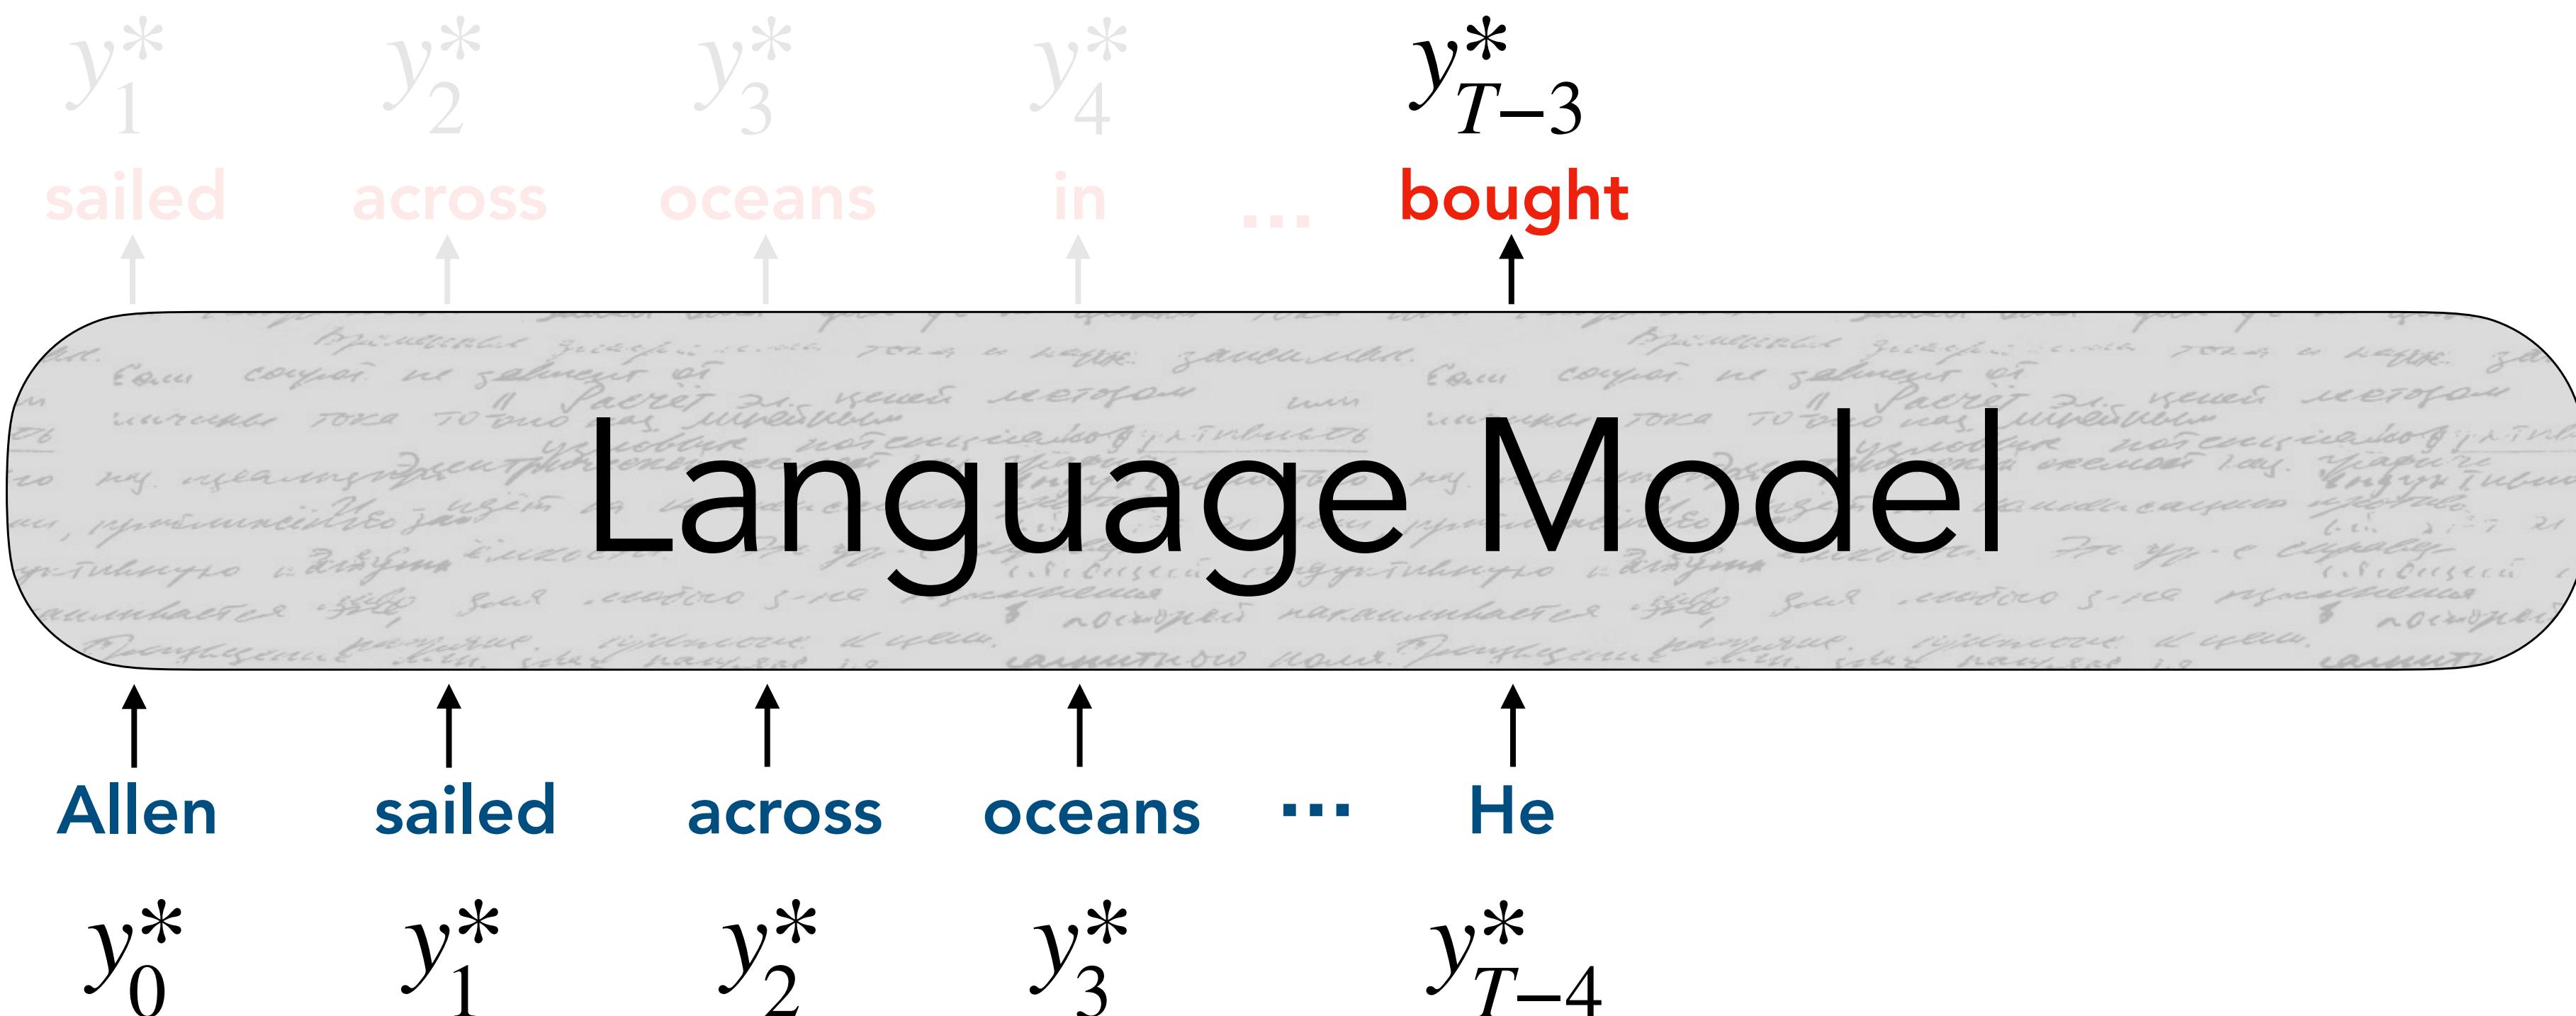
Maximizing Likelihood

- Trained to generate the next word given a set of preceding words



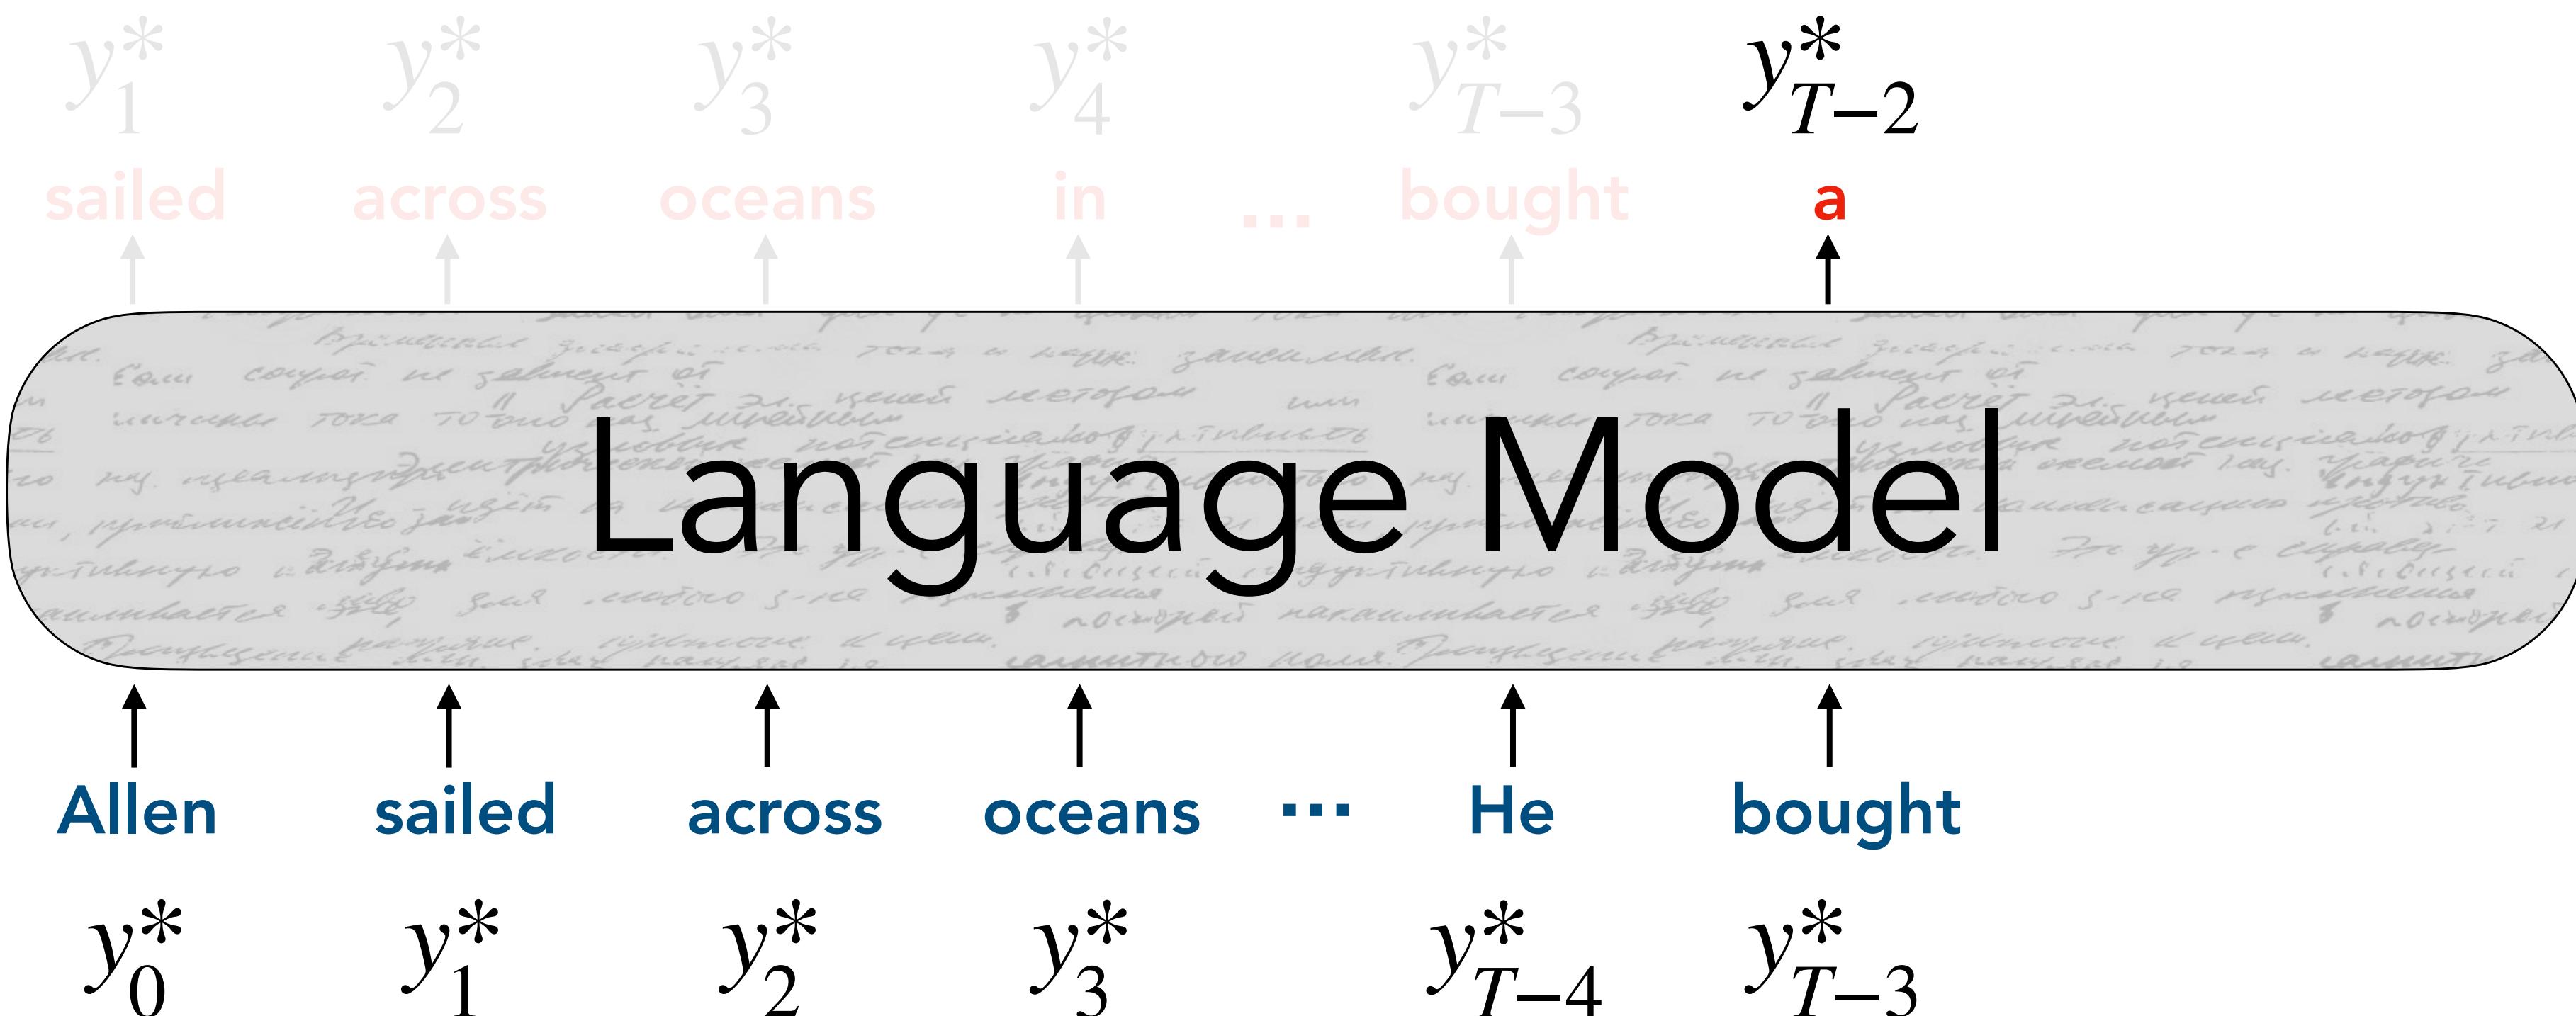
Maximizing Likelihood

- Trained to generate the next word given a set of preceding words



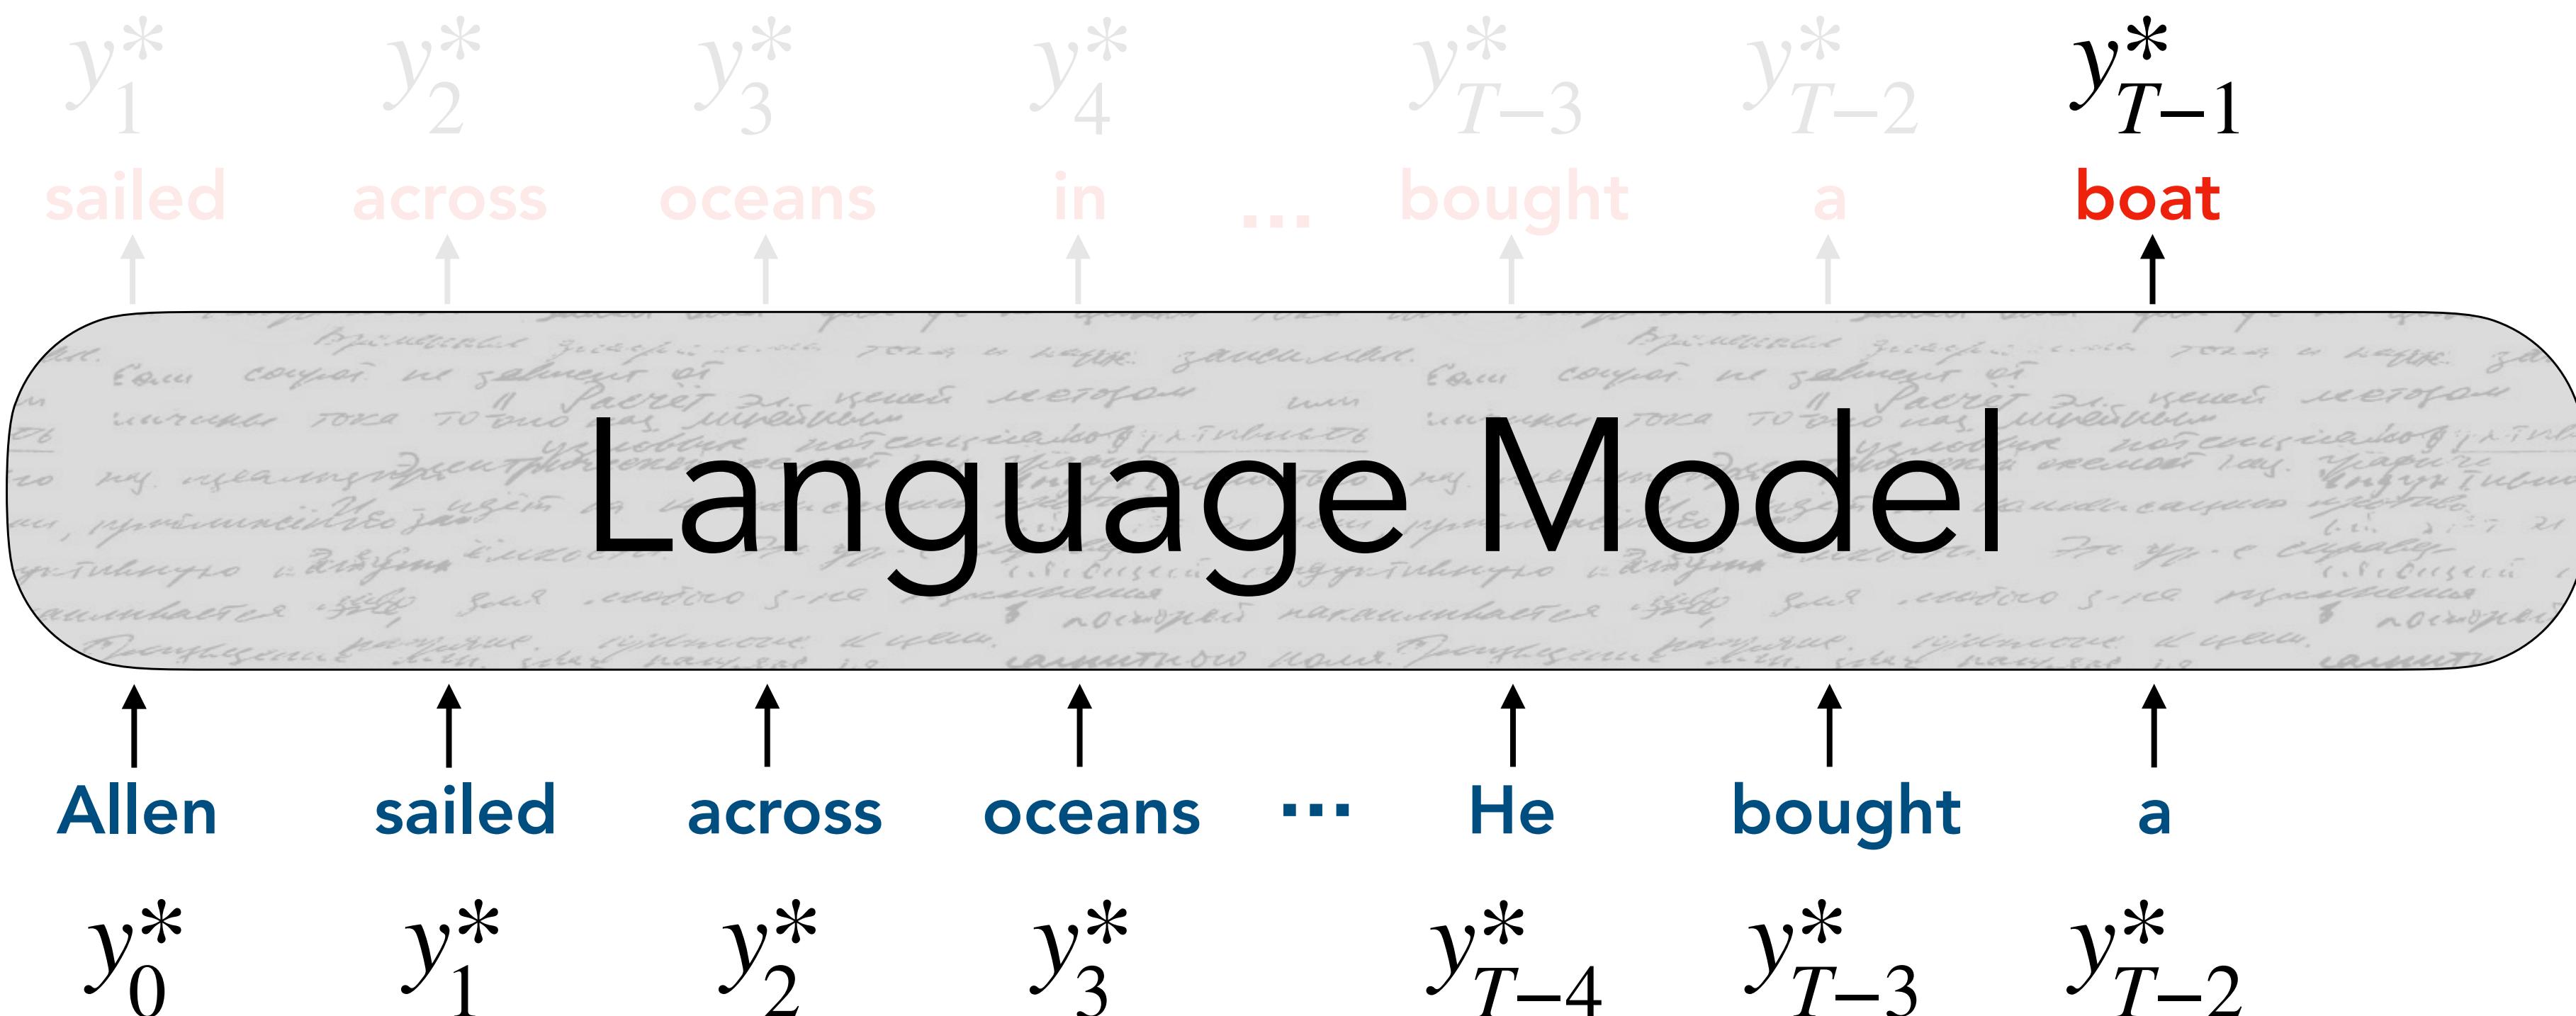
Maximizing Likelihood

- Trained to generate the next word given a set of preceding words



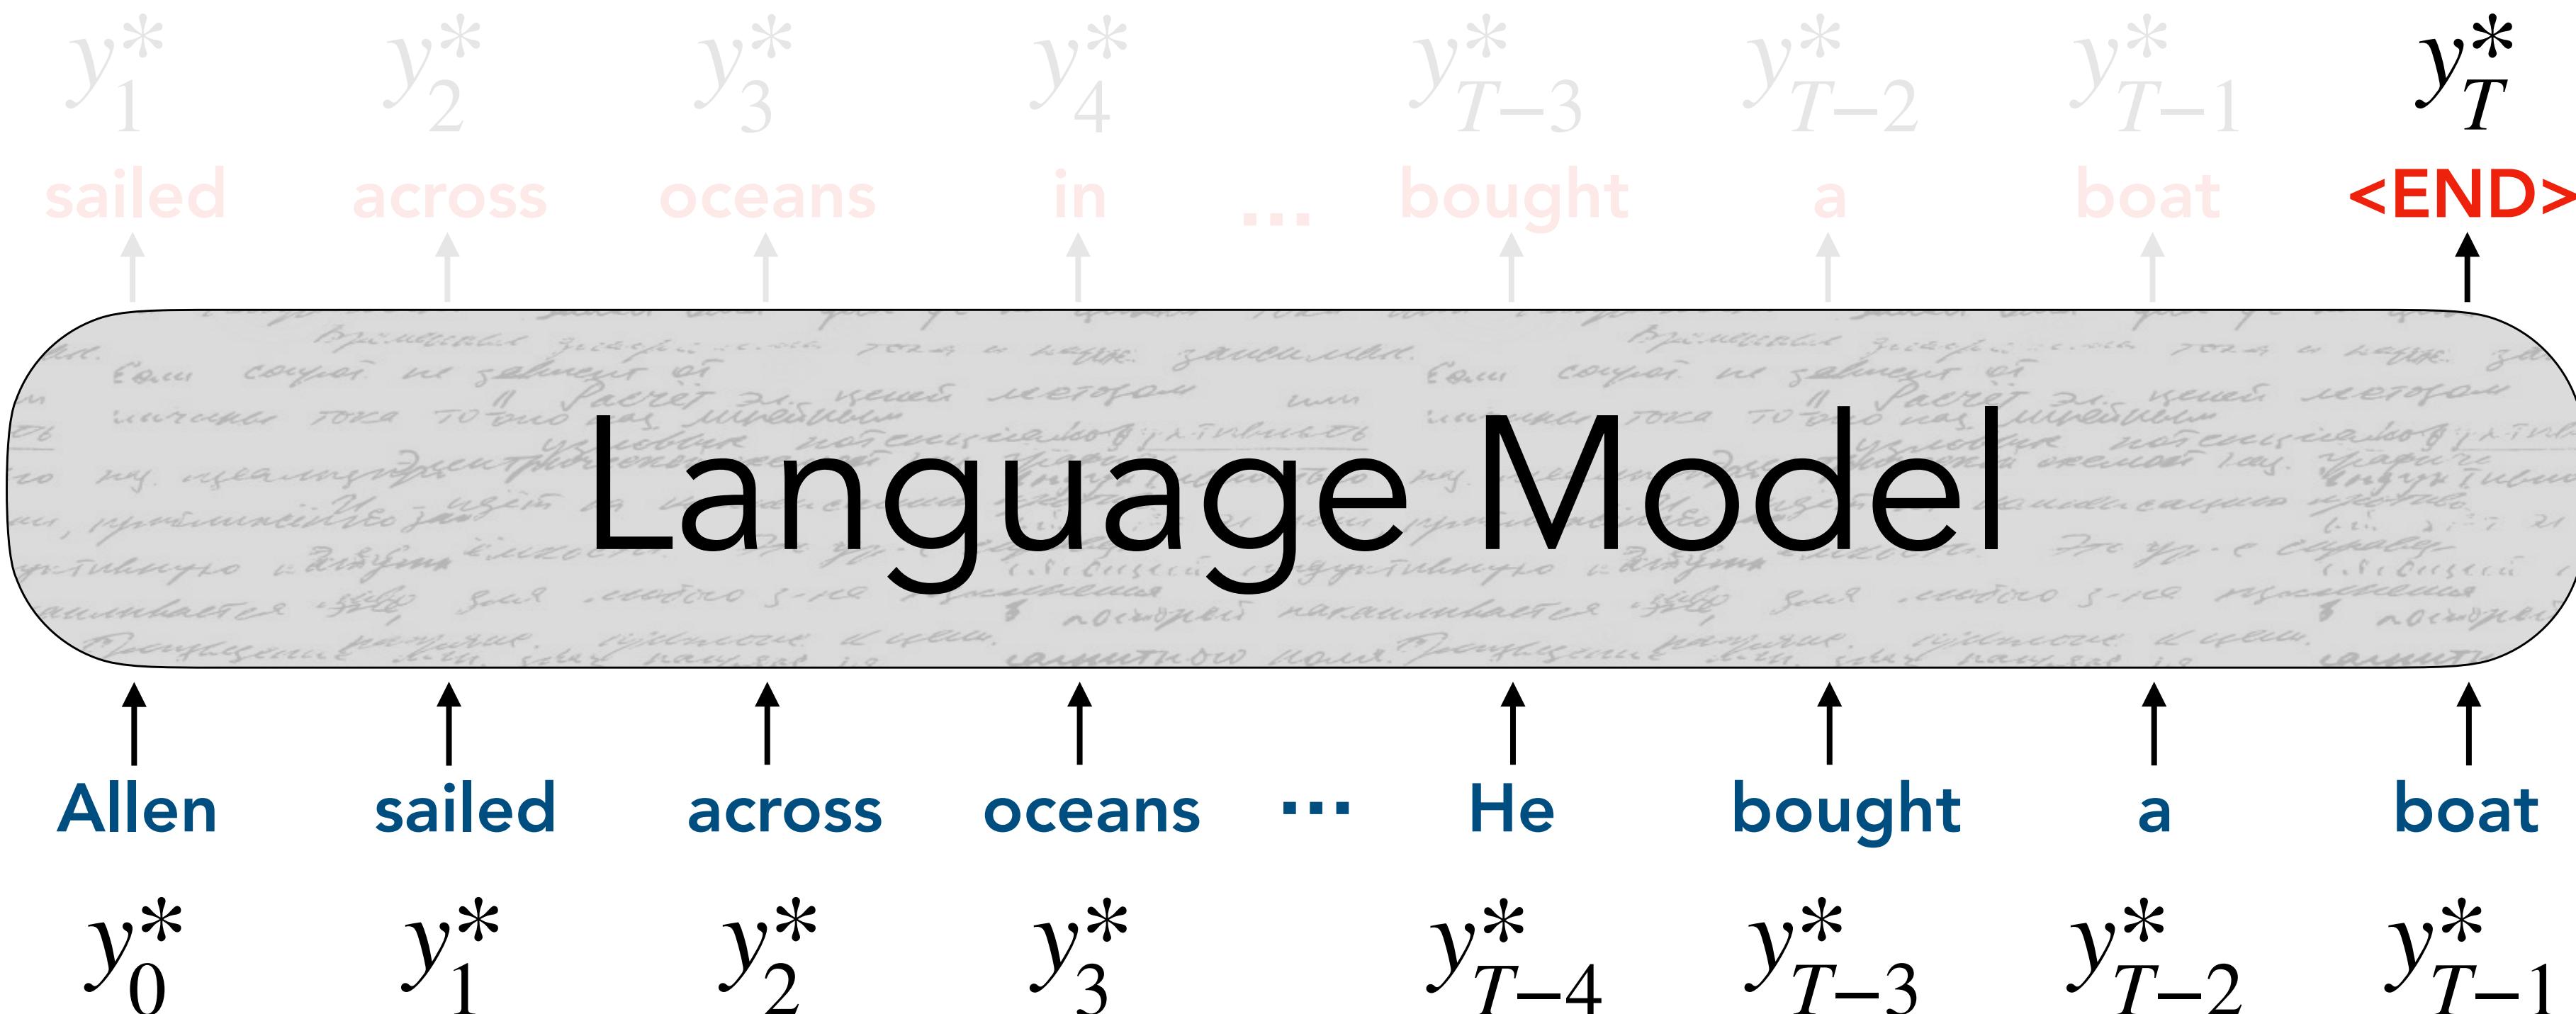
Maximizing Likelihood

- Trained to generate the next word given a set of preceding words



Maximizing Likelihood

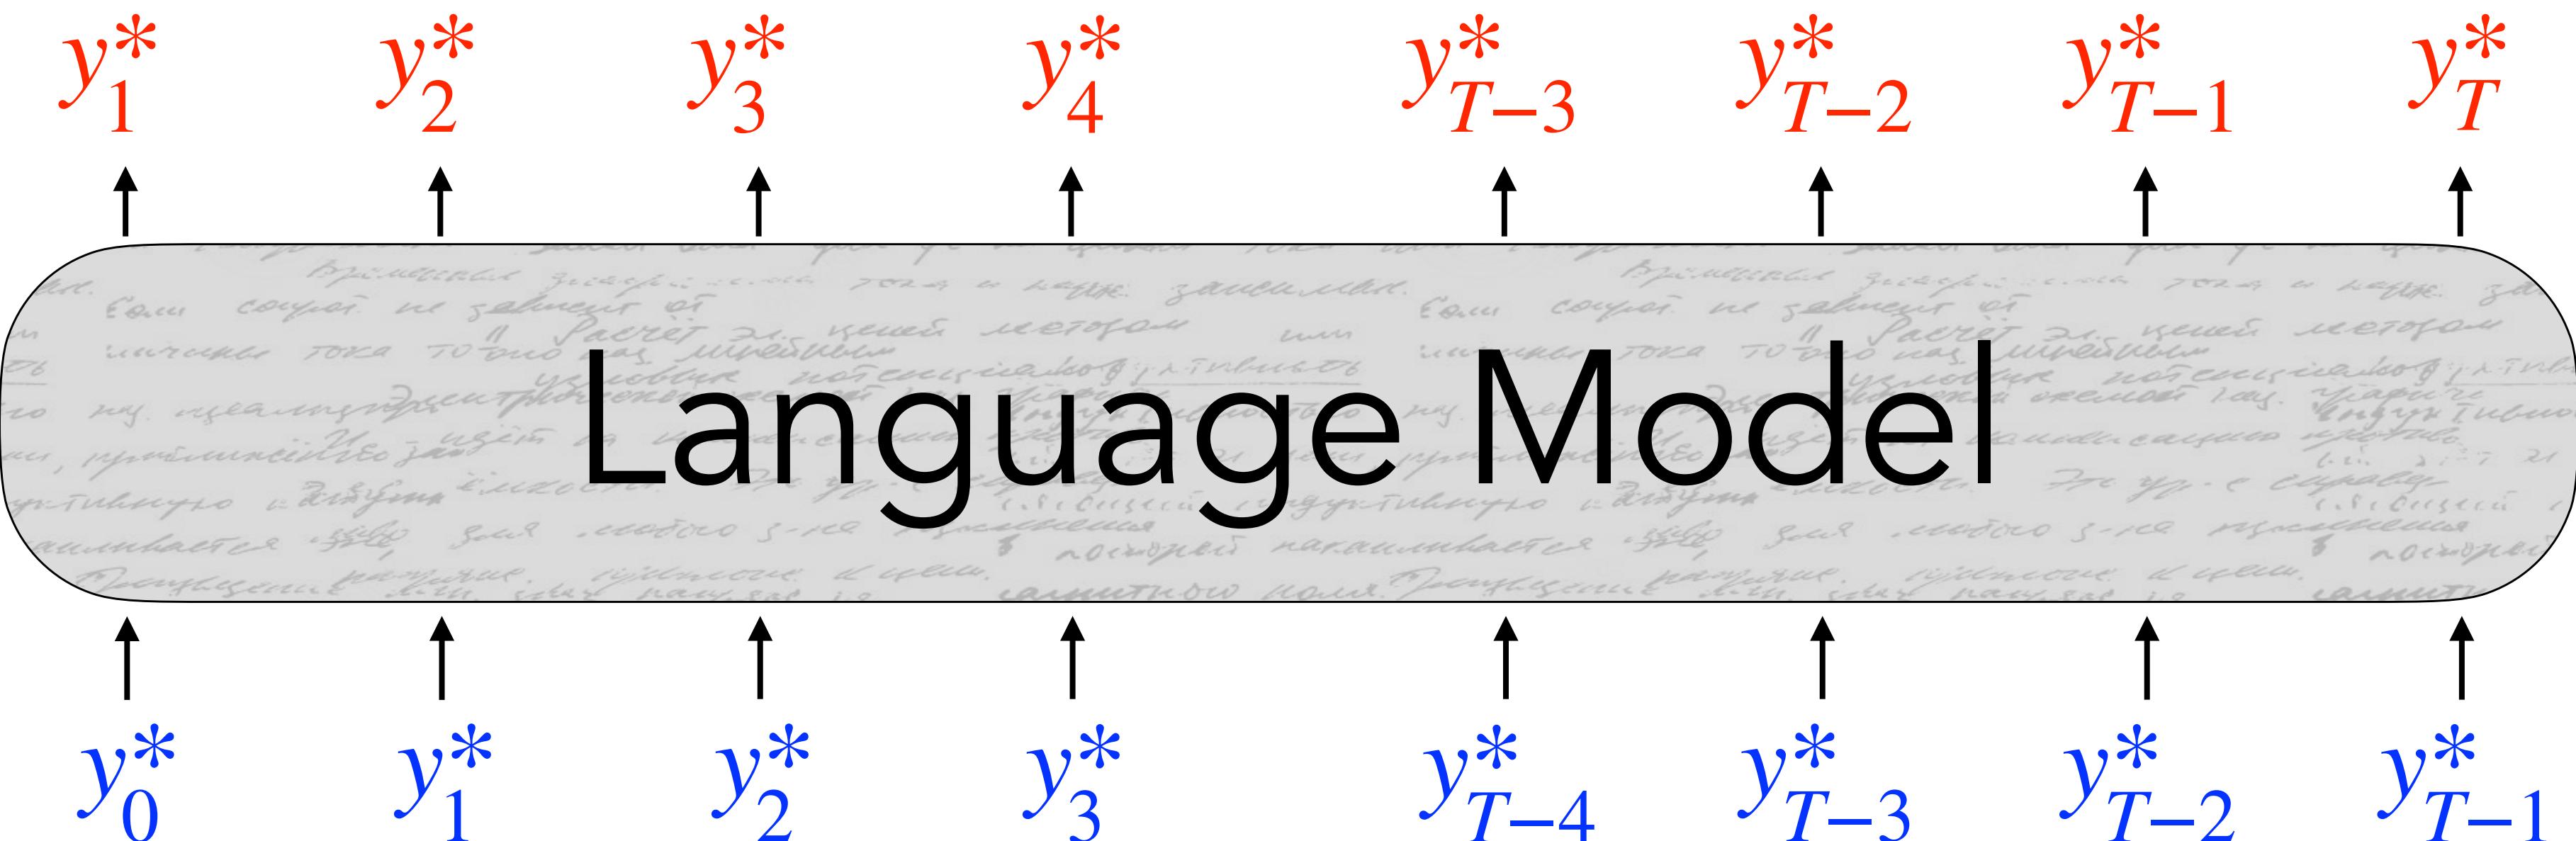
- Trained to generate the next word given a set of preceding words



Maximizing Likelihood

- Trained to generate the next word given a set of preceding words

$$\mathcal{L} = - \sum_{t=1}^T \log P(y_t^* | \{y^*\}_{<t})$$

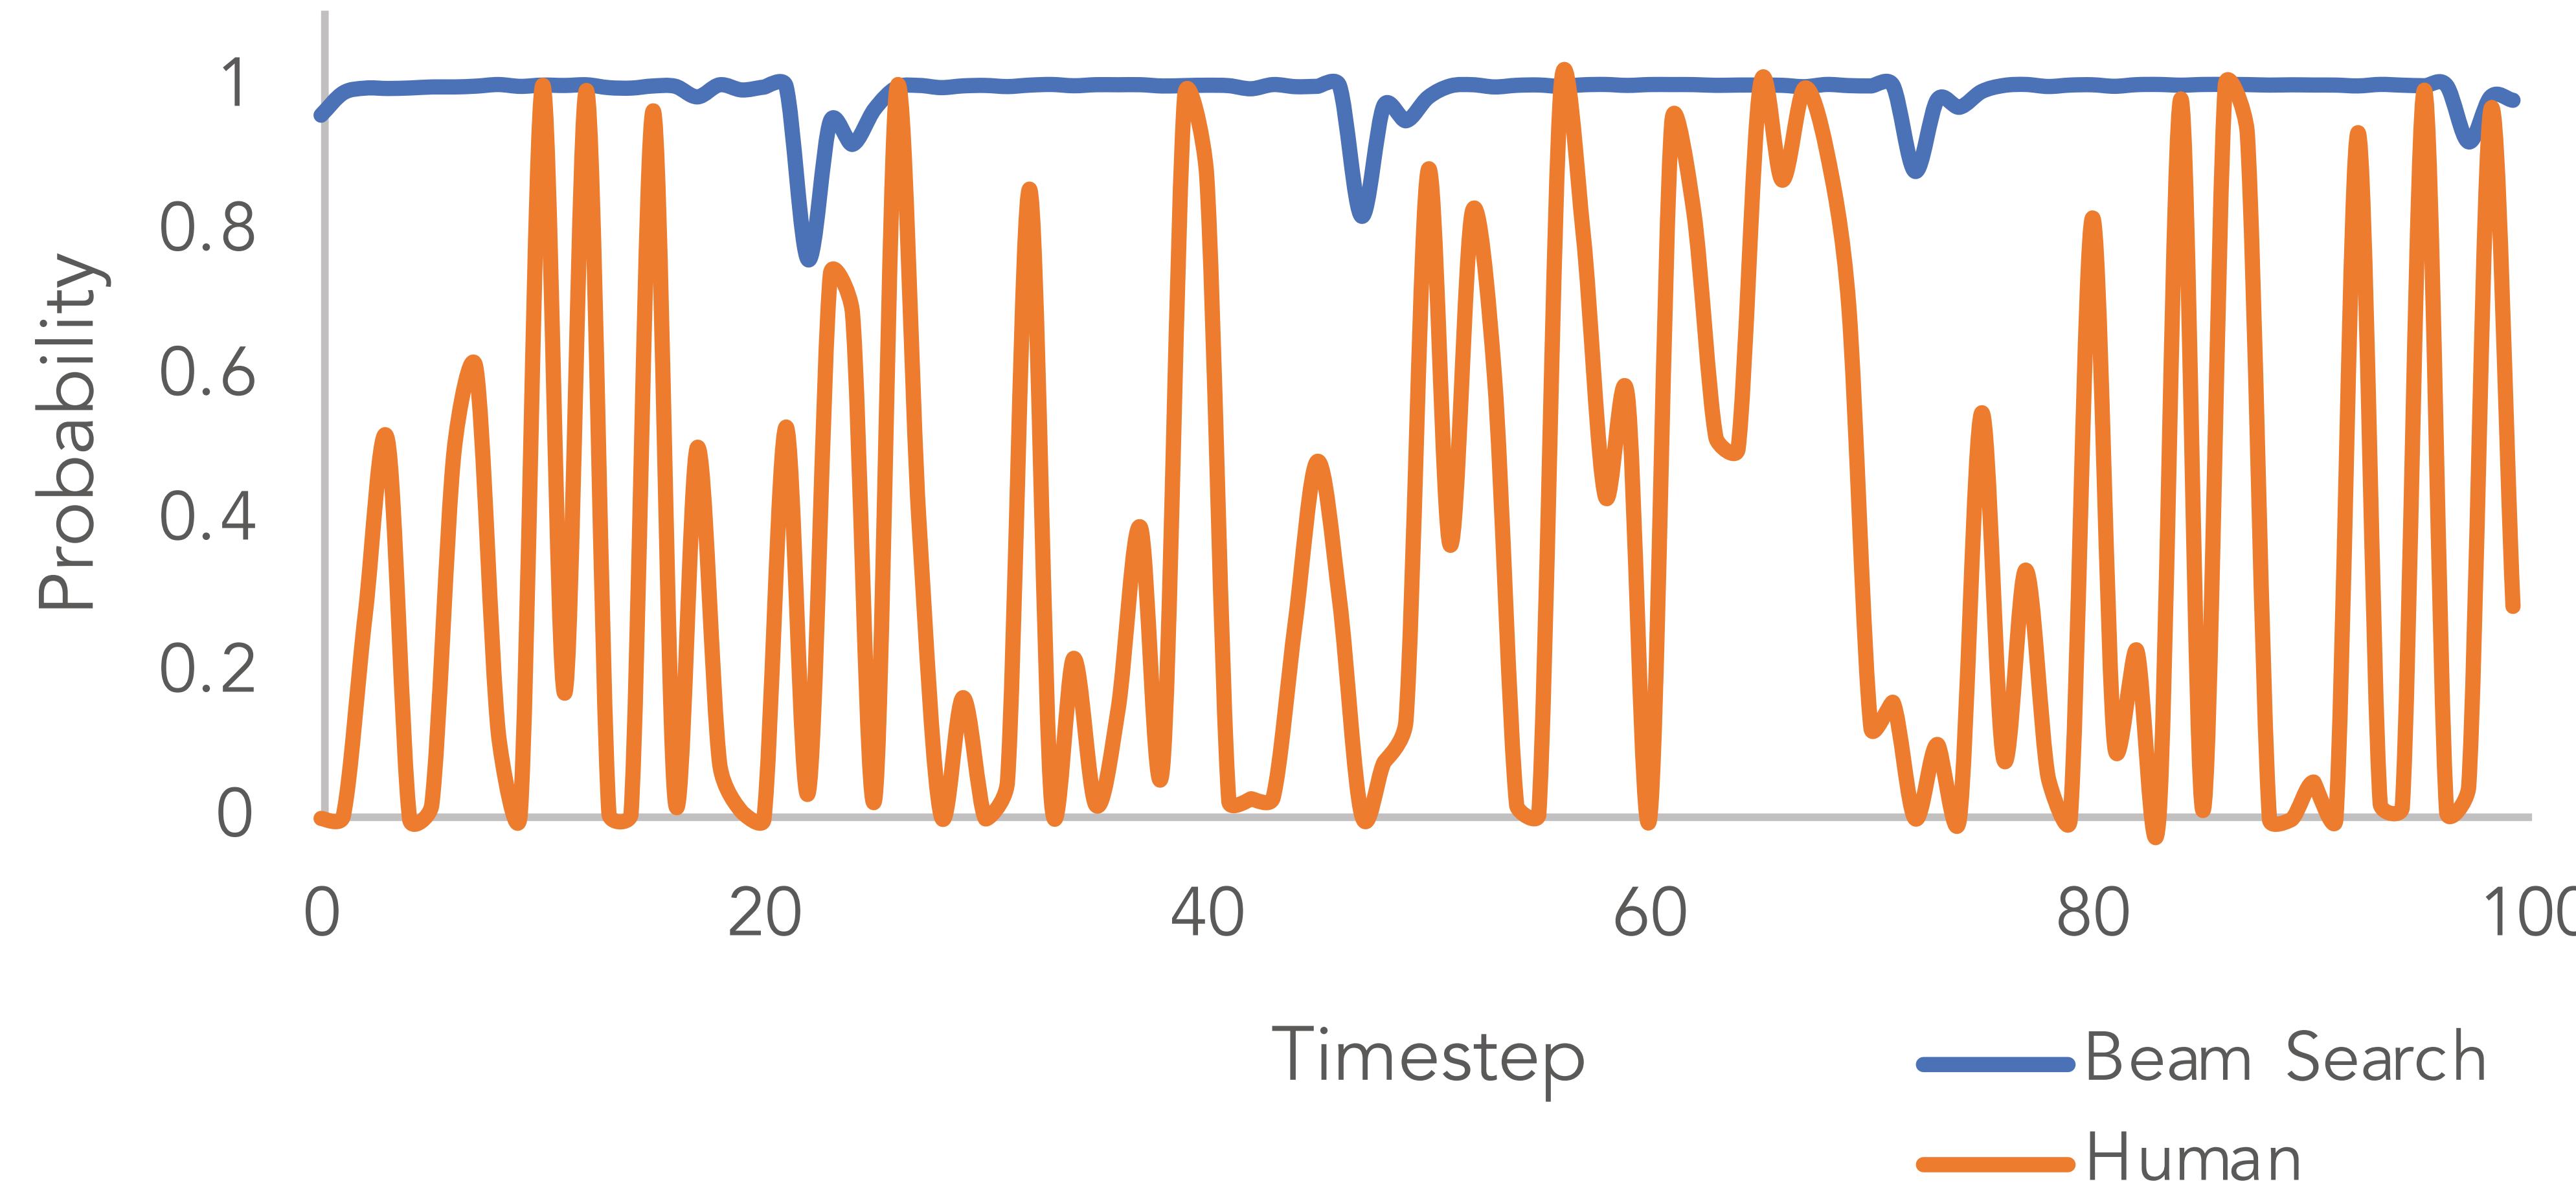


Issue #1: MLE discourages diversity

Context: In a shocking finding, scientist discovered a herd of unicorns living in a remote, previously unexplored valley, in the Andes Mountains. Even more surprising to the researchers was the fact that the unicorns spoke perfect English.

Continuation: The study, published in the Proceedings of the National Academy of Sciences of the United States of America (PNAS), was conducted by researchers from the **Universidad Nacional Autónoma de México (UNAM)** and **the Universidad Nacional Autónoma de México (UNAM/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México/Universidad Nacional Autónoma de México...)**

Issue #1: MLE discourages diversity



Alternatives: Unlikelihood Training

- Sequence-level Unlikelihood Training
- Given a set of undesired tokens \mathcal{C} , lower their likelihood in context

$$\mathcal{L}_{MLE}^t = -\log P(y_t^* \mid \{y^*\}_{<t})$$

$$\mathcal{L}_{UL}^t = - \sum_{y_{neg} \in \mathcal{C}} \log(1 - P(y_{neg} \mid \{y^*\}_{<t}))$$

Alternatives: Unlikelihood Training

- Sequence-level Unlikelihood Training
- Given a set of undesired tokens \mathcal{C} , lower their likelihood in context

$$\mathcal{L}_{MLE}^t = -\log P(y_t^* | \{y^*\}_{<t})$$

$$\mathcal{L}_{UL}^t = - \sum_{y_{neg} \in \mathcal{C}} \log(1 - P(y_{neg} | \{y^*\}_{<t}))$$

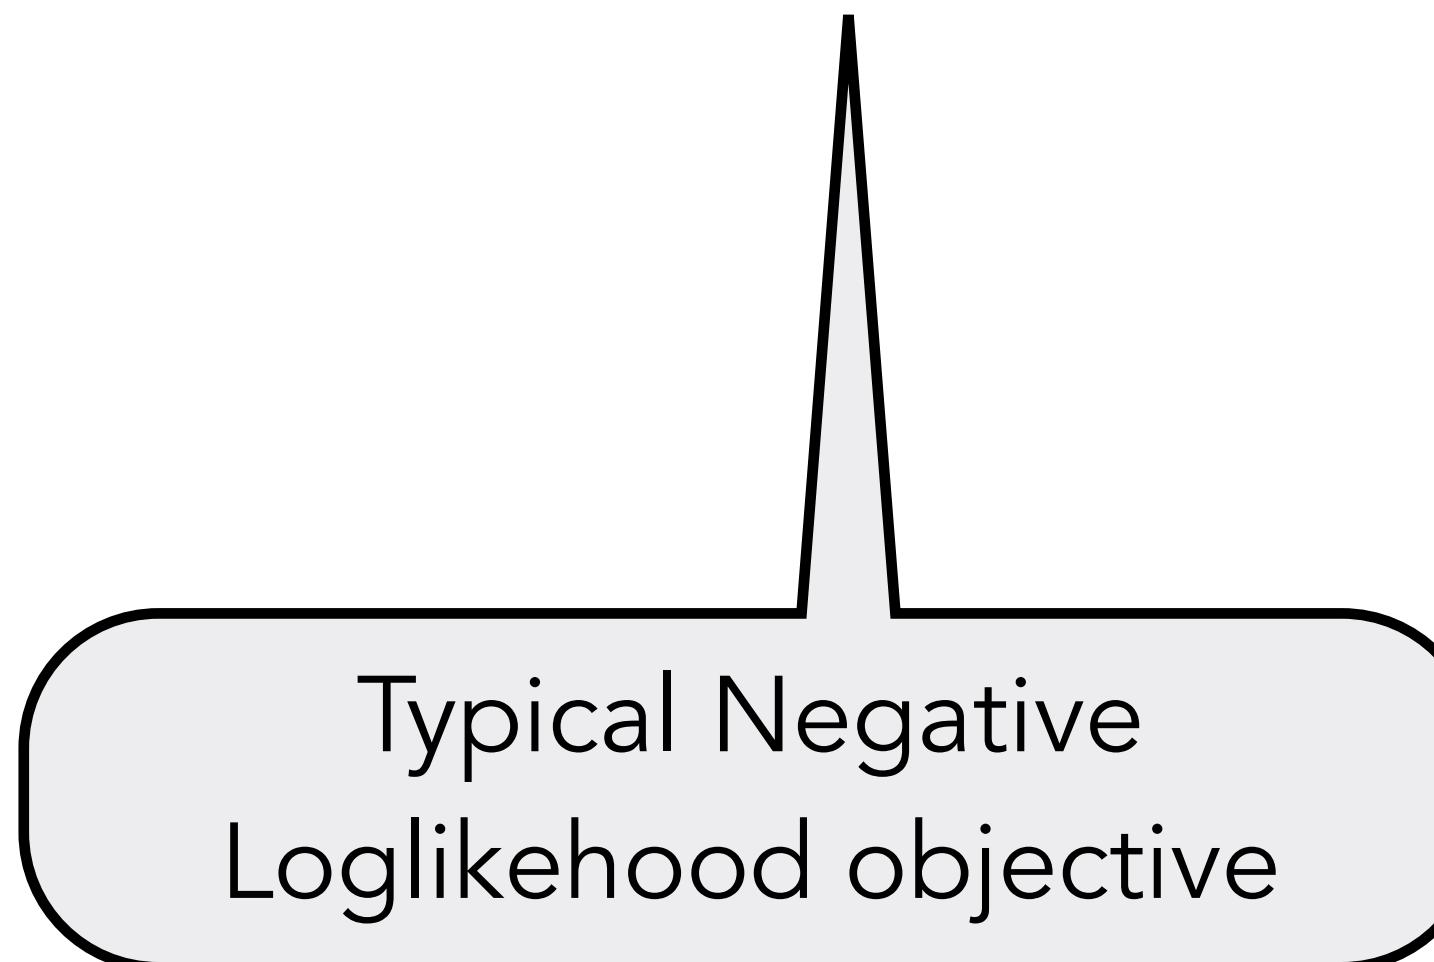


Typical Negative
Loglikelihood objective

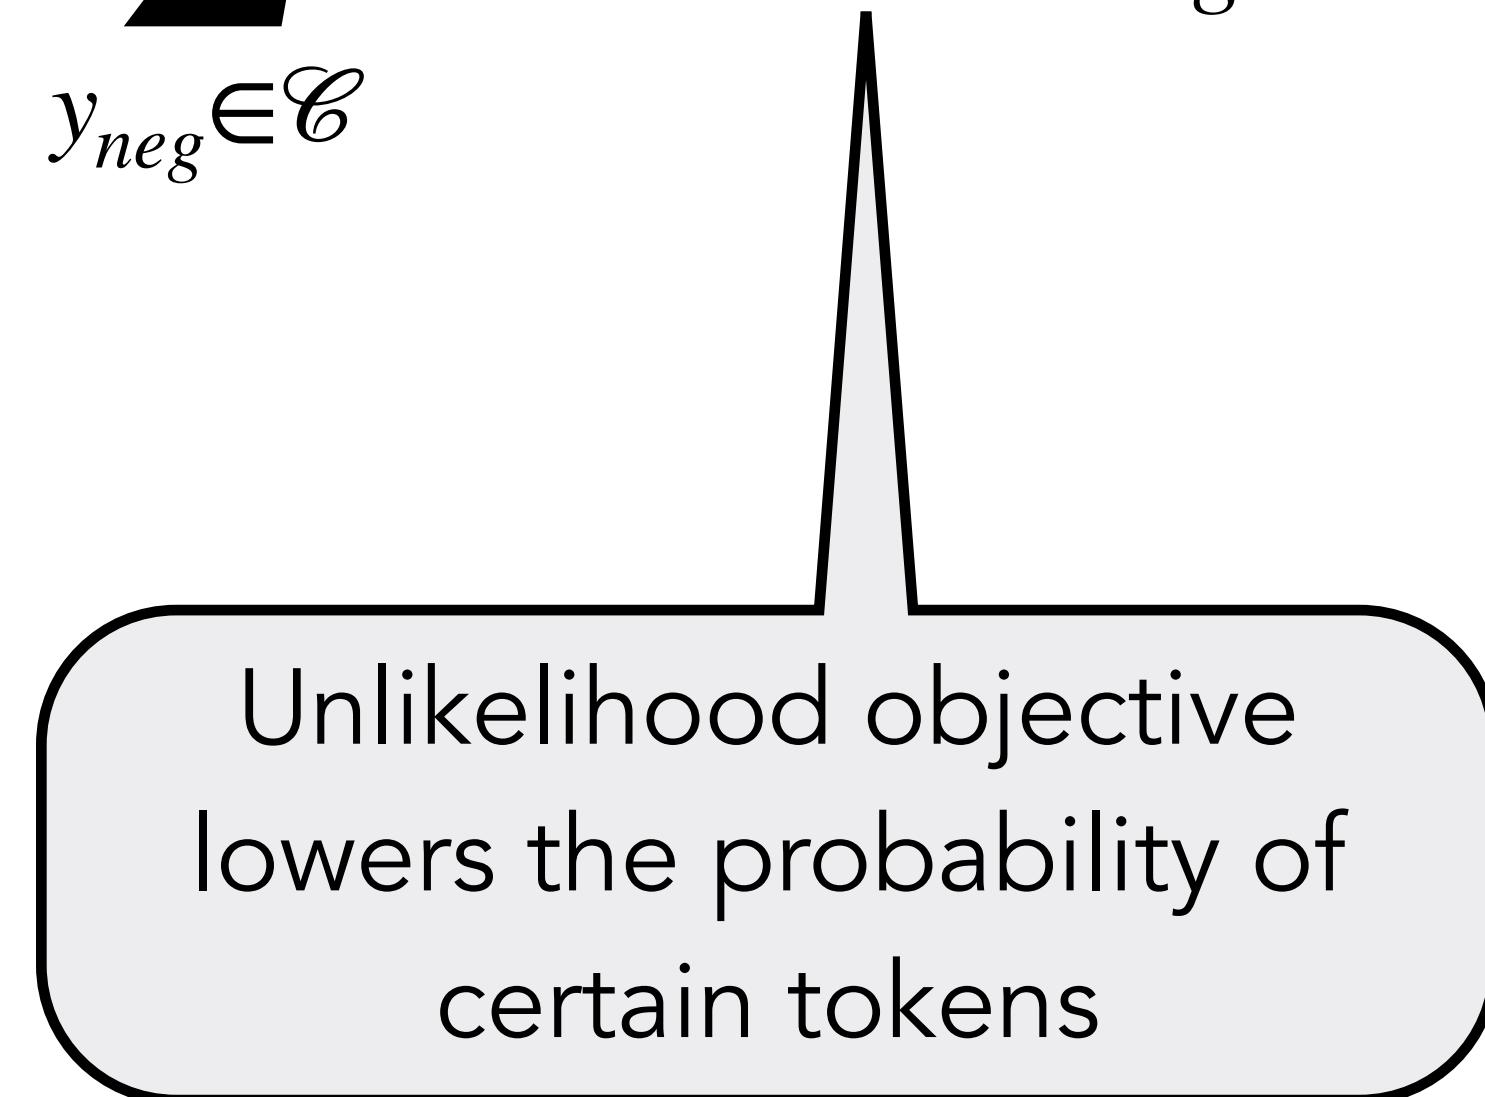
Alternatives: Unlikelihood Training

- Sequence-level Unlikelihood Training
- Given a set of undesired tokens \mathcal{C} , lower their likelihood in context

$$\mathcal{L}_{MLE}^t = -\log P(y_t^* | \{y^*\}_{<t})$$



$$\mathcal{L}_{UL}^t = - \sum_{y_{neg} \in \mathcal{C}} \log(1 - P(y_{neg} | \{y^*\}_{<t}))$$



Alternatives: Unlikelihood Training

- Sequence-level Unlikelihood Training
- Given a set of undesired tokens \mathcal{C} , lower their likelihood in context

$$\mathcal{L}_{MLE}^t = -\log P(y_t^* | \{y^*\}_{<t})$$

$$\mathcal{L}_{UL}^t = - \sum_{y_{neg} \in \mathcal{C}} \log(1 - P(y_{neg} | \{y^*\}_{<t}))$$

$$\mathcal{L}_{ULE}^t = \mathcal{L}_{MLE}^t + \alpha \mathcal{L}_{UL}^t$$

Combine them for full
unlikelihood training

Alternatives: Unlikelihood Training

- Sequence-level Unlikelihood Training
- Given a set of undesired tokens \mathcal{C} , lower their likelihood in context

$$\mathcal{L}_{MLE}^t = -\log P(y_t^* | \{y^*\}_{<t})$$

$$\mathcal{L}_{UL}^t = - \sum_{y_{neg} \in \mathcal{C}} \log(1 - P(y_{neg} | \{y^*\}_{<t}))$$

$$\mathcal{L}_{ULE}^t = \mathcal{L}_{MLE}^t + \alpha \mathcal{L}_{UL}^t$$

But wait, what's \mathcal{C} ?

Alternatives: Unlikelihood Training

- Sequence-level Unlikelihood Training
- Given a set of undesired tokens \mathcal{C} , lower their likelihood in context

$$\mathcal{L}_{MLE}^t = -\log P(y_t^* | \{y^*\}_{<t})$$

$$\mathcal{L}_{UL}^t = - \sum_{y_{neg} \in \mathcal{C}} \log(1 - P(y_{neg} | \{y^*\}_{<t}))$$

$$\mathcal{L}_{ULE}^t = \mathcal{L}_{MLE}^t + \alpha \mathcal{L}_{UL}^t$$

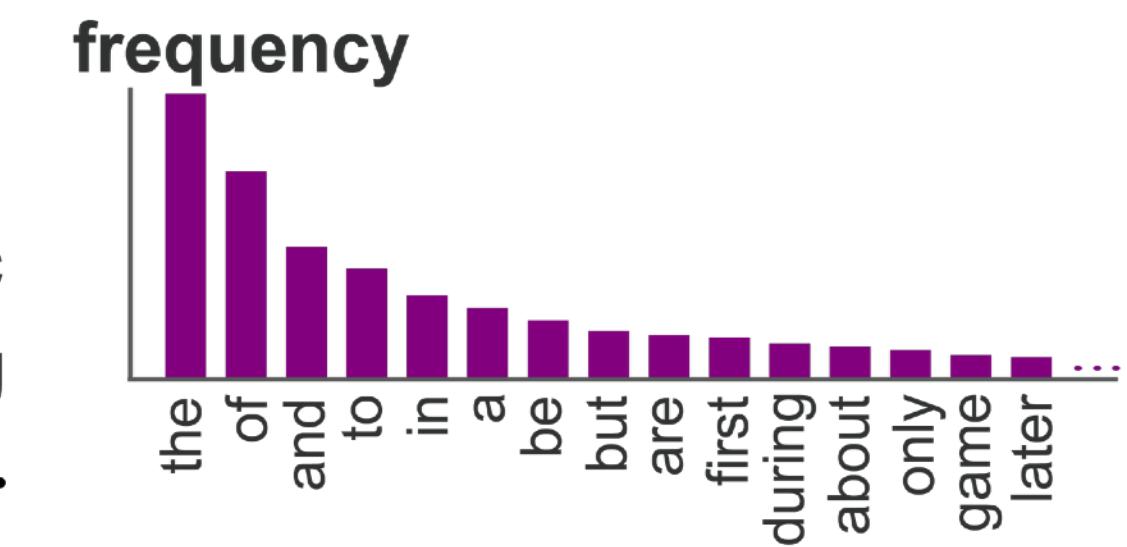
$$\mathcal{C} = \{y^*\}_{<t}$$

Alternatives: F² Softmax

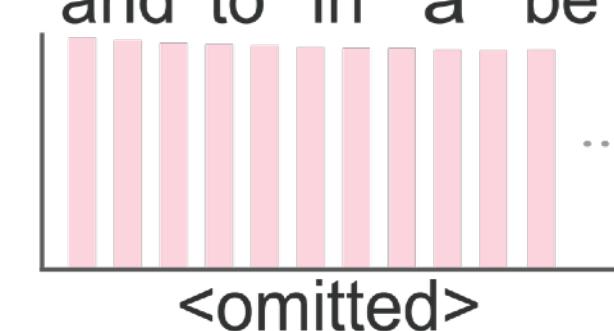
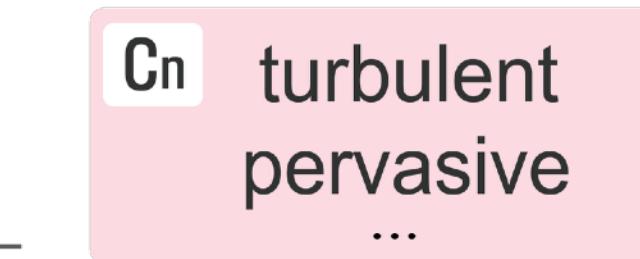
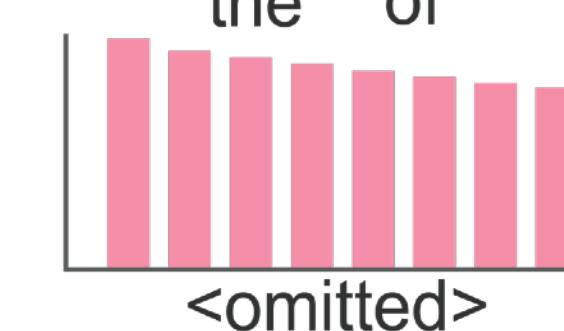
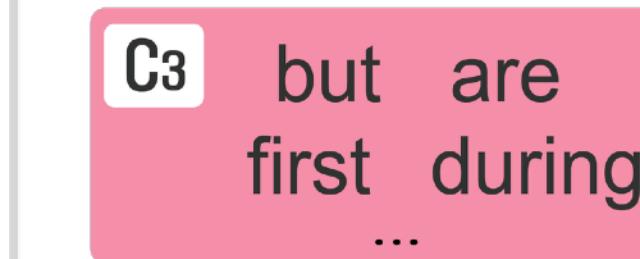
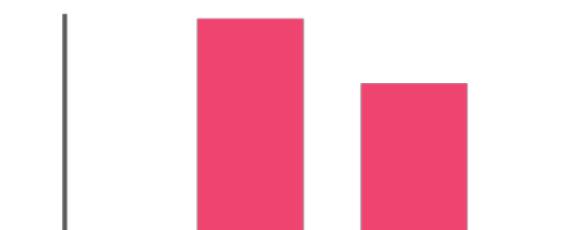
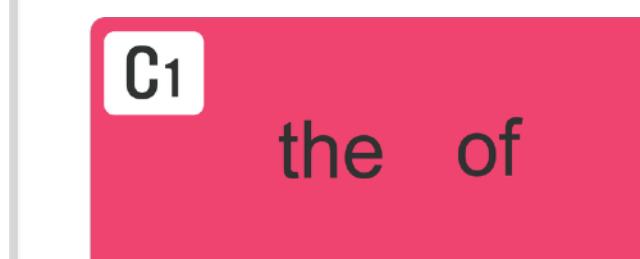
- Avoid likelihood issues by factorizing the softmax
- Initialize C frequency classes
- Distribute vocabulary into classes so that token frequency uniformly distributed across **and** between classes

(i) Unique tokens, sorted by frequency

the of and to in a be but are first during
only about game later three found music
role match way common sometimes decision king
mission organize scoring castle property curb ...



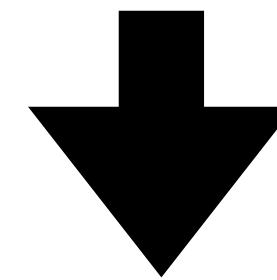
(ii) Assigning frequency class



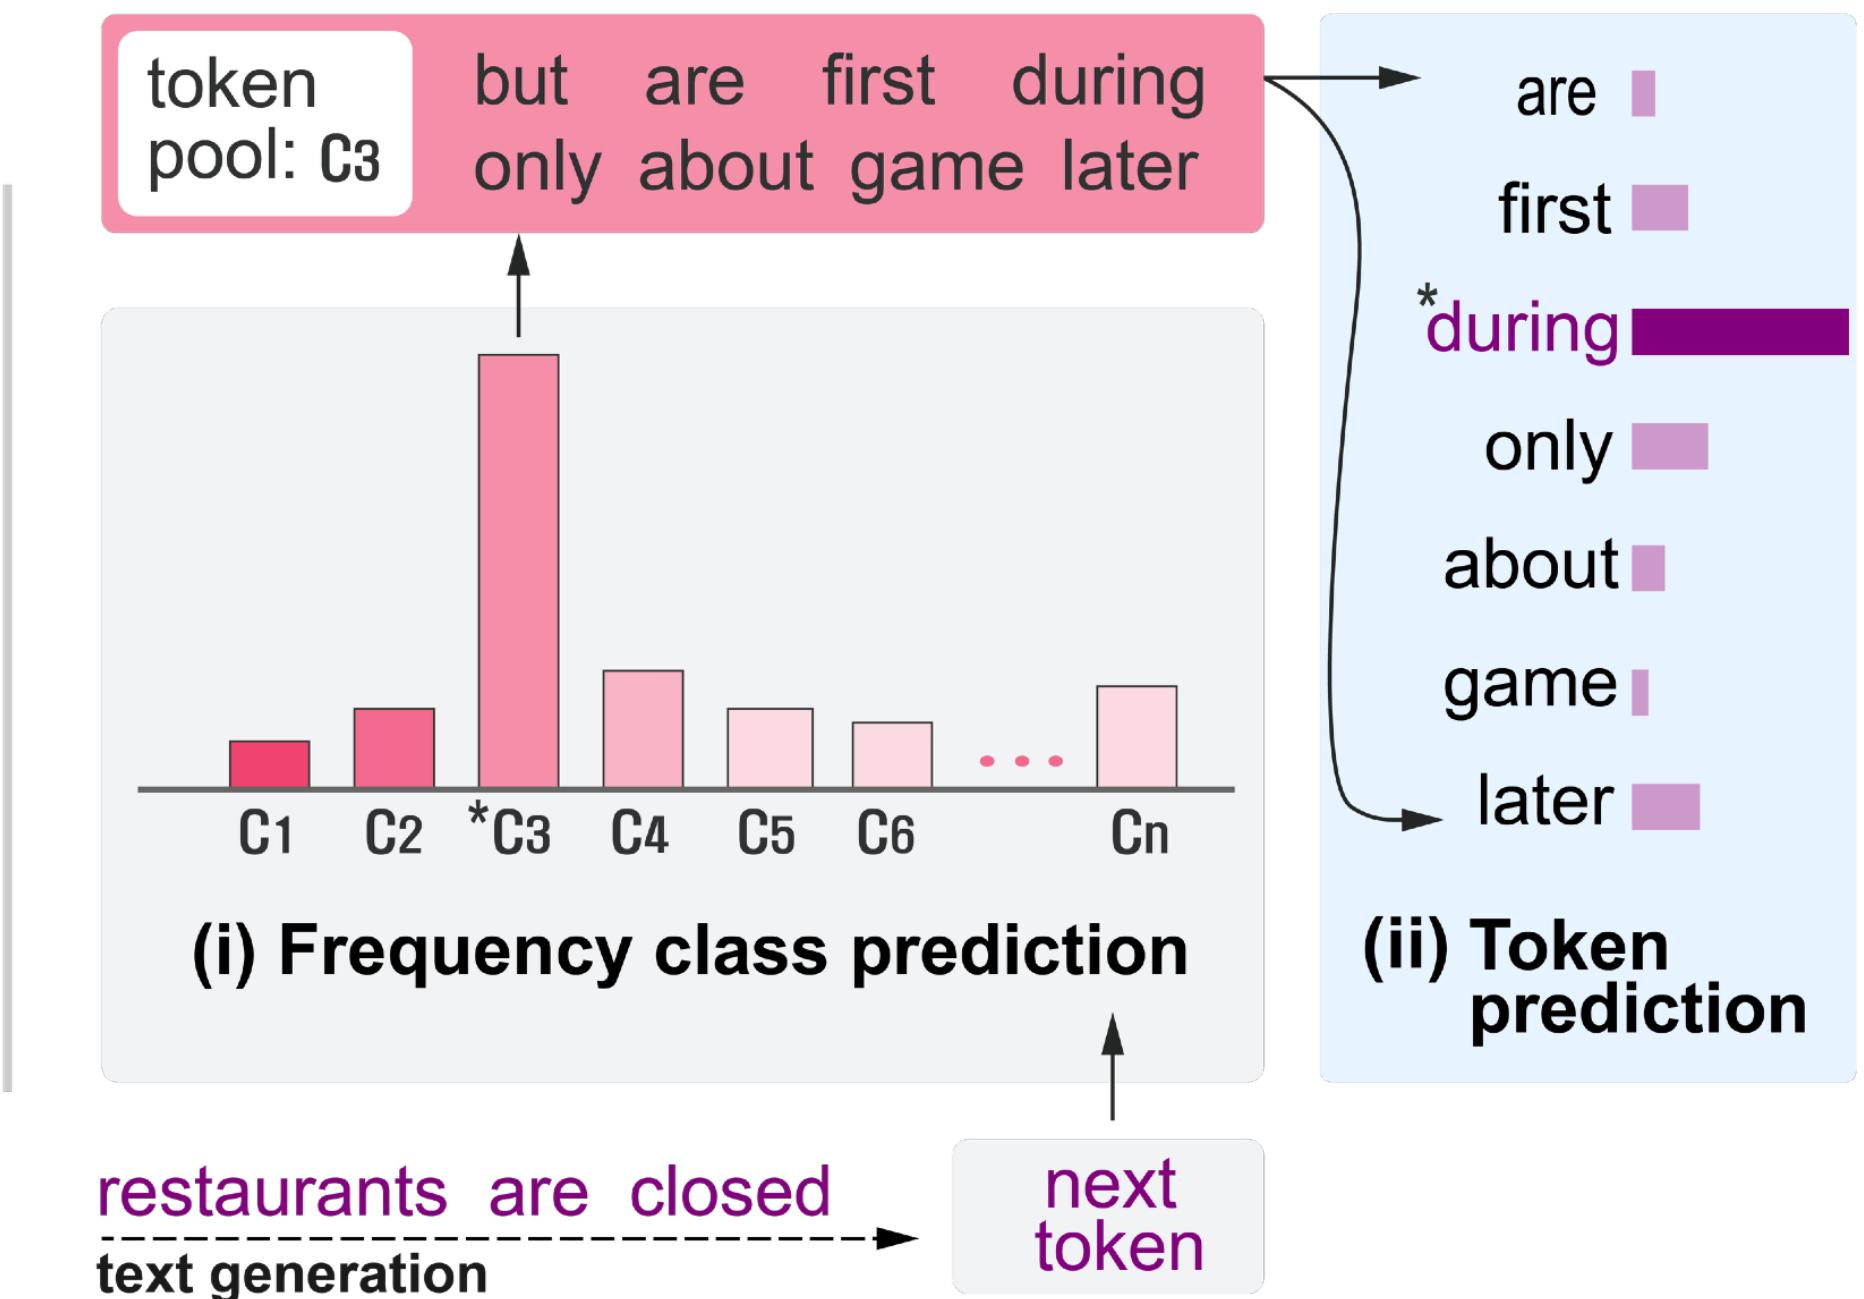
Alternatives: F² Softmax

- Learn to select both frequency class and vocabulary token during training

$$P(y_t = w_n | \{y\}_{<t}) = \frac{e^{U_n h}}{\sum_{m=1}^M e^{U_m h}}$$



$$P(y_t = w_n | \{y\}_{<t}) = \left(\frac{e^{V_f h}}{\sum_{c=1}^C e^{V_c h}} \right) \left(\frac{e^{U_n h}}{\sum_{m=1}^{M_f} e^{U_m h}} \right)$$



Issue #2: Exposure Bias

- During training, we condition on gold context tokens that are real human-generated text

$$\mathcal{L}_{MLE} = - \log P(y_t^* | \{y^*\}_{<t})$$

- During inference, we decode from distributions conditioned on previously generated tokens

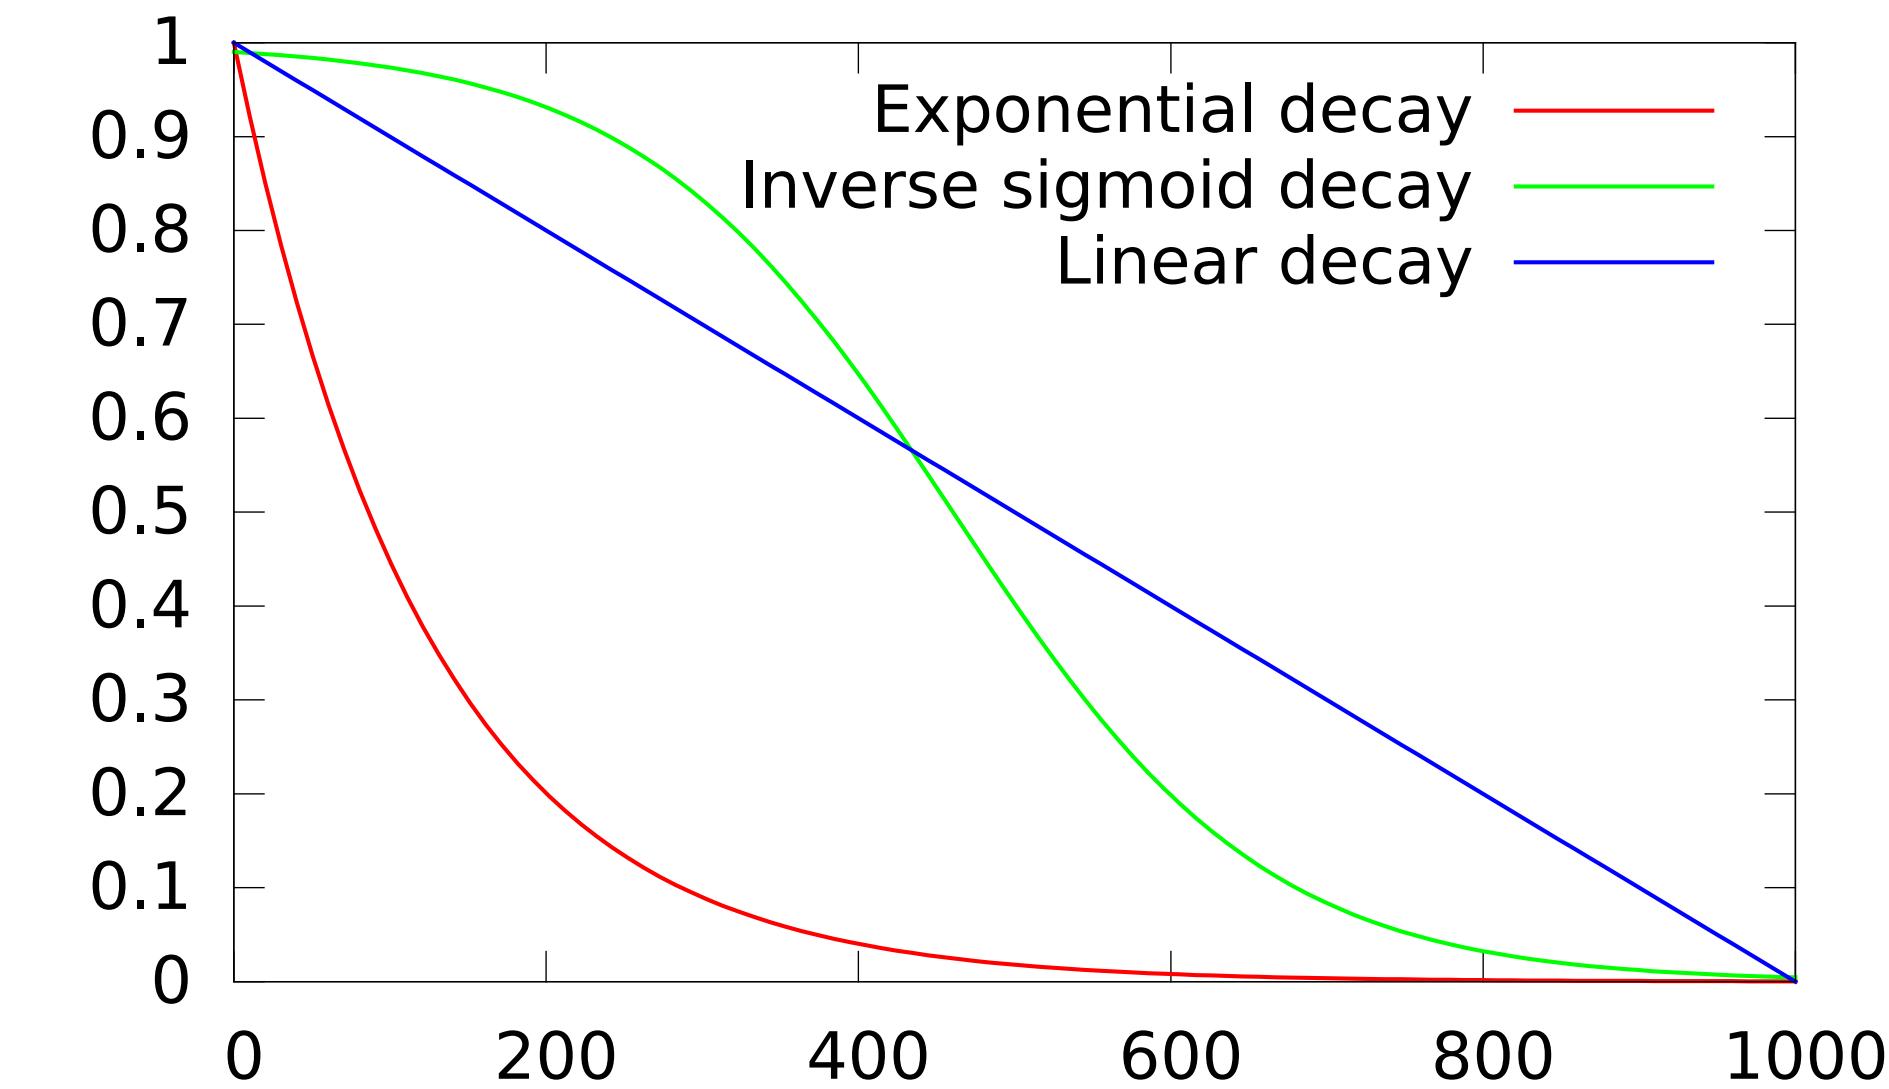
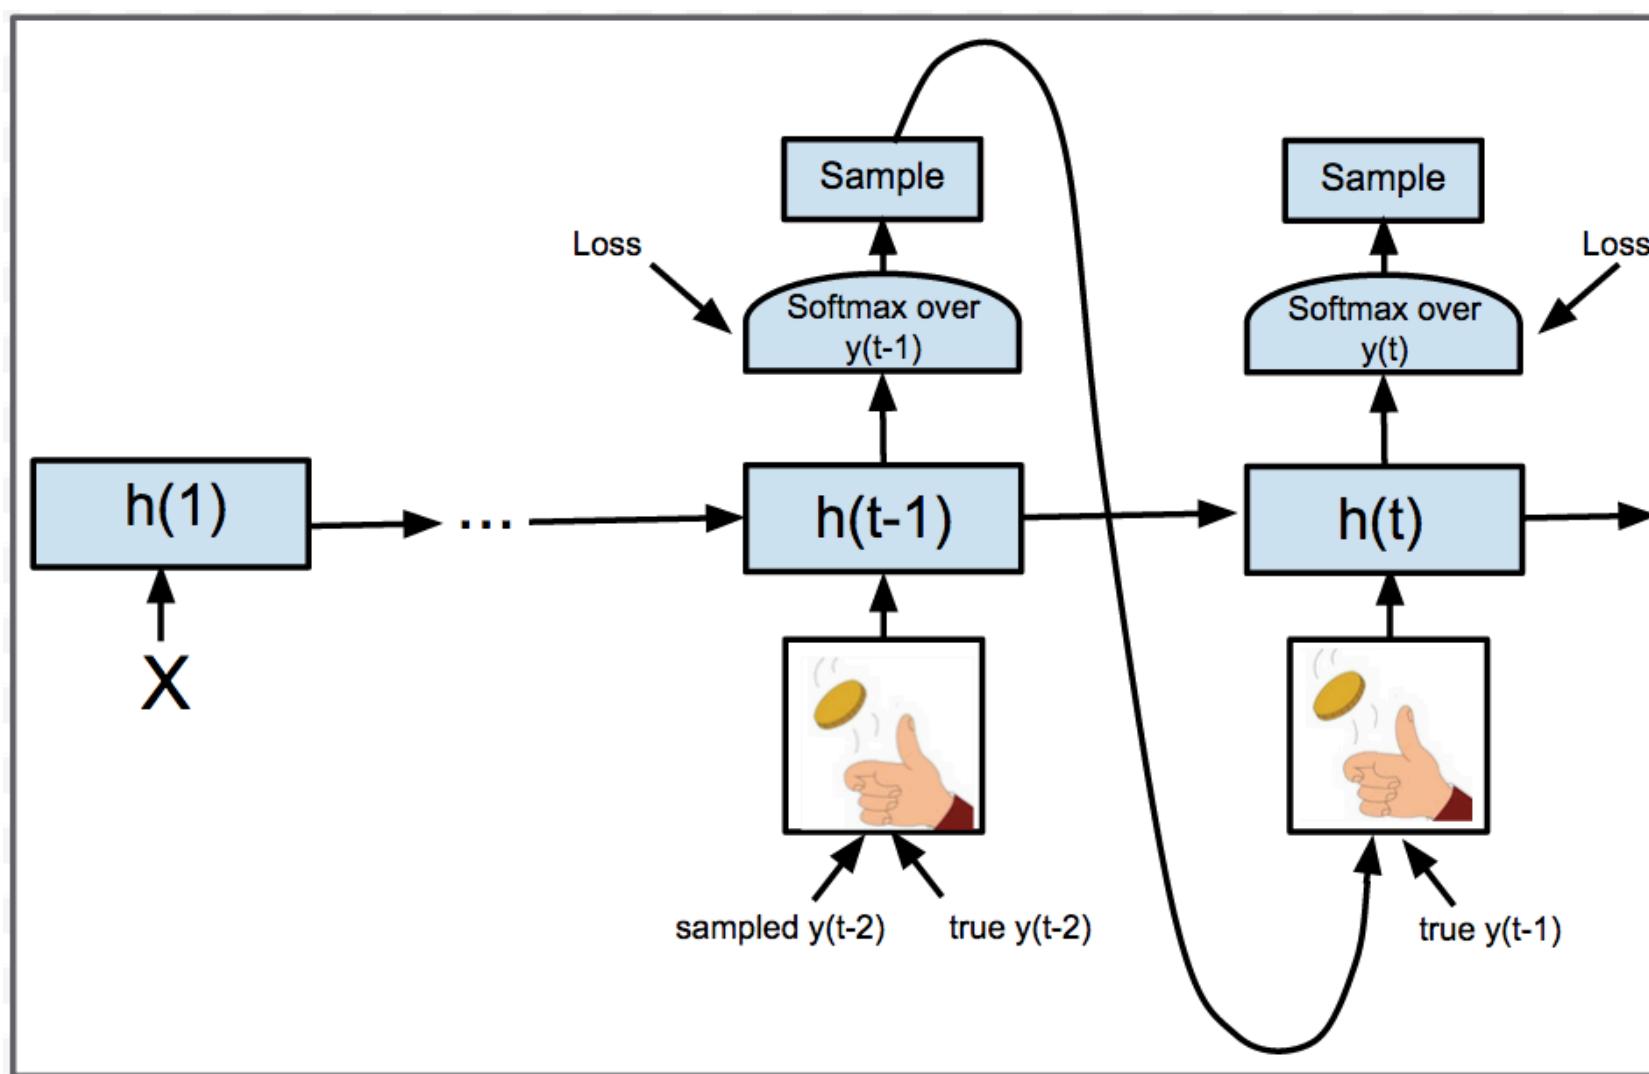
$$\mathcal{L}_{dec} = - \log P(\hat{y}_t | \{\hat{y}\}_{<t})$$

Alternative: Online Learning!

- Scheduled Sampling
- Dataset Aggregation (DAGGER)
- Reinforcement Learning
- Adversarial Learning

Alternative: Scheduled Sampling

- Barely used in newer transformer architectures, but relevant for RNNs
- After enough epochs of training with maximum likelihood estimation, interleave gold tokens and sampled tokens



Dataset Aggregation

1. Sample sequences according to your model
2. Add these sequences back to your training set
3. Re-train your model
4. Repeat

REINFORCE

- Sample sequences
- Use REINFORCE trick to scale losses by reward
- Define scoring metric as reward
- Set appropriate baseline
- Mix with

Adversarial Learning

- Similar to REINFORCE, but using discriminator as the reward function
-

Issue #3: Can we learn other behaviors?

- Can we tie discriminators to desired behaviors? Yes! If we can formalize as reward functions
- Cross-modal consistency (Ren et al., CVPR 2017)
- Simplicity (Zhang and Lapata, EMNLP 2017)
- Temporal consistency (Bosselut et al., NAACL 2018)
- Politeness (Tan et al., TACL 2018)
- Paraphrasing (Li et al., EMNLP 2018)
- Sentiment (Gong et al., NAACL 2019)
- Formality (Gong et al., NAACL 2019)