



HUGGING FACE



The Amazing World of Neural Language Generation

EMNLP 2020

November 20, 2020



Yangfeng Ji
University Of Virginia



Antoine Bosselut
Stanford University



Thomas Wolf
Huggingface



Asli Celikyilmaz
Microsoft Research

Natural Language Generation

- Sub-field of natural language processing (NLP)
- Building software systems to produce:
 - ***coherent, readable*** and **useful** written or spoken text.
- Produces explanations, summaries, narratives, etc.

Machine Translation



[User] We will discuss several issues today!

[System] Nous discuterons plusieurs questions aujourd'hui.

[System] 我们今天将讨论几个问题！

[System] Bugun çok sayıda sorunu tartışacağız.

Conversational Dialog Systems



[USER] Where is my next appointment and am I free for lunch?

[Agent] Your next meeting is at 10:30 at City Center. Did you want me to book a place for lunch in downtown ?

Understanding Visually Grounded Language



Start chopping
lettuce

Don't chop too
thin!

Not yet, start
with lettuce



When do I add
the fish ?

Image or Video Captioning



a man is showing how to make a drink



A car drives really fast

Multi Document Summarization



Summary

High Quality Content by WIKIPEDIA articles! Multi-document summarization is an automatic procedure aimed at extraction of information from multiple texts written about the same topic. Resulting summary report allows individual users, so as professional information consumers, to quickly familiarize themselves with information contained in a large cluster of documents. In such a way, multi-document summarization systems are complementing the news aggregators performing the next step down the road of coping with information overload.

Data-to-Text Generation

TEAM	WIN	LOSS	PTS	FG_PCT	RB	AST	...
Pacers	4	6	99	42	40	17	...
Celtics	5	4	105	44	47	22	...
PLAYER	H/V	AST	RB	PTS	FG	CITY	...
Jeff Teague	H	4	3	20	4	Indiana	...
Miles Turner	H	1	8	17	6	Indiana	...
Isaiah Thomas	V	5	0	23	4	Boston	...
Kelly Olynyk	V	4	6	16	6	Boston	...
Amir Johnson	V	3	9	14	4	Boston	...
...

PTS: points, FT_PCT: free throw percentage, RB: rebounds, AST: assists, H/V: home or visiting, FG: field goals, CITY: player team city.

The **Boston Celtics** defeated the host **Indiana Pacers 105-99** at Bankers Life Fieldhouse on Saturday. In a battle between two injury-riddled teams, the Celtics were able to prevail with a much needed road victory. The key was shooting and defense, as the **Celtics** outshot the **Pacers** from the field, from three-point range and from the free-throw line. Boston also held Indiana to **42 percent** from the field and **22 percent** from long distance. The Celtics also won the rebounding and assisting differentials, while tying the Pacers in turnovers. There were 10 ties and 10 lead changes, as this game went down to the final seconds. Boston (**5-4**) has had to deal with a glut of injuries, but they had the fortunate task of playing a team just as injured here. **Isaiah Thomas** led the team in scoring, totaling **23 points and five assists on 4-of-13** shooting. He got most of those points by going 14-of-15 from the free-throw line. **Kelly Olynyk** got a rare start and finished second on the team with his **16 points, six rebounds and four assists**.

Meeting Summarization



C: Looking at what we've got, we we want an LCD display with a spinning wheel.

B: You have to have some push-buttons, don't you?

C: Just spinning and not scrolling, I would say.

B: I think the spinning wheel is definitely very now.

A: but since LCDs seems to be uh a definite yes,

C: We're having push-buttons on the outside

C: and then on the inside an LCD with spinning wheel,

Decision Abstract (Summary):

The remote will have push buttons outside, and an LCD and spinning wheel inside.

A: and um I'm not sure about the buttons being in the shape of fruit though.

D: Maybe make it like fruity colours or something.

C: The power button could be like a big apple or something.

D: Um like I'm just thinking bright colours.

Problem Abstract (Summary):

How to incorporate a fruit and vegetable theme into the remote.

Conversational Dialog Summarization

[Nurse] Hi Mr.#name#, you were discharged on #date#. There are some questions i'd like to check with you.

[Patient] Ok, Ok.

[Nurse] Well, have you been experiencing swelling recently?

[Patient] Swelling? It comes and go, comes and go.

[Nurse] Comes and go ... I see .. **#repetition#**

[Nurse] ... **#pause#**... When did it start?

[Patient] Let me see, started from three weeks ago.

... ..

[Nurse] Are you experiencing any headache right now as we speak?

[Patient] Umm ... **#back-channel#**

[Nurse] Let me check, the last time you told me is sometimes at night.

[Patient] Oh, right, only a bit.

... ..

[Nurse] Still feel some chest pain or chest discomfort?

[Patient] Yes, my head is... **#false-start#** no, the pain is much better.

Still feel headache though ... **#topic-drift#**

... ..

[Nurse] Any giddiness or palpitation?

[Patient] Palpitation? Do not have-- **#interruption#**

[Nurse] Well ... Do you-- **#interruption#**

[Patient] and no giddiness, no, nothing.

... ..

[Nurse] Ok, you need to check your heartrate everyday.

[Nurse] Do you know how to use the device?

[Patient] Yes, yes, no problem.

... ..

Swelling: started from three weeks ago, comes and go.

Headache: sometimes, at night, only a bit.

Chest pain: much better.

Dizziness: none.

Other Text Generation Tasks



Question
Generation



Paraphrase
Generation



Poetry
Generation



Long Question
Answering



Visual Dialog
Systems



Document/Article
Generation



Program
Synthesis



Search Snippet
Generation

Why Automatic Text Generation?



Efficiency

Why Automatic Text Generation?



Efficiency



Education

Why Automatic Text Generation?



Efficiency



Education



Productivity

History of Natural Language Generation Systems

Template based systems:

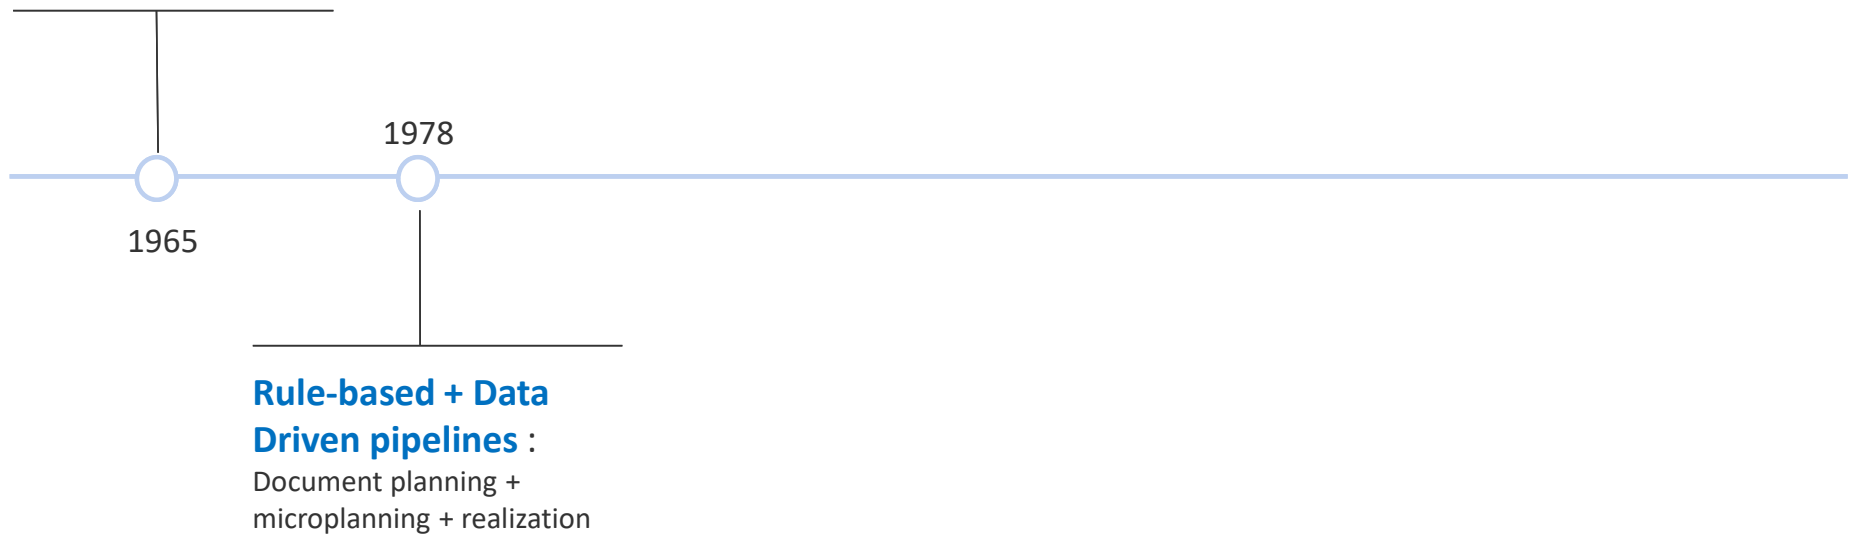
Uses rules and
templates



History of Natural Language Generation Systems

Template based systems:

Uses rules and templates



History of Natural Language Generation Systems

Template based systems:


Uses rules and templates

1965

Modeling Discourse Structures :

Relation learning,
Rhetorical Structure Theory

1985



1978

Rule-based + Data Driven pipelines :

Document planning +
microplanning + realization

History of Natural Language Generation Systems

Template based systems:

Uses rules and templates

1965

1978


Rule-based + Data Driven pipelines :

Document planning + microplanning + realization

1985

Modeling Discourse Structures :


Relation learning, Rhetorical Structure Theory



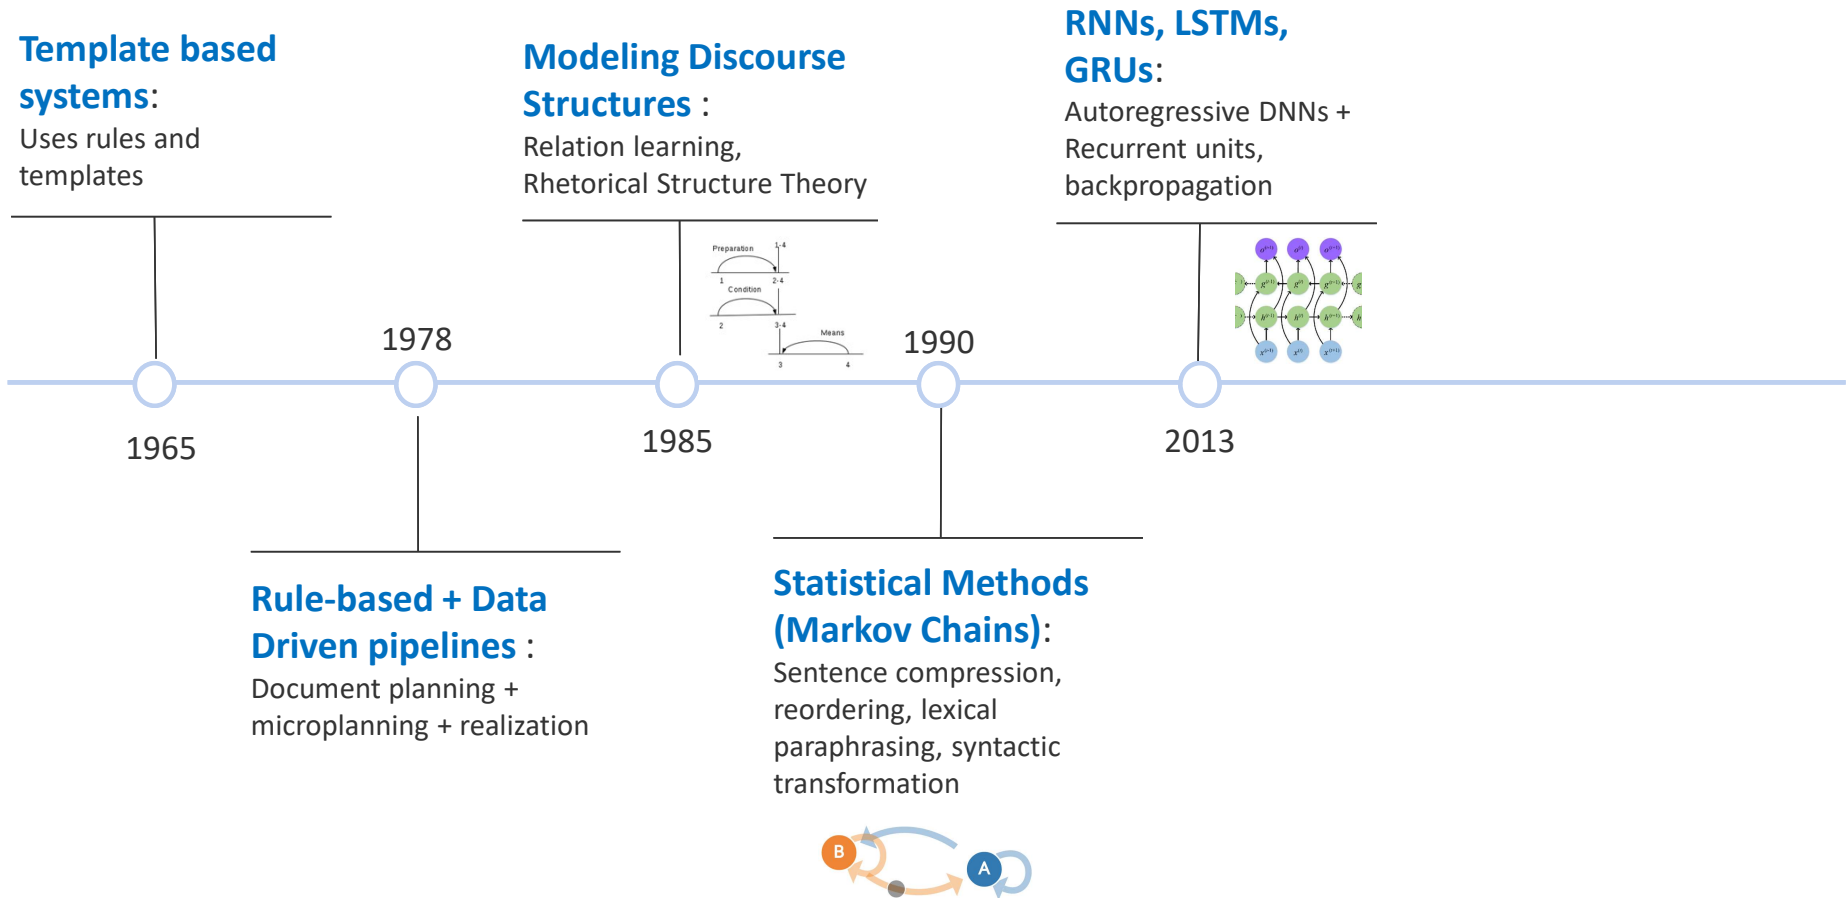
1990

Statistical Methods (Markov Chains):

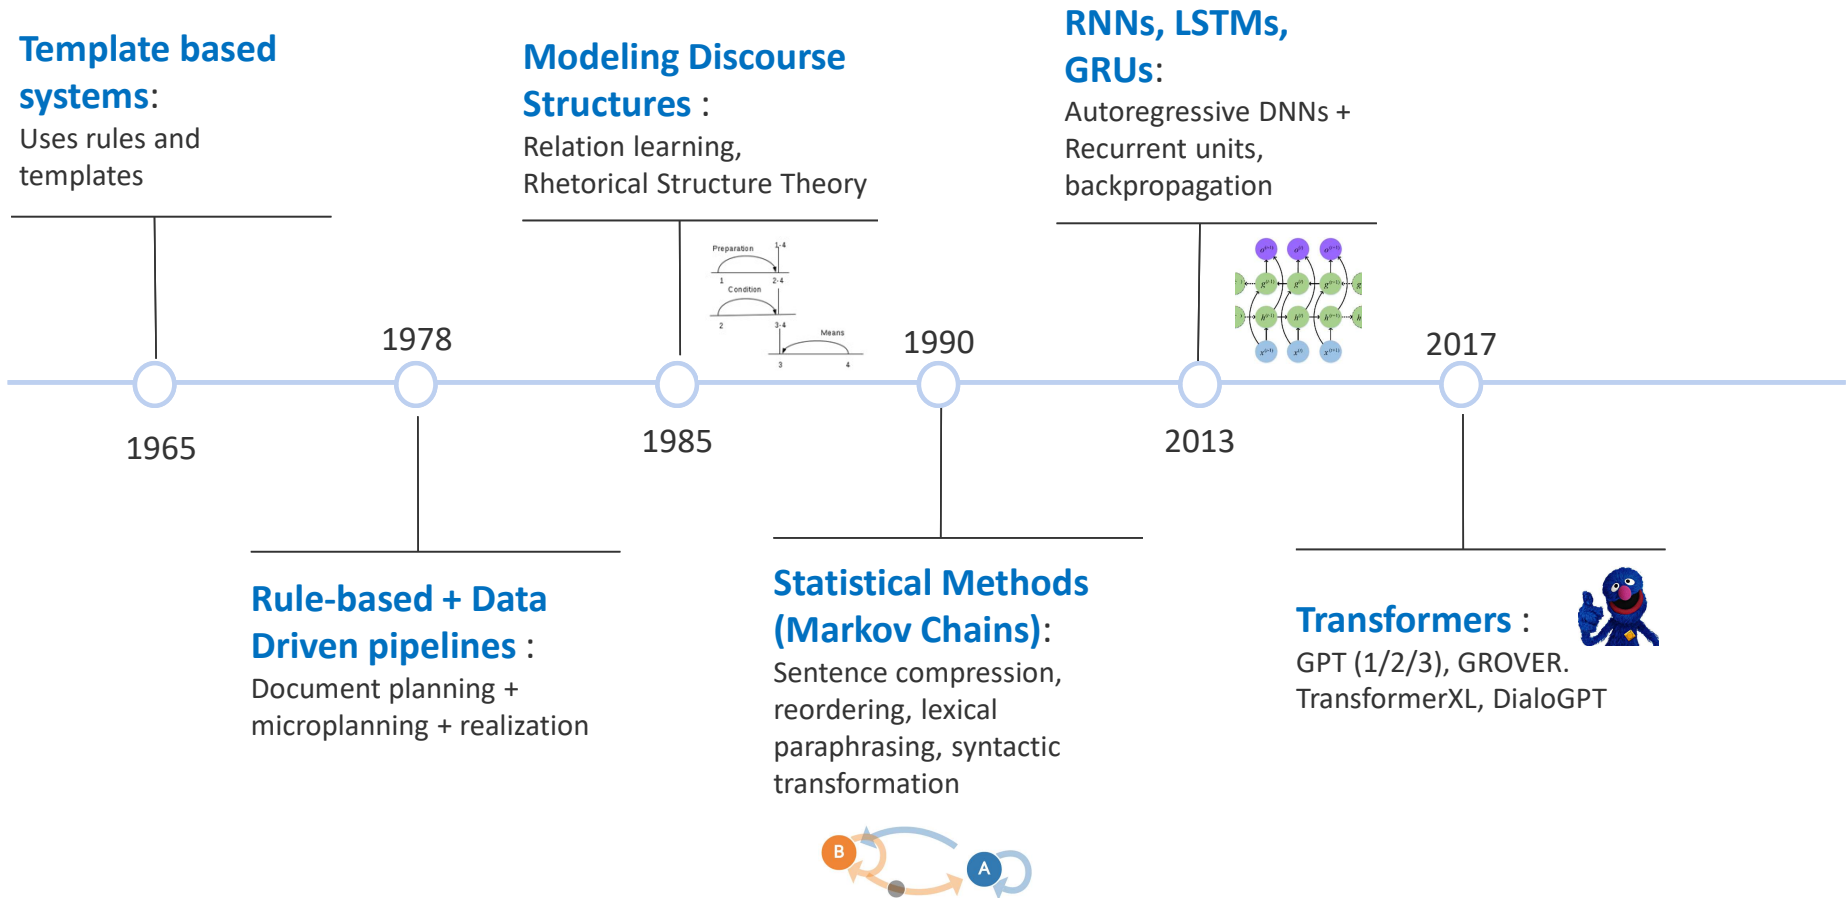
Sentence compression, reordering, lexical paraphrasing, syntactic transformation



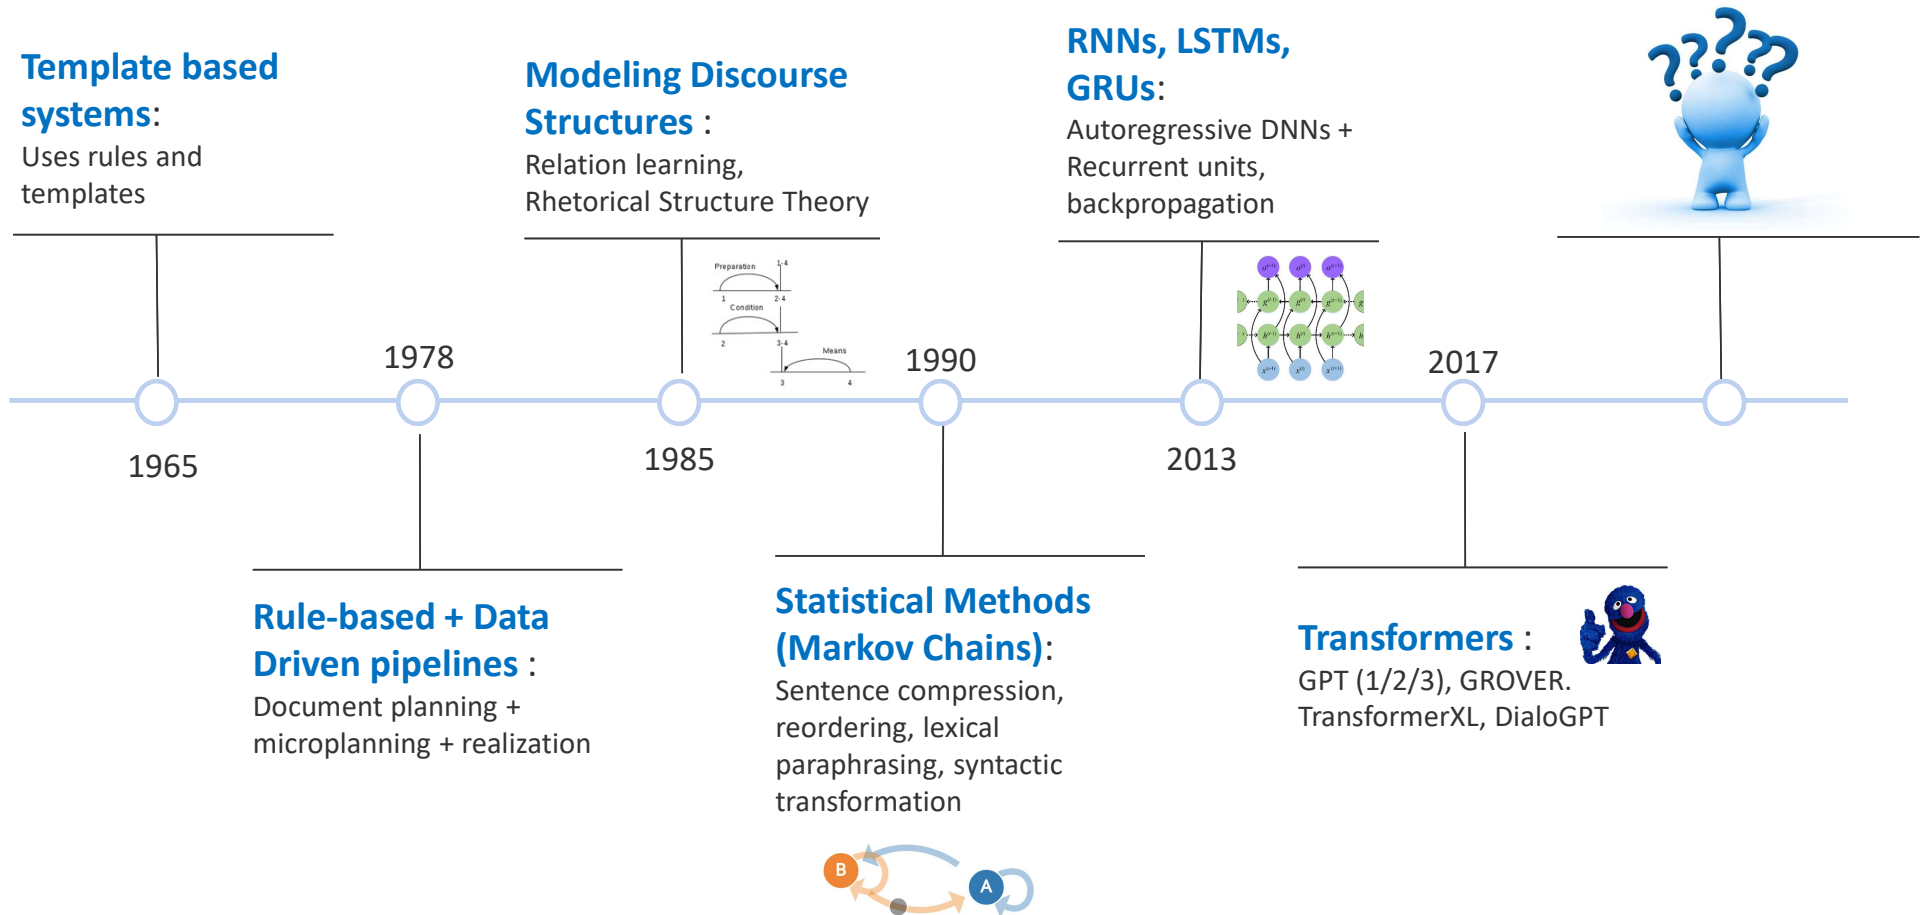
History of Natural Language Generation Systems



History of Natural Language Generation Systems



History of Natural Language Generation Systems



Language Modeling



Conditional Language Modelling

$$P(w_1, \dots, w_n) = \prod_{i=1}^n P(w_i | w_{i-1}, \dots, w_1)$$

Conditional Language Modelling

$$P(w_1, \dots, w_n) = \prod_{i=1}^n P(w_i | w_{i-1}, \dots, w_1)$$

- Speech Recognition
 - $P(\text{I saw a man}) \gg P(\text{eyes awe of an})$

Conditional Language Modelling

$$P(w_1, \dots, w_n) = \prod_{i=1}^n P(w_i | w_{i-1}, \dots, w_1)$$

- Speech Recognition
 - P(I saw a man) >> P(eyes awe of an)
- Spell Correction
 - P(about fifteen minutes from) > P(about fifteen minuets from)

Conditional Language Modelling

$$P(w_1, \dots, w_n) = \prod_{i=1}^n P(w_i | w_{i-1}, \dots, w_1)$$

- Speech Recognition
 - $P(\text{I saw a man}) \gg P(\text{eyes awe of an})$
- Spell Correction
 - $P(\text{about fifteen minutes from}) > P(\text{about fifteen minuets from})$
- Machine translation, Question Answering, Paraphrasing, Image captioning, Summarization, others...

Long Text Generation – Image Story Telling



Two men are riding bicycles down a street. One man is behind the other riding on the side of a paved road. Both bicycles have black baskets on the front attached to the handlebars. Both men are holding umbrellas. The first man is wearing blue pants and a white and black striped shirt. His umbrella is blue. The second man is wearing black pants and a light purple shirt. His umbrella is aqua blue colored. Behind them, along the road are shops. One of the shops is closed with a silver metal fence covering the entrance. The other shop is a clothing store and there are four mannequins with clothing outside the door. There is also a pot-ted plant with red flowers outside one of the shops.

Long Text Generation – Image Story Telling



Two men are riding bicycles down a street. One man is behind the other riding on the side of a paved road. Both bicycles have black baskets **on the front** attached to the handlebars. Both men are holding umbrellas. The first man is wearing blue pants and a white and black striped shirt. **His** umbrella is blue. **The second man** is wearing black pants and a light purple shirt. **His** umbrella is aqua blue colored. **Behind** them, along the road are shops. **One of** the shops is closed with a silver metal fence covering the entrance. **The other** shop is a clothing store and there are four mannequins with clothing **outside** the door. There is also a pot-ted plant with red flowers **outside** one of the shops.

Long Text Generation with RNN-LM




two men are **riding a bicycle** on the road. it is a **sunny day**. a man is **riding a blue bicycle** on the street. the men with blue umbrella is **riding a bicycle** on a **rainy day**. the woman black dress is standing on the road.

Neural Text Generation Issues – Model Selection


Neural Text Generation Issues – Model Selection

Exposure Bias



Neural Text Generation Issues – Model Selection


Exposure Bias



Compounding Errors / Label Bias


Gold: The cat purrs

Pred: The **dog** barks



Neural Text Generation Issues – Model Selection


Exposure Bias



Compounding Errors / Label Bias

Gold: The cat purrs

Pred: The **dog** barks



Coherence and Narrative Flow

“do you like animals?”


yes, I have three cats.

how many cats do you have ?

I don't have cats.

Neural Text Generation Issues – Model Selection


Exposure Bias



Compounding Errors / Label Bias

Gold: The cat purrs

Pred: The **dog** barks



Coherence and Narrative Flow

“do you like animals?”

yes, I have three cats.

how many cats do you have ?

I don't have cats.

Surrogate objective function

$$P^*(x|c) \rightarrow P_\theta(x|c)$$

Perplexity

vs

BLEU/ROUGE/Neural Rewards

Weaknesses of MEGA Language Models for **GENERATION!**

Inconsistent
output

Maintaining **coherence**
between **paragraphs**

MLE: cannot capture high-level semantics

Crippled by length

Biased pre-trained models

coreference issues

No **real** understanding,
commonsense, **factual**
correctness

Domain transfer is hard

How to learn discourse?

Single path generative flow

Longer string that are
repeated many times in
the dataset

Lack of implicit
“**planning**”


Softmax Bottleneck issues!

Unnecessarily
repeating entities

sub-optimal evaluation metrics
“perplexity”

Surrogate-loss functions

Auto-regressive!



exposure-bias

Degenerative sampling
methods

word-by-word
generation: can't see
global context!


Research on Evaluation: Now more than ever!



Detecting machine generated text



Standards in Text Generation Evaluations




Effective Evaluation




Evaluating Ethical Issues


Epilepsy Detection Model with Brain MRI Data



Brain MRI data



Complex ML model



Making Evaluation Explainable

Tutorial Schedule

Approximate Time	Segment Title	Speaker	Description
15 mins	Introduction	Asli Celikyilmaz	Why is it important today to discuss neural text generation?
20 mins	Modeling: NN Architecture	Yangfeng Ji	Journey of NNs architectures that are used for text generation up-to-date
20 mins	Modeling: Generation With Rich context	Yangfeng Ji	How to efficiently use context in language generation?
20 mins	Training Text Generation Models	Antoine Bosselut	What are the best practices in training neural text generation systems today?
20 mins	Decoding Algorithms	Antoine Bosselut	Amazing world of neural decoding algorithms and beyond!
35 mins	Evaluation and Benchmarks	Asli Celikyilmaz	How are neural text generation models evaluated today? Future directions!
40 mins	Building NLG Systems	Thomas Wolf	Adventure into building an NLG system using state-of-the art tools and libraries.

Introduction and Evaluation



“Let’s discuss the Evaluation of Text Generation Systems, which is more important today than ever!”

Time:

Modeling



Training and Decoding



Building of NLG Systems

